

# Implementing Threads 3

# Blocking Locks

```
void blocking_lock(mutex_t *mut) {  
    if (mut->holder != 0) {  
        enqueue(mut->wait_queue,  
                CurrentThread);  
        uthread_switch();  
    } else  
        mut->holder = CurrentThread;  
}
```

```
void blocking_unlock(mutex_t *mut) {  
    if (queue_empty(mut->wait_queue))  
        mut->holder = 0;  
    else {  
        mut->holder =  
            dequeue(mut->wait_queue);  
        enqueue(RunQueue, mut->holder);  
    }  
}
```

Does it work?

# Working Blocking Locks (?)

```
void blocking_lock(mutex_t *mut) {
    spin_lock(&mut->spinlock);
    if (mut->holder != 0) {
        enqueue(mut->wait_queue,
                CurrentThread);
        spin_unlock(&mut->spinlock);
        uthread_switch();
    } else {
        mut->holder = CurrentThread;
        spin_unlock(&mut->spinlock);
    }
}
```

## Quiz 1

**This**

- a) always works
- b) sometimes doesn't work
- c) never works

```
void blocking_unlock(mutex_t *mut) {
    spin_lock(&mut->spinlock);
    if (queue_empty(
        mut->wait_queue)) {
        mut->holder = 0;
    } else {
        mut->holder =
            dequeue(mut->wait_queue);
        enqueue(RunQueue,
                mut->holder);
    }
    spin_unlock(&mut->spinlock);
}
```

# Futexes

- **Safe, *efficient* kernel conditional queueing in Linux**
- **All operations performed atomically**
  - `futex_wait(futex_t *futex, int val)`
    - **if `futex->val` is equal to `val`, then sleep**
    - **otherwise return**
  - `futex_wake(futex_t *futex)`
    - **wake up one thread from `futex`'s wait queue, if there are any waiting threads**

# Ancillary Functions

- `int atomic_inc(int *val)`
  - add 1 to `*val`, return its original value
- `int atomic_dec(int *val)`
  - subtract 1 from `*val`, return its original value

# Attempt 1

```
void lock(futex_t *futex) {  
    int c;  
    while ((c = atomic_inc(&futex->val)) != 0)  
        futex_wait(futex, c+1);  
}
```

```
void unlock(futex_t *futex) {  
    futex->val = 0;  
    futex_wake(futex);  
}
```

# Attempt 2

```
void lock(futex_t *futex) {
    int c;
    if ((c = CAS(&futex->val, 0, 1) != 0)
        do {
            if (c == 2 || (CAS(&futex->val, 1, 2) != 0))
                futex_wait(futex, 2);
            while ((c = CAS(&futex->val, 0, 2)) != 0))
        }

void unlock(futex_t *futex) {
    if (atomic_dec(&futex->val) != 1) {
        futex->val = 0;
        futex_wake(futex);
    }
}
```

**Quiz 2**  
**Does it work?**

- a) No**
- b) Yes**

# Blocking Locks in MThreads

- We could use futexes, but don't
- *uthread\_switch* gets an additional argument
  - a POSIX mutex (representing a spin lock)
  - unlock it after getting out of the context of the calling thread




# Actual Code

```
uthread_mtx_lock(uthread_mtx_t *mtx) {  
    uthread_nopreempt_on();  
    pthread_mutex_lock(&mtx->m_pmut);  
    if (mtx->m_owner == NULL) {  
        mtx->m_owner = ut_curthr;  
        pthread_mutex_unlock(&mtx->m_pmut);  
        uthread_nopreempt_off();  
    } else {  
        ut_curthr->ut_state = UT_WAIT;  
        uthread_switch(&mtx->m_waiters, 0, &mtx->m_pmut);  
        uthread_nopreempt_off();  
    }  
}
```


# MP Memory Issues

- **Naive view is that all processors in MP system see same memory contents at all times**
  - they don't

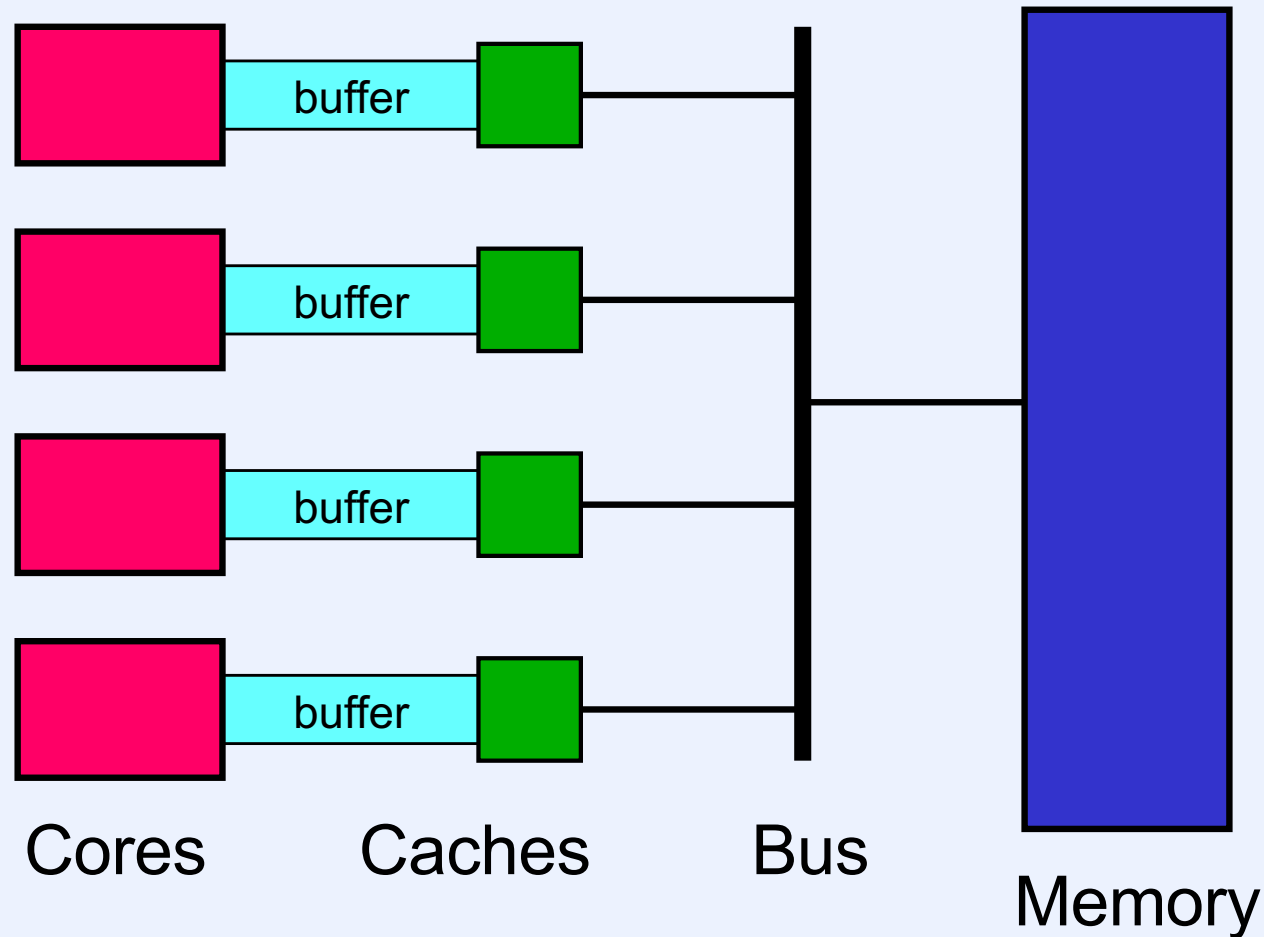
# Multi-Core Processor: Simple View



# Multi-Core Processor: More Realistic View



# Multi-Core Processor: Even More Realistic



# Concurrent Reading and Writing

**Thread 1:**

```
i = shared_counter;
```

**Thread 2:**

```
shared_counter++;
```

# Mutual Exclusion w/o Mutexes

```
void peterson(long me) {  
    static long loser;           // shared  
    static long active[2] = {0, 0}; // shared  
    long other = 1 - me;        // private  
    active[me] = 1;  
    loser = me;  
    while (loser == me && active[other])  
        ;  
    // critical section  
    active[me] = 0;  
}
```

## Quiz 3

**With delayed stores**

- a) works
- b) doesn't work

# Busy-Waiting Producer/Consumer

```
void producer(char item) {  
  
    while(in - out == BSIZE)  
        ;  
  
    buf[in%BSIZE] = item;  
  
    in++;  
}
```

## Quiz 4

**With re-ordered stores**

- a) works**
- b) doesn't work**


```
char consumer( ) {  
    char item;  
    while(in - out == 0)  
        ;  
  
    item = buf[out%BSIZE];  
  
    out++;  
  
    return(item);  
}
```



# Coping


- **Use what's available in the architecture to make sure all cores have the same view of memory (when necessary)**
  - lock prefix on x86
  - mfence x86 instruction
- **Use the synchronization primitives**
  - presumably the implementers knew what they were doing

# One-Level Model

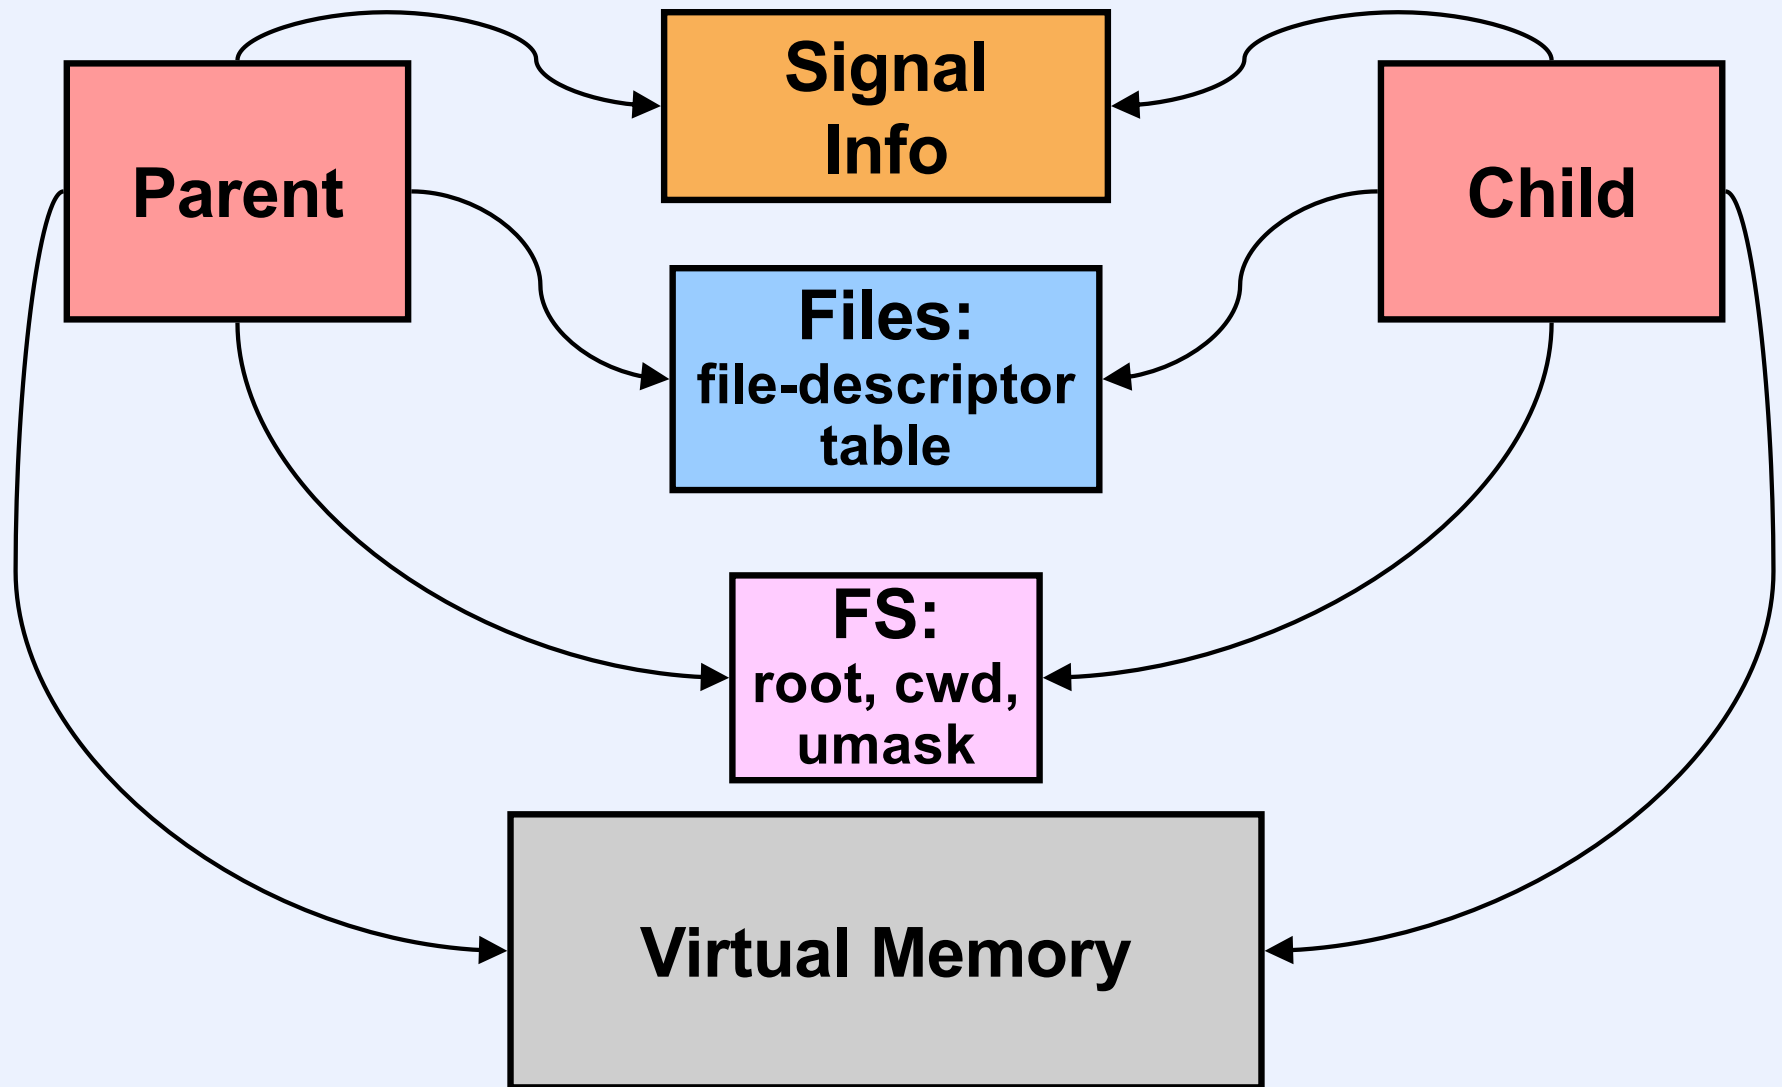


# Variable-Weight Processes


- Variant of one-level model
- Portions of parent process selectively *copied* into or *shared* with child process
- Children created using *clone* system call



# Cloning



# Linux Threads (pre 2.6)




# NPTL in Linux 2.6+


- **Native POSIX-Threads Library**
  - full POSIX-threads semantics on improved variable-weight processes
    - threads of a “process” form a *thread group*
      - *getpid()* returns process ID of first thread in group
      - any thread in group can wait for any other to terminate
      - signals to process delivered by kernel to any thread in group

# Two-Level Model

## One Kernel Thread



# Two-Level Model: Multiple Kernel Threads






# Quiz 5

**One kernel thread for each user thread is clearly a sufficient number of kernel threads in the two-level model. Is it necessary for maximum concurrency?**

- a) there are no situations in which that number of threads is necessary, as long as there are at least as many kernel threads as processors.**
- b) there must always be that number of kernel threads for the two-level model to work well.**
- c) there are situations in which that number is necessary, but they occur rarely.**


# Deadlock



# MThreads

- **Two-level threads implementation of Uthreads**
  - kernel-supported threads are POSIX threads
  - user threads based on your implementation of Uthreads
- **Effectively a multiprocessor implementation**
  - use POSIX mutexes rather than spin locks
  - use POSIX condition variables rather than the idle loop

# Two-Level Model: MThreads



# Synchronizing LWPs

```
uthread_switch(...) {  
    uthread_mtx_lock(&runq_mtx)  
    volatile int first = 1;  
    getcontext(&ut_curthr->ut_ctx);  
    if (!first) {  
        ...  
    }  
    setcontext(&curlwp->lwp_ctx);  
}  
lwp_switch() {  
    ...  
    ut_curthr = top_priority_thread(&runq);  
    uthread_mtx_unlock(&runq_mtx);  
    setcontext(&ut_curthr->ut_ctx);  
    ...  
}
```

# Synchronizing LWPs (2)

```
uthread_switch(...) {  
    spin_lock(&runq_mtx)  
    volatile int first = 1;  
    getcontext(&ut_curthr->ut_ctx);  
    if (!first) {  
        ...  
    }  
    setcontext(&curlwp->lwp_ctx);  
}  
lwp_switch() {  
    ...  
    ut_curthr = top_priority_thread(&runq);  
    spin_unlock(&runq_mtx);  
    setcontext(&ut_curthr->ut_ctx);  
    ...  
}
```

---

# Synchronizing LWPs (3)

```
uthread_switch(...) {  
    pthread_mutex_lock(&runq_mtx)  
    volatile int first = 1;  
    getcontext(&ut_curthr->ut_ctx);  
    if (!first) {  
        ...  
    }  
    setcontext(&curlwp->lwp_ctx);  
}  
lwp_switch() {  
    ...  
    ut_curthr = top_priority_thread(&runq);  
    pthread_mutex_unlock(&runq_mtx);  
    setcontext(&ut_curthr->ut_ctx);  
    ...  
}
```

---

# POSIX Mutexes and MThreads

- **POSIX mutexes used to synchronize activity among LWPs**
- **Problem case**
  - **uthread (running on LWP) locks mutex**
  - **clock interrupt occurs, uthread yields LWP to another uthread**
  - **that uthread (running on same LWP) locks same mutex**
  - **deadlock: LWP attempting to lock mutex it currently has locked**
- **Solution**
  - **mask interrupts while thread has mutex locked**



# Example

```
void uthread_wake(uthread_t *uthr) {  
  
    pthread_mutex_lock(&runq_mtx);  
  
    ...  
  
    // wake up thread, put it on runq  
  
    ...  
  
    pthread_mutex_unlock(&runq_mtx);  
  
}
```

# Example: Fixed

```
void uthread_wake(uthread_t *uthr) {  
    uthread_nopreempt_on();  
    pthread_mutex_lock(&runq_mtx);  
  
    ...  
  
    // wake up thread, put it on runq  
  
    ...  
  
    pthread_mutex_unlock(&runq_mtx);  
    uthread_nopreempt_off();  
}
```

# Thread-Local Storage in Mthreads

- `__thread thread_t *ut_curthr;`
  - **reference to the current uthread**
- `__thread lwp_t *curlwp`
  - **reference to the current LWP (POSIX thread)**
- **Thread-Local Storage accesses are not async-signal safe!**
- **Must turn off preemption while using TLS**
  - **otherwise thead could be preempted and later resumed on another LWP**
  - **TLS pointer refers to the wrong item!**