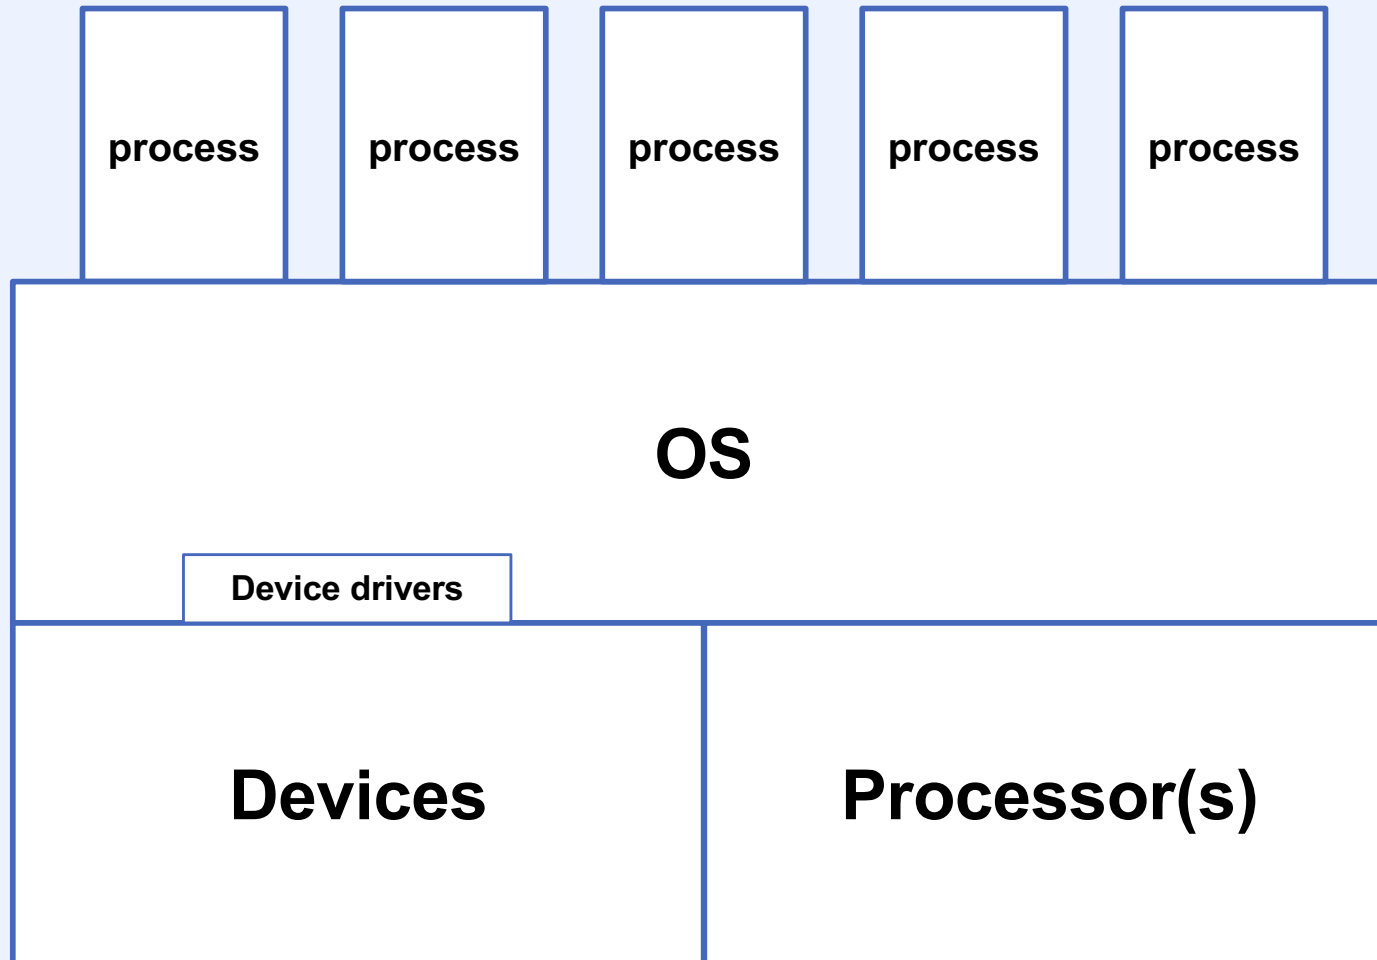


Virtual Machines

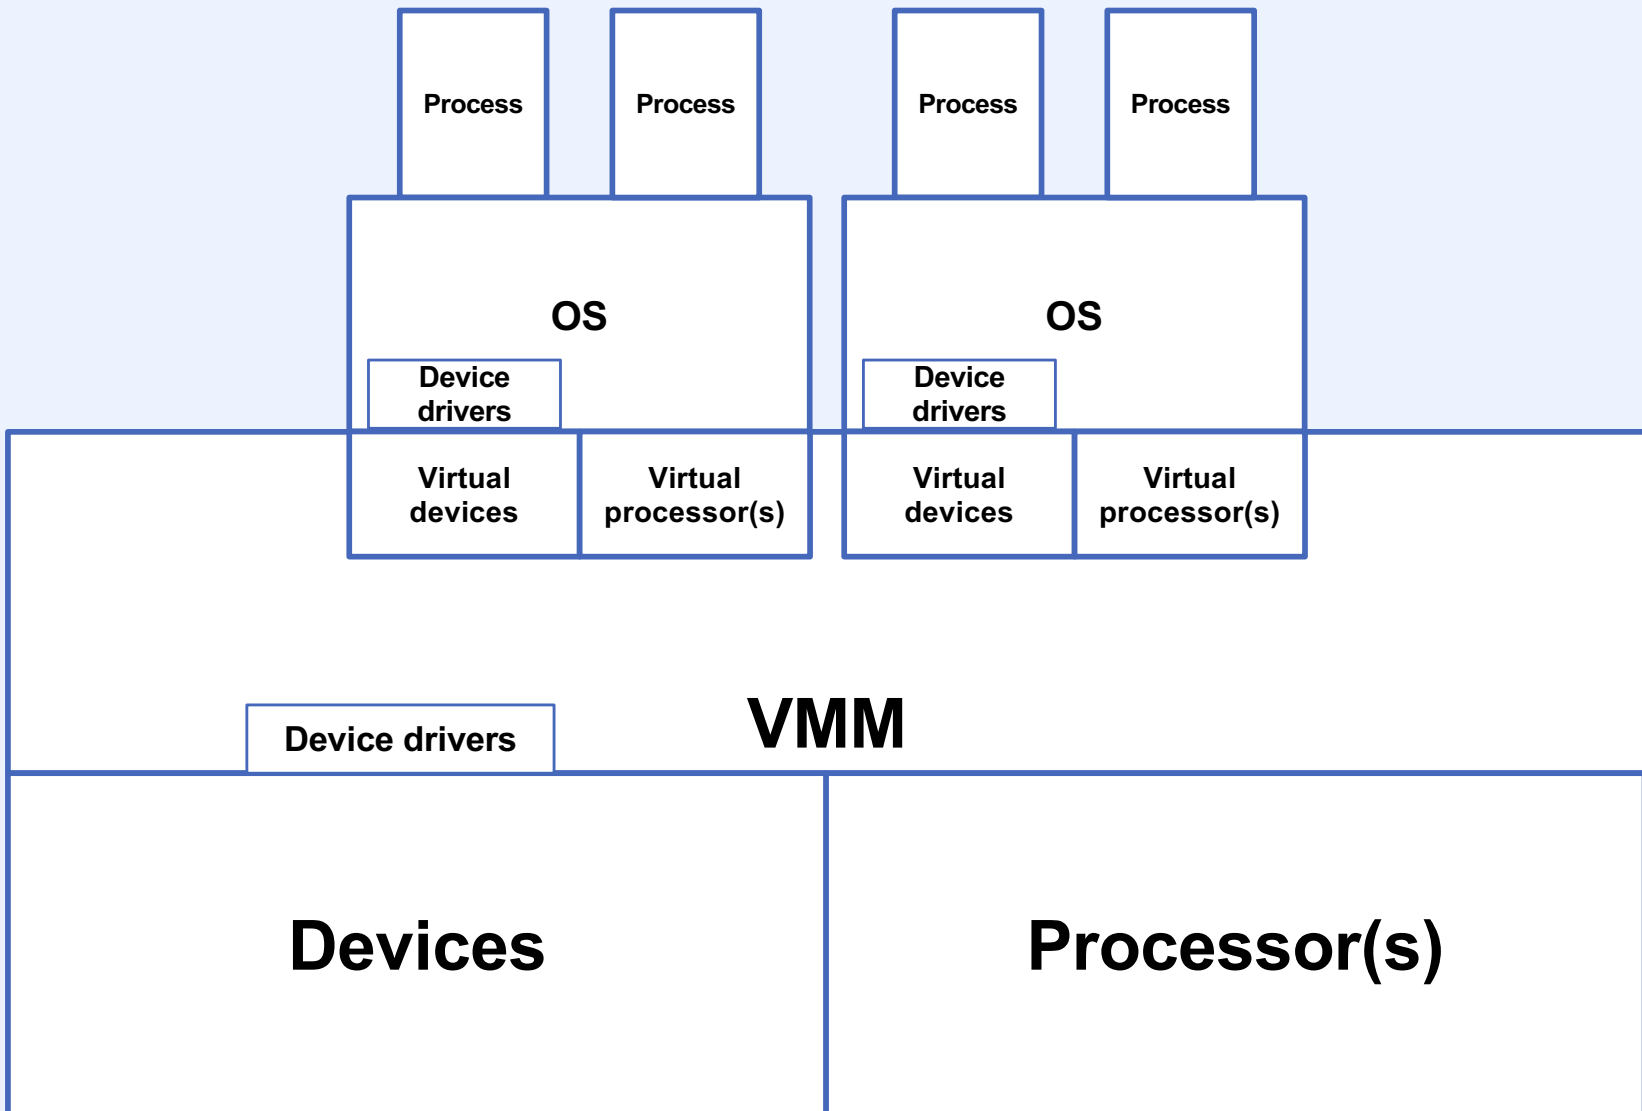
Part 2: starting ~20 years ago

(continued)

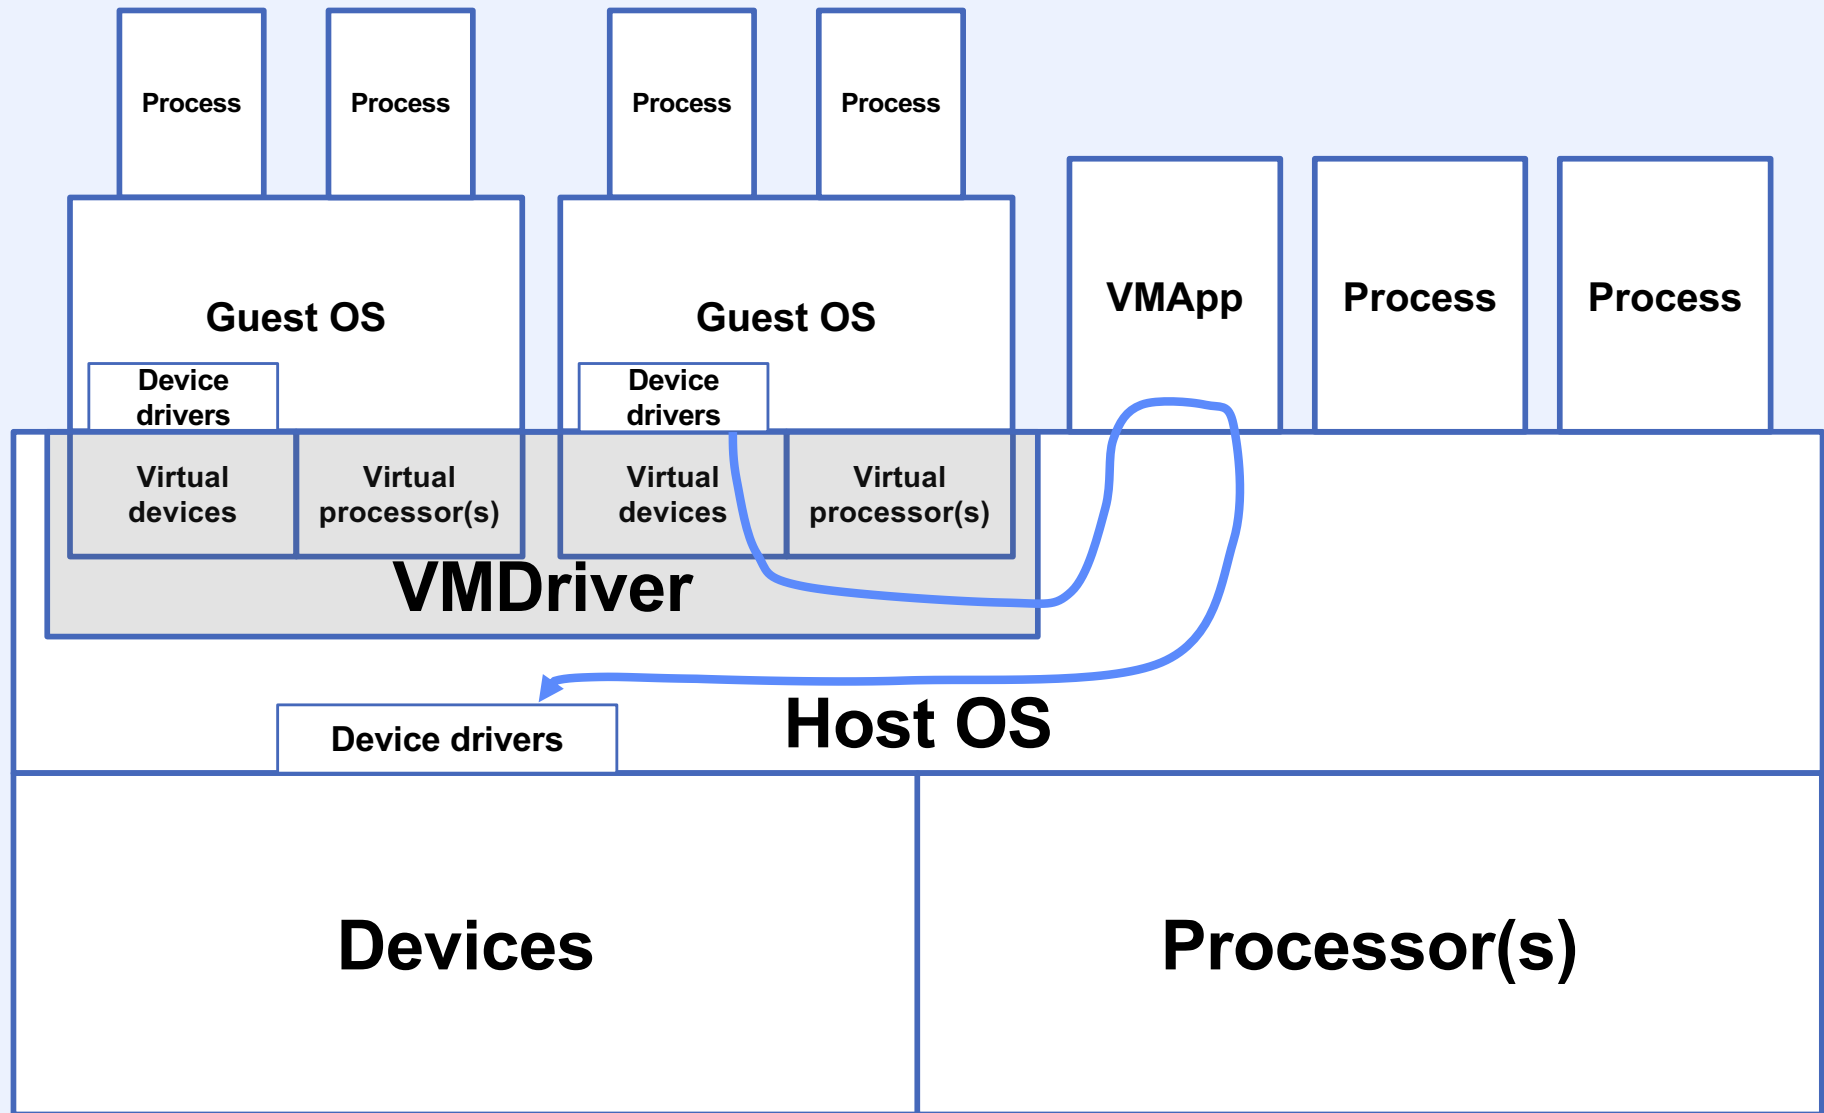
Real-Machine OS Structure



On a Virtual Machine ...



VMware Workstation



Quiz 1

The host OS contains a scheduler that multiplexes the execution of threads on its processors. Each guest OS contains a scheduler that does the same thing.

**Assume all threads have the same priority.
Assume just one real processor.**

- a) Threads on the guest OS compete equally for the real processor as do threads on the host OS**
 - b) Threads on the guest OS effectively have higher priority than host OS threads**
 - c) Threads on the guest OS effectively have lower priority than host OS threads**
-

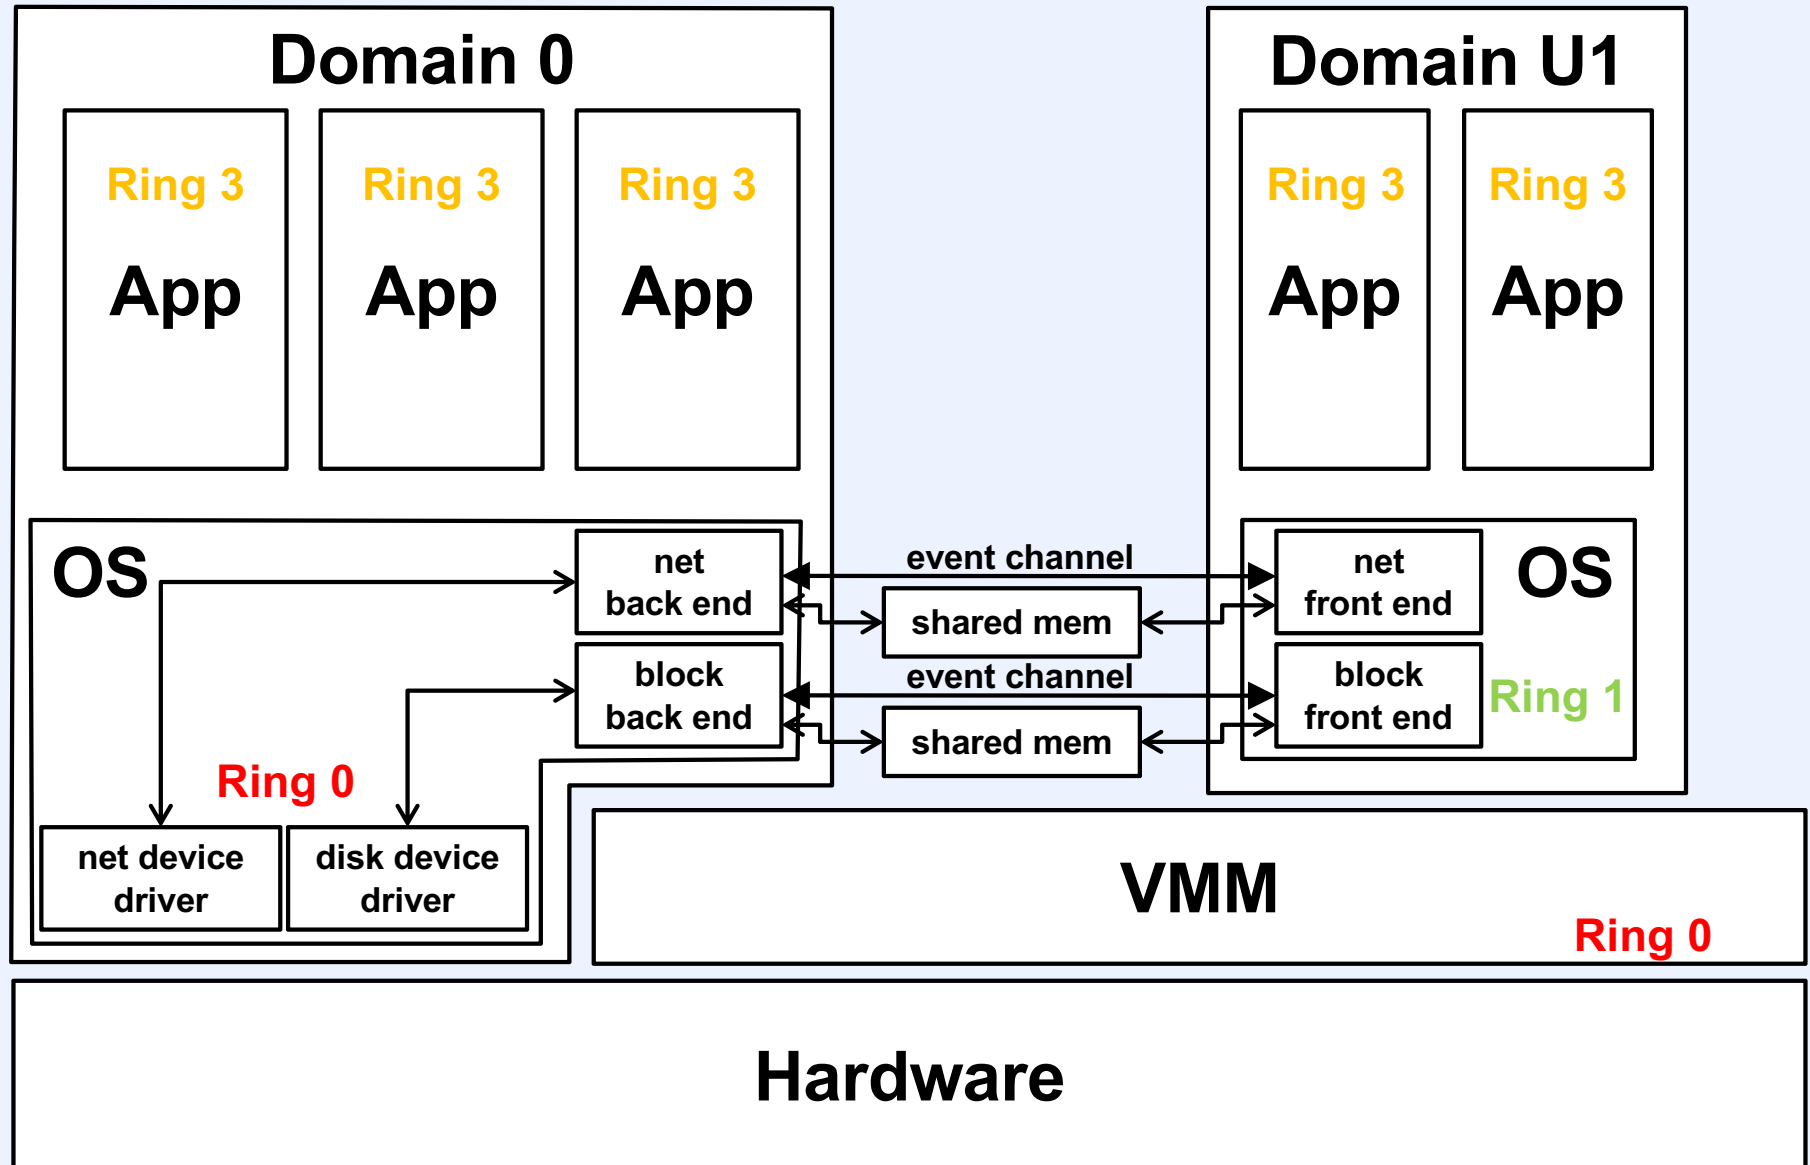
KVM/QEMU

- **KVM**
 - kernel virtual machine monitor for Linux
 - uses VMX technology (or AMD equivalent)
- **QEMU**
 - generic and open source machine emulator and virtualizer
 - does binary rewriting and caching as does VMware
 - emulates I/O devices as well
- **KVM/QEMU**
 - code executes natively until VM-exit
 - user-space QEMU code does I/O emulation

Paravirtualization

- **Sensitive instructions replaced with hypervisor calls**
 - traps to VMM
- **Virtual machine provides higher-level device interface**
 - guest machine has no device drivers

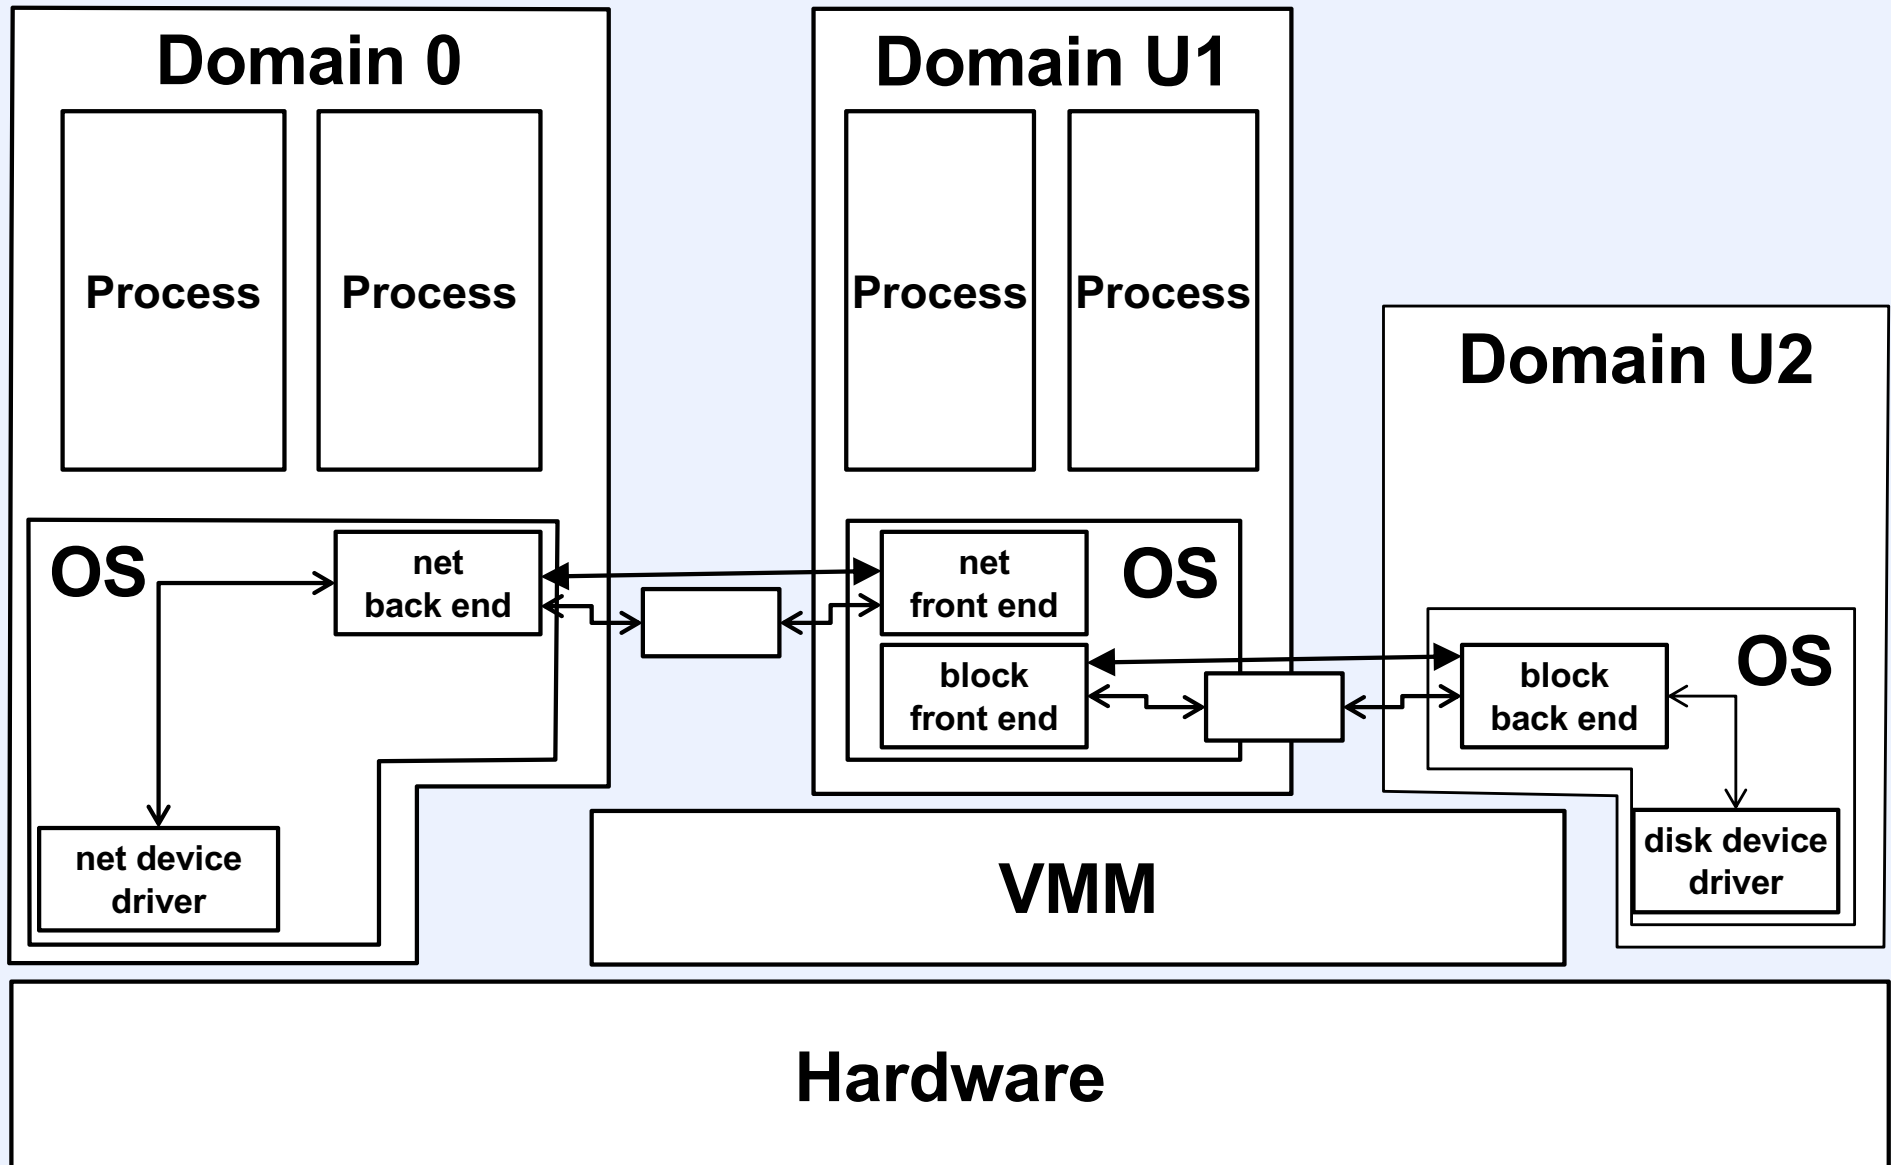
Xen






Additional Applications

- **Sandboxing**
 - isolate web servers
 - isolate device drivers
- **Migration**
 - VM not tied to particular hardware
 - easy to move from one (real) platform to another





Xen with Isolated Driver



Process Migration




Approaches: Before




Approaches: After

Communication



Communication



Virtual-Machine Migration

- **Virtual machines are isolated**
 - by definition!
- **State is well defined**
 - thus easy to identify and move
 - possible exception of virtual memory

Transferring Virtual Memory

- **Eager**
 - all
 - dirty
 - (clean pages come from common source)
- **Lazy**
 - copy on reference
- **Straightforward**
 - flush everything to file system on source, then access file system on target
- **Weird**
 - precopy

Eager–Dirty

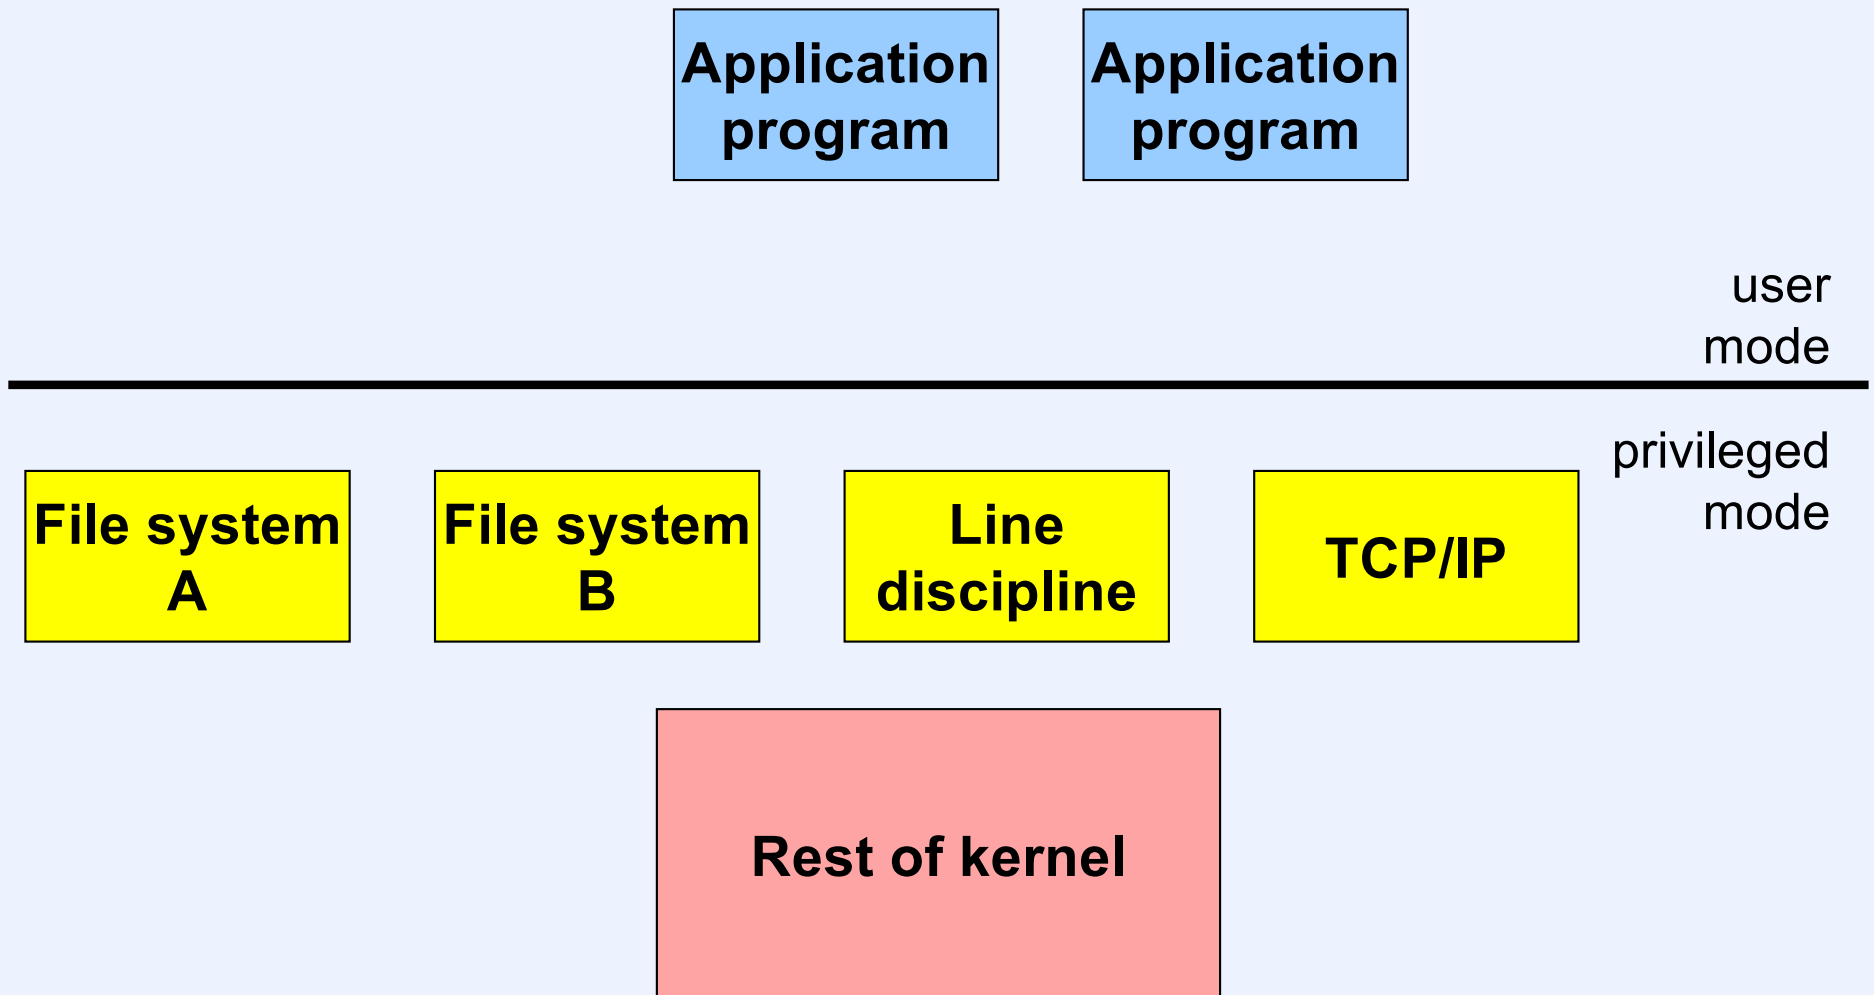
- Freeze process on source
- Transfer all dirty pages to target
- Resume process on target

Precopy

- **While process still running on source**
 - transfer dirty pages to target (eager-dirty)
- **While more than x pages dirty on source**
 - transfer newly dirtied pages to target
- **Freeze process on source**
- **Transfer remaining dirty pages to target**
- **Resume process on target**

Microkernels

Traditional OS Organization

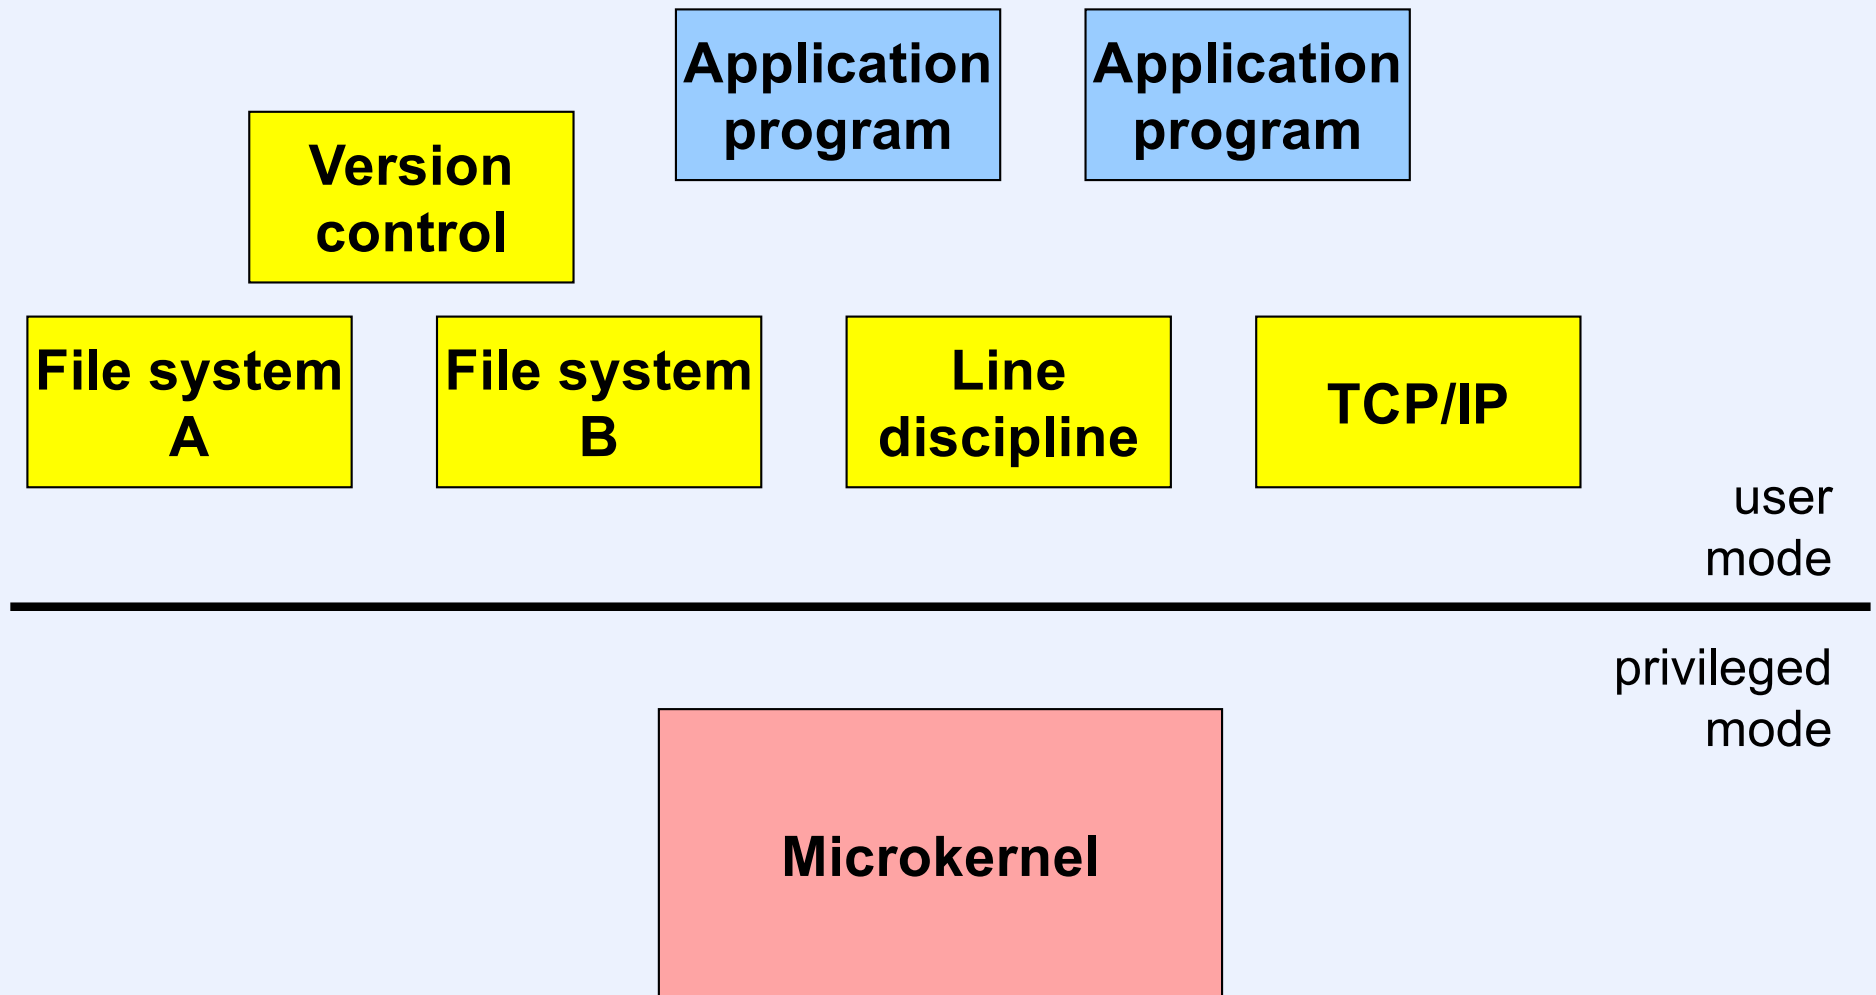


Quiz 2

In the previous slide, assume each of the two application programs runs as a separate process. What's in the slide employs:

- a) two address spaces: one for each process, with the kernel existing in a shared portion of the two process address spaces**
- b) three address spaces: one for each process, and one for the kernel**
- c) six address spaces: one for each process, one for each of the four kernel components, and one for the rest of the kernel**

OS Services as User Apps



Why?

- **It's cool ...**
- **Assume that OS coders are incompetent, malicious, or both ...**
 - OS components run as protected user-level applications
- **Extensibility**
 - easier to add, modify, and extend user-level components than kernel components

Implementation Issues

- **What are the building blocks?**
- **What is run in privileged mode?**

Mach


- **Developed at CMU, then Utah**
- **Early versions shared kernel with Unix**
 - **basis of NeXT OS**
 - **basis of Apple OS X**
- **Later versions still shared kernel with Unix**
 - **basis of OSF/1**
- **Even later versions actually functioned as working microkernel**
 - **basis of GNU/HURD project**
 - **HURD: HIRD of Unix-replacing daemons**
 - **HIRD: HURD of interfaces representing depth**

Mach's Building Blocks

- **Tasks**
 - represent services/objects
 - holders of access rights
- **Threads**
 - represent virtual processors
- **Ports**
 - communication channels and access rights
- **Messages**
 - carriers of data and access rights

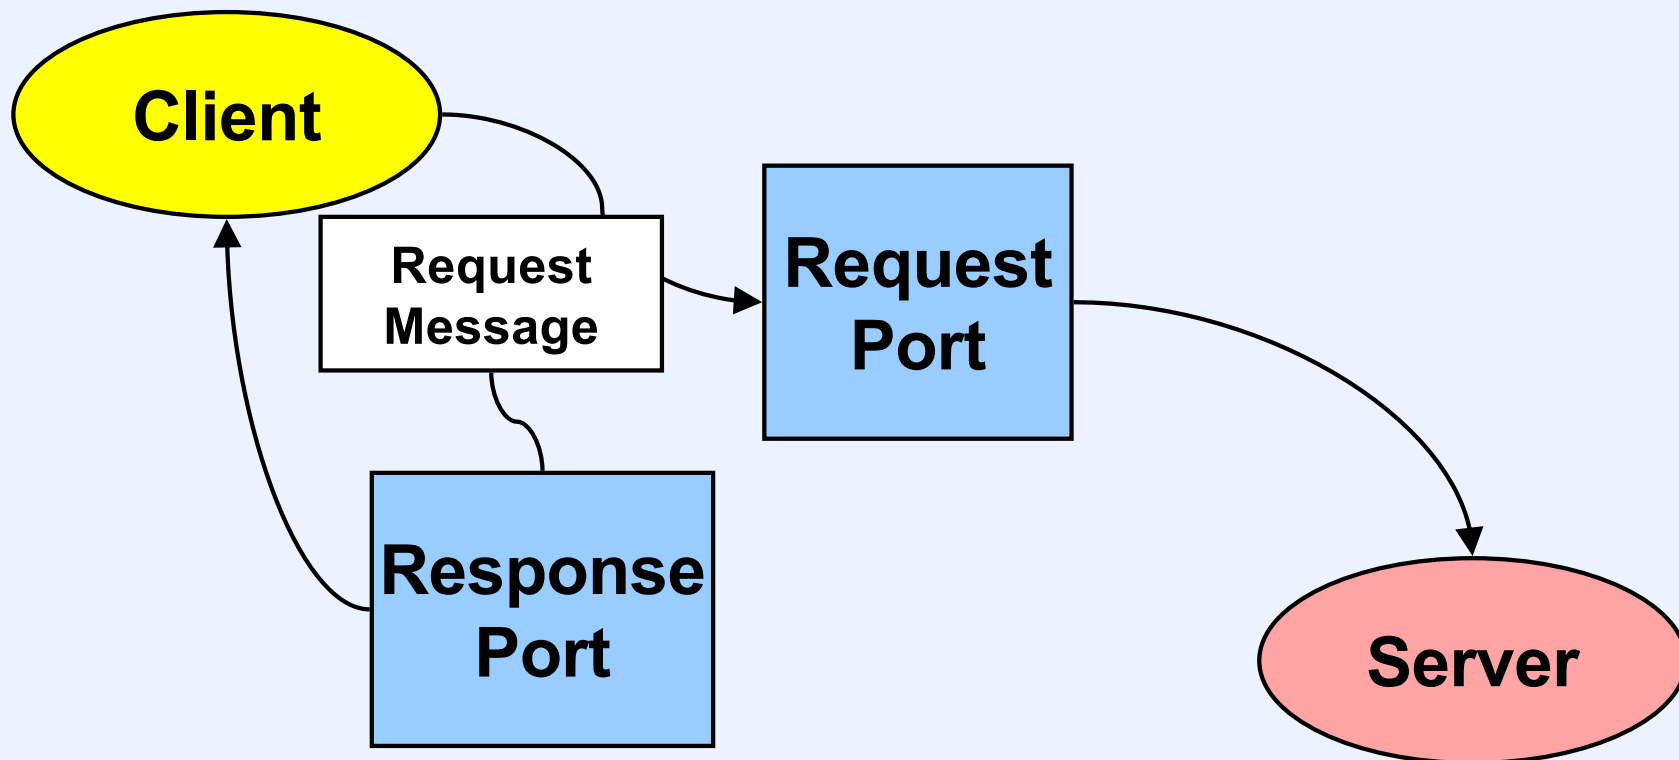
Mach Ports (1)

- Access rights




Mach Ports (2)

- Communication construct



Mach Ports (3)

- Communication construct



Method Invocation


- ***Tasks*** implement objects
- **Access rights to *ports* are secure object references**
- ***Messages* are method invocations and responses**

Messages

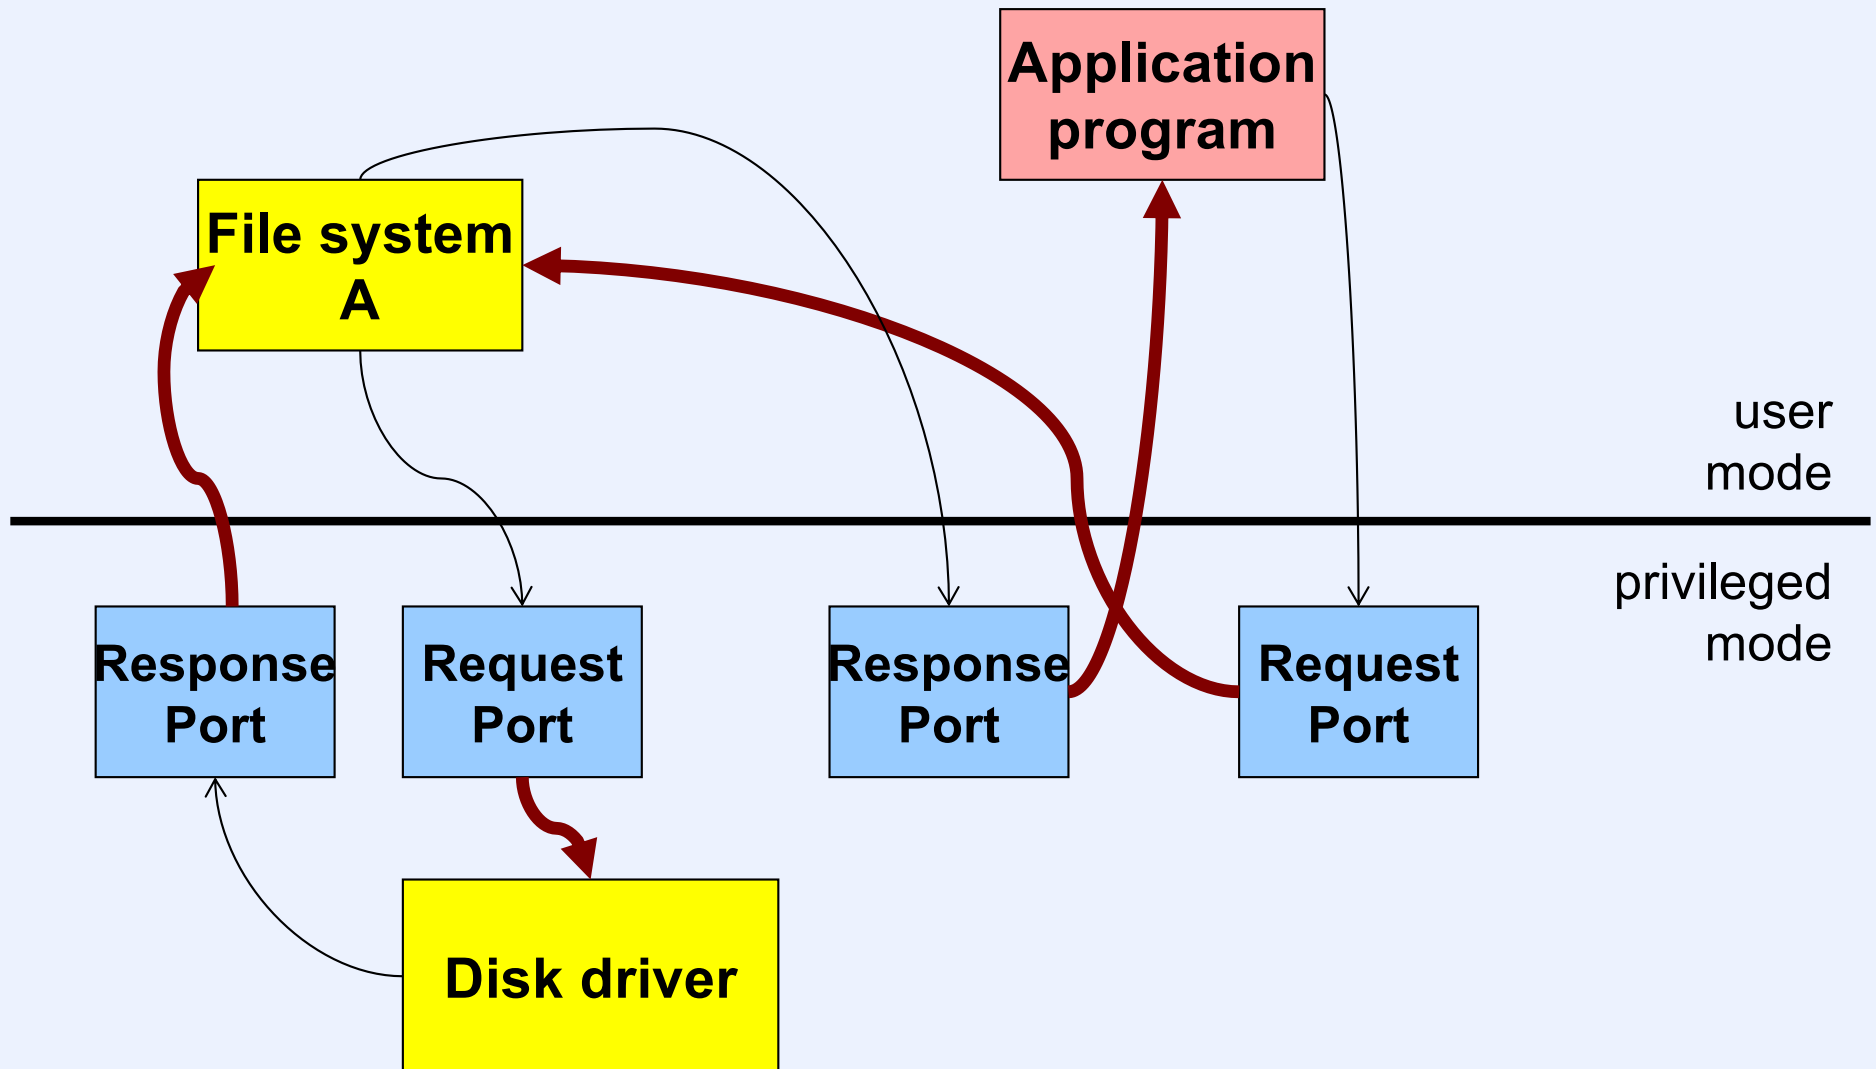
- **Cost of passing messages is critical factor**
- **Small messages are copied**
- **Large messages within single address space passed by reference**
- **What about messages across address spaces?**

Implementing Port Rights

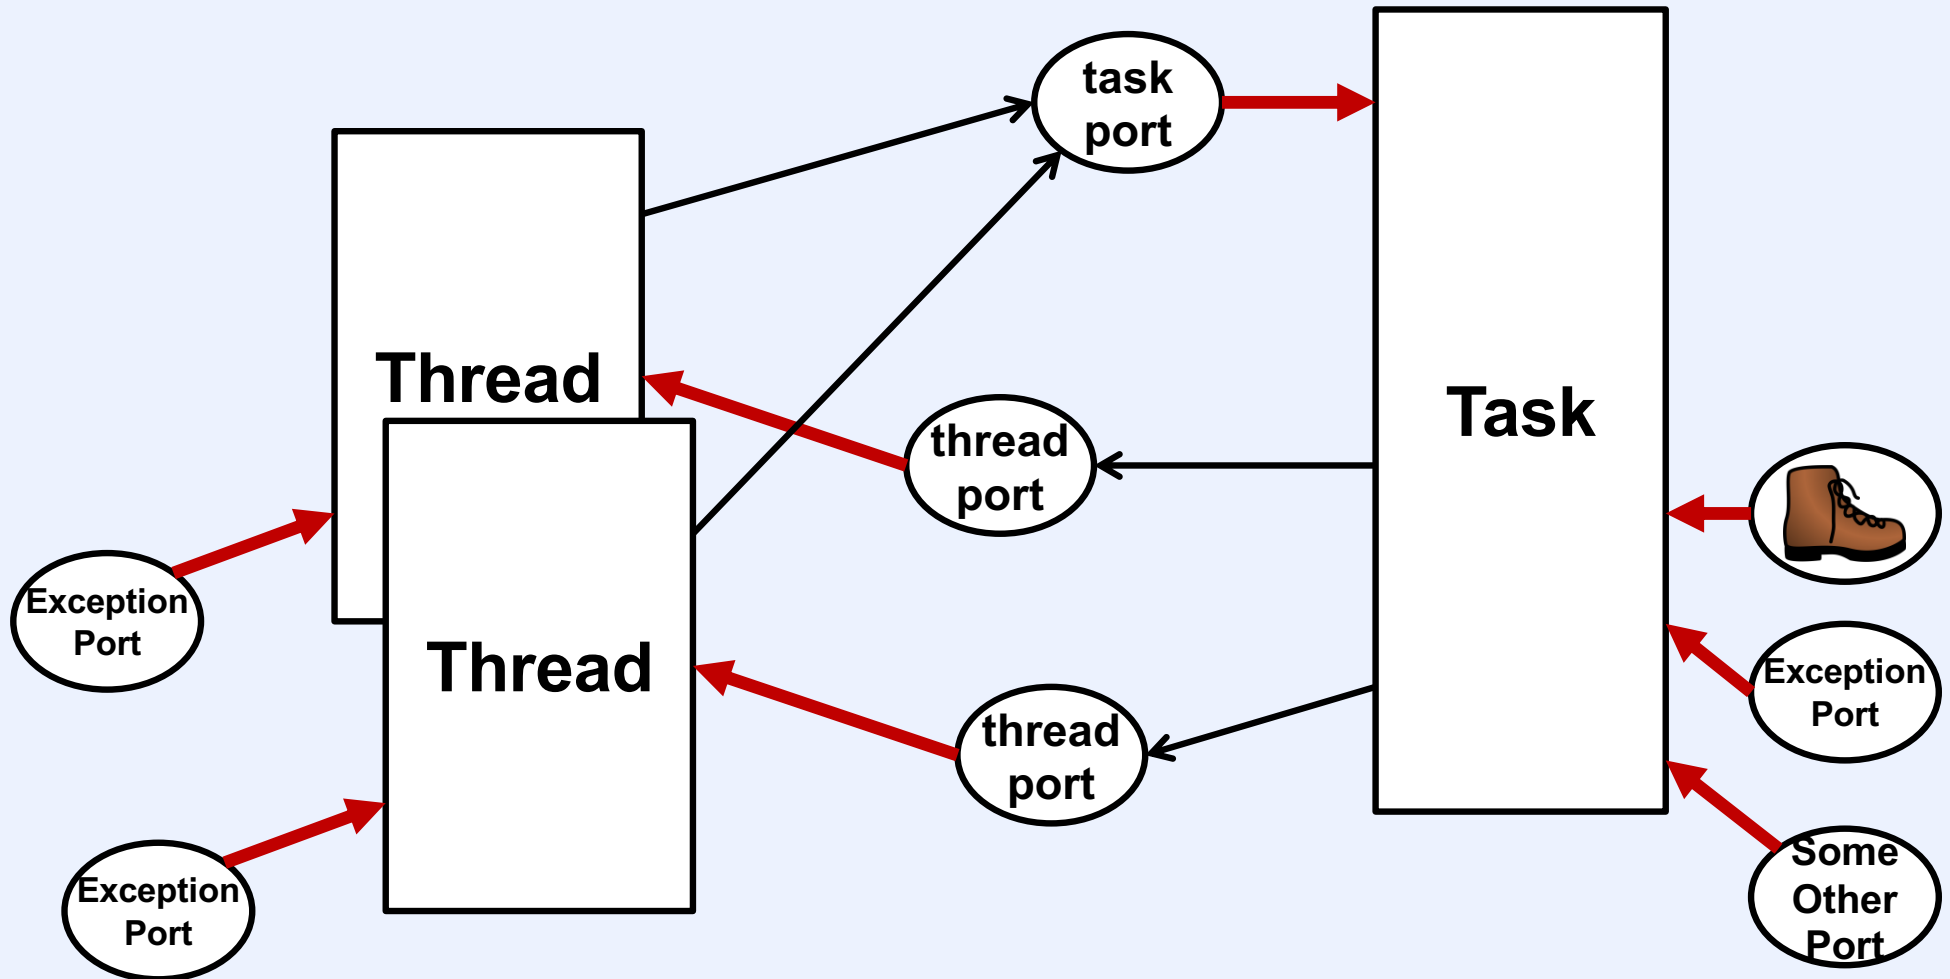
- References to ports must be secure and unforgeable
 - how are they done?



Example



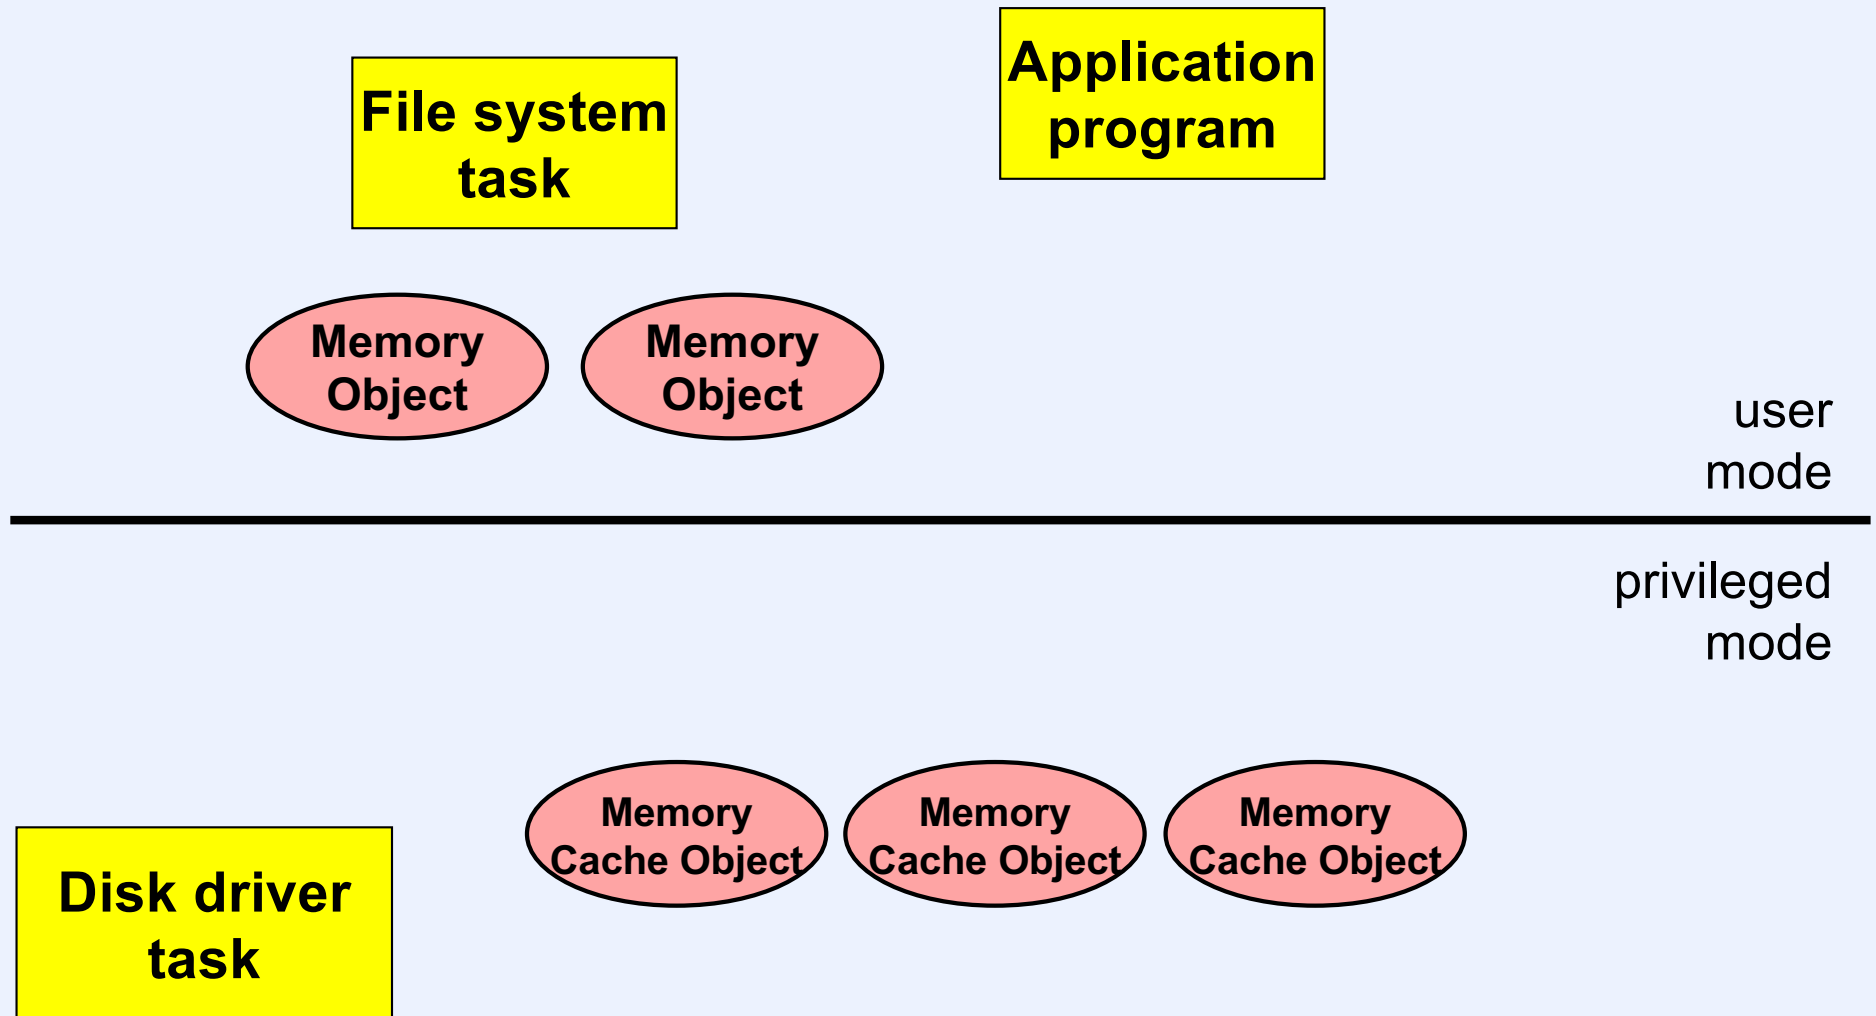
Task and Thread Objects



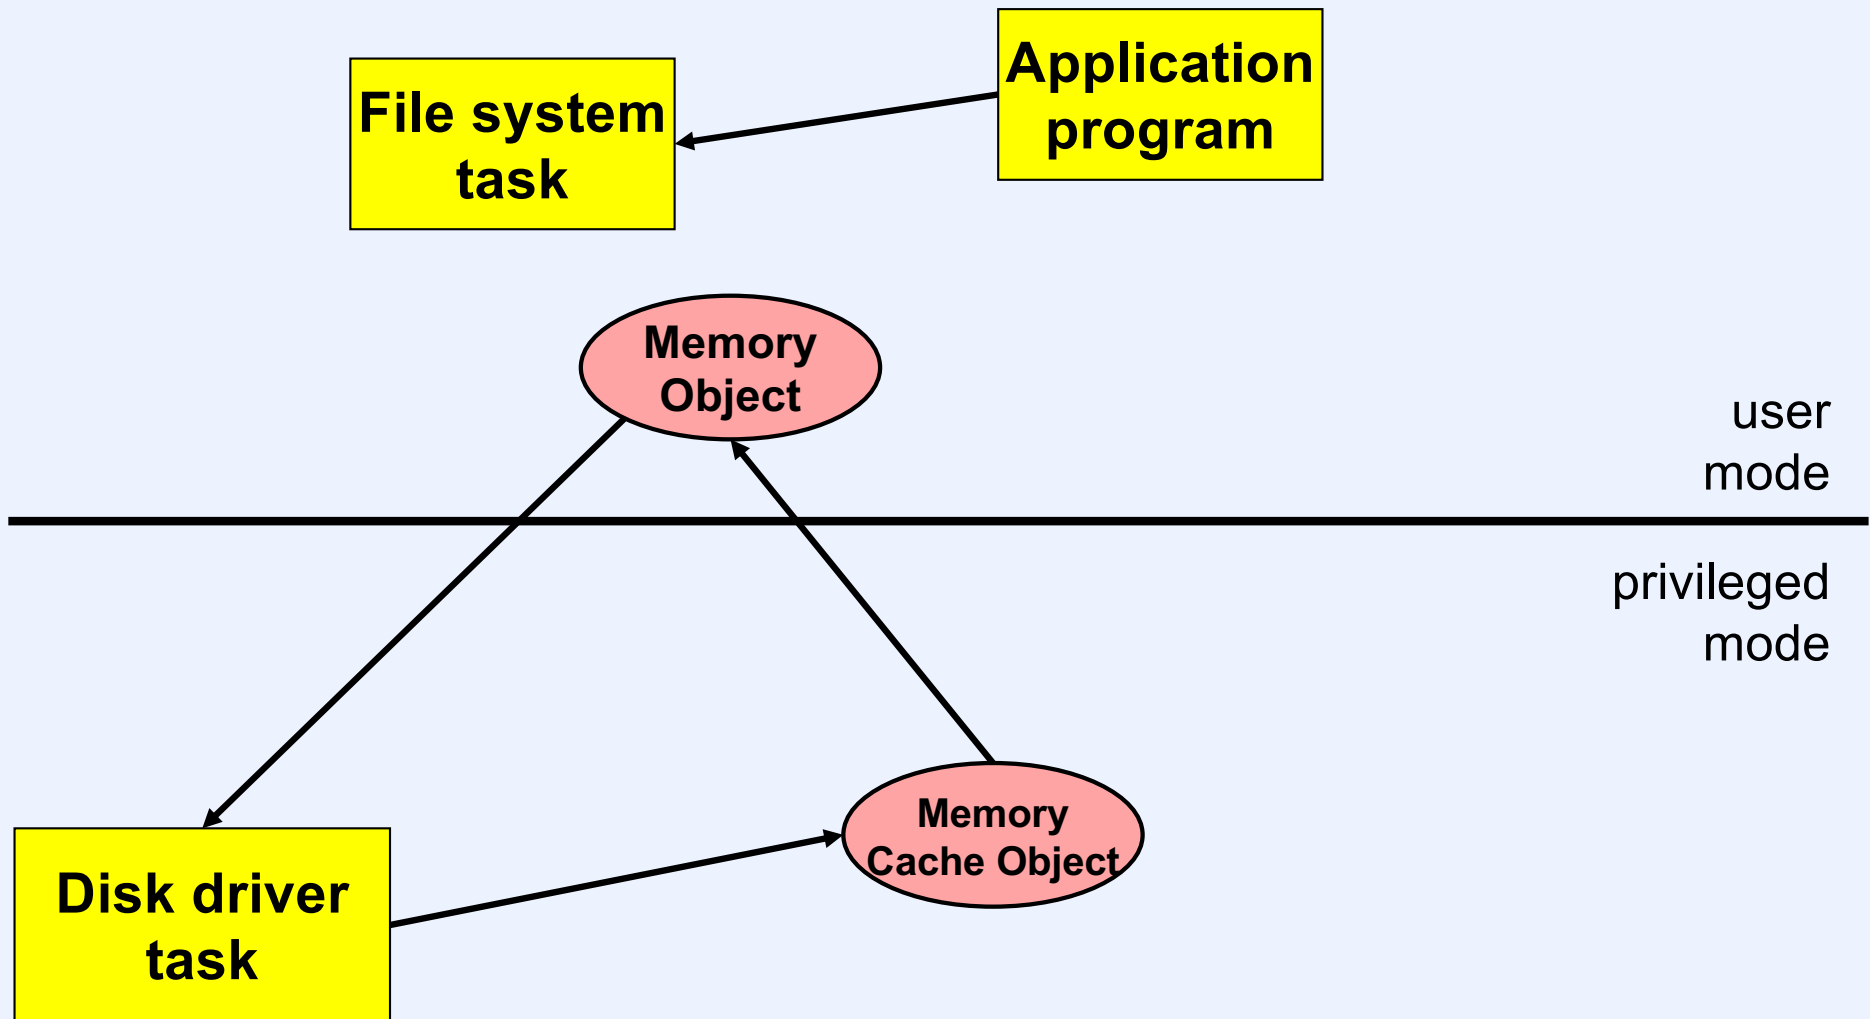
Virtual Memory

- **Memory cache objects**
 - implemented in kernel
 - represent what's in real memory
- **Memory objects**
 - implemented in kernel or as user tasks
 - represent what's mapped into real memory

Memory-Object Example



Memory-Object Example



Quiz 3

The application thread attempts to read from memory. The page containing the desired bytes is not present and must be fetched from disk. How many threads are involved?

- a) one (just the application thread)**
- b) two**
- c) three or more**

Devices

- **Device master port exported by kernel**
- **Tasks holding send rights may request access to any device**
 - **send rights given for device port**

Successful Microkernel Systems

-
-
- ...

Attempts

- **Windows NT 3.1**
 - graphics subsystem ran as user-level process
 - moved to kernel in 4.0 for performance reasons
- **Apple OS X**
 - based on Mach
 - all services in kernel for performance reasons
- **HURD**
 - based on Mach
 - services implemented as user processes
 - no one used it, for performance reasons