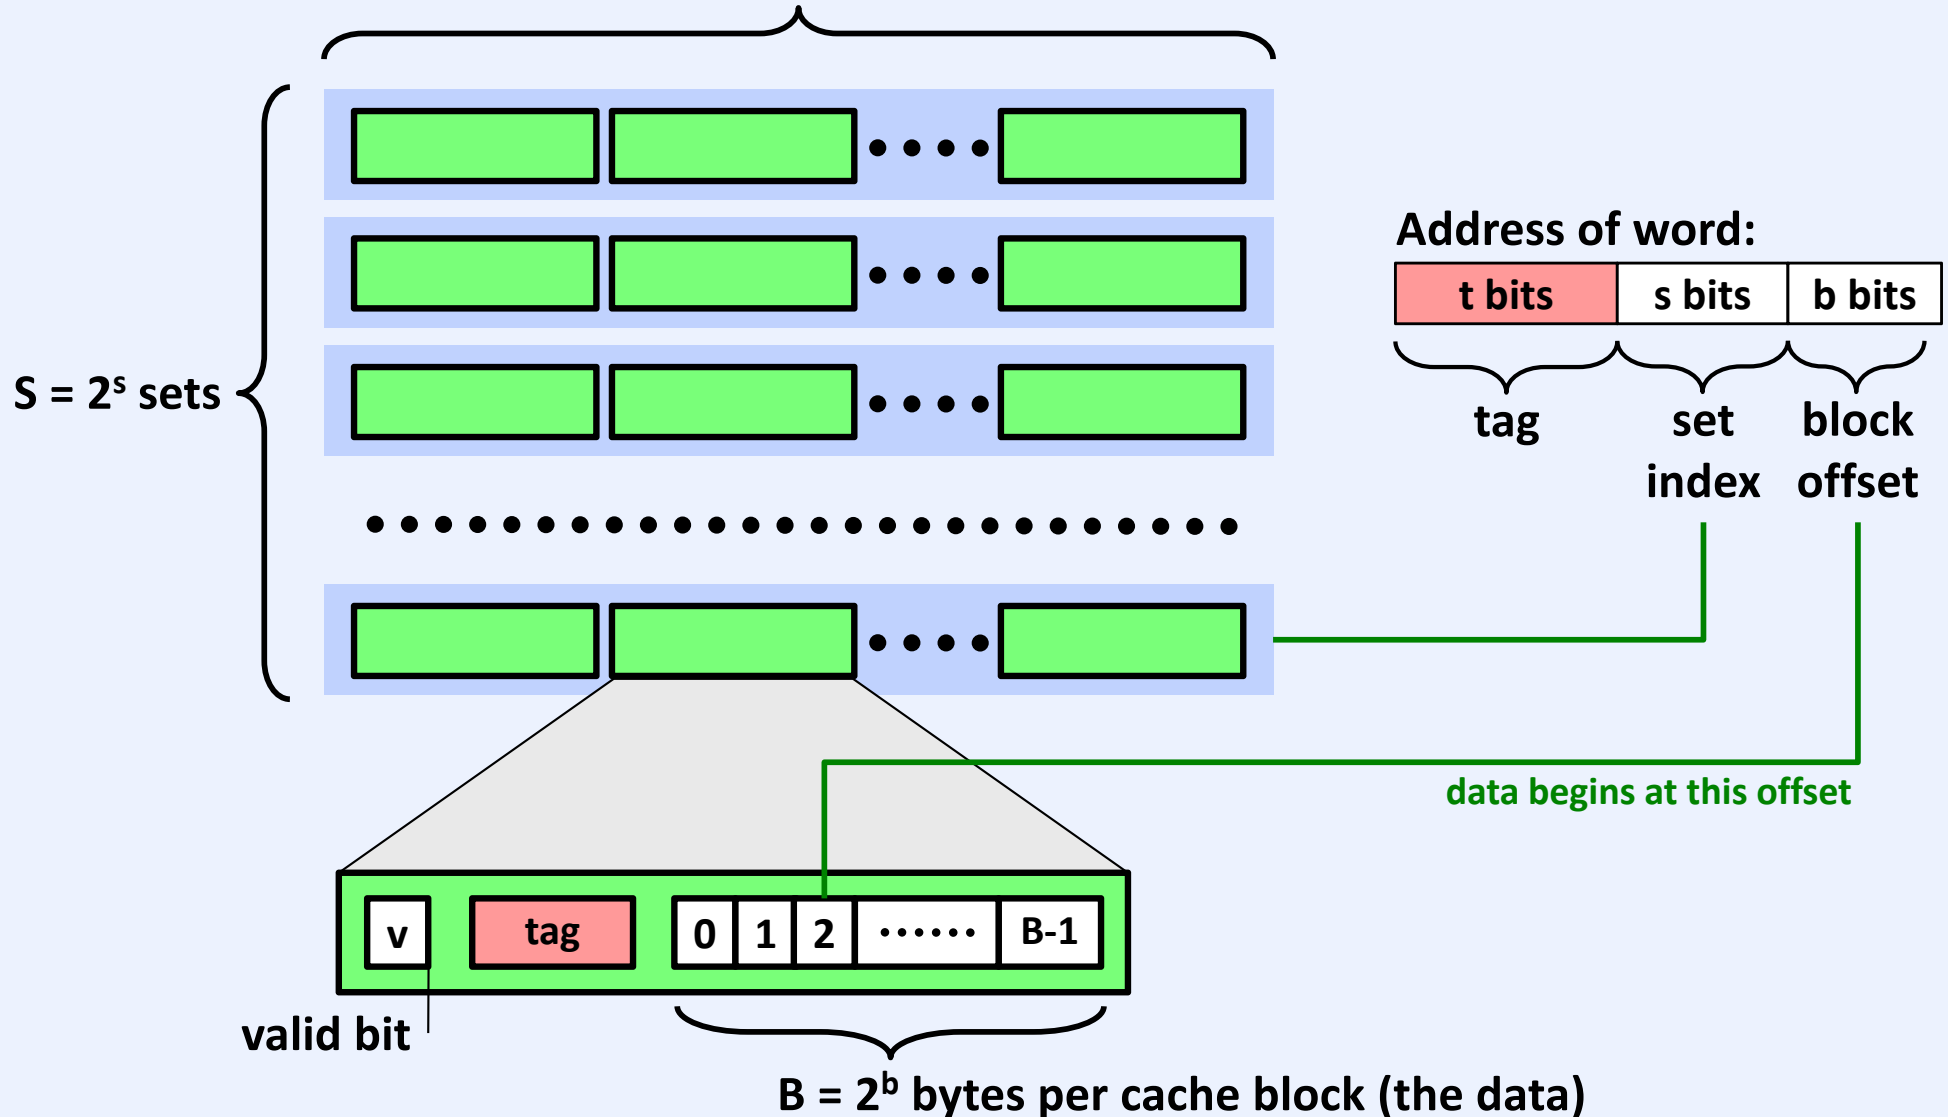


Memory Management Part 5

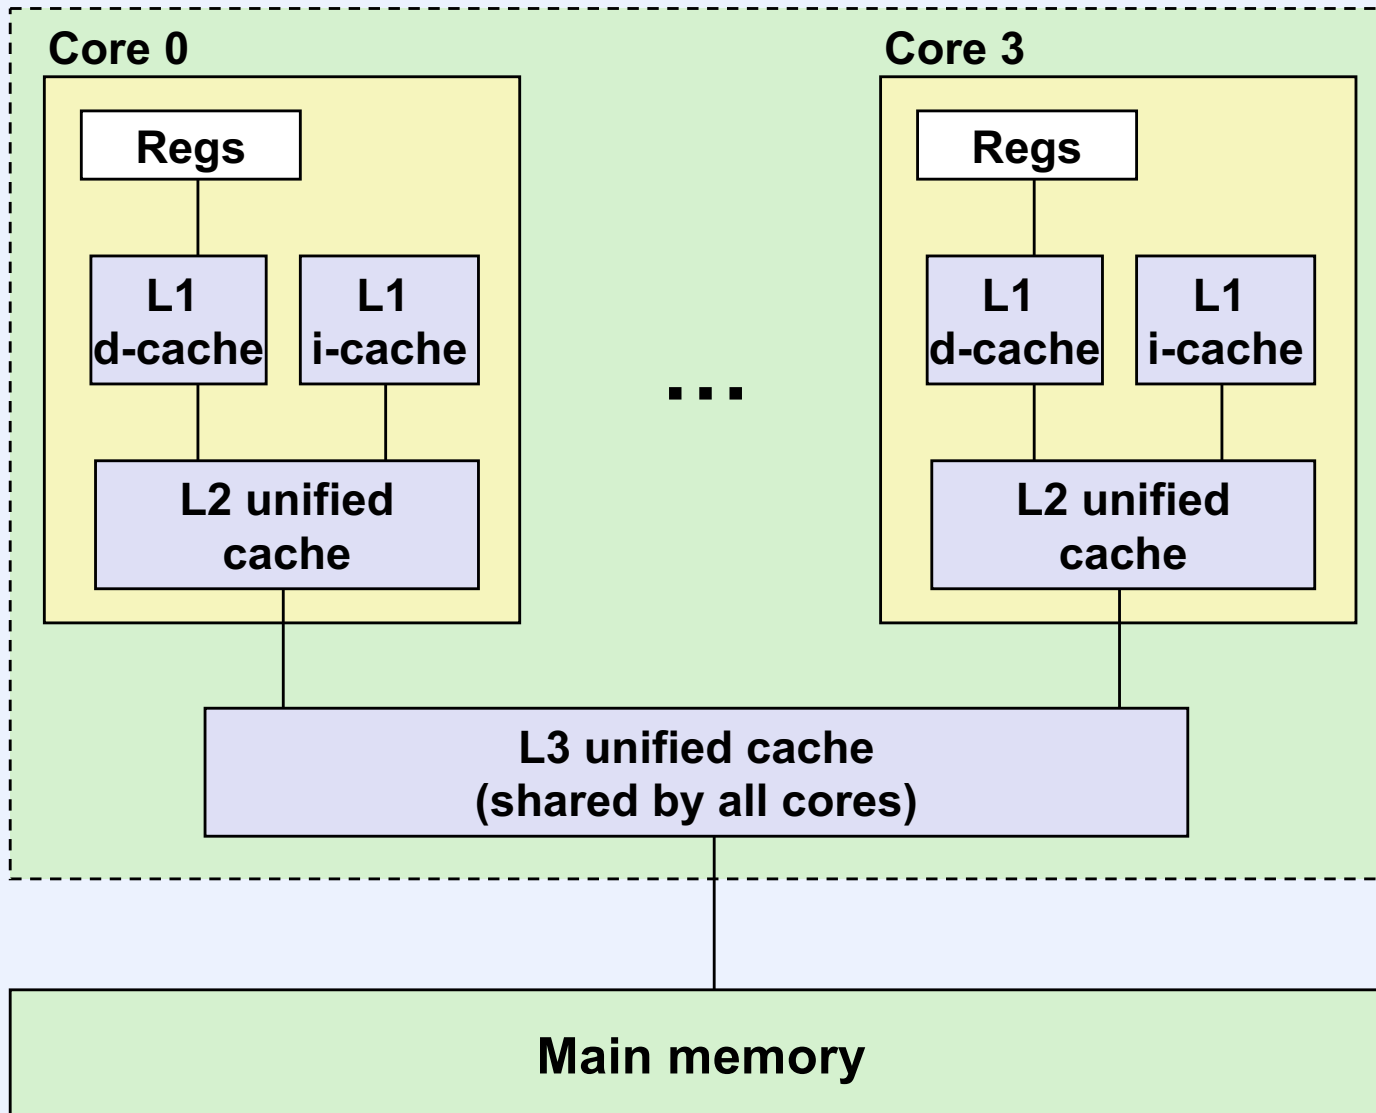
E-Way Set-Associative Cache

$E = 2^e$ lines per set



Intel Core i5 and i7 Cache Hierarchy

Processor package



L1 i-cache and d-cache:

32 KB, 8-way,
Access: 4 cycles

L2 unified cache:

256 KB, 8-way,
Access: 11 cycles

L3 unified cache:

8 MB, 16-way,
Access: 30-40 cycles


Block size: 64 bytes for
all caches

Quiz 1


You're designing the algorithm for allocating an often-used and -allocated kernel data structure that fits within a cache line. We'd like to make sure that a number of these data structures can coexist in the hardware caches. Which one of the following would help make this happen (and is doable)?

- a) Rounding the size of the data structure up to a power of 2**
- b) Making sure all reside in the same cache set**
- c) Making sure they are distributed across cache sets**
- d) Nothing would help**

Slab Allocation



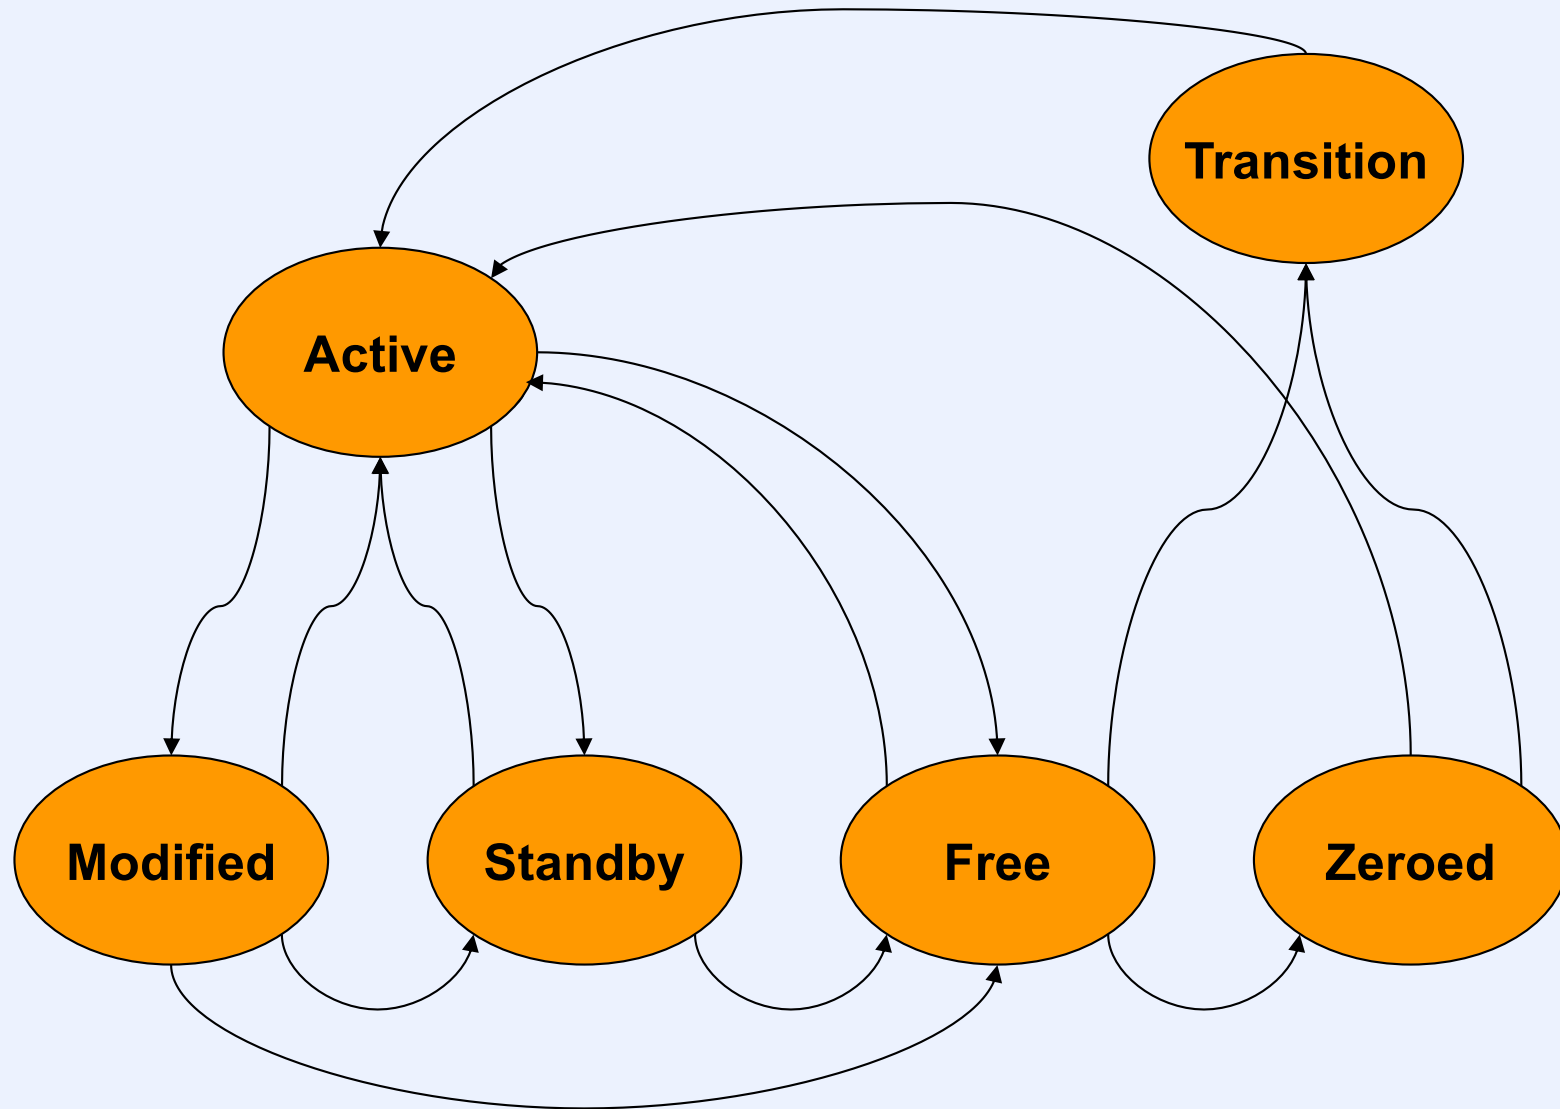
Windows x86 Layout



Windows Paging Strategy

- **All processes guaranteed a “working set”**
 - lower bound on page frames
- **Competition for additional page frames**
- **“Balance-set” manager thread maintains working sets**
 - one-handed clock algorithm
- **Swapper thread swaps out idle processes**
 - first kernel stacks
 - then working set
- **Some of kernel memory is paged**
 - page faults are possible

Windows Page-Frame States



Unix and Virtual Memory: The *fork/exec* Problem

- Naive implementation:
 - fork actually makes a copy of the parent's address space for the child
 - child executes a few instructions (setting up file descriptors, etc.)
 - child calls exec
 - result: a lot of time wasted copying the address space, though very little of the copy is actually used

vfork

- **Don't make a copy of the address space for the child; instead, give the address space to the child**
 - the parent is suspended until the child returns it
 - **The child executes a few instructions, then does an *exec***
 - as part of the *exec* (or *exit*), the address space is handed back to the parent
 - **Advantages**
 - very efficient
 - **Disadvantages**
 - works only if child does an *exec* (or *exit*)
 - child shouldn't do anything to the address space
-

Quiz 2

Will the assertion evaluate to true?

```
volatile int A = 6;  
...  
if (vfork() == 0) {  
    A = 7;  
    exit(0);  
}  
sleep(1); // sleep for one second  
assert(A == 7);  
...
```

a) yes

b) no

Lazy Evaluation

- Always put things off as long as possible
- If you wait long enough, you might not have to do them

A Better *fork*

- Parent and child share the pages comprising their address spaces
 - if either party attempts to modify a page, the modifying process gets a copy of just that page
- Advantages
 - semantically equivalent to the original *fork*
 - usually faster than the original *fork*
- Disadvantages
 - slower than *vfork*


Quiz 3

How many pages of virtual memory must be copied from the parent to the child in the following code?


```
if (fork() == 0) {  
    close(0);  
    dup(open("input_file", O_RDONLY));  
    execv("newprog", 0);  
}
```

- a) 0**
 - b) 1-2**
 - c) 4-8**
 - d) lots**
-

Copy on Write (1)



Copy on Write (2)




Quiz 4


We have a file that contains one billion 64-bit integers. We are writing a program to read in the file and add up all the integers. Which approach will be fastest:

- a) read the file 8 bytes at a time, adding to a running total what is read in**
- b) read the file 8k bytes at a time, then add each of the integers contained in that block to the running total**
- c) *mmap* the file into the process's address space, then sum up all the integers in this mapped region of memory**

The *mmap* System Call




Share-Mapped Files




Data = 17;

Private-Mapped Files



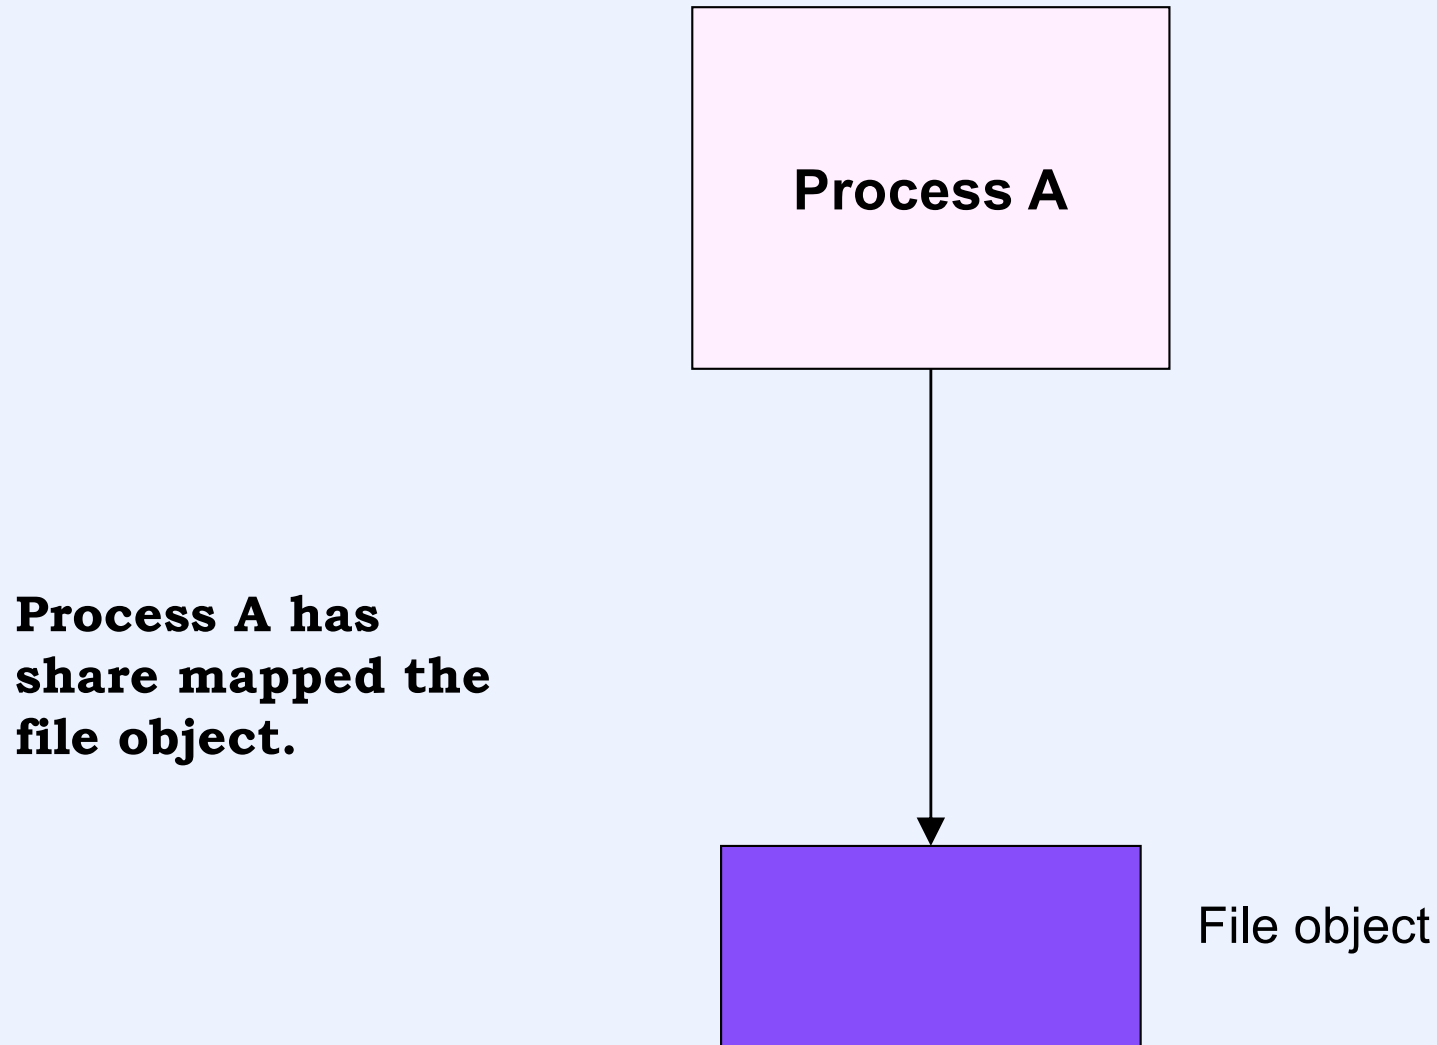
A Private-Mapped File Changes



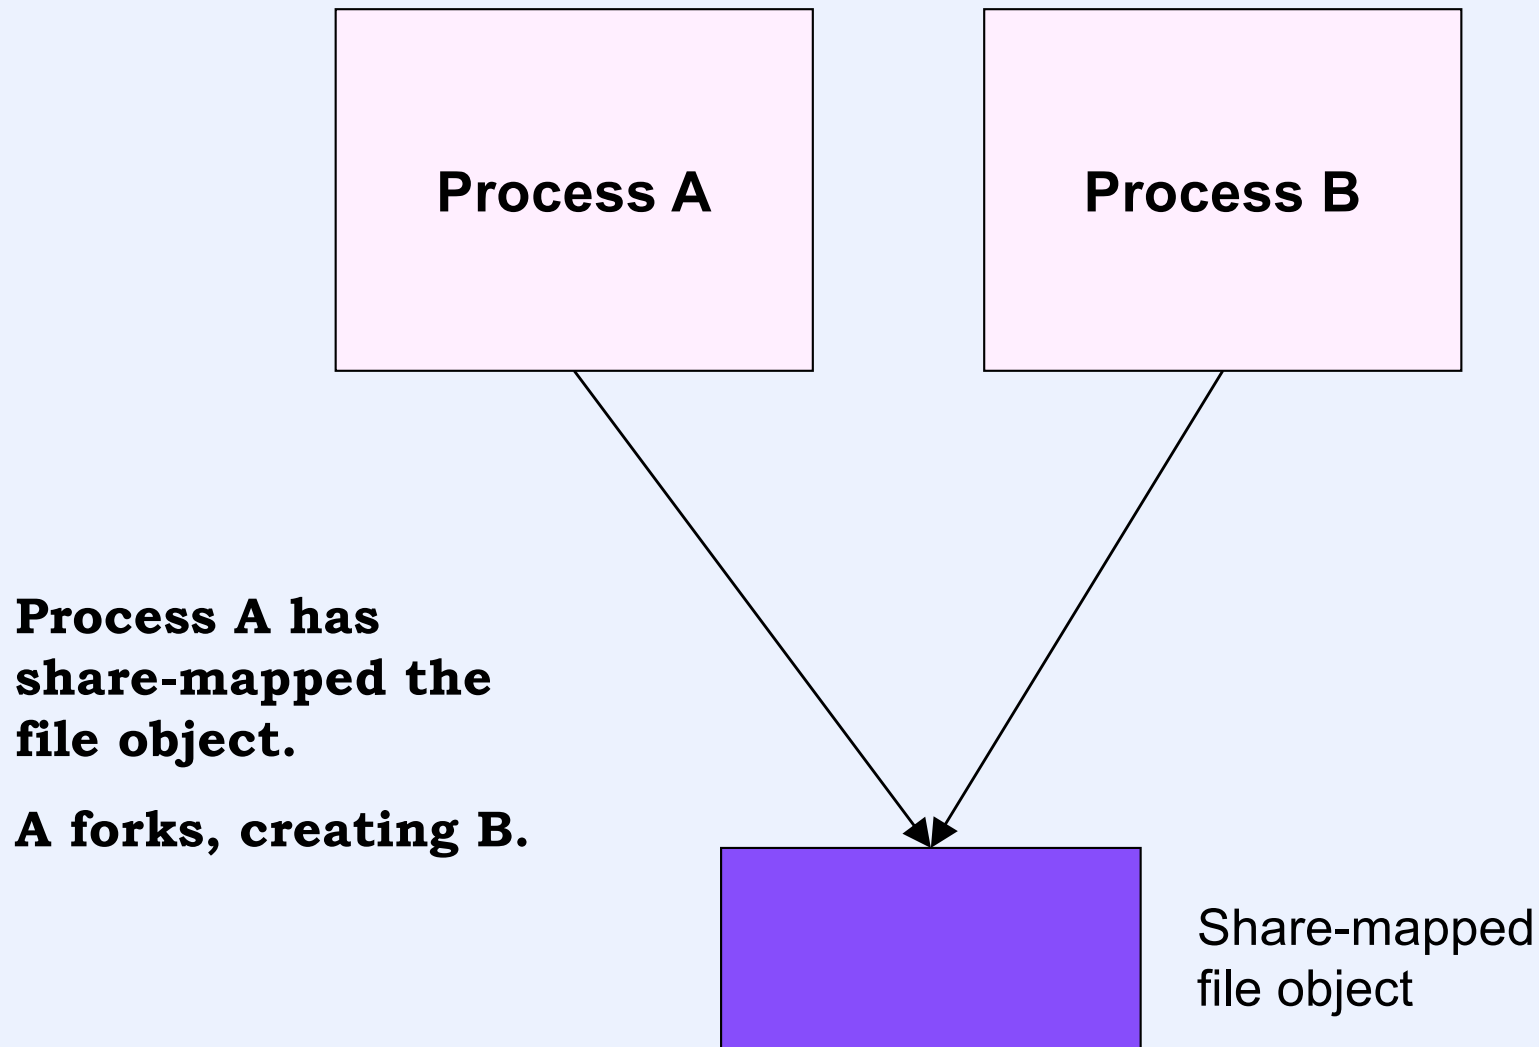
Virtual Copy

- **Local RPC**
 - **“copy” arguments from one process to another**
 - **assume arguments are page-aligned and page-sized**
 - **map pages into both caller and callee, copy-on-write**

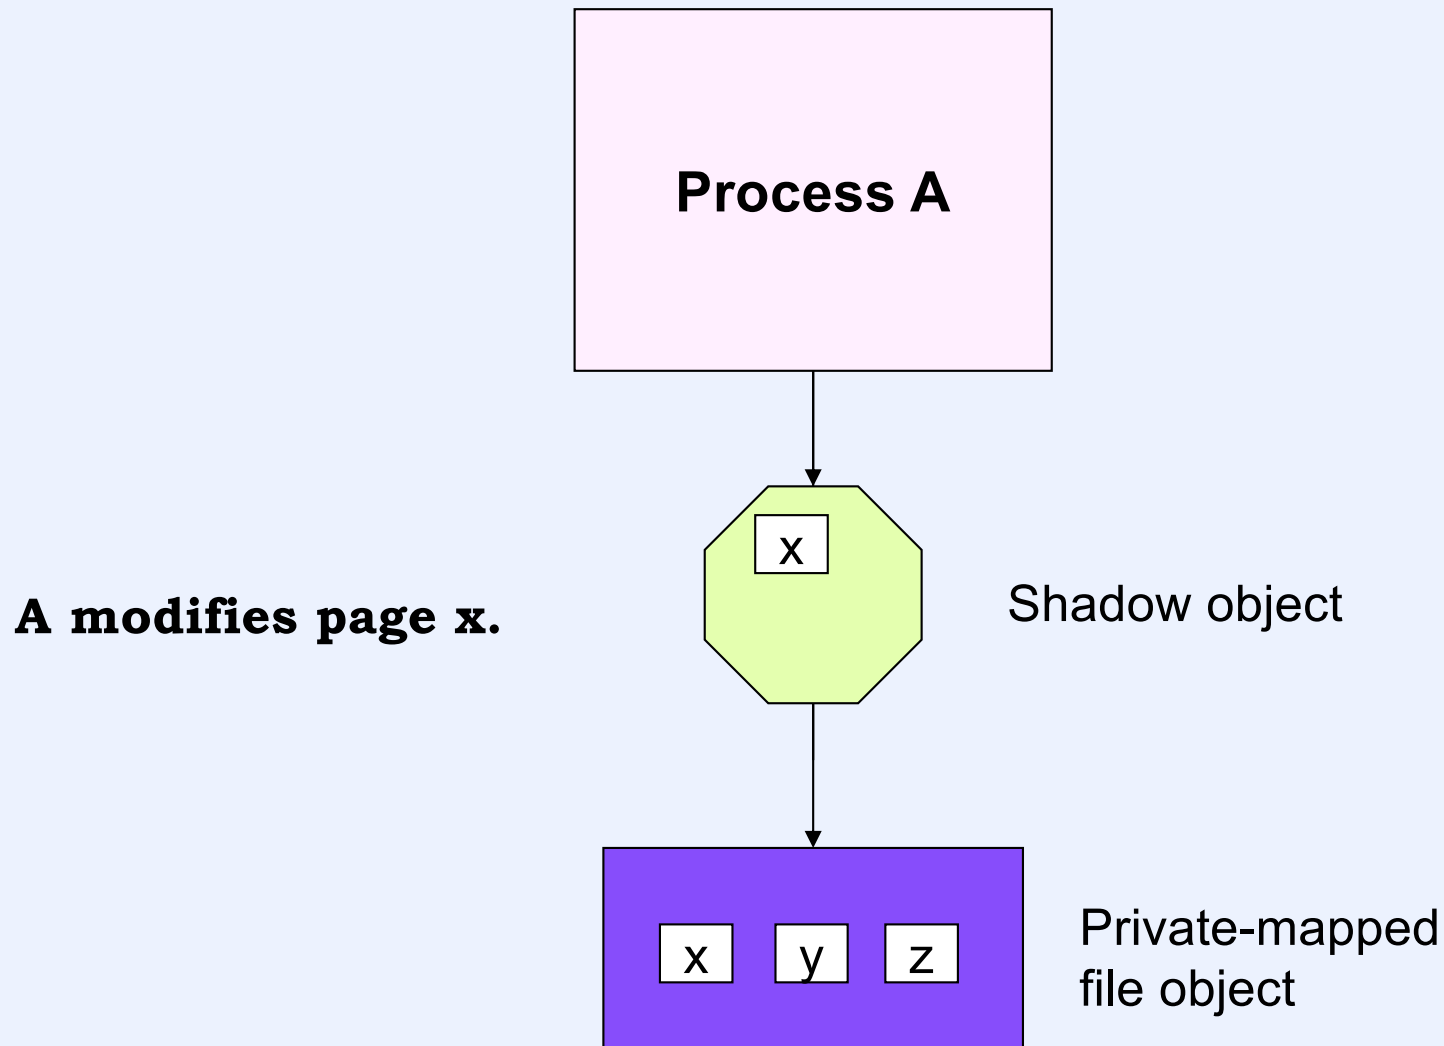
Share Mapping (1)



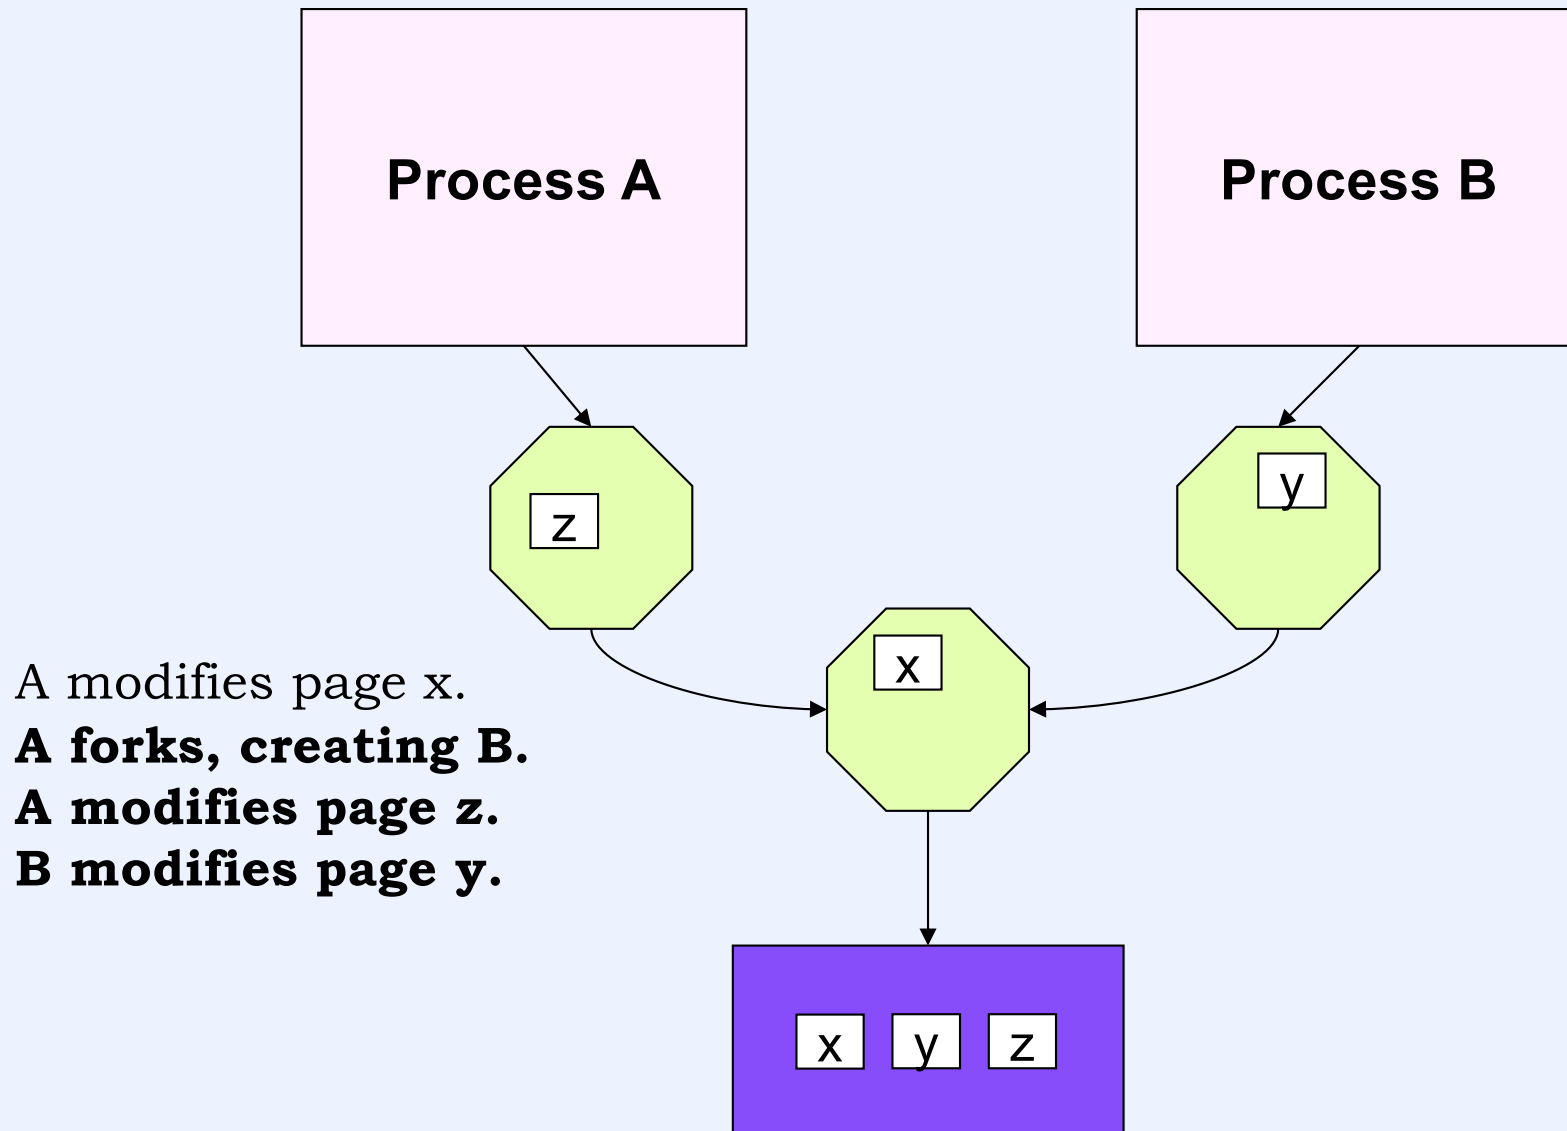
Share Mapping (2)



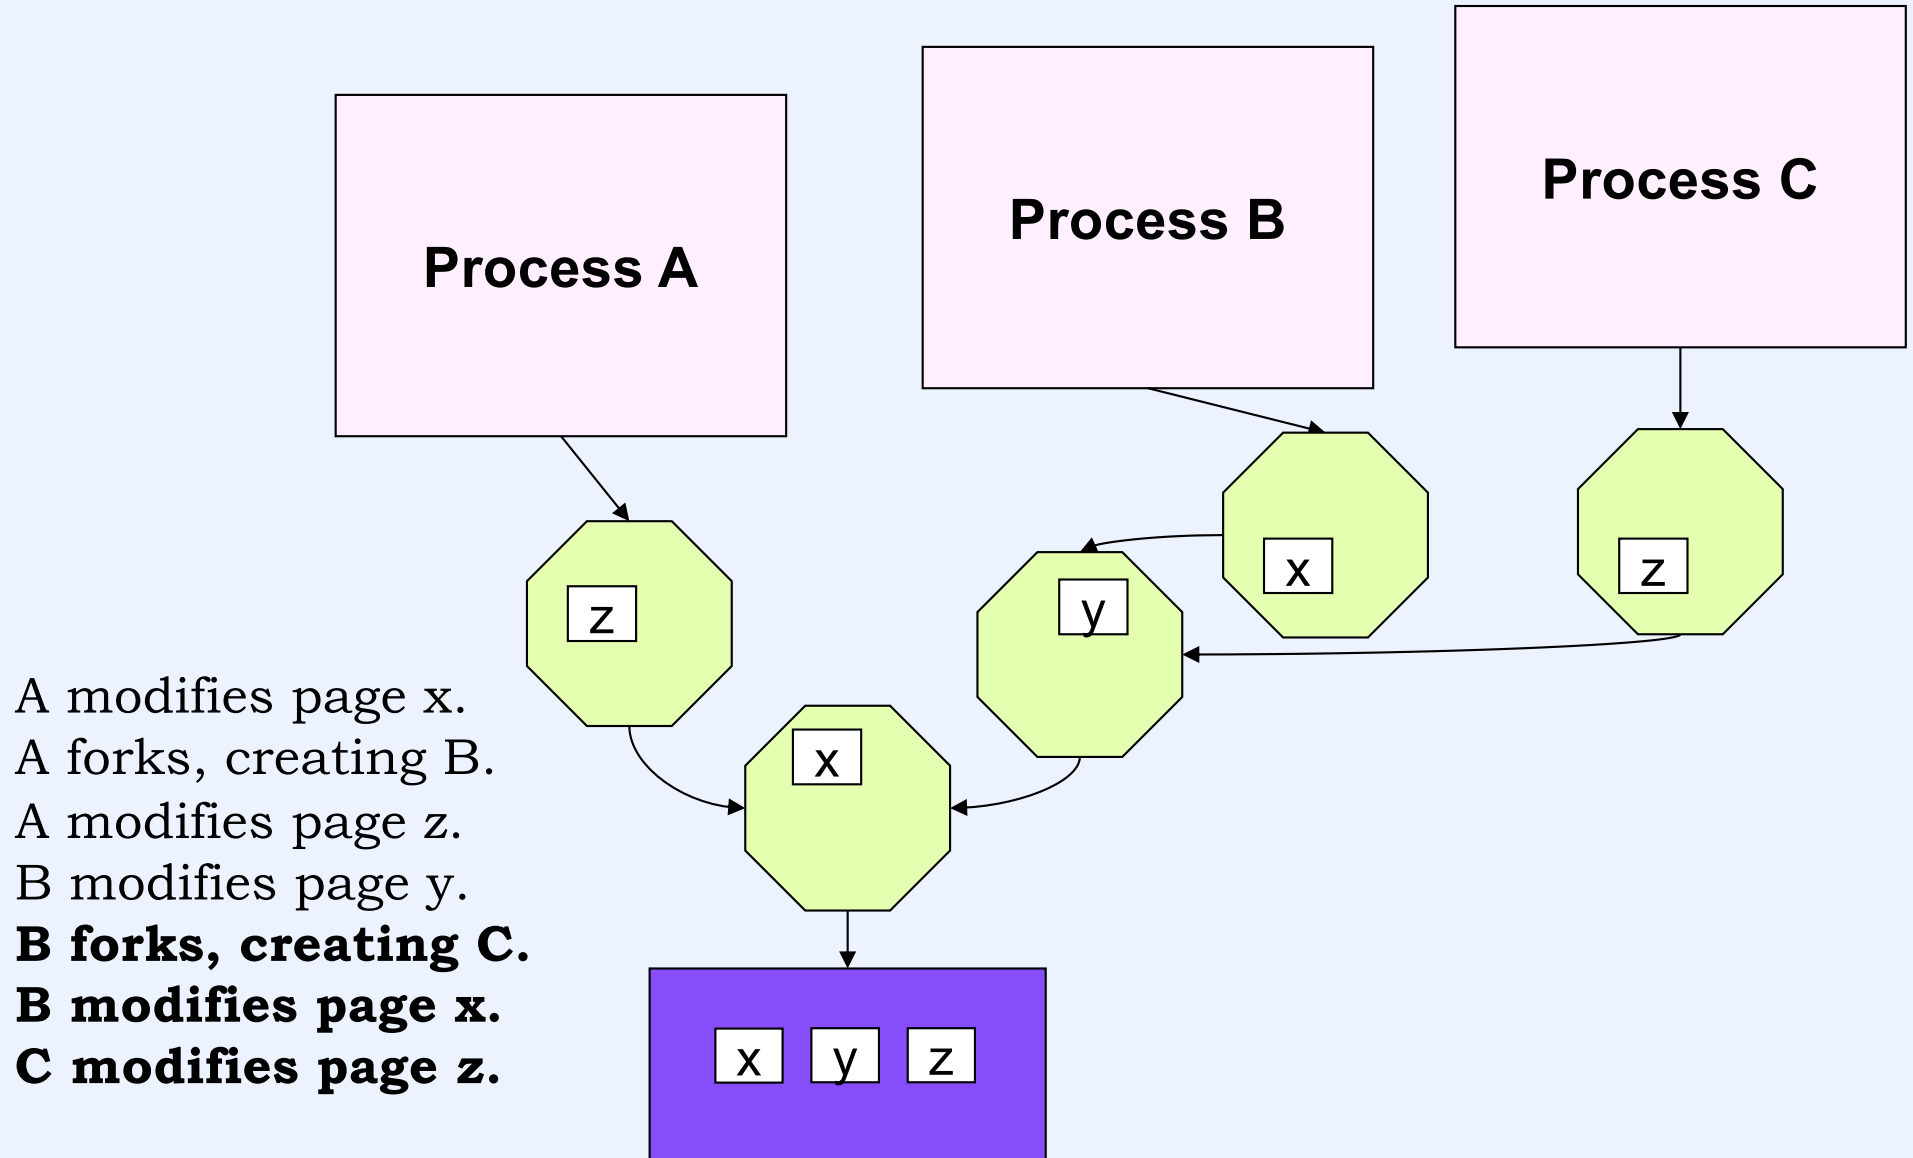
Private Mapping (1)



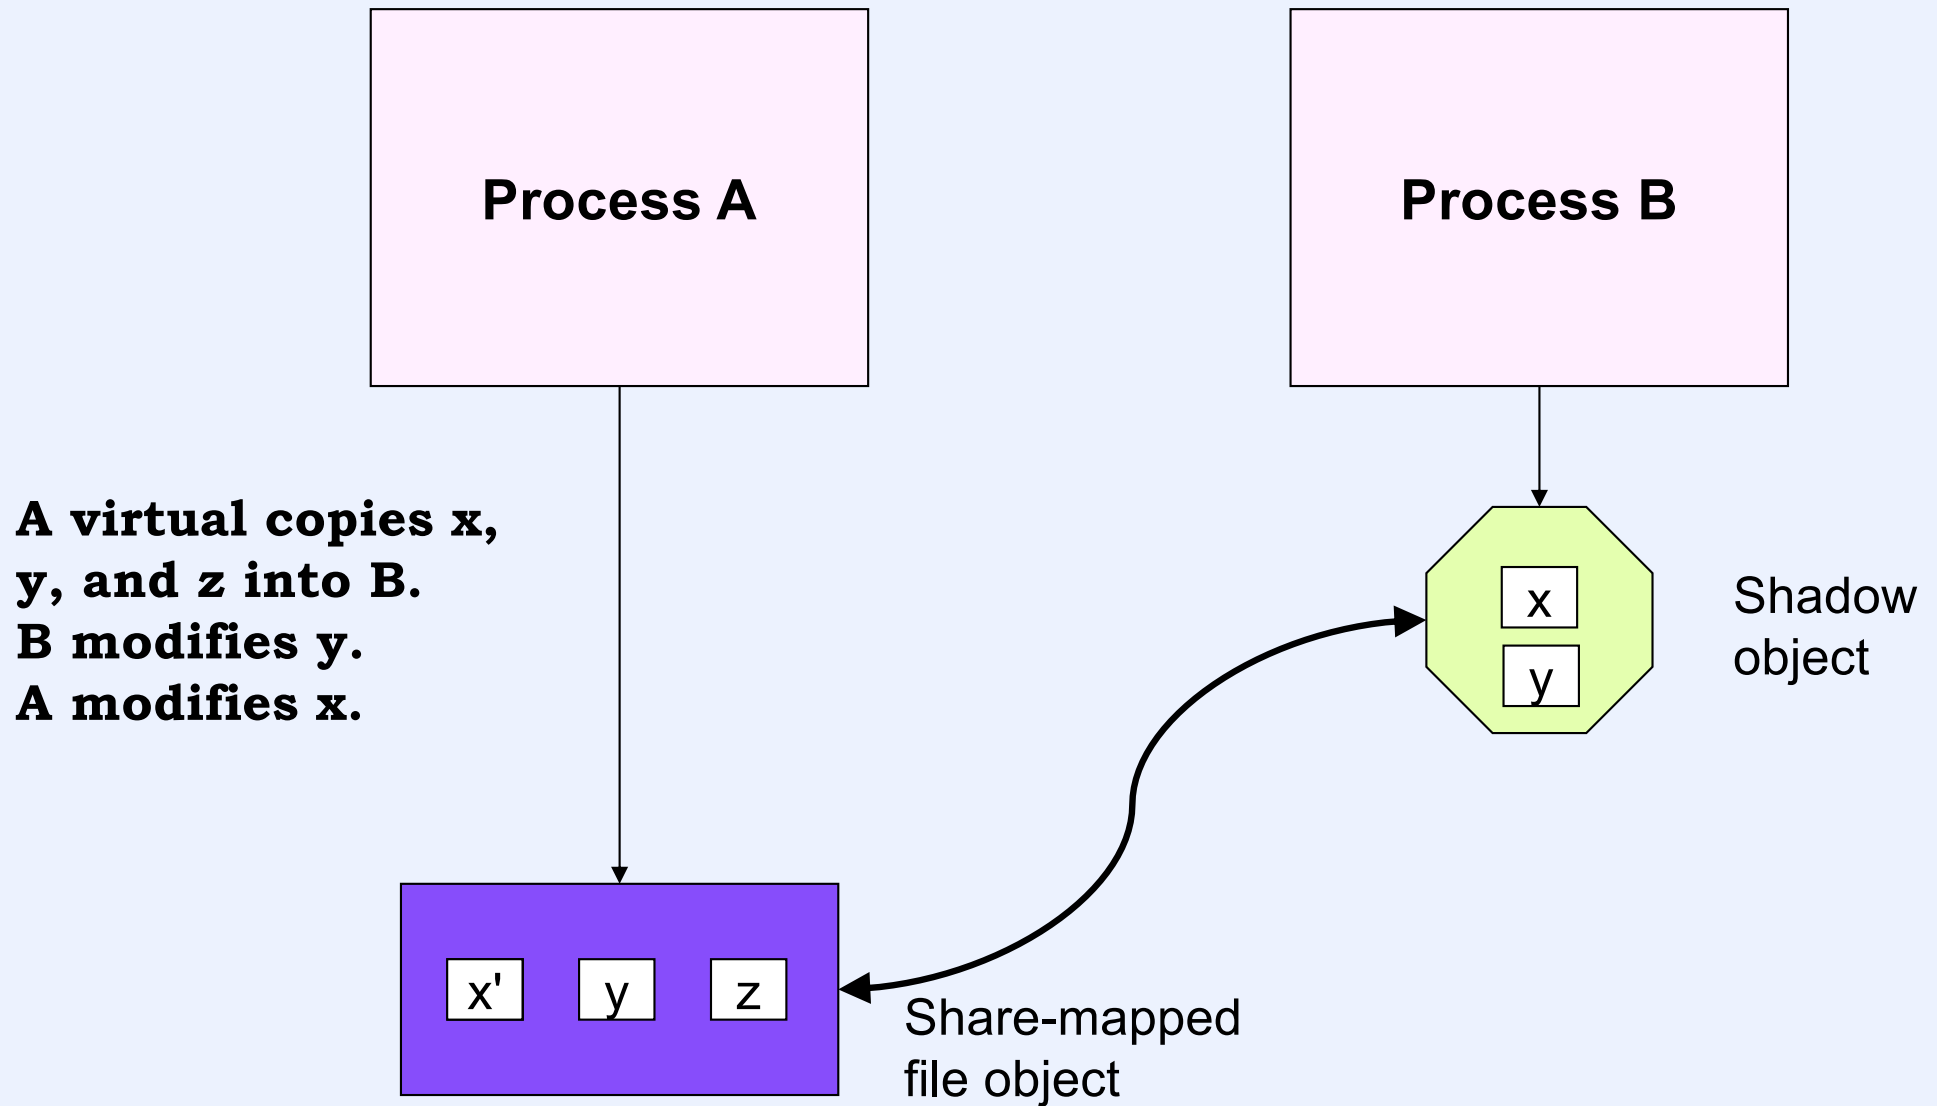
Private Mapping (2)



Private Mapping (3)



Share and Private Mapping



Quiz 5

Unix process X has private-mapped a file into its address space. Our system has one-byte pages and the file consists of four pages. The pages are mapped into locations 100 through 103. The initial values of these pages are all zeroes.

- 1) X stores a 1 into location 100
- 2) X forks, creating process Y
- 3) X stores a 1 into location 101
- 4) Y stores a 2 into location 102
- 5) Y forks, creating process Z
- 6) X stores 111 into location 100
- 7) Y stores 222 into location 103
- 8) Z sums the contents of locations 100, 101, and 102, and stores them into location 103

Answer:

- a) 0
- b) 3
- c) 4
- d) 113

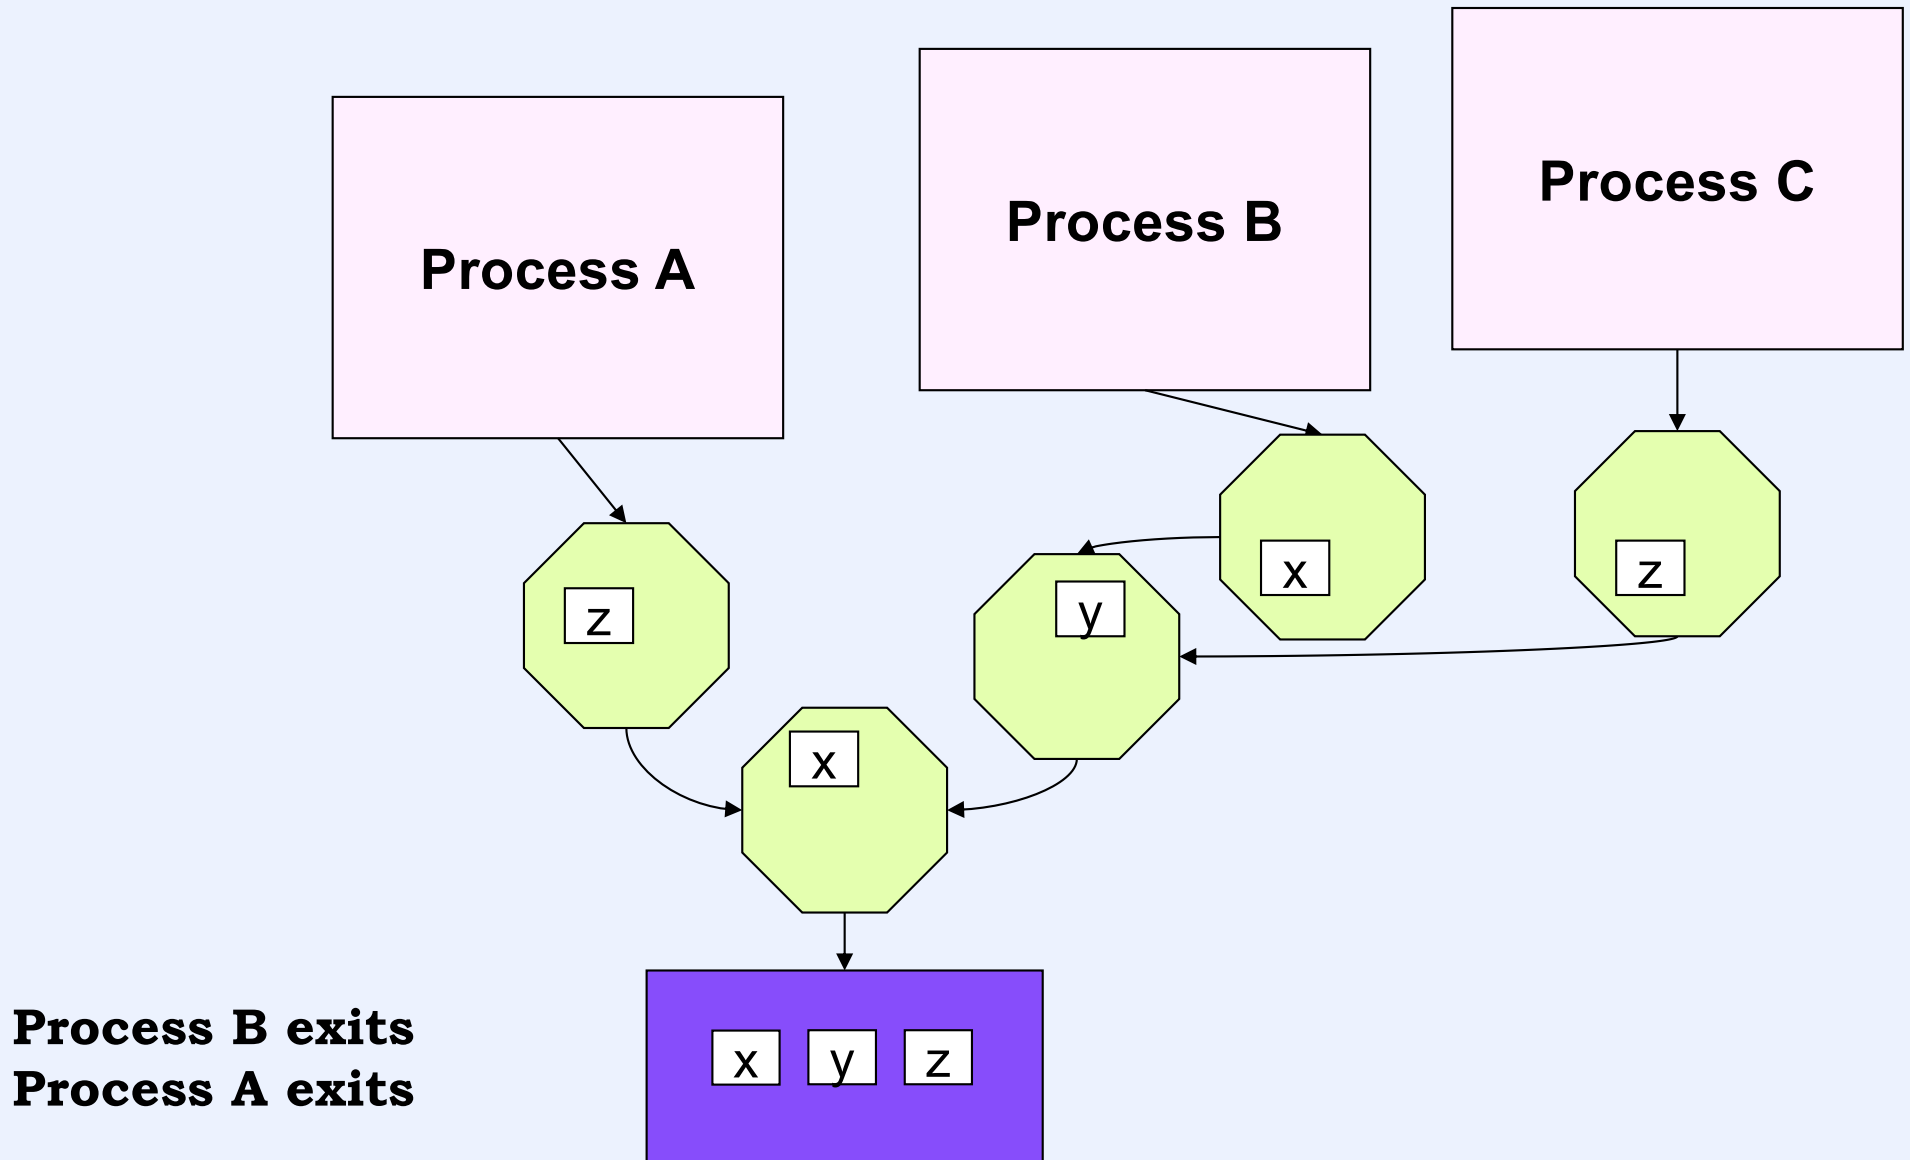
What value did Z store into 103?

Fork Bomb!


```
int main() {  
    while (1) {  
        if (fork() <= 0)  
            exit(0);  
    }  
    return 0;  
}
```

```
int main() {  
    while (1) {  
        if (fork() > 0)  
            exit(0);  
    }  
    return 0;  
}
```

Private Mapping (Continued)



The Backing Store



Backing Up Pages (1)

- **Read-only mapping of a file (e.g. text)**
 - pages come from the file, but, since they are never modified, they never need to be written back
- **Read-write shared mapping of a file (e.g. via *mmap* system call)**
 - pages come from the file, modified pages are written back to the file

Backing Up Pages (2)

- Read-write private mapping of a file (e.g. the data section as well as memory mapped private by the *mmap* system call)
 - pages come from the file, but modified pages, associated with shadow objects, must be backed up in swap space
- Anonymous memory (e.g. bss, stack, and shared memory)
 - pages are created as *zero fill on demand*; they must be backed up in swap space

Swap Space

- **Space management possibilities**
 - **radical-conservative approach: pre-allocation**
 - **backing-store space is allocated when virtual memory is allocated**
 - **page outs always succeed**
 - **might need to have much more backing store than needed**
 - **radical-liberal approach: lazy evaluation**
 - **backing-store space is allocated only when needed**
 - **page outs could fail because of no space**
 - **can get by with minimal backing-store space**

Space Allocation in Linux

- Total memory = primary + swap space
- System-wide parameter:
overcommit_memory
 - three possibilities
 - maybe (default)
 - always
 - never
- mmap has MAP_NORESERVE flag
 - don't worry about over-committing

Space Allocation in Windows

- **Space reservation**
 - allocation of virtual memory
- **Space commitment**
 - reservation of physical resources
 - paging space + physical memory
- ***MapViewOfFile* (sort of like *mmap*)**
 - no over-commitment
- **Thread creation**
 - creator specifies both reservation and commitment for stack pages