

File Systems Part 3

Transactions

- **“ACID” property:**
 - **atomic**
 - **all or nothing**
 - **consistent**
 - **take system from one consistent state to another**
 - **isolated**
 - **have no effect on other transactions until committed**
 - **durable**
 - **persists**

How?

- **Journaling**
 - before updating disk with steps of transaction:
 - record previous contents: *undo journaling*
 - record new contents: *redo journaling*
- **Shadow paging**
 - steps of transaction written to disk, but old values remain
 - single write switches old state to new

Data vs. Metadata

- **Metadata**
 - system-maintained data pertaining to the structure of the file system
 - inodes
 - indirect, doubly indirect, triply indirect blocks
 - directories
 - free space description
 - etc.
- **Data**
 - data written via write system calls

Journaling

- **Journaling options**
 - **journal everything**
 - **everything on disk made consistent after crash**
 - **last few updates possibly lost**
 - **expensive**
 - **journal metadata only**
 - **metadata made consistent after a crash**
 - **user data not**
 - **last few updates possibly lost**
 - **relatively cheap**

Committing vs. Checkpointing

- **Checkpointed updates**
 - written to file system and are thus permanent
- **Committed updates**
 - not necessarily written to file system, but guaranteed to be written eventually (checkpointed), even if there is a crash
- **Uncommitted updates**
 - not necessarily written to file system (yet), may disappear if there is a crash

Ext3

- **A journaled file system used in Linux**
 - **same on-disk format as Ext2 (except for the journal)**
 - **(Ext2 is an FFS clone)**
 - **supports both full journaling and metadata-only**
 - **does redo journaling**

Full Journaling in Ext3

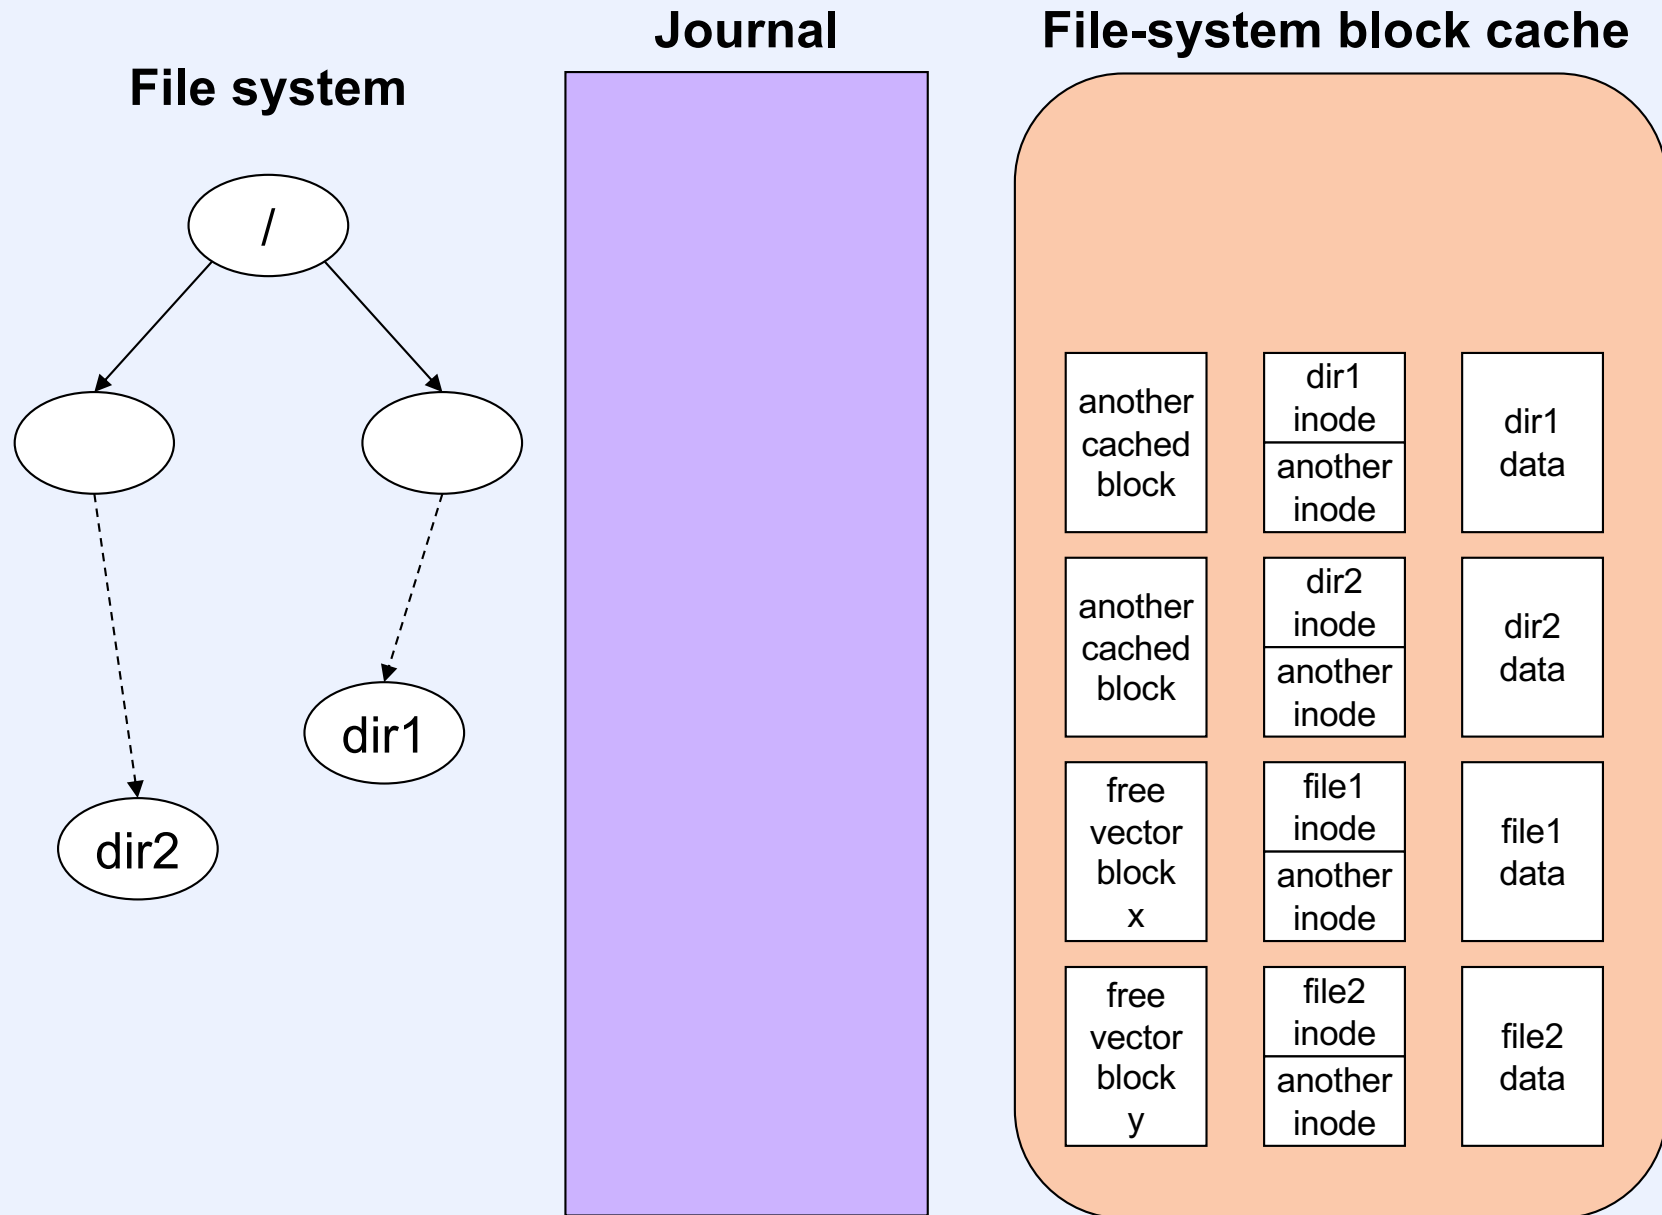
- **File-oriented system calls divided into subtransactions**
 - updates go to cache only
 - subtransactions grouped together
- **When sufficient quantity collected or 5 seconds elapsed, *commit* processing starts**
 - updates (new values) written to journal
 - once entire batch is journaled, end-of-transaction record is written
- **Cached updates are then *checkpointed* — written to file system**
 - journal cleared after checkpointing completes

Quiz 1

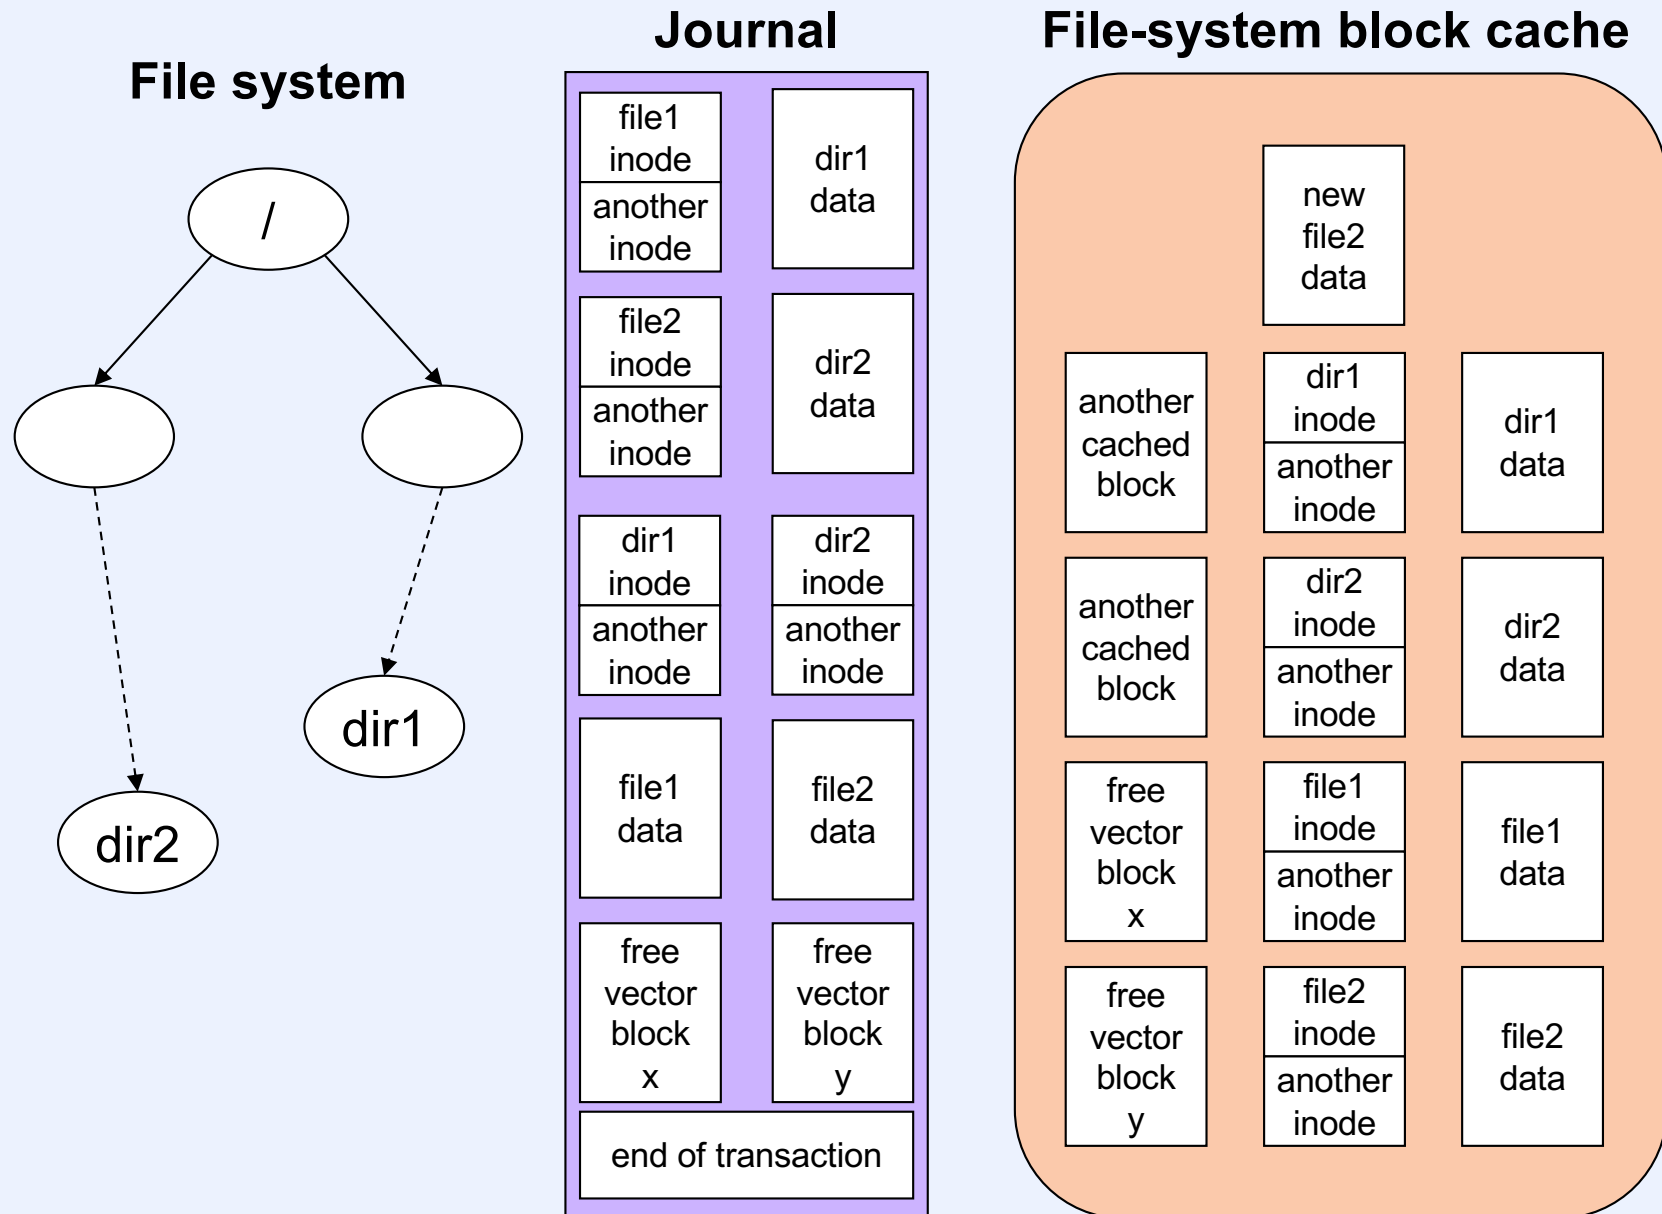
You have a Linux system with an Ext3 file system with full journaling. You run a script that deletes some files, creates some new files and writes data to them, then renames the new files. This takes two seconds. Immediately after the script finishes, there's a power failure and the system crashes. When it comes back up, after crash recovery:

- a) some later parts of the script may have completed, even though earlier parts did not**
 - b) it will appear as if the script ran to some point and then terminated (this allows for both none of it and all of it)**
 - c) it will definitely appear as if the script ran to completion**
-

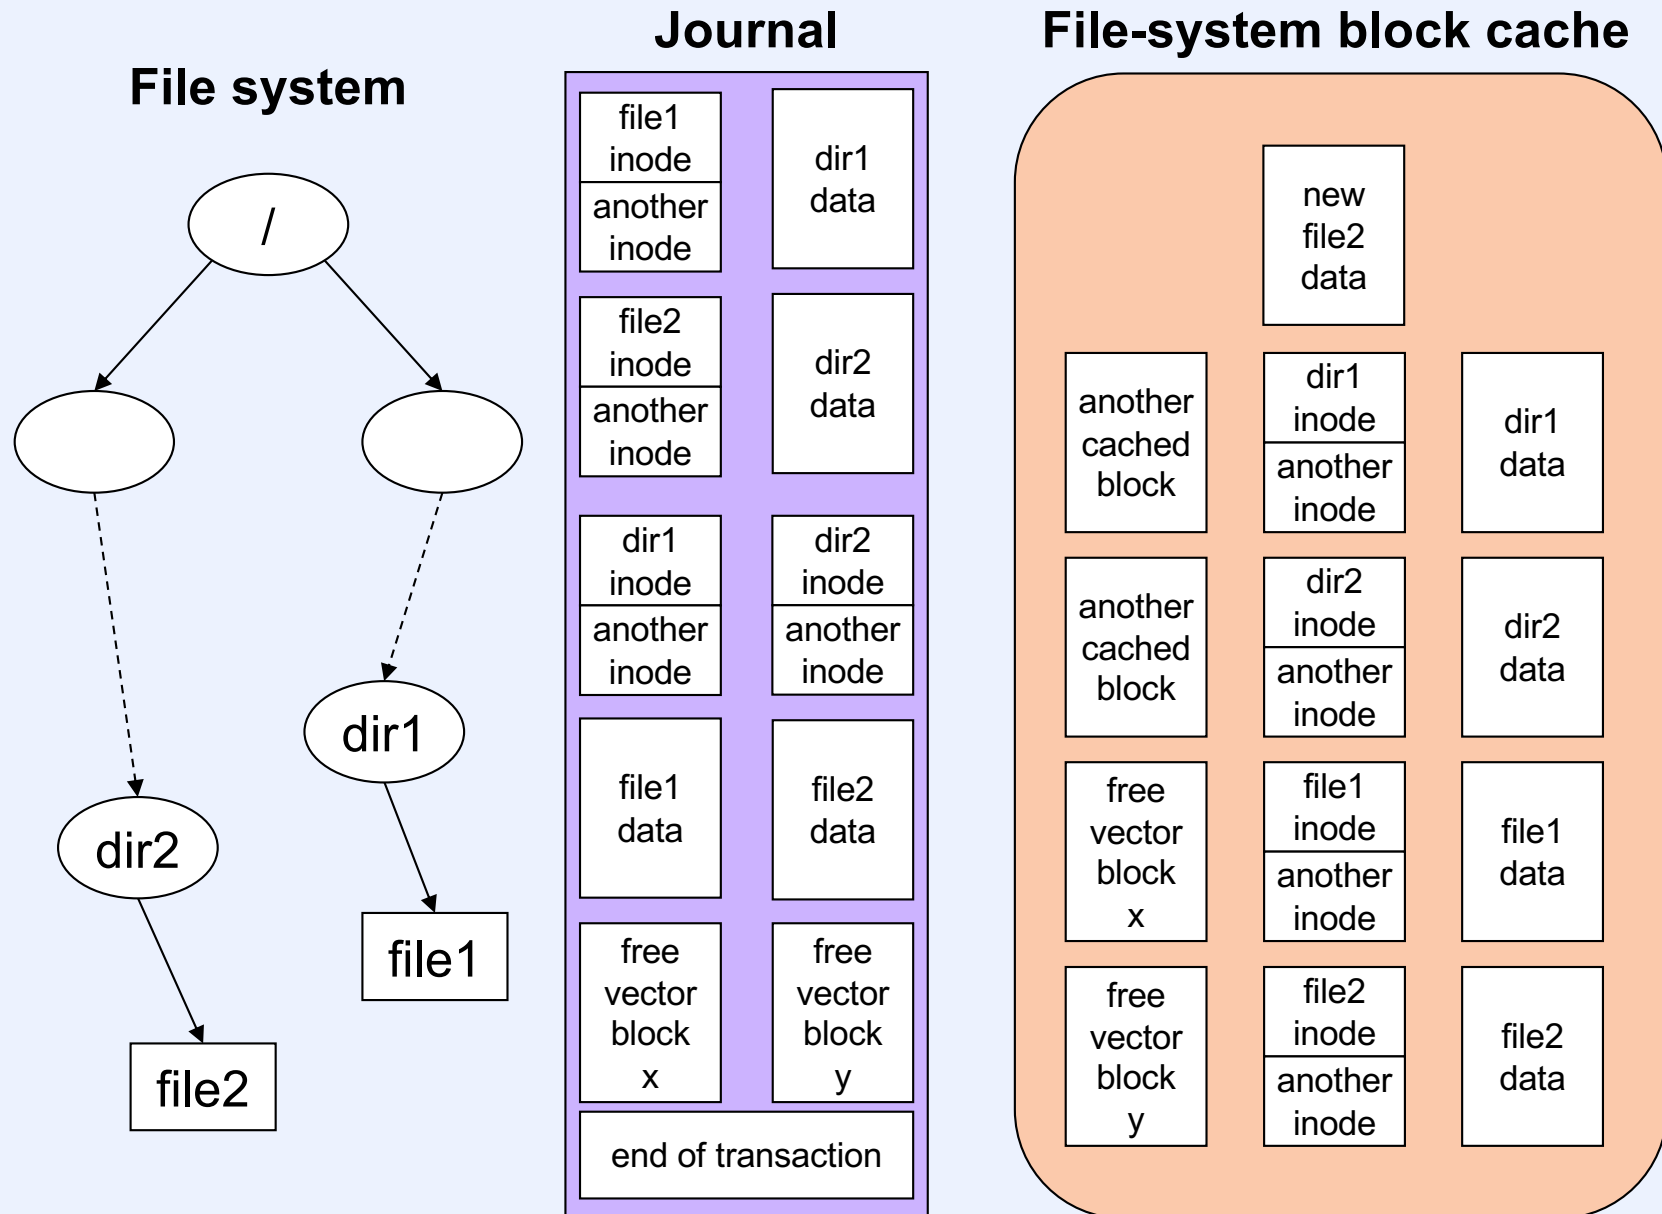
Journaling in Ext3 (part 1)



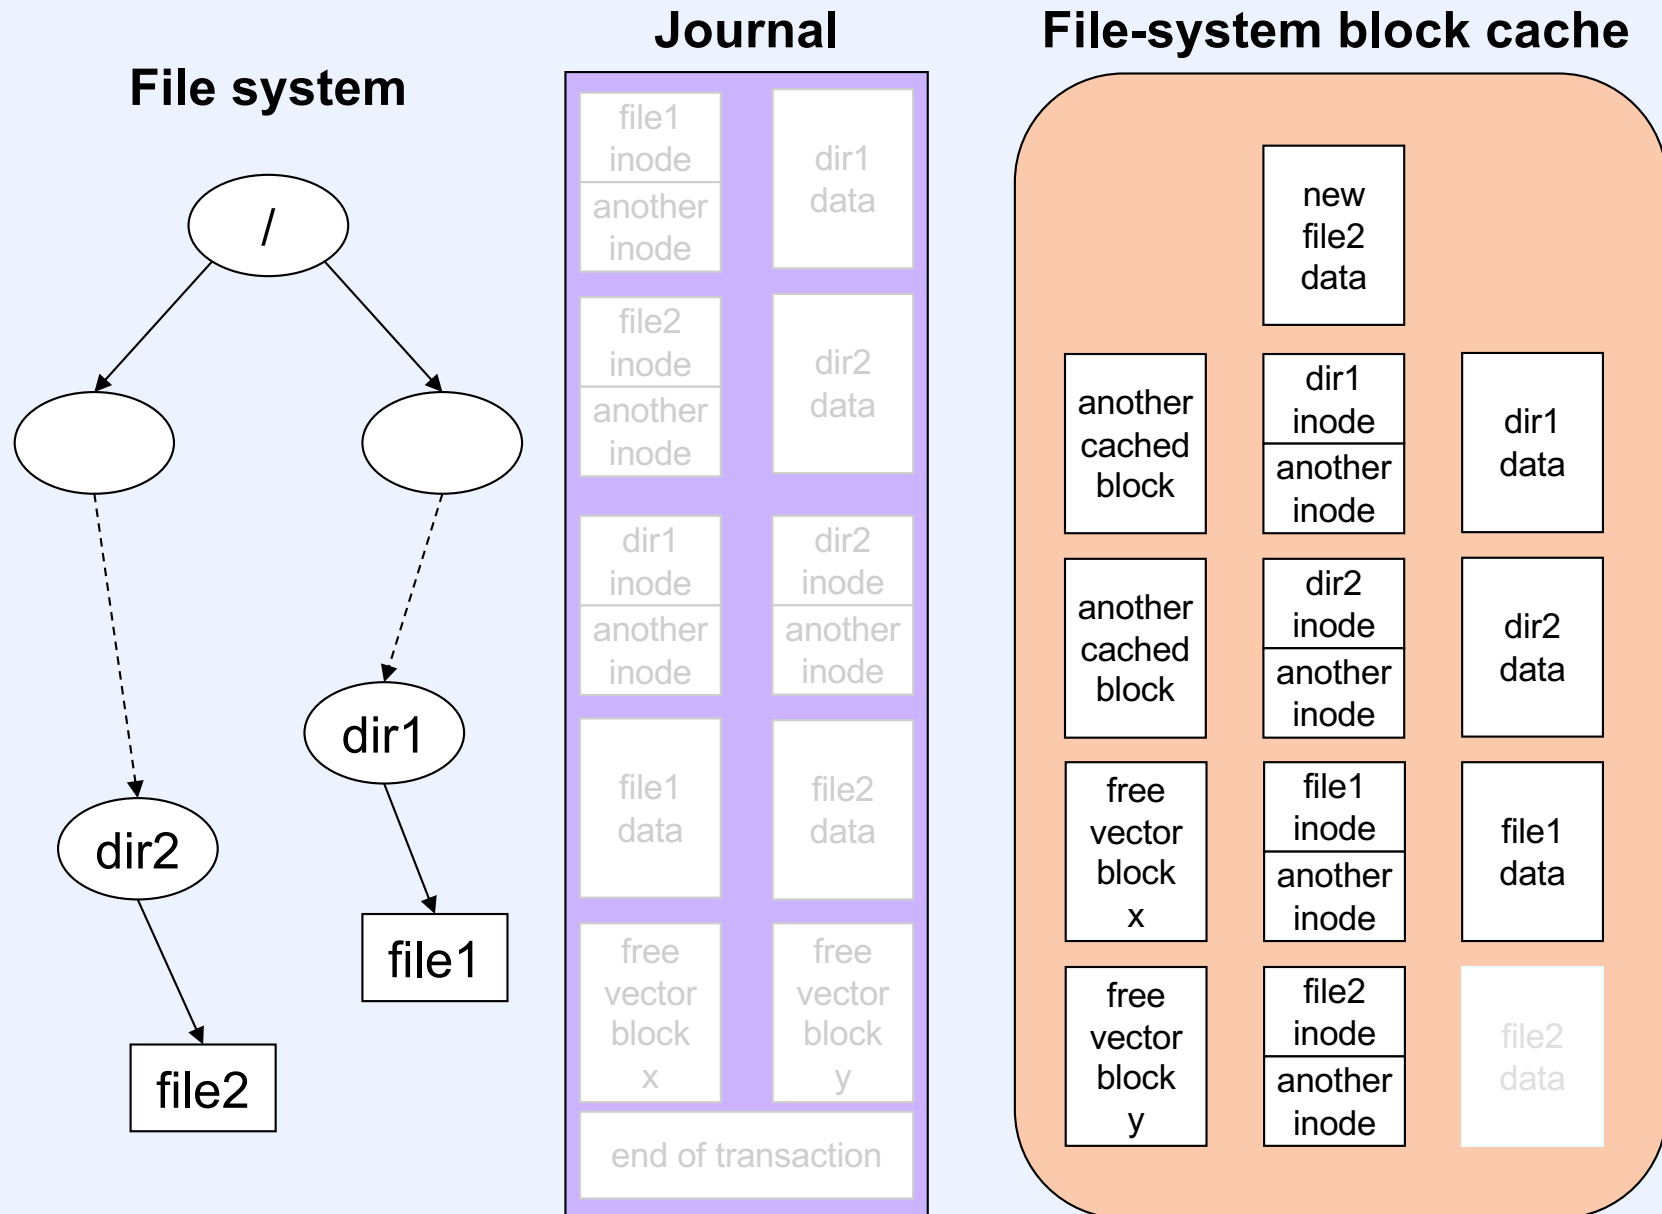
Journaling in Ext3 (part 2)



Journaling in Ext3 (part 3)



Journaling in Ext3 (part 4)



Quiz 2

You have a Linux system with an Ext3 file system with metadata-only journaling. You run a script that deletes some files, creates some new files and writes data to them, then renames the new files. This takes two seconds. Immediately after the script finishes, there's a power failure and the system crashes. When it comes back up, after crash recovery:

- a) there may be data written to the new files, but no files were deleted**
 - b) it will definitely appear as if the script ran to some point and then terminated (this allows for both none of it and all of it)**
 - c) the script may appear to have completed, though there's no data written to the new files**
-

Metadata-Only Journaling in Ext3

- **File-oriented system calls divided into subtransactions**
 - updates to metadata go to cache only
 - updates to data go to cache; may be written to disk at any time
- **When sufficient quantity collected or 5 seconds elapsed, *commit* processing starts**
 - metadata updates written to journal
 - once entire batch is journaled, end-of-transaction record is written
- **Cached metadata updates are then *checkpointed* — written to file system**

Metadata-Only Issues

- **Scenario (one of many):**
 - you create a new file and write data to it
 - transaction is committed
 - metadata is in journal
 - user data still in cache
 - system crashes
 - system reboots; journal is recovered
 - new file's metadata are in file system
 - user data are not
 - metadata refer to disk blocks containing other users' data

Coping

- **Zero all disk blocks as they are freed**
 - done in “secure” operating systems
 - expensive
- **Ext3 approach**
 - write newly allocated data blocks to file system before committing metadata to journal
 - fixed?

Yes, but ...

- **Mary deletes file A**
 - A's data block *x* added to free vector
 - **Ted creates file B**
 - **Ted writes to file B**
 - block *x* allocated from free vector
 - new data goes into *x*
 - system writes newly allocated *x* to file system in preparation for committing metadata, but ...
 - **System crashes**
 - metadata did not get journaled
 - A still exists; B does not
 - B's data is in A
-

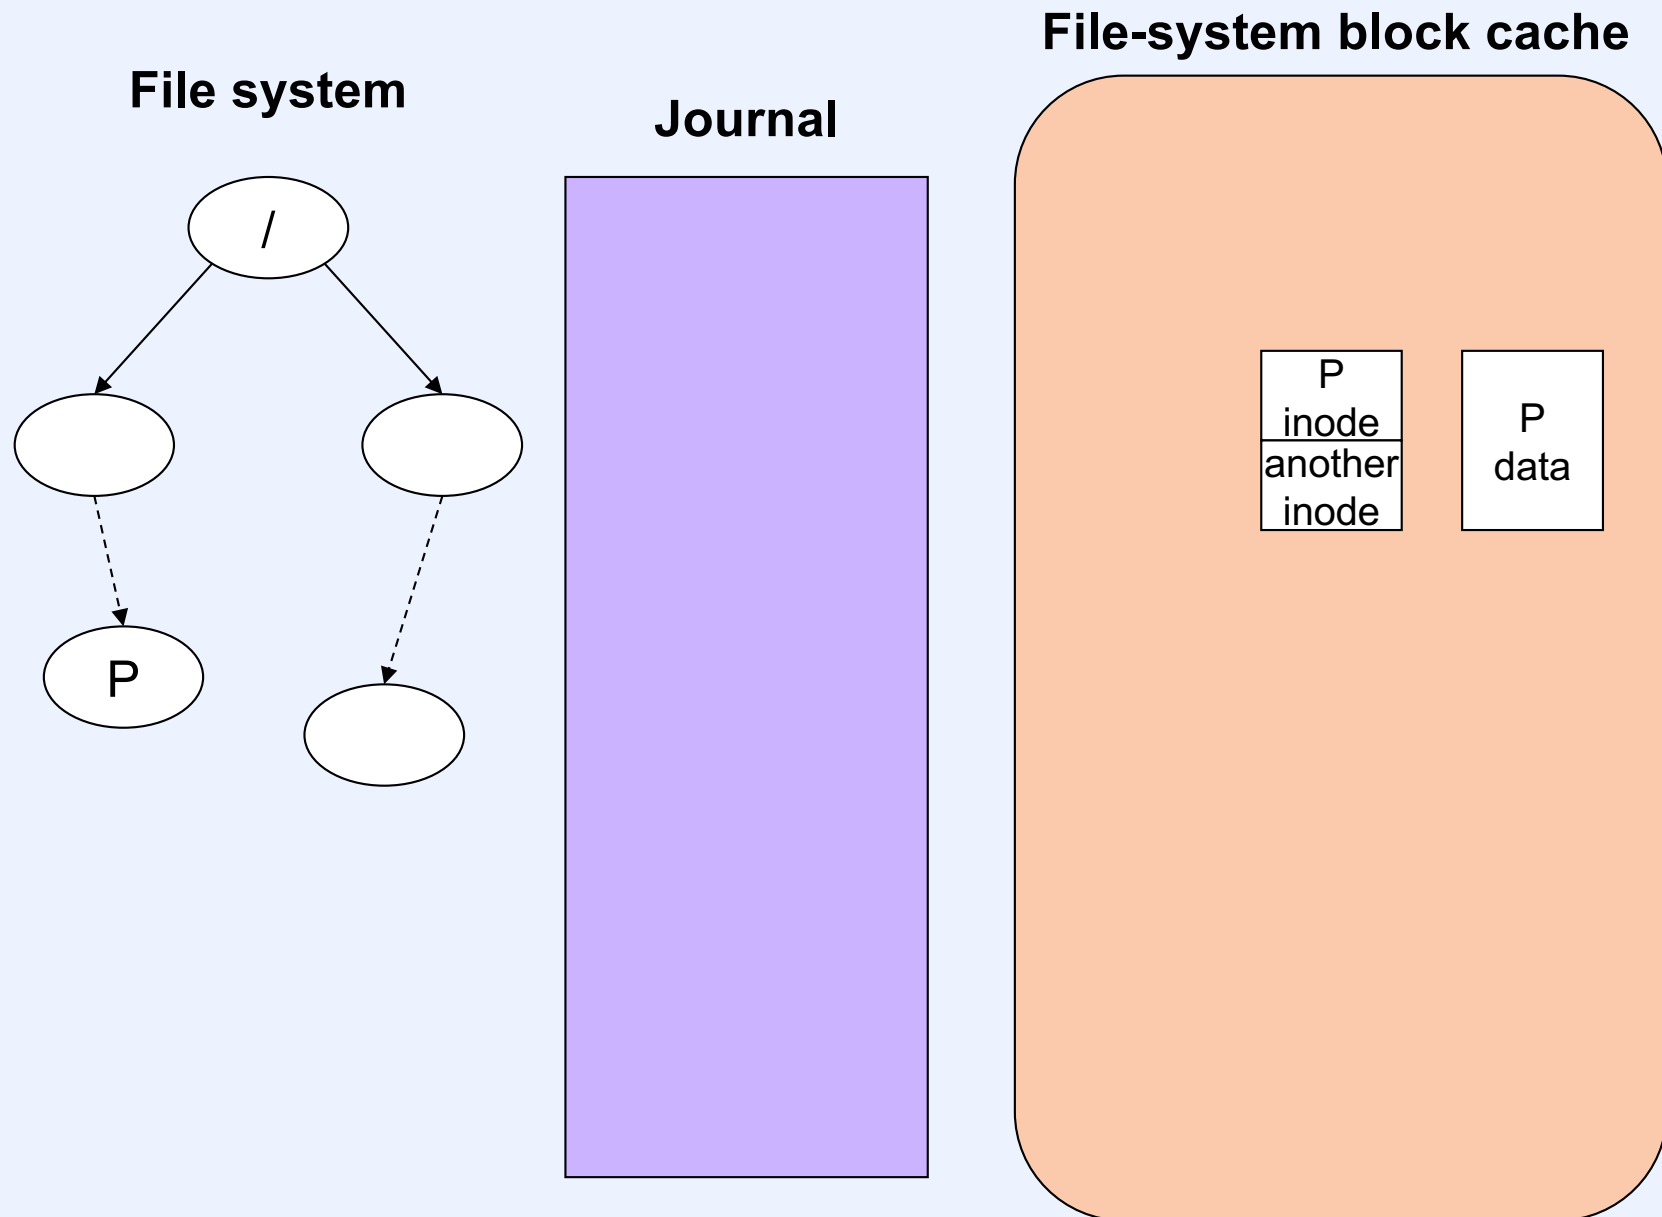
Fixing the Fix

- **Don't reuse a block until transaction freeing it has been committed**
 - keep track of most recently committed free vector
 - allocate from it

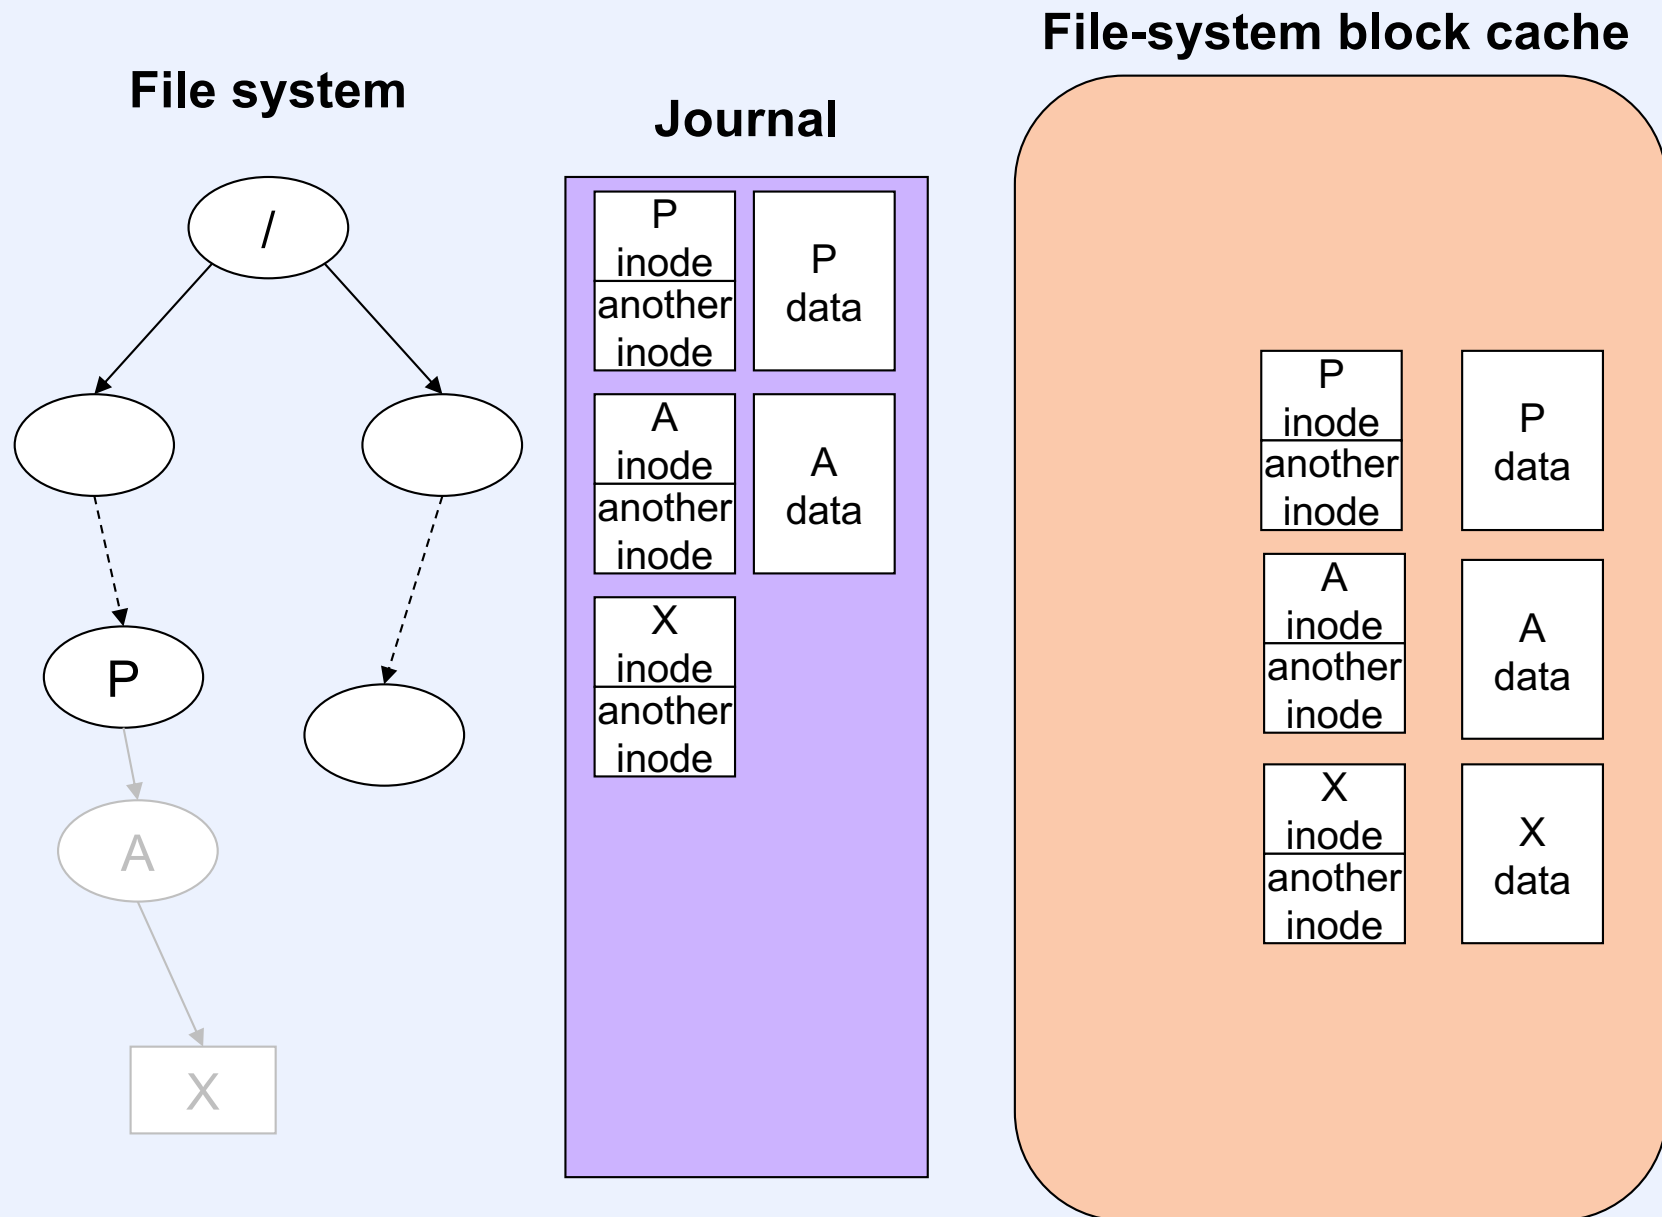
Fixed Now?

- No ...

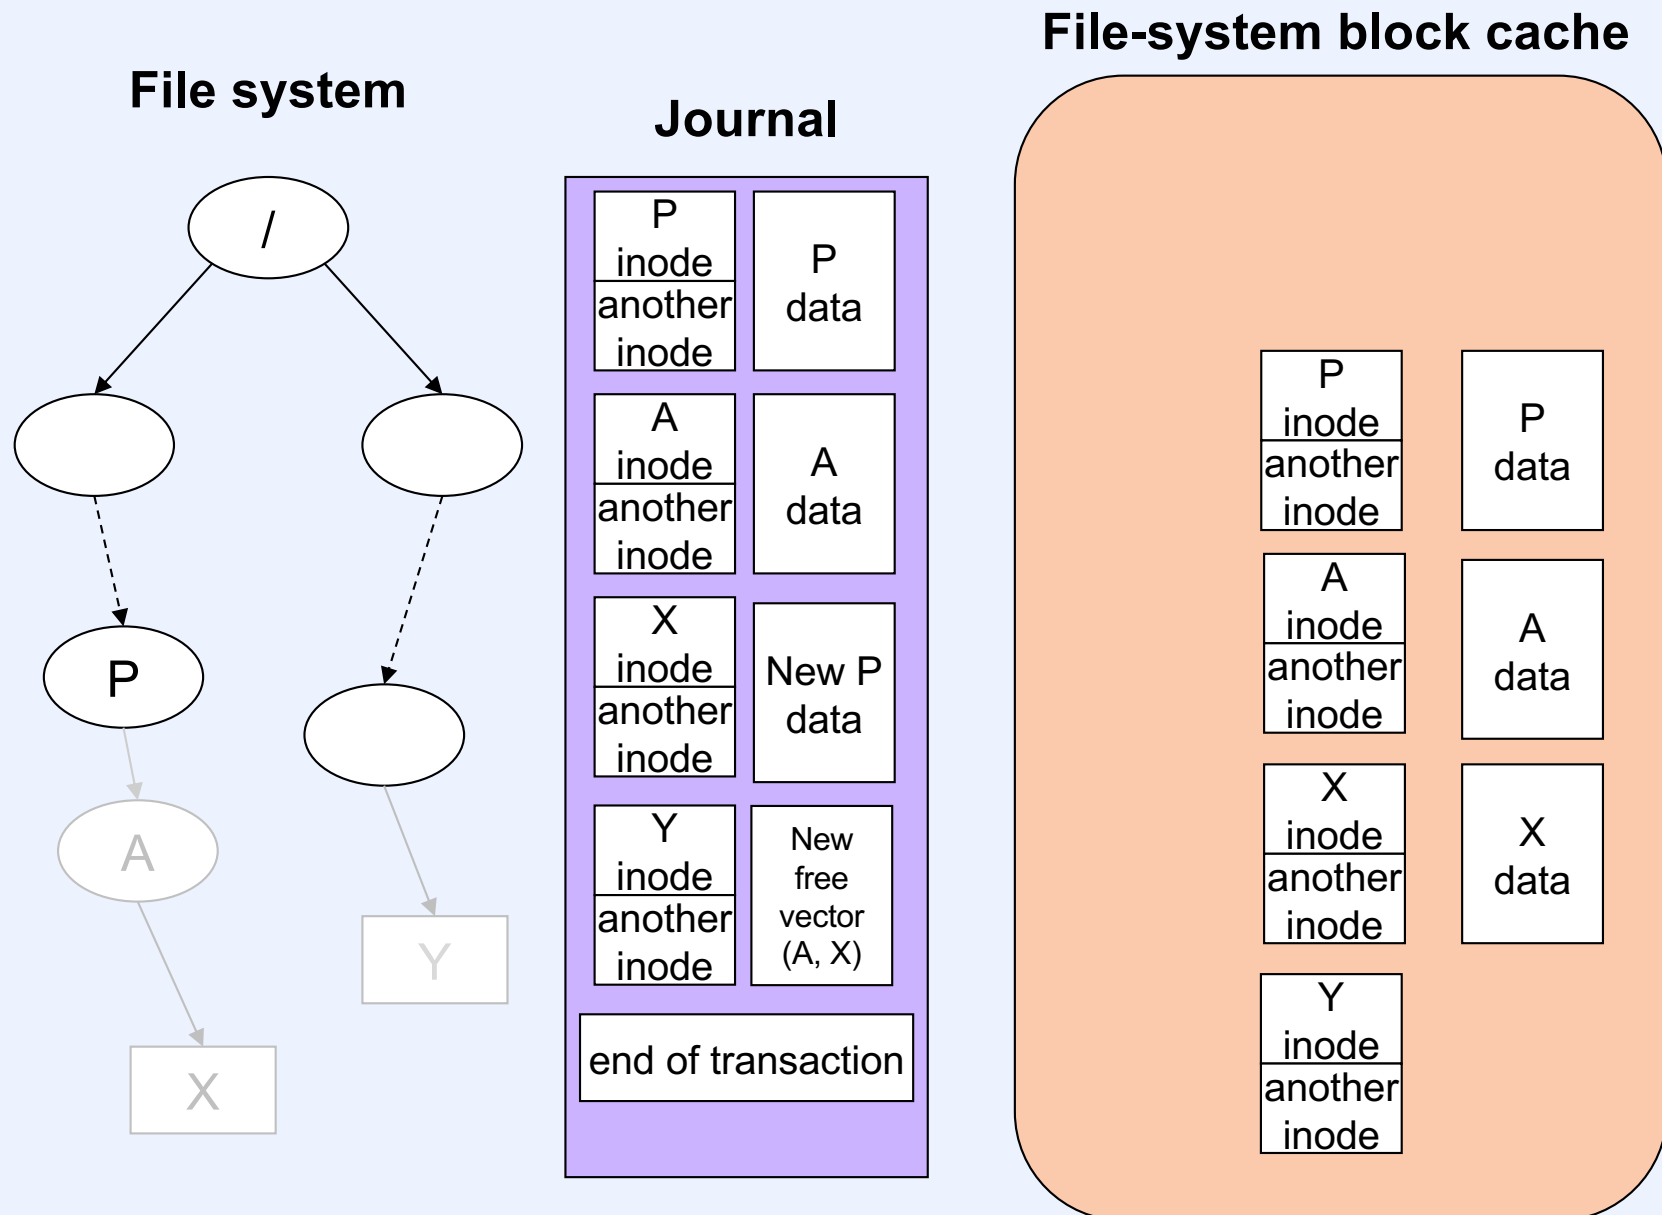
Yet Another Problem (part 1)



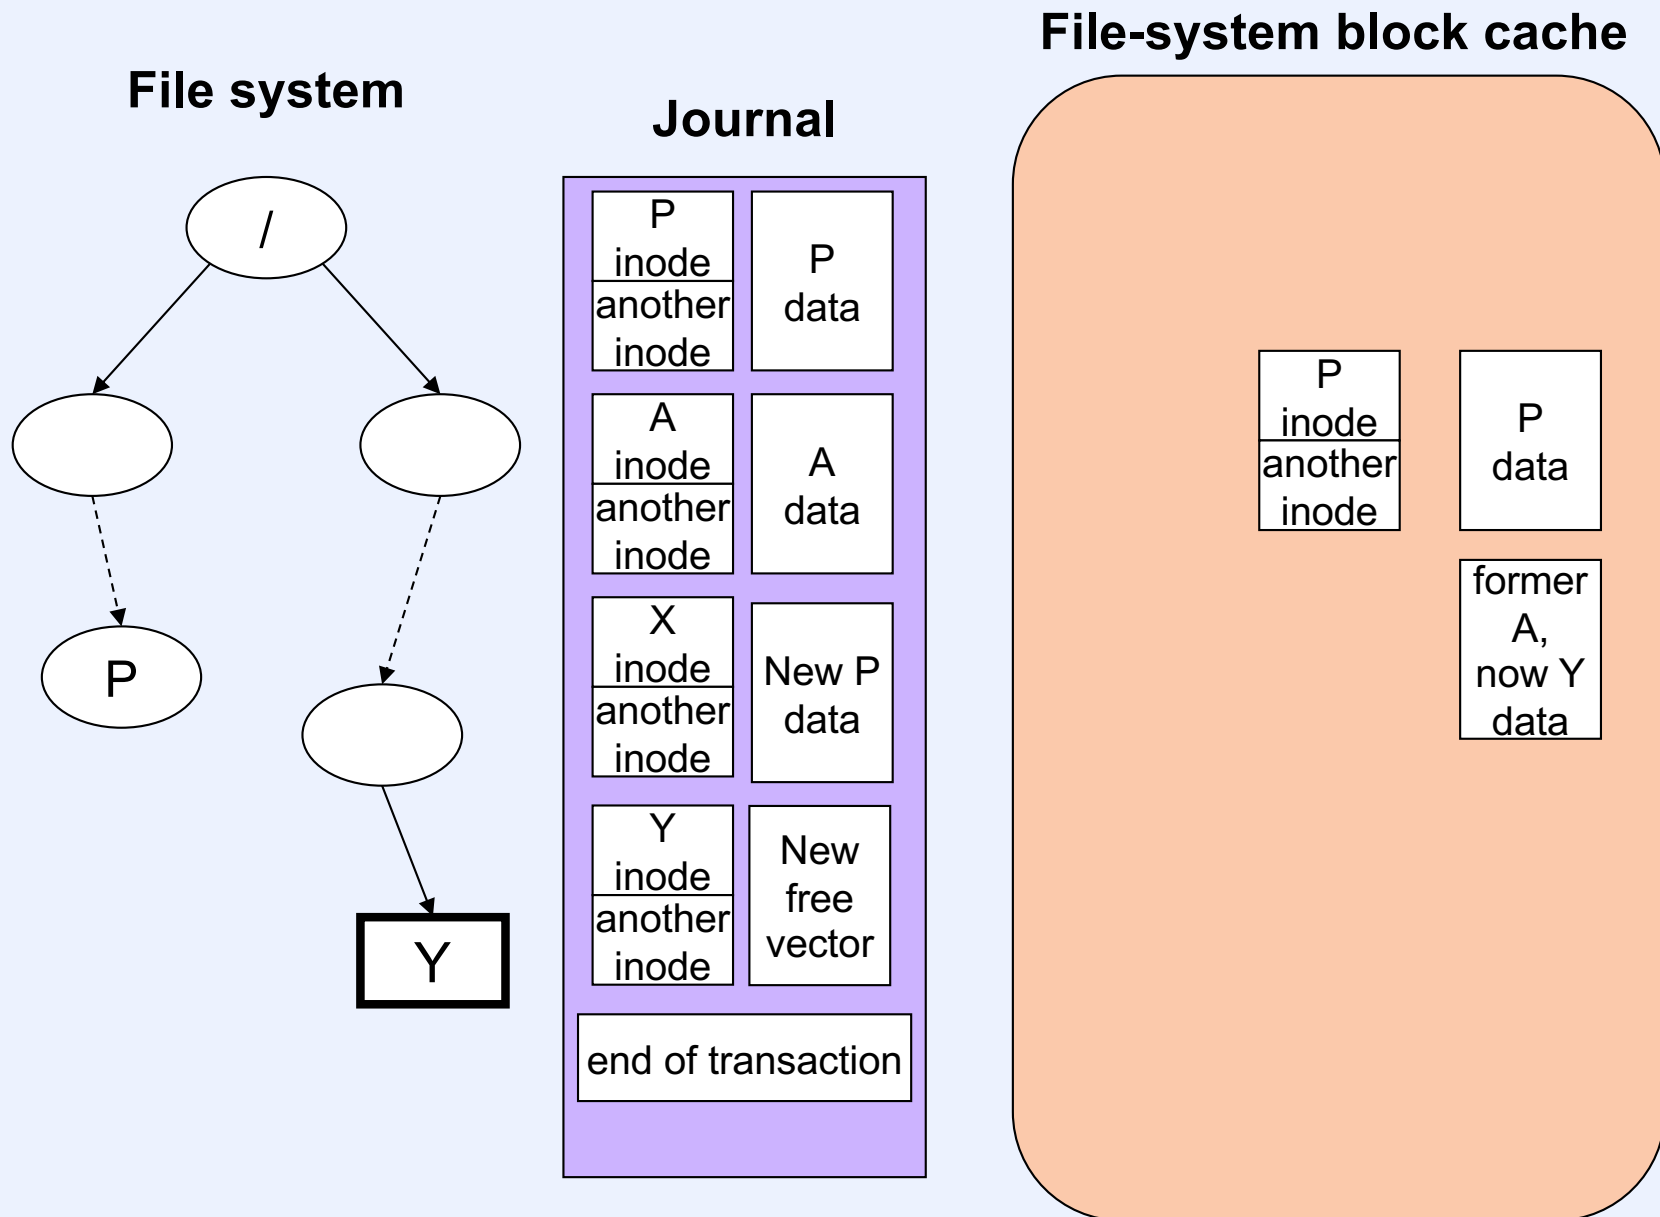
Yet Another Problem (part 2)



Yet Another Problem (part 3)

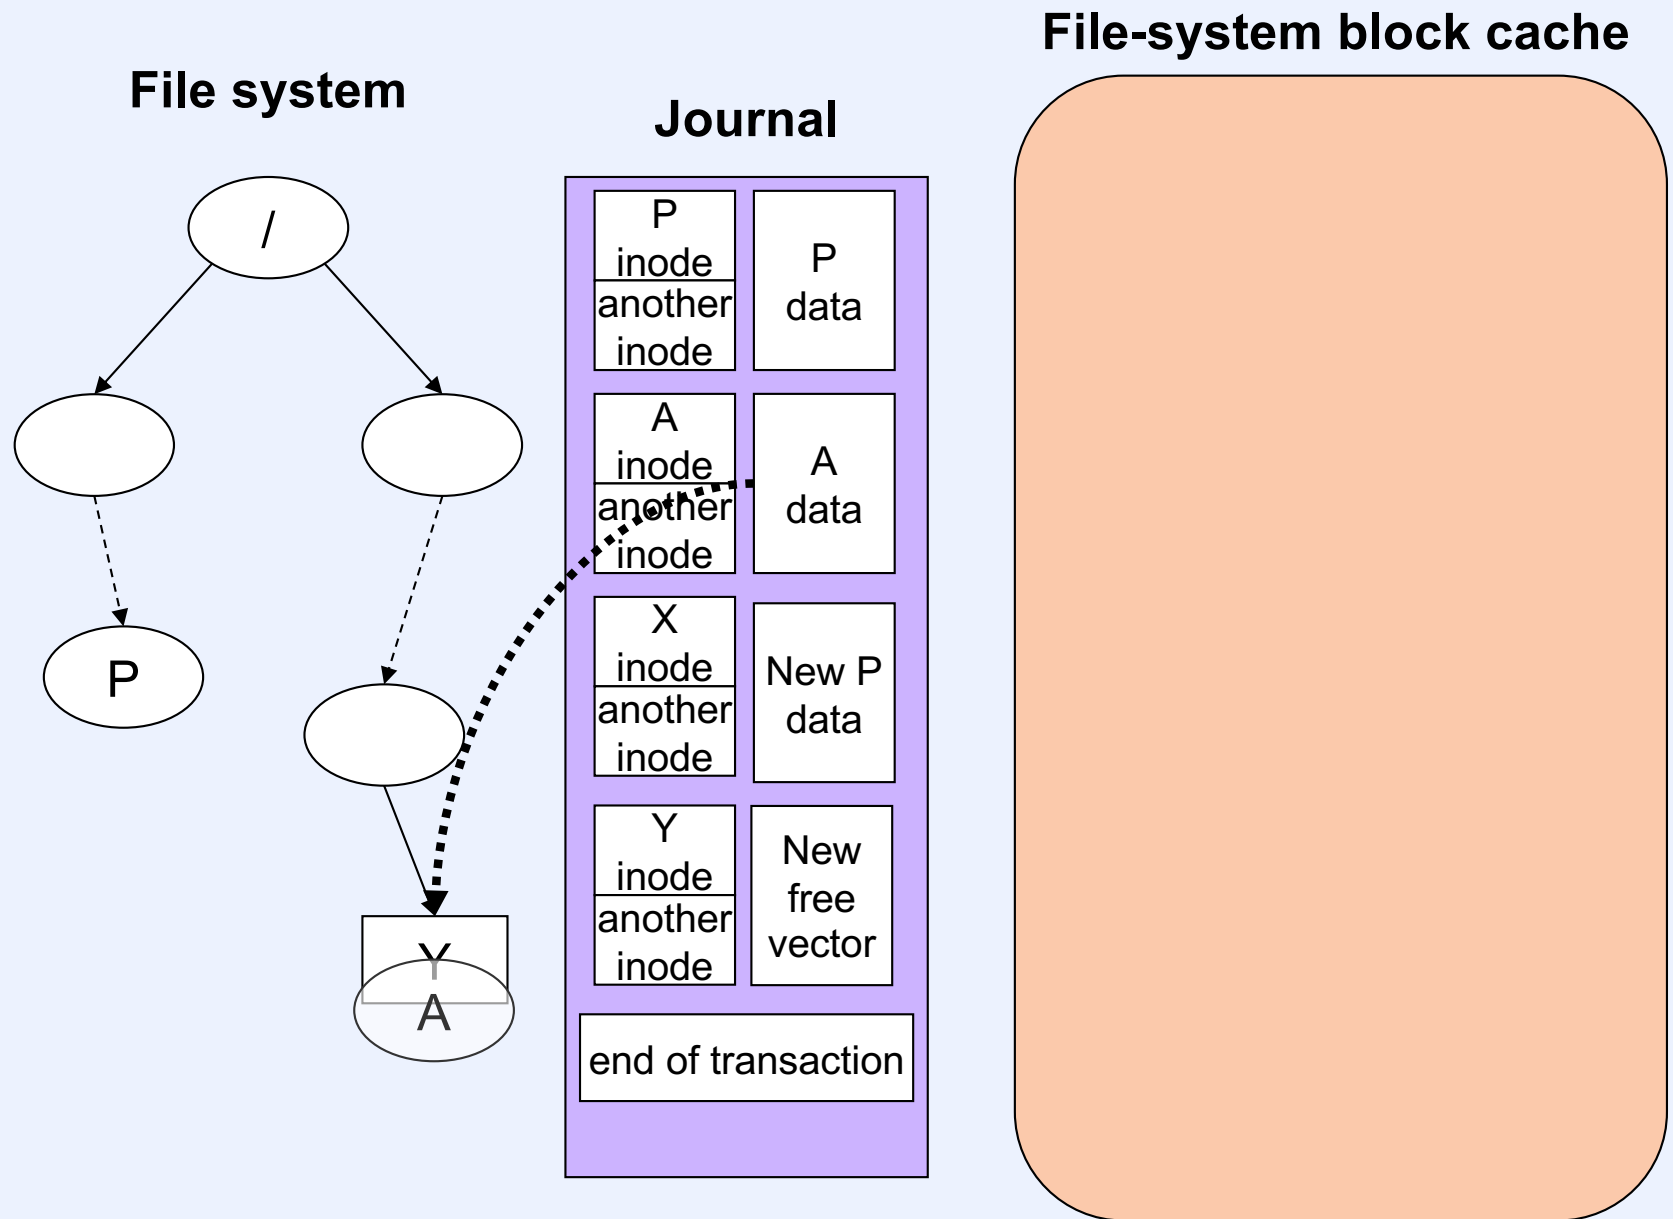


Yet Another Problem (part 4)





Yet Another Problem (part 5)



The Fix

- The problem occurs because metadata is modified, then deleted.
- Don't blindly do both operations as part of crash recovery
 - no need to modify the metadata!
 - Ext3 puts a “revoke” record in the journal, which means “never mind ...”

Fixed Now?

- **Yes!**
 - (or, at least, it seems to work ...)

Ext4

- **Latest Linux file system**
 - used at Brown CS as local FS on Linux systems
- **Retains much of Ext3**
 - journaling
 - meta-data only used at Brown CS
 - inodes
- **Adds extents**
 - four extents in inode
 - if more needed, B-tree is used

Undo Journaling

- **Old data written to journal (from cache)**
- **New data written to cache; written to file system when convenient**
- **Committed after checkpointing**
 - new data is on disk
 - (undo) journal entries removed
- **Crash recovery**
 - if transaction not committed, all changes to file system are undone by playing back the journal

Undo vs. Redo

- **Redo**
 - data written to cache
 - data copied to journal
 - committed when entire transaction is in journal
 - data stays in cache until it is checkpointed
- **Undo**
 - old data written to journal (from cache)
 - new data written to cache
 - new data written to disk
 - committed when all data on disk
 - data may be removed from cache prior to committing

Quiz 3


If undo journaling is being used, when can a transaction be committed?

- a) when the new contents of all blocks involved in the transaction have been written to the cache**
- b) when the new contents of all blocks involved in the transaction have been checkpointed to the file system on disk**
- c) when the old contents of all blocks involved in the transaction have been written to the journal**


Shadow Paging

- Refreshingly simple
- Provides historical snapshots
- Examples
 - WAFL (Network Appliance)
 - ZFS (Sun)


Shadow-Page Tree



Shadow-Page Tree: Modifying a Node



Shadow-Page Tree: Propagating Changes



Benefits

- Each update results in a new shadow-page tree (having much in common with the previous one)
- The current root identifies the current tree
- If the system crashes after an update has been made, but before changes are propagated to the new root, the update is lost
 - a single write (to the root) effectively serves as a commit
- Older roots refer to previous states of the file system – snapshots

Quiz 4

When the shadow-page tree is updated:

- a) file-system data may be cached (and written asynchronously) as long as the root is written last**
- b) file-system data may be cached (and written asynchronously) only if all lower parts of the tree are written to disk before upper parts**
- c) all file-system data must be written to disk synchronously: writes may be cached for the sake of reads, but write system calls may not return until the data is on disk**