

# **File Systems Part 7 Continued**

# s5fs\_get\_pframe (1)

```
1 static long s5fs_get_pframe(vnode_t *vnode,  
    uint64_t pagenum, long forwrite, pframe_t **pfp) {  
2     if (vnode->vn_len <= pagenum * PAGE_SIZE)  
3         return -EINVAL;  
4     mobj_find_pframe(&vnode->vn_mobj, pagenum, pfp);  
5     if (*pfp) {  
6         // block is cached  
7         (*pfp)->pf_dirty |= forwrite;  
8         return 0;  
9     }
```

## s5fs\_get\_pframe (2)

```
10     int new;
11     long loc = s5_file_block_to_disk_block(
        VNODE_TO_S5NODE(vnode), pagenum, forwrite, &new);
12     if (loc < 0) return loc;
13     if (loc) {
14         if (new) {
15             *pfp = s5_cache_and_clear_block(
                &vnode->vn_mobj, pagenum, loc);
16         } else
17             s5_get_file_disk_block(vnode, pagenum, loc,
                forwrite, pfp);
18     return 0;
19 }
```

# s5fs\_get\_pframe (3)

```
20     else {
21         KASSERT(!forwrite);
22         return mobj_default_get_pframe(
23             &vn->vn_mobj, pagenum, forwrite, pfp);
24     }
25 }
```

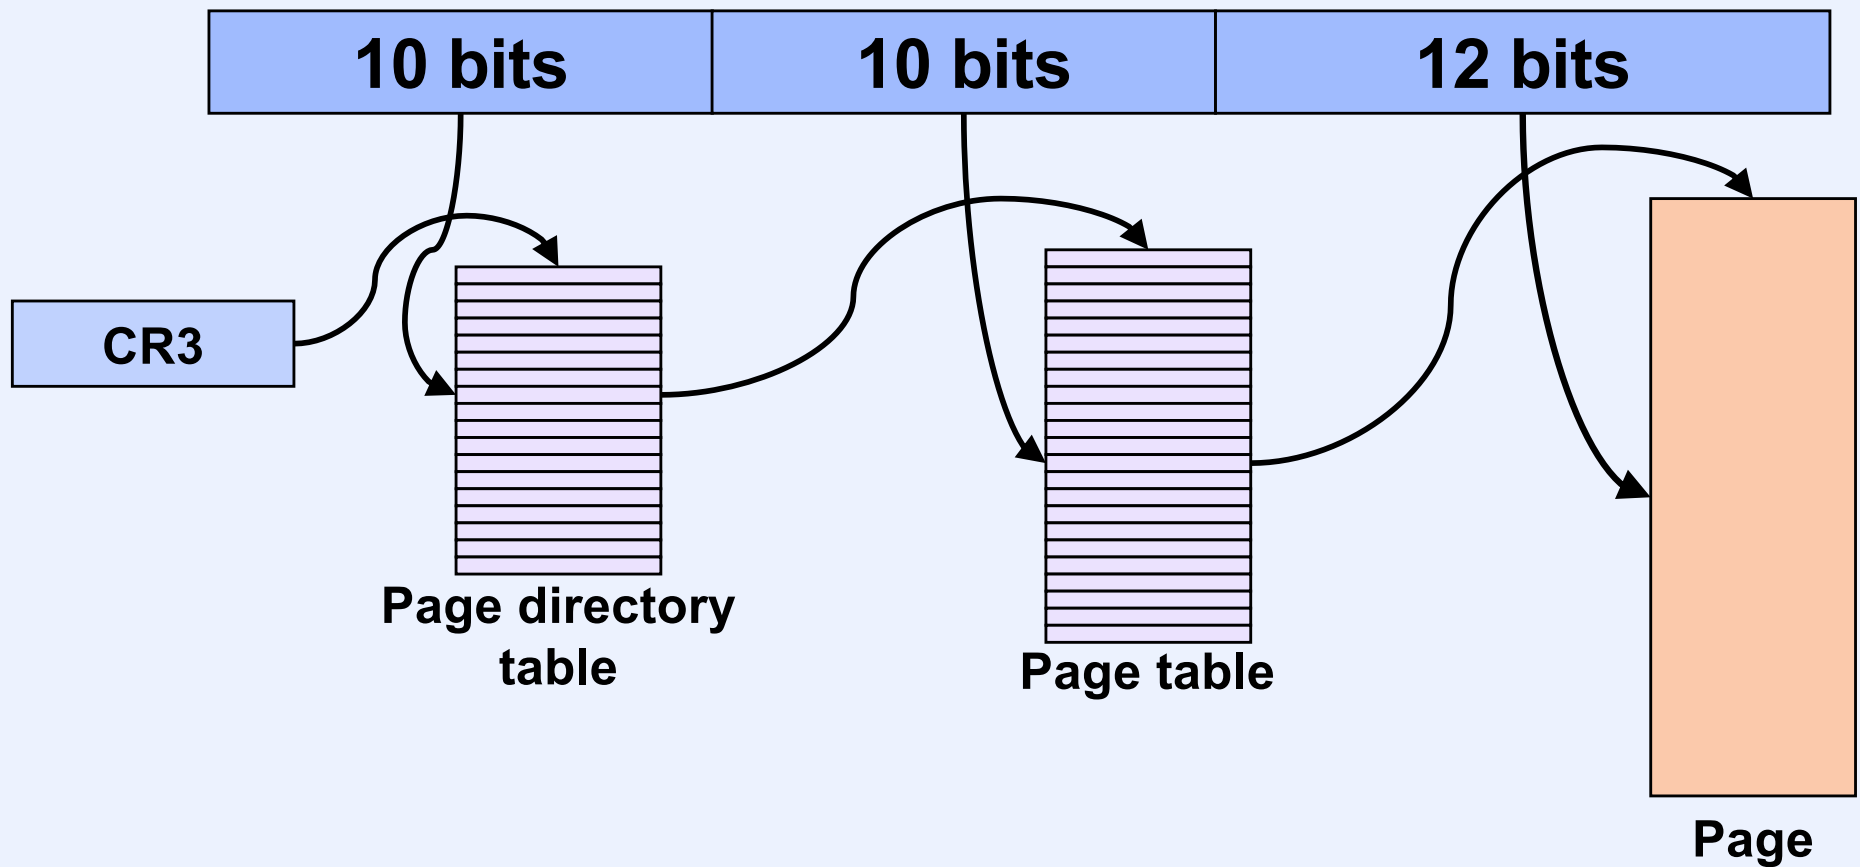
# Quiz 1

Suppose a thread does a *read* system call (which calls *s5fs\_get\_pframe*) to read a portion of a block that is sparse. It then writes data to the block, using the *write* system call. Will, as part of handling this *write*, *mobj\_default\_get\_pframe* be called?

- a) no, because the block was originally a sparse block
- b) no, the block doesn't need to be zeroed and the caller of *s5fs\_get\_pframe* will have put data into it
- c) yes, since the block is sparse, *mobj\_default\_get\_pframe* must be called to zero the block, then modify a portion of it
- d) yes, for some other reason

# Memory Management Part 2

# IA32 Paging




# Quiz 2

**Suppose a process on an IA32 has exactly one page residing in real memory. What is the total number of combined pages of page-directory table and page tables required to map this page?**


- a) 1**
- b) 2**
- c) 4**
- d) 8**



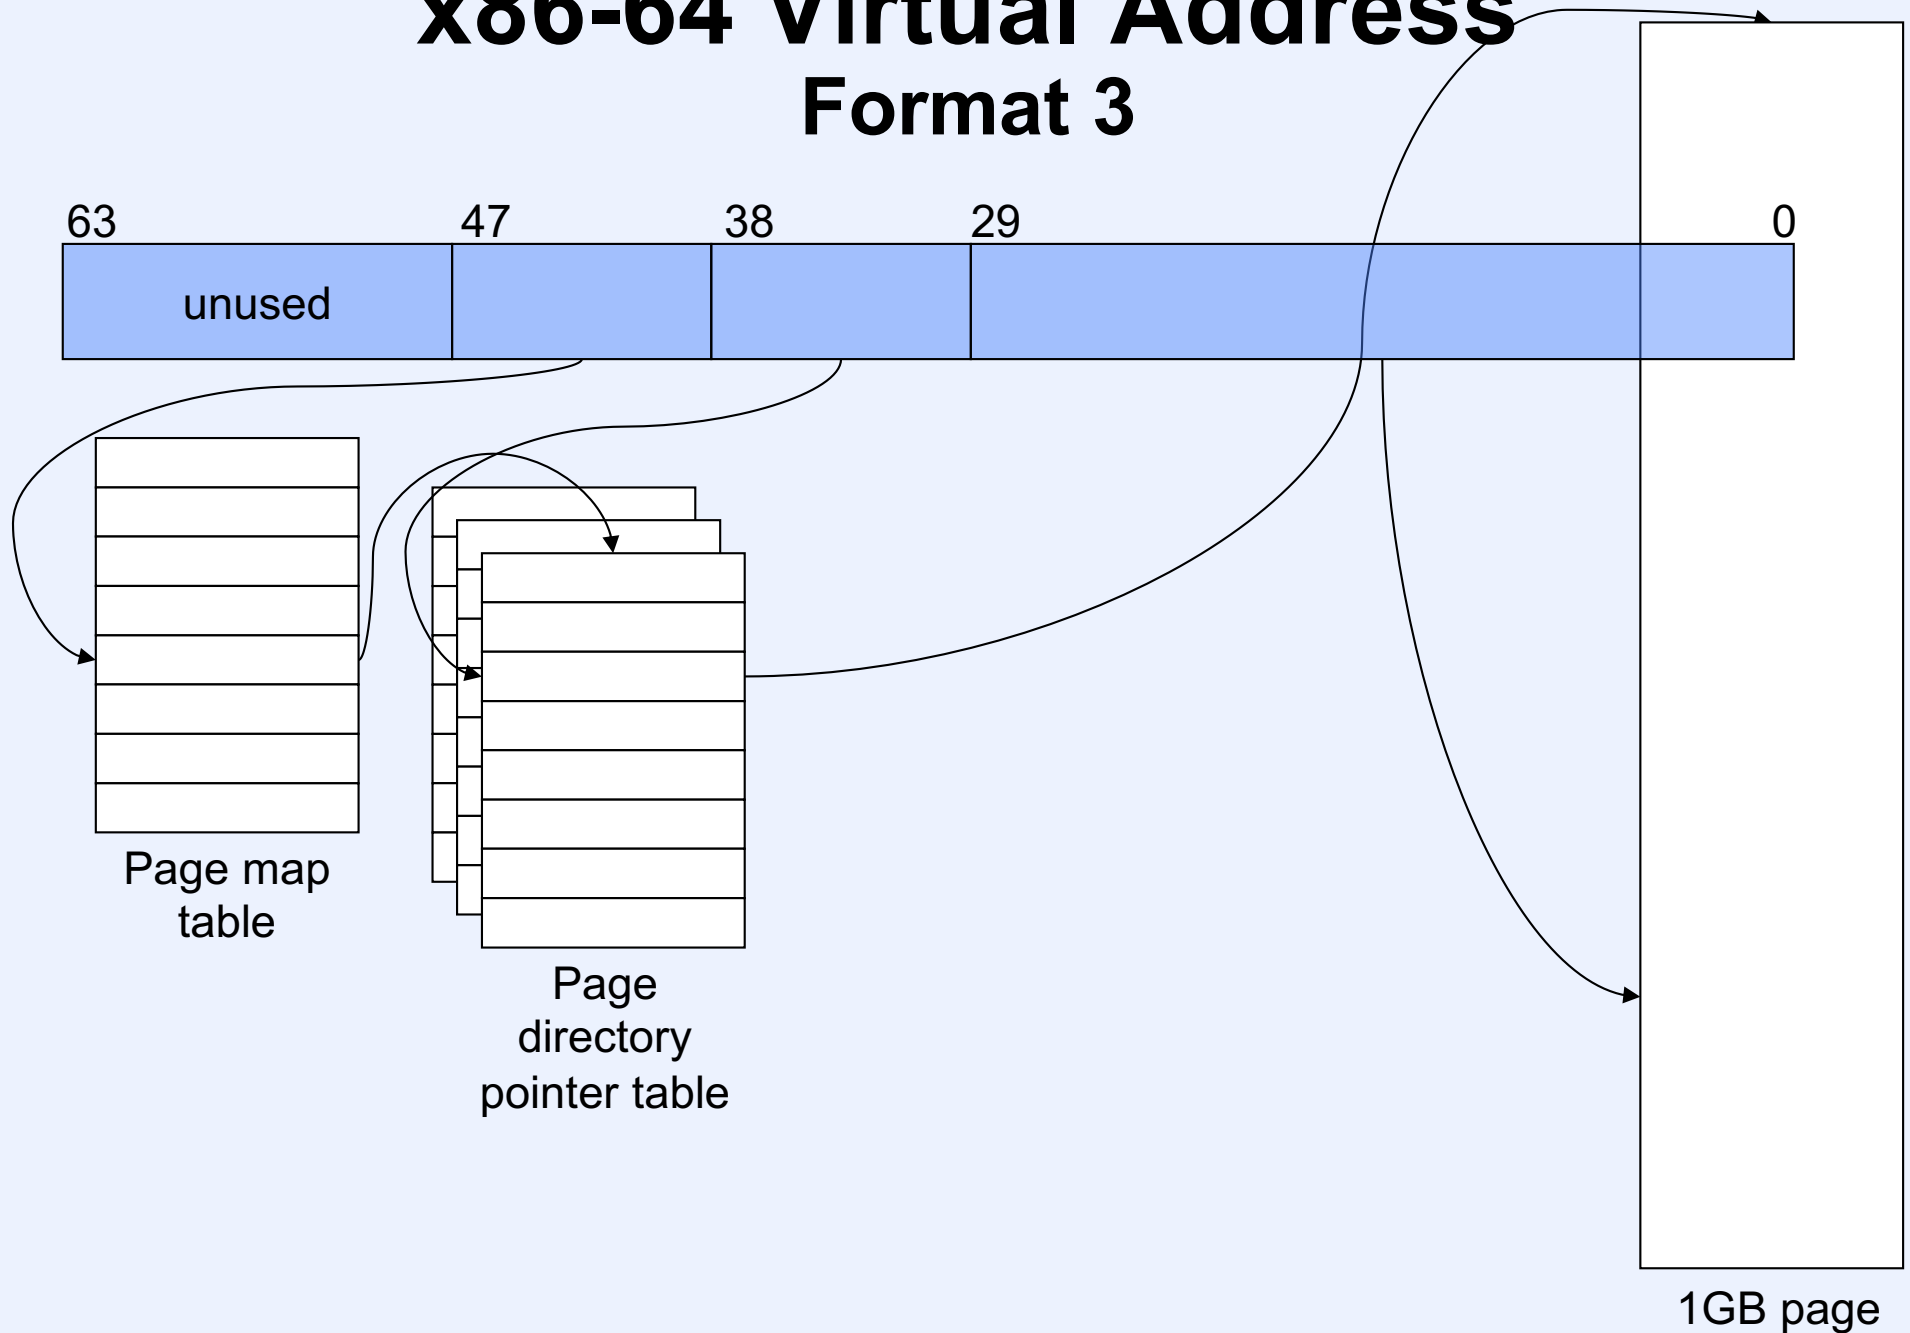
# x86-64 Virtual Address Format 1



# x86-64 Virtual Address Format 2



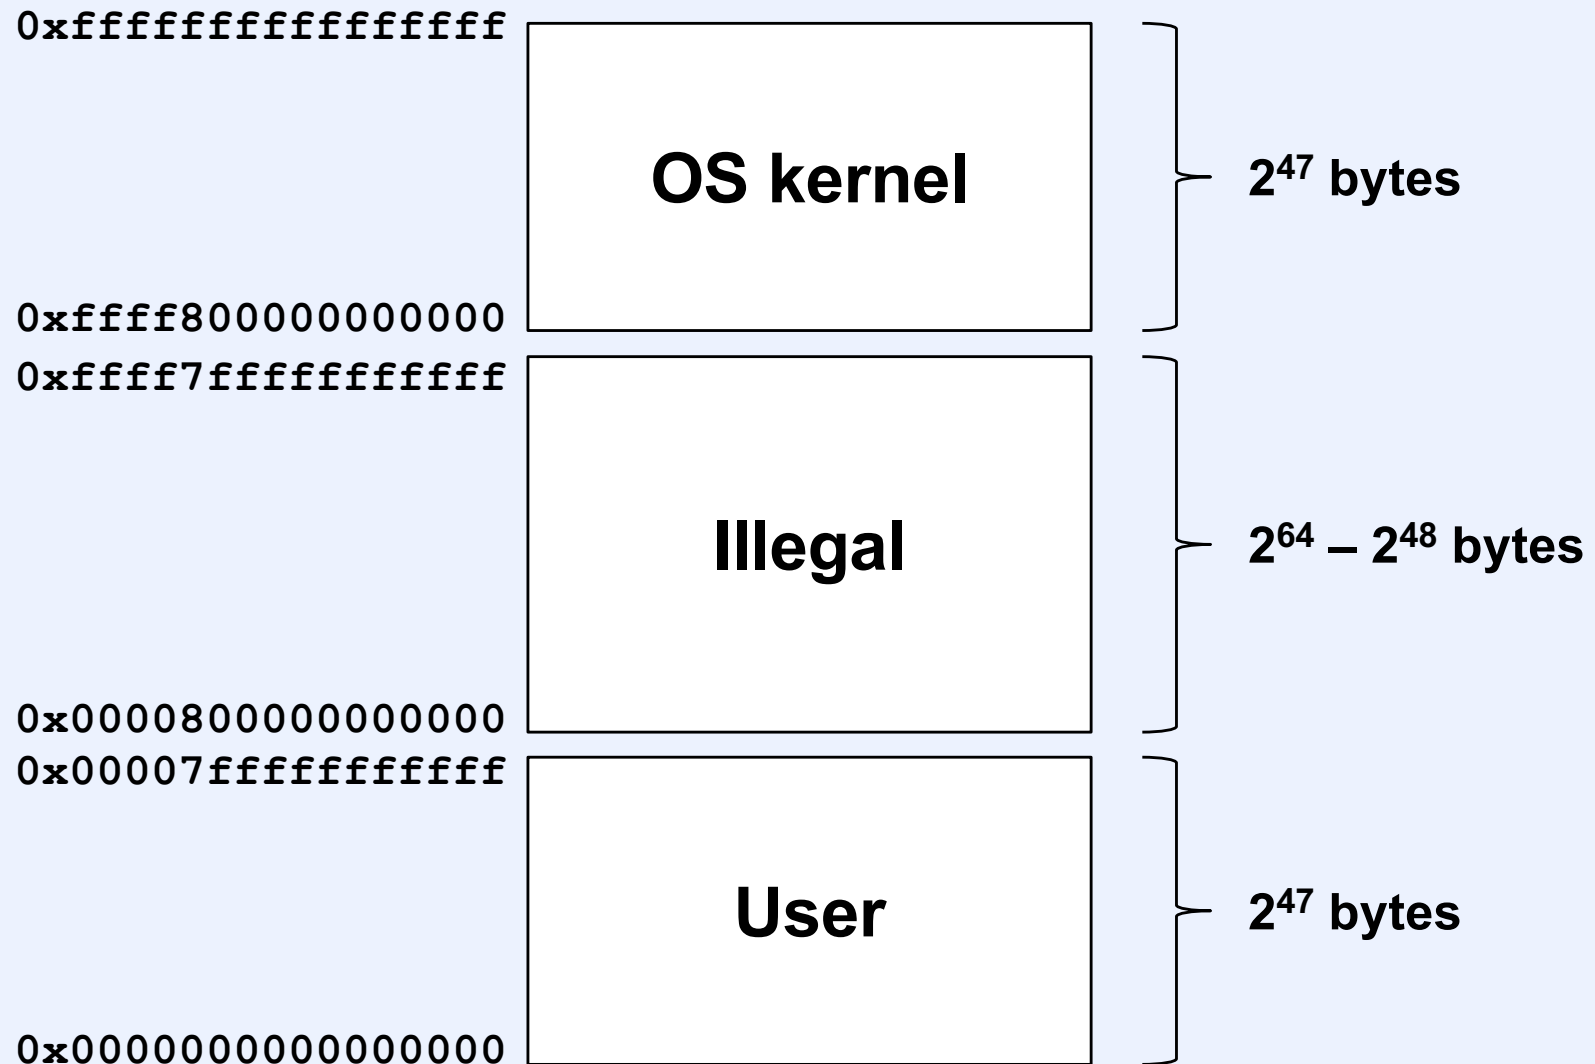
# x86-64 Virtual Address Format 3




# Why Multiple Page Sizes?

- **Internal fragmentation**
  - for region composed of 4KB pages, average internal fragmentation is 2KB
  - for region composed of 1GB pages, average internal fragmentation is 512MB
- **Page-table overhead**
  - larger page sizes have fewer page tables
    - less overhead in representing mappings
      - both in memory and in cache

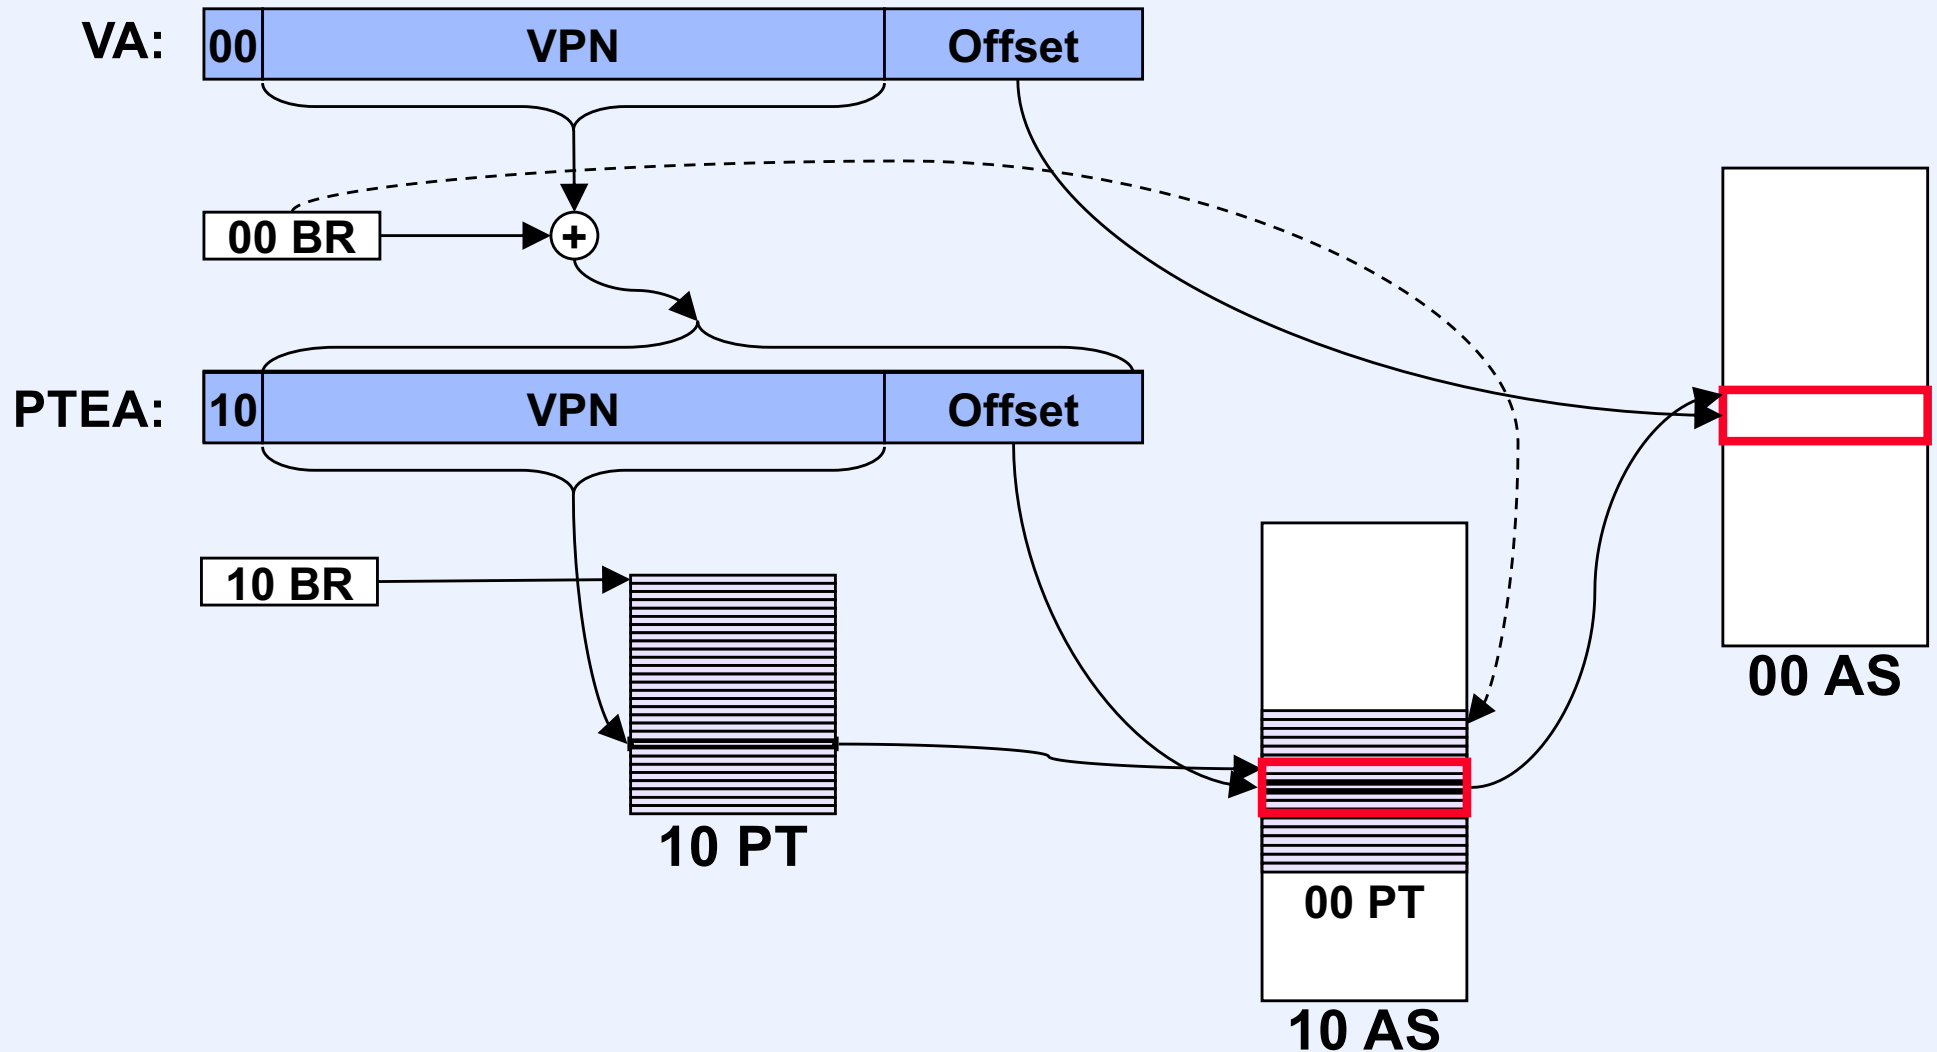
# Address Space



# Linear Page Table



# VAX Linear Page Translation



# \$

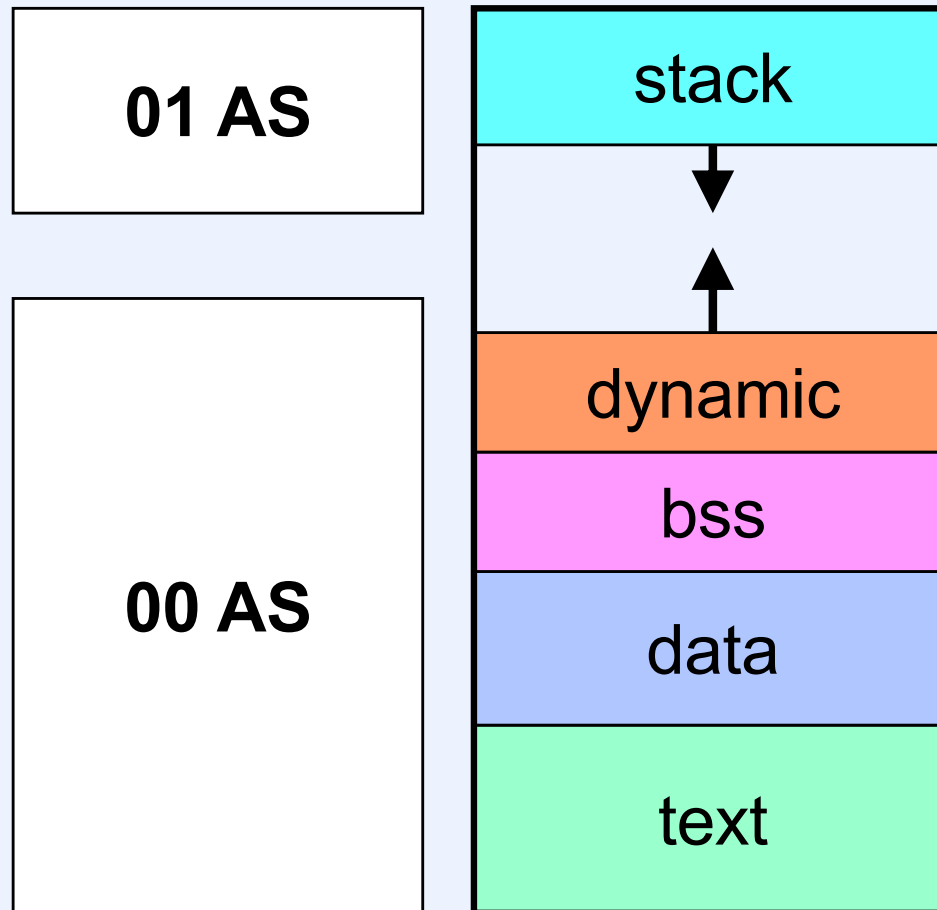
- **VAX architecture introduced in 1978**
  - **memory cost \$40,000/MB**
    - **3.8¢/byte**  
(.475¢/bit)



# Linear Page-Table Management

- **00 and 01 page tables each require contiguous locations in 10 space**
  - **with 512-byte pages, 8MB each:**
    - **maximum of 128 such page tables**
    - **(need room for other things, e.g. OS)**
- **Reduce size requirements with partial page tables**
  - **length registers constrain size of each space**

# Traditional Unix with Linear PTs



# Quiz 3

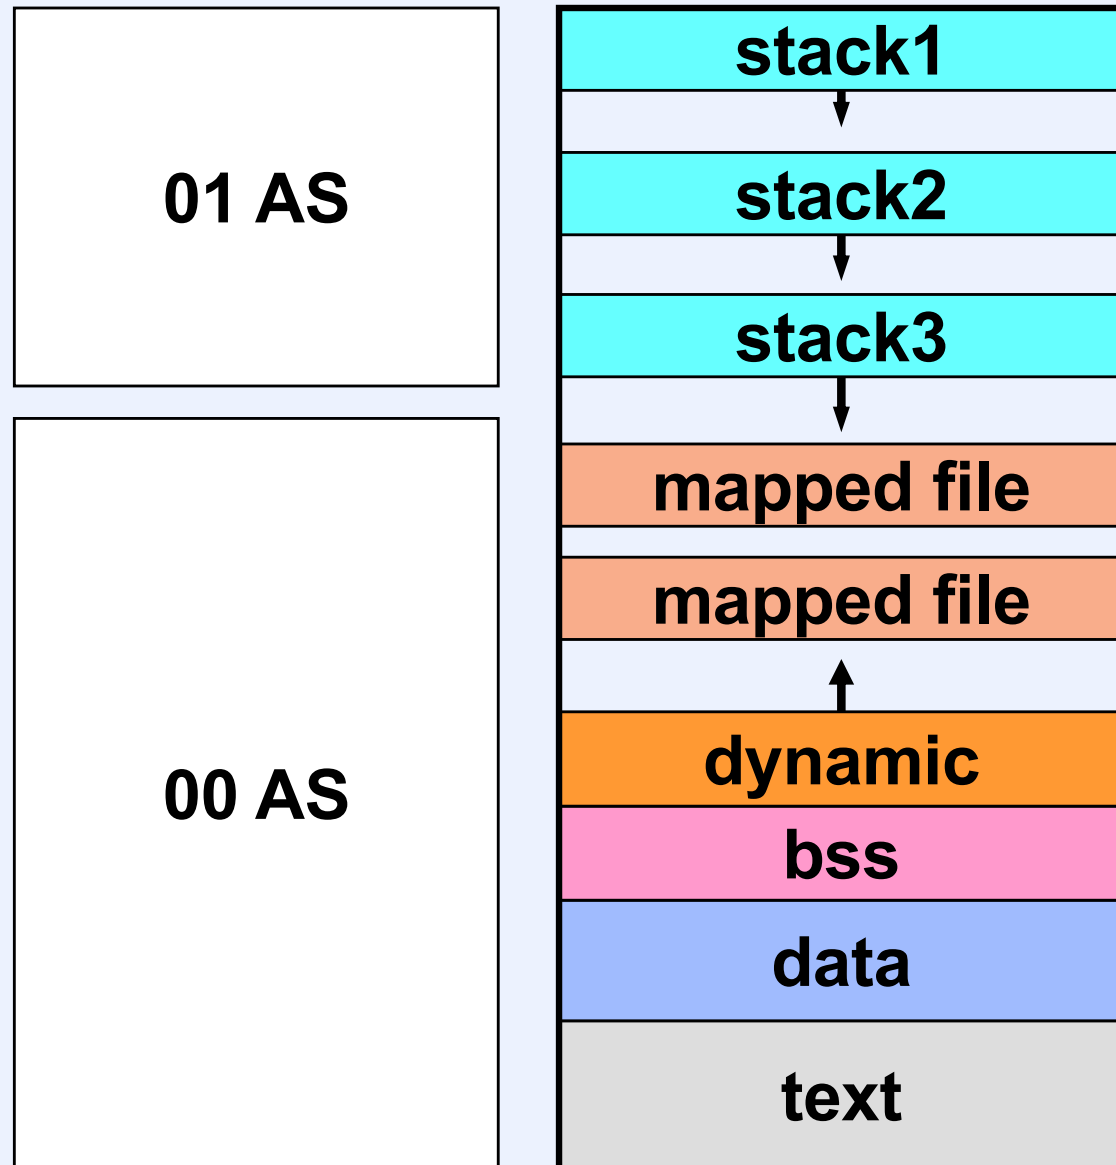
**Suppose the page size is 512 bytes ( $2^9$ ) and each page-table entry requires 4 bytes. How many pages of page-table entries are required to map 1 megabyte ( $2^{20}$ ) of address space?**

- a) 4**
- b) 8**
- c) 16**
- d) 32**

# \$

- **Limit size of 00 space to 1 MB**
  - requires 16-page 00 page table in 10 virtual memory
    - requires 16 entries in 10 page table
- **Same requirements if 01 space limited to 1 MB**
- **What are real-memory requirements?**
  - 10 page table resides in real memory
  - at least one page of real memory must be allocated for each of 00 and 01 page tables
  - minimum real memory is 1152 bytes
    - \$43.95 in 1978

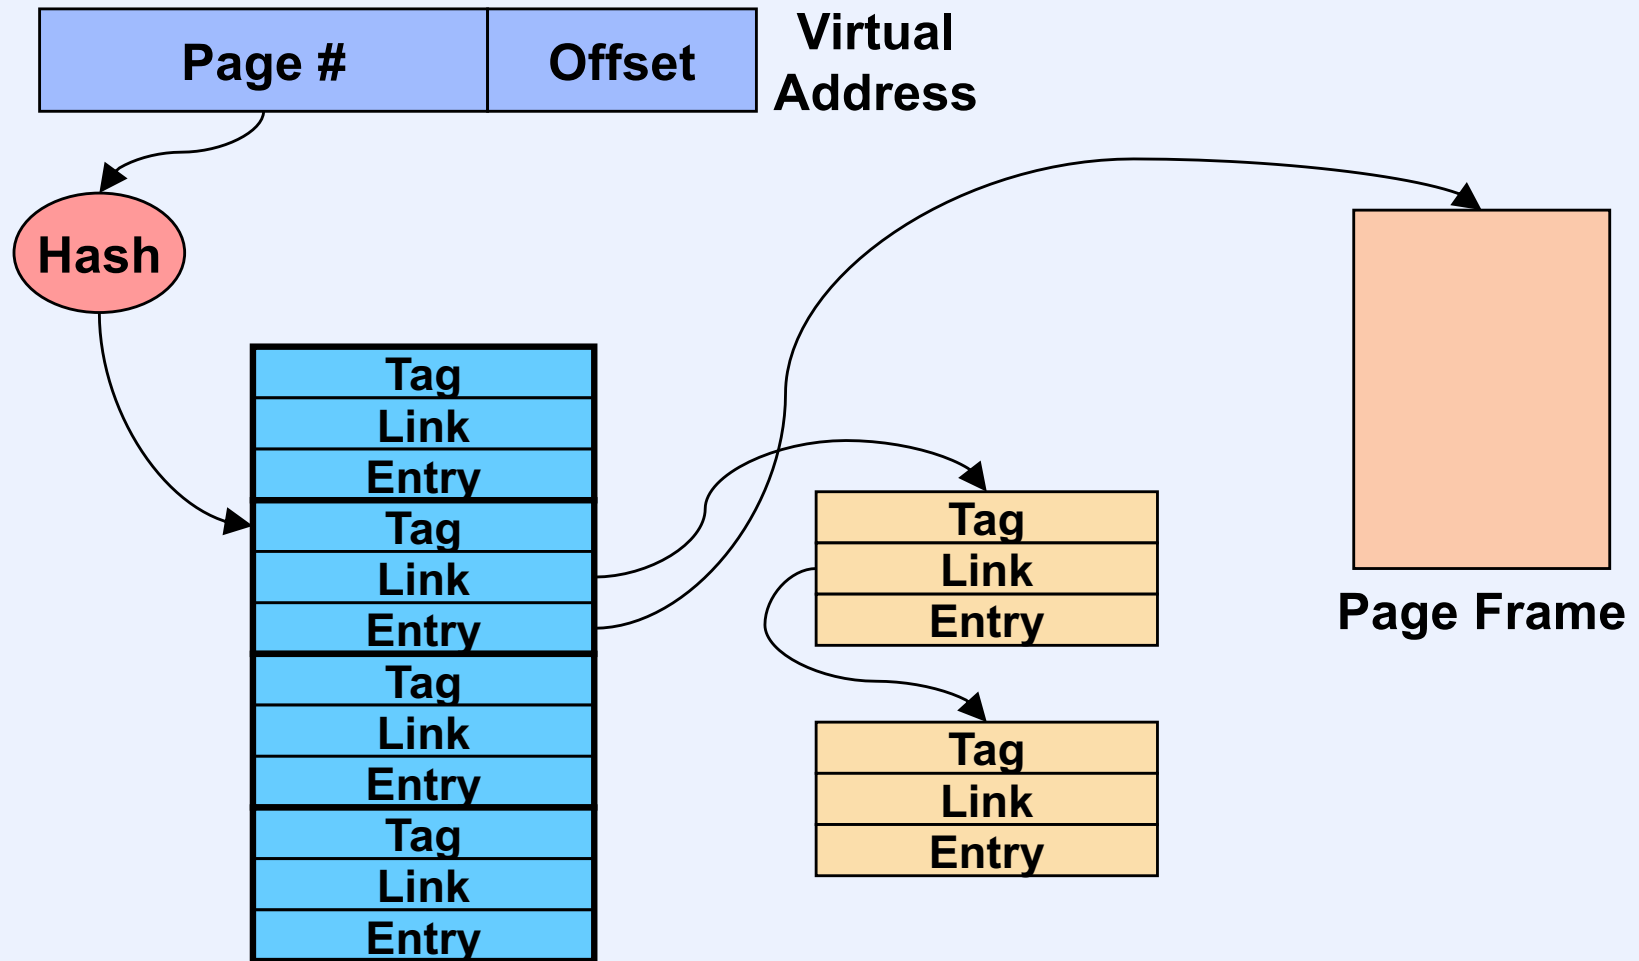
# Modern Unix



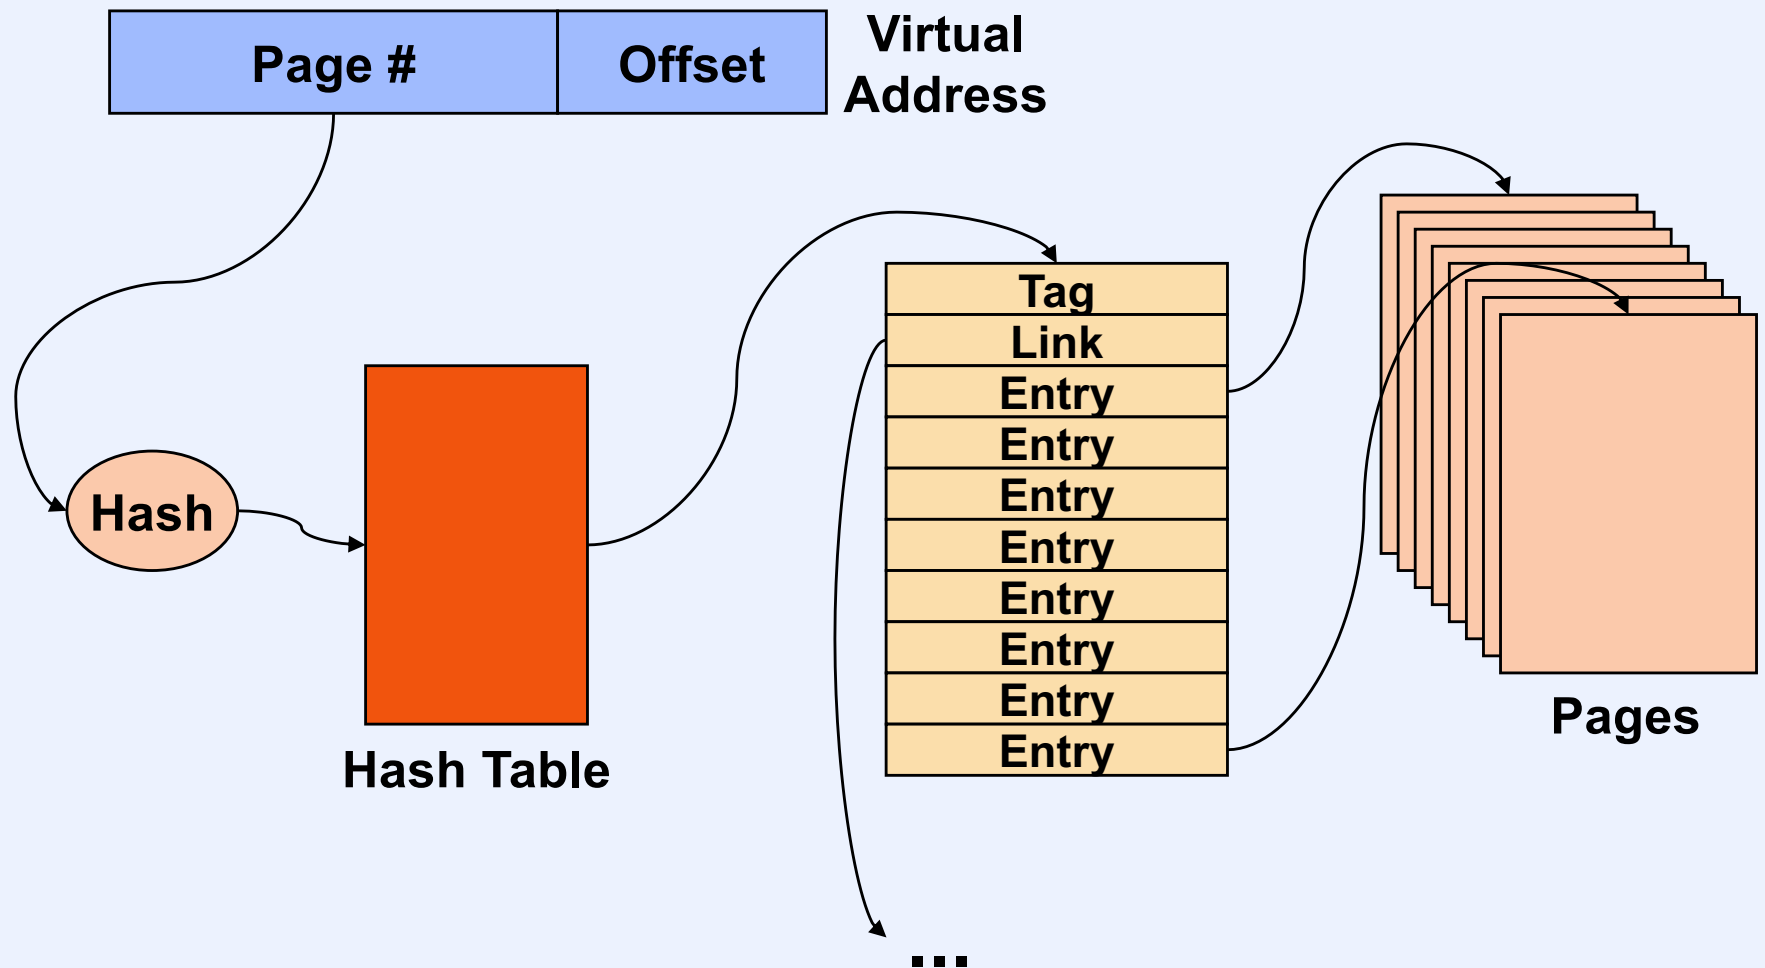
# \$

- **Requires sufficient 10 page-table entries to map almost all of 00 and 01 space**
  - **$2^{14}$  10 page-table entries for each space**
    - **requiring 64KB each, 128KB total**
    - **\$5000 in 1978**
  - **<1¢/process today**
    - **who cares?**
    - **increase address space from  $2^{32}$  to  $2^{64}$** 
      - **4,294,967,296-fold increase**
      - **significant ...**

# Hashed Page Tables

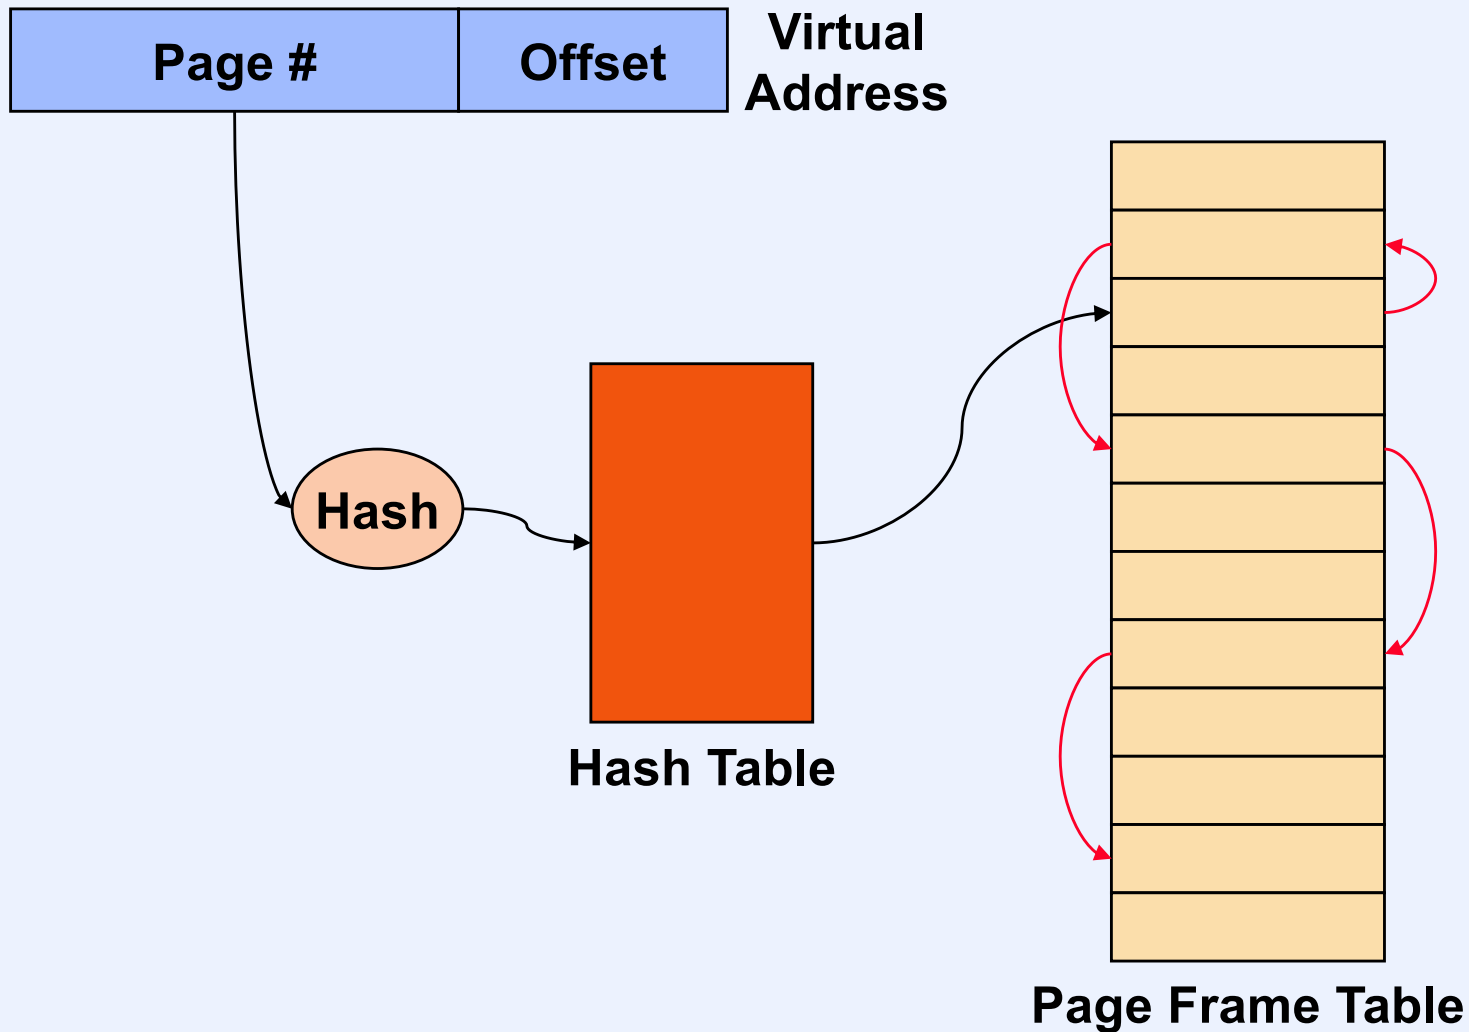


# Clustered Page Tables





# Inverted Page Tables

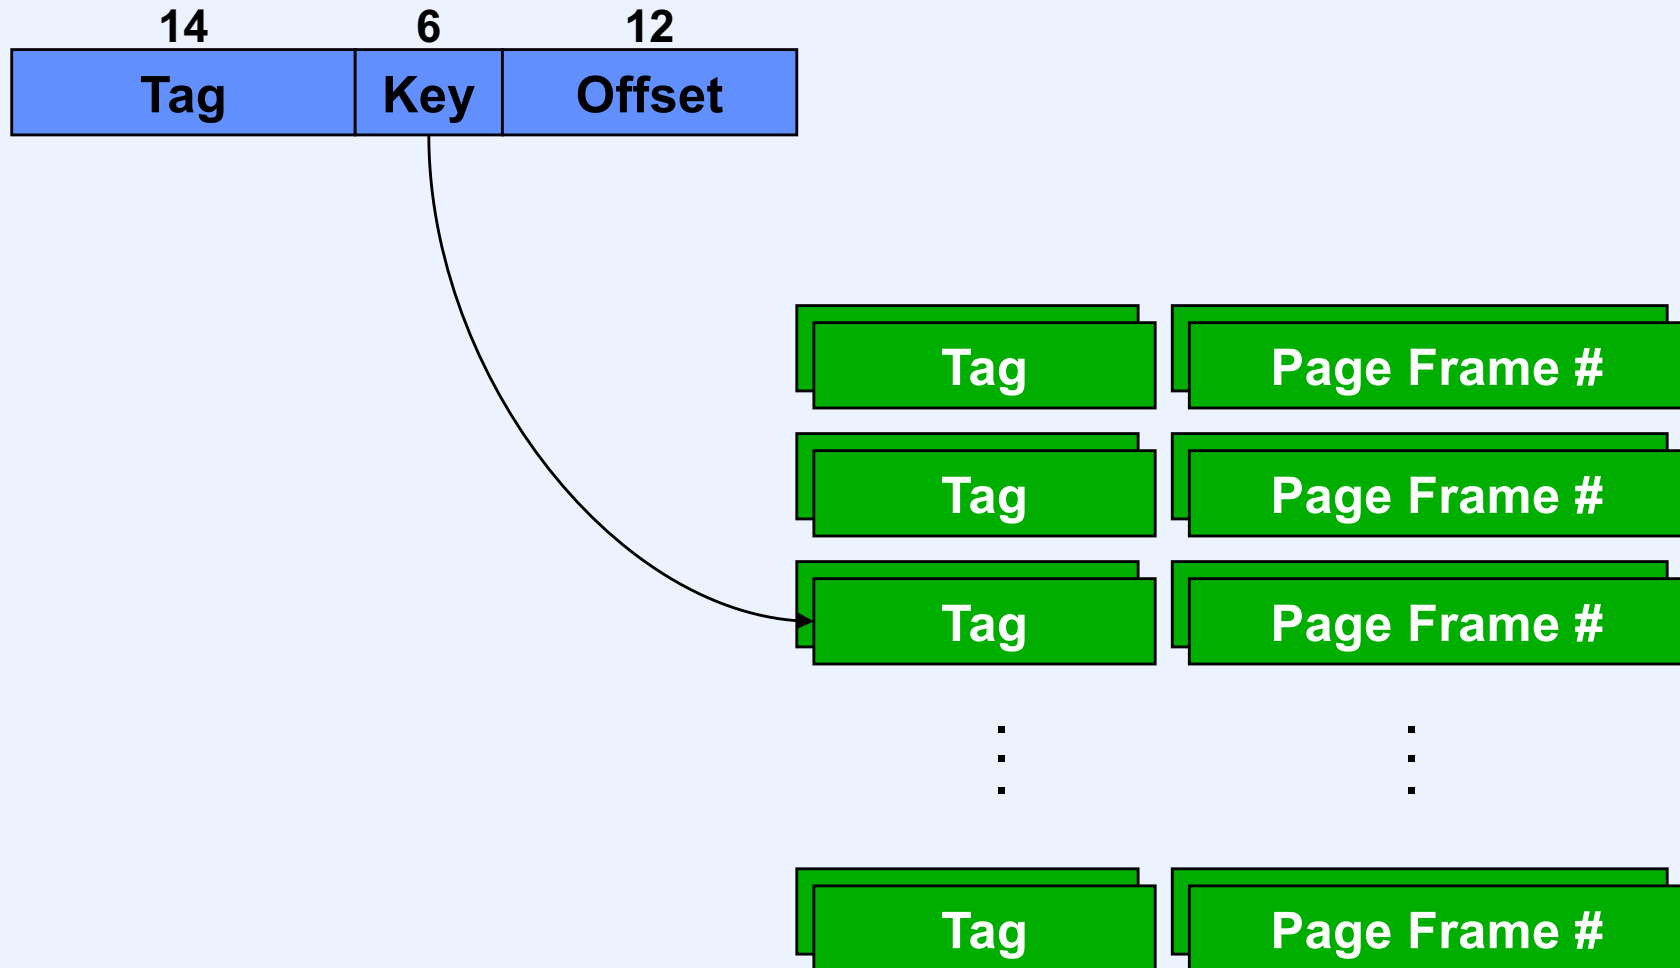


# Quiz 4

**Normal page tables map virtual memory to real memory. More precisely, they map an address space and a location within that address space to real memory. Inverted page tables do the inverse mapping: given an address space ID and a location in real memory, they produce the corresponding virtual location.**

- a) Inverted page tables work with all Unix systems**
- b) They don't work with any Unix system**
- c) They don't work with Unix systems that support *mmap* with shared mappings**


# Translation Lookaside Buffers



# TLBs and Mappings

- **TLBs provide a cache for mappings from virtual addresses to real addresses**
- **Mappings change when**
  - pages are removed (unmapped) from memory
  - when the address space is switched from one process's to another's
- **OS must explicitly flush old contents of TLB**
  - either individual entries or all of it

# TLBs and Multiprocessors



# TLB Shutdown Algorithm

```
// shooter code
for all processors i sharing
    address space
    interrupt(i);
for all processors i sharing
    address space
    while (noted[i] == 0)
        ;
modify_page_table();
update_or_flush_tlb();
done[me] = 1;
```

```
// shootee i interrupt handler
receive_interrupt_from_
    processor j
noted[i] = 1
while (done[j] == 0)
    ;
flush_tlb()
```