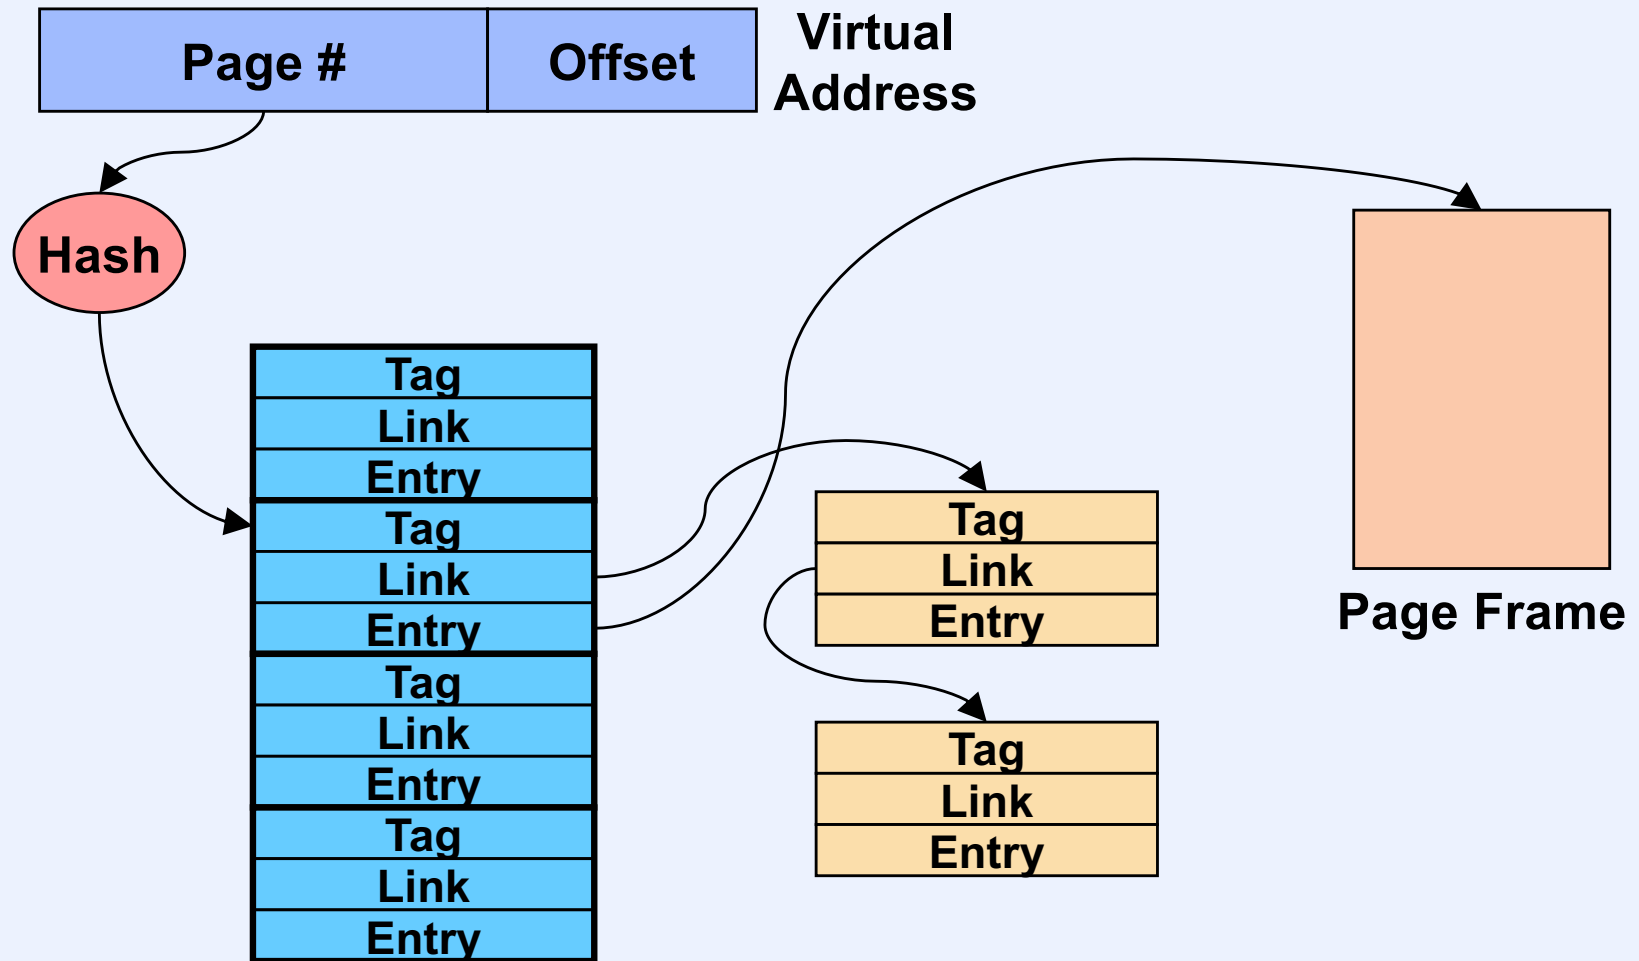
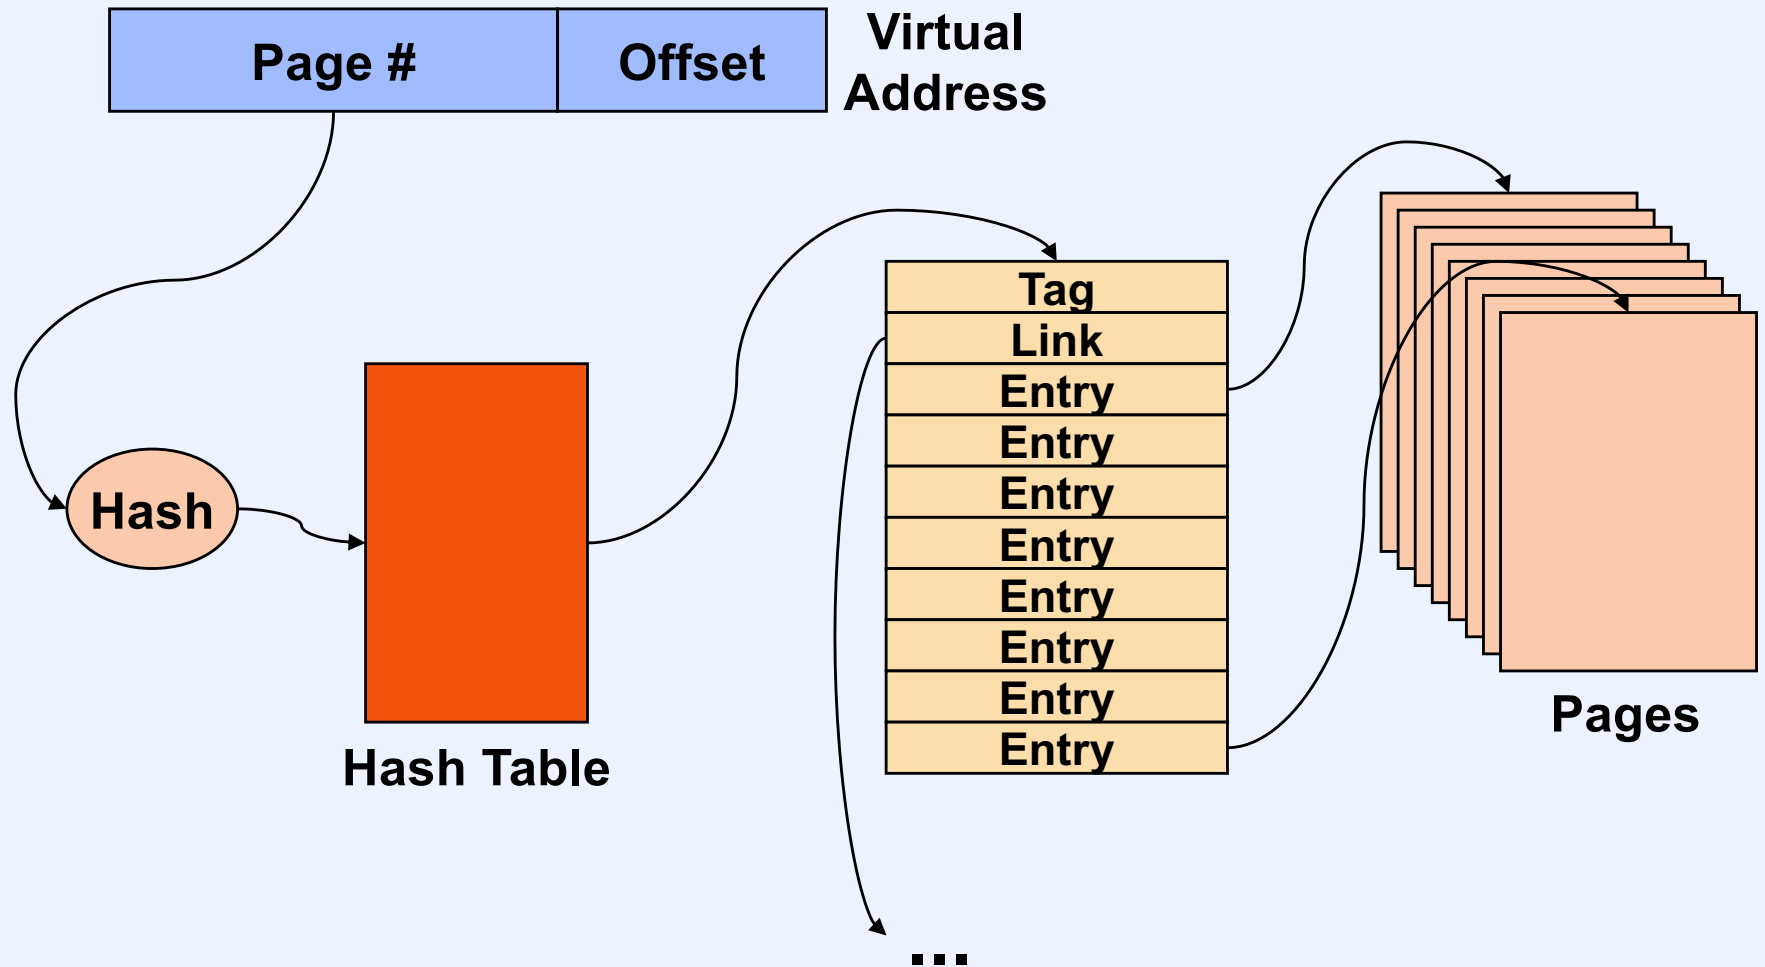


Memory Management Part 3

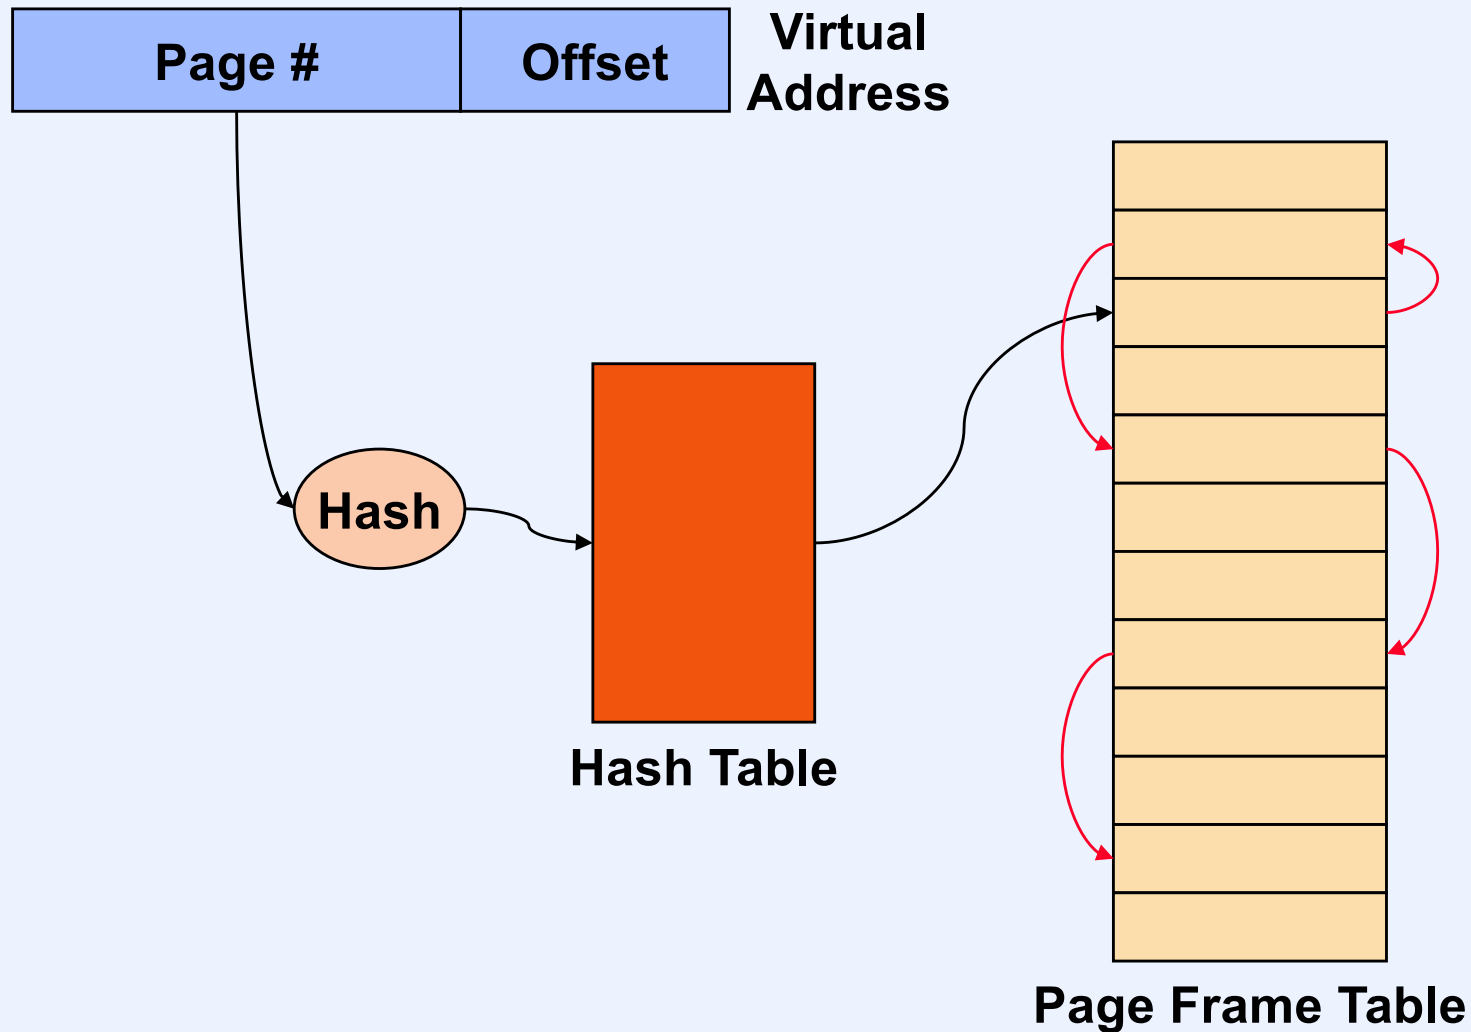
Hashed Page Tables



Clustered Page Tables



Inverted Page Tables

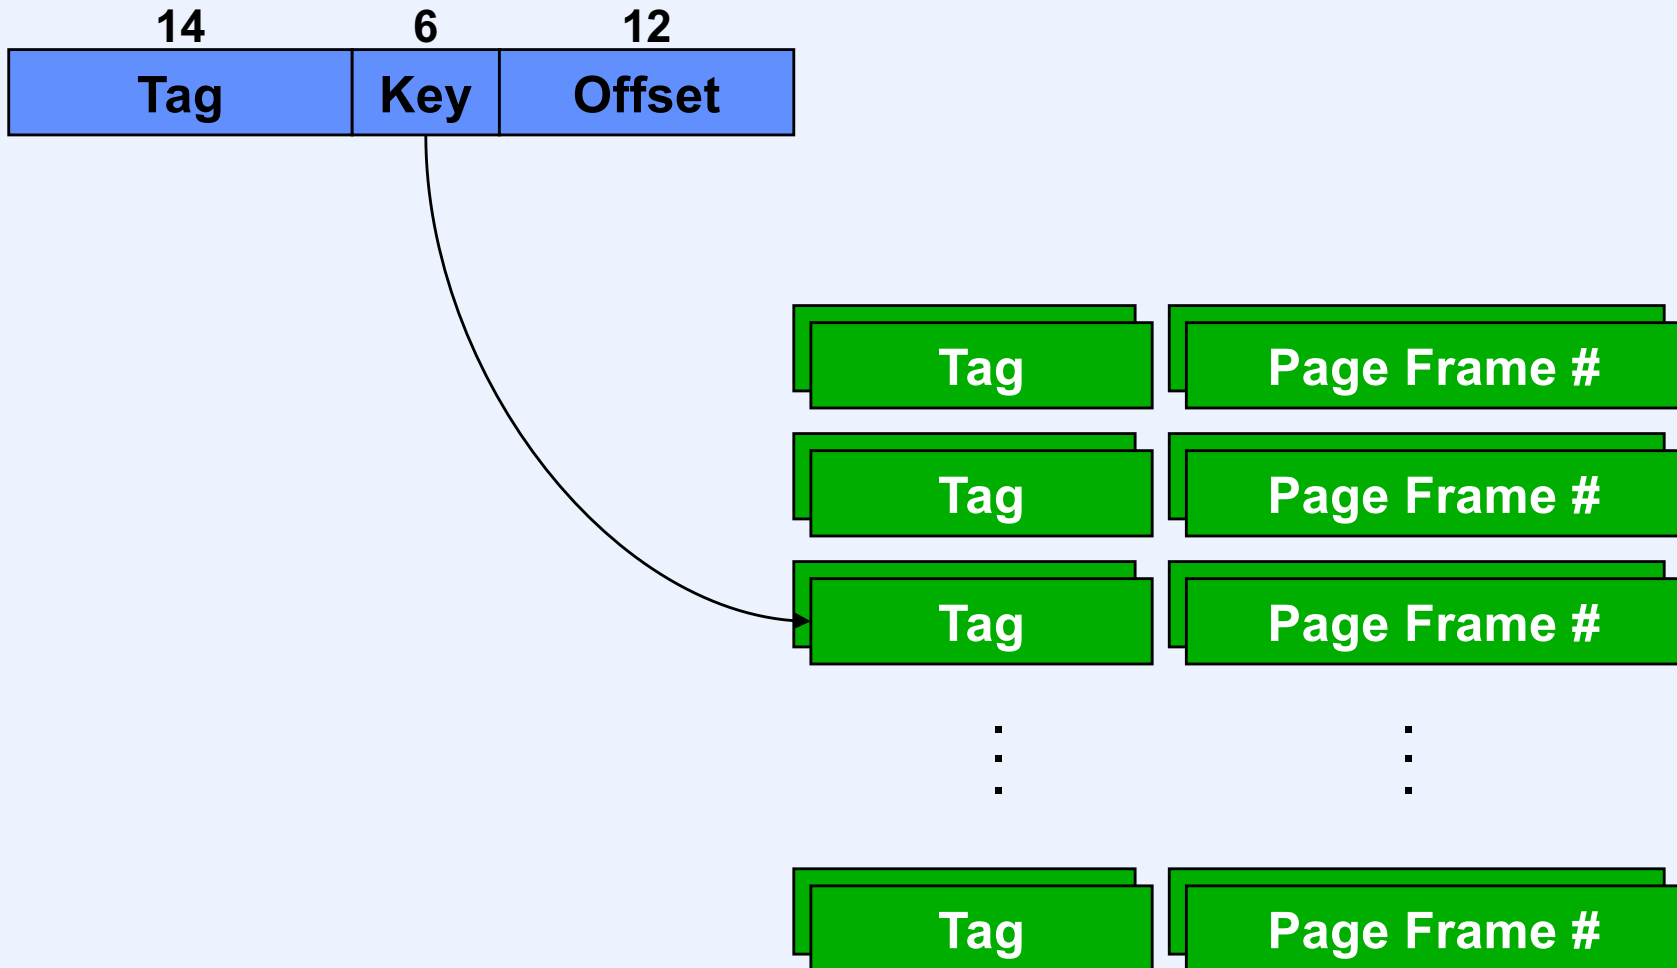


Quiz 1

Normal page tables map virtual memory to real memory. More precisely, they map an address space and a location within that address space to real memory. Inverted page tables do the inverse mapping: given an address space ID and a location in real memory, they produce the corresponding virtual location.

- a) Inverted page tables work with all Unix systems**
- b) They don't work with any Unix system**
- c) They don't work with Unix systems that support *mmap* with shared mappings**

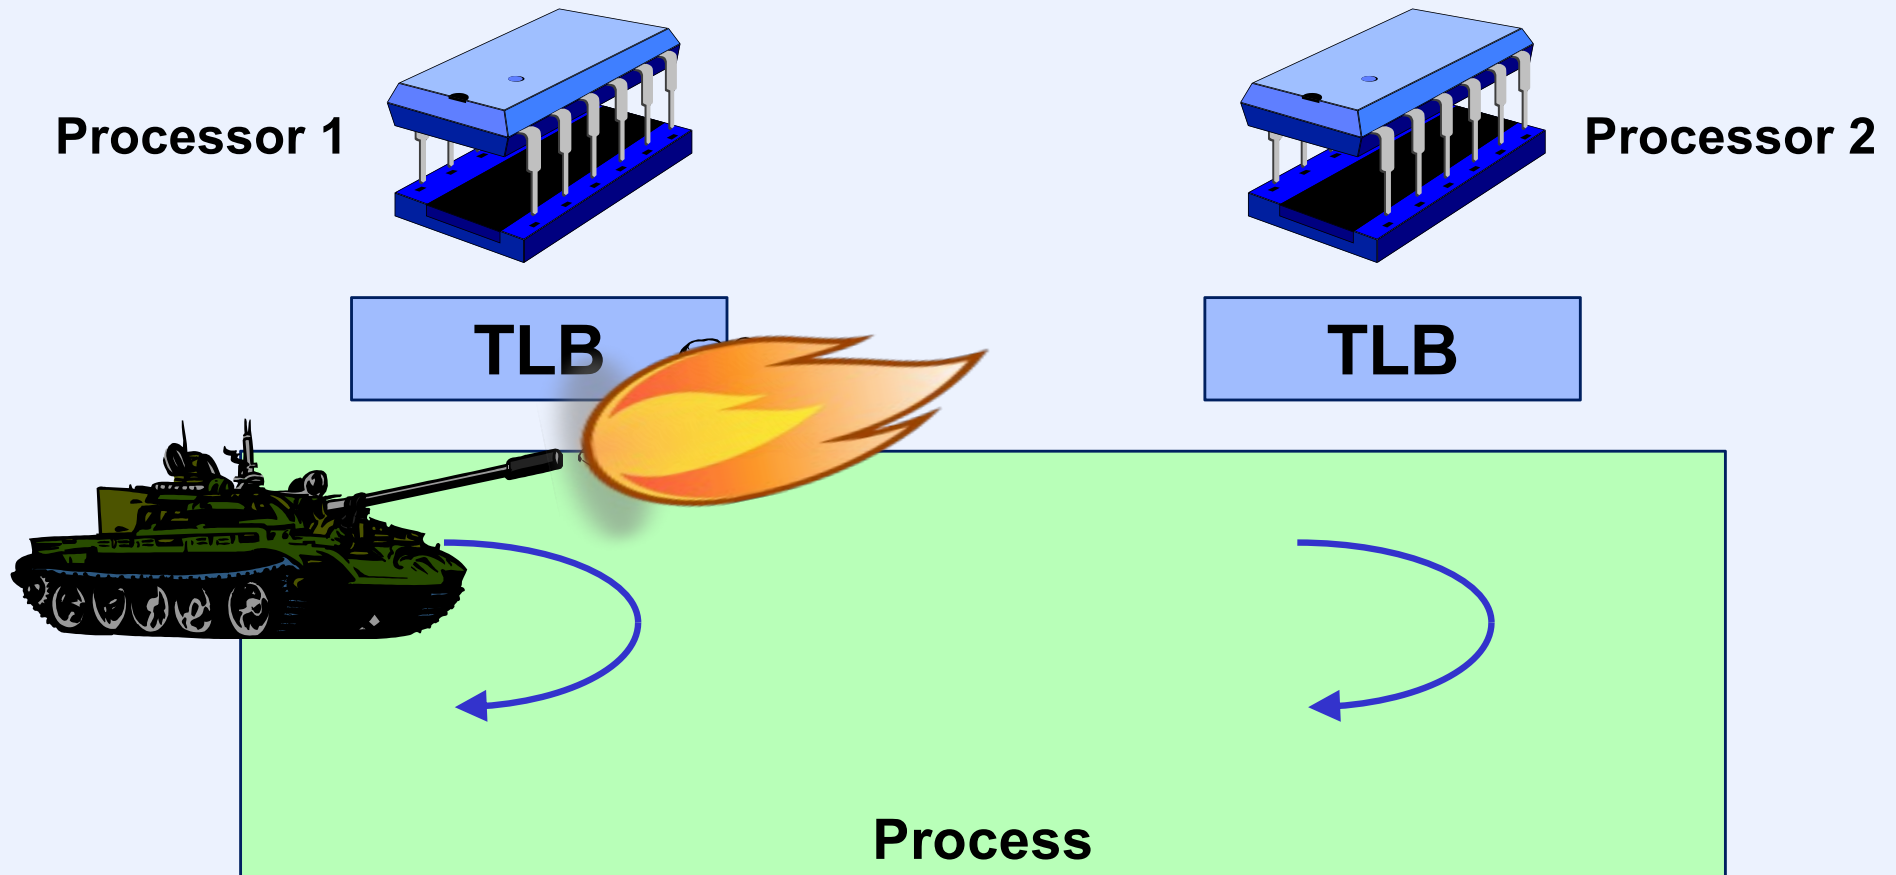
Translation Lookaside Buffers



TLBs and Mappings

- **TLBs provide a cache for mappings from virtual addresses to real addresses**
- **Mappings change when**
 - pages are removed (unmapped) from memory
 - when the address space is switched from one process's to another's
- **OS must explicitly flush old contents of TLB**
 - either individual entries or all of it

TLBs and Multiprocessors

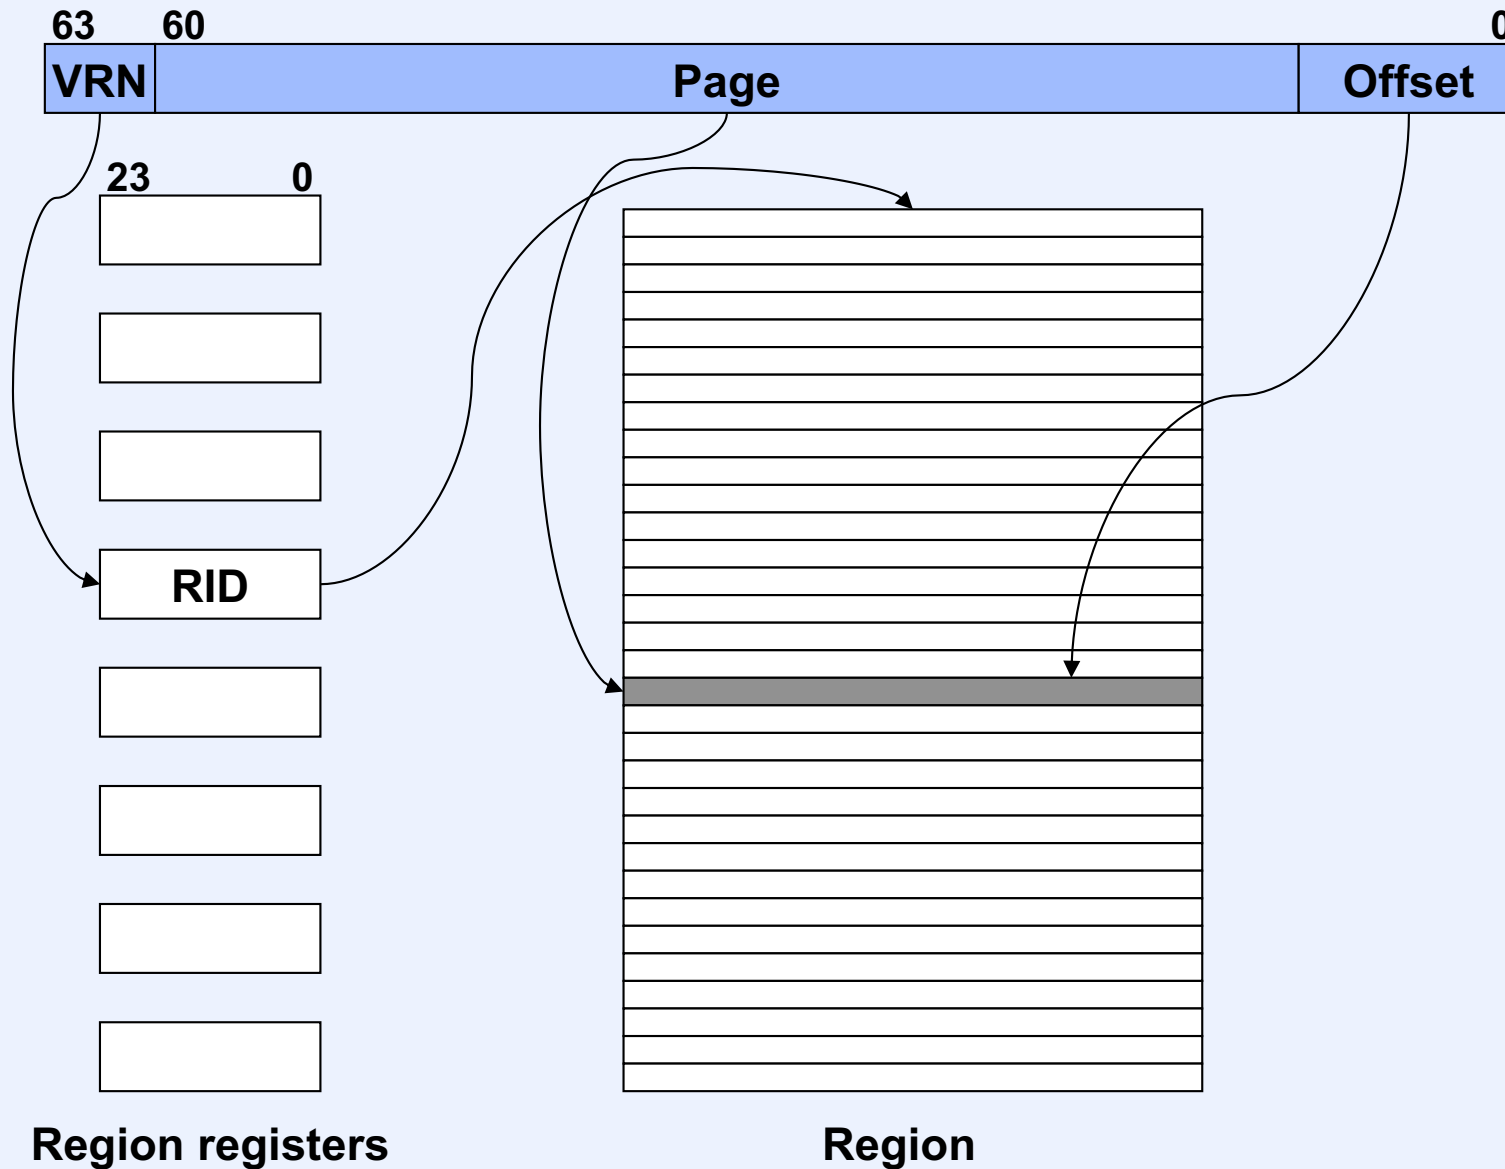


TLB Shutdown Algorithm

```
// shooter code
for all processors i sharing
    address space
    interrupt(i);
for all processors i sharing
    address space
    while (noted[i] == 0)
        ;
modify_page_table();
update_or_flush_tlb();
done[me] = 1;
```

```
// shootee i interrupt handler
receive_interrupt_from_
    processor j
noted[i] = 1
while (done[j] == 0)
    ;
flush_tlb()
```

Intel IA-64



Quiz 2

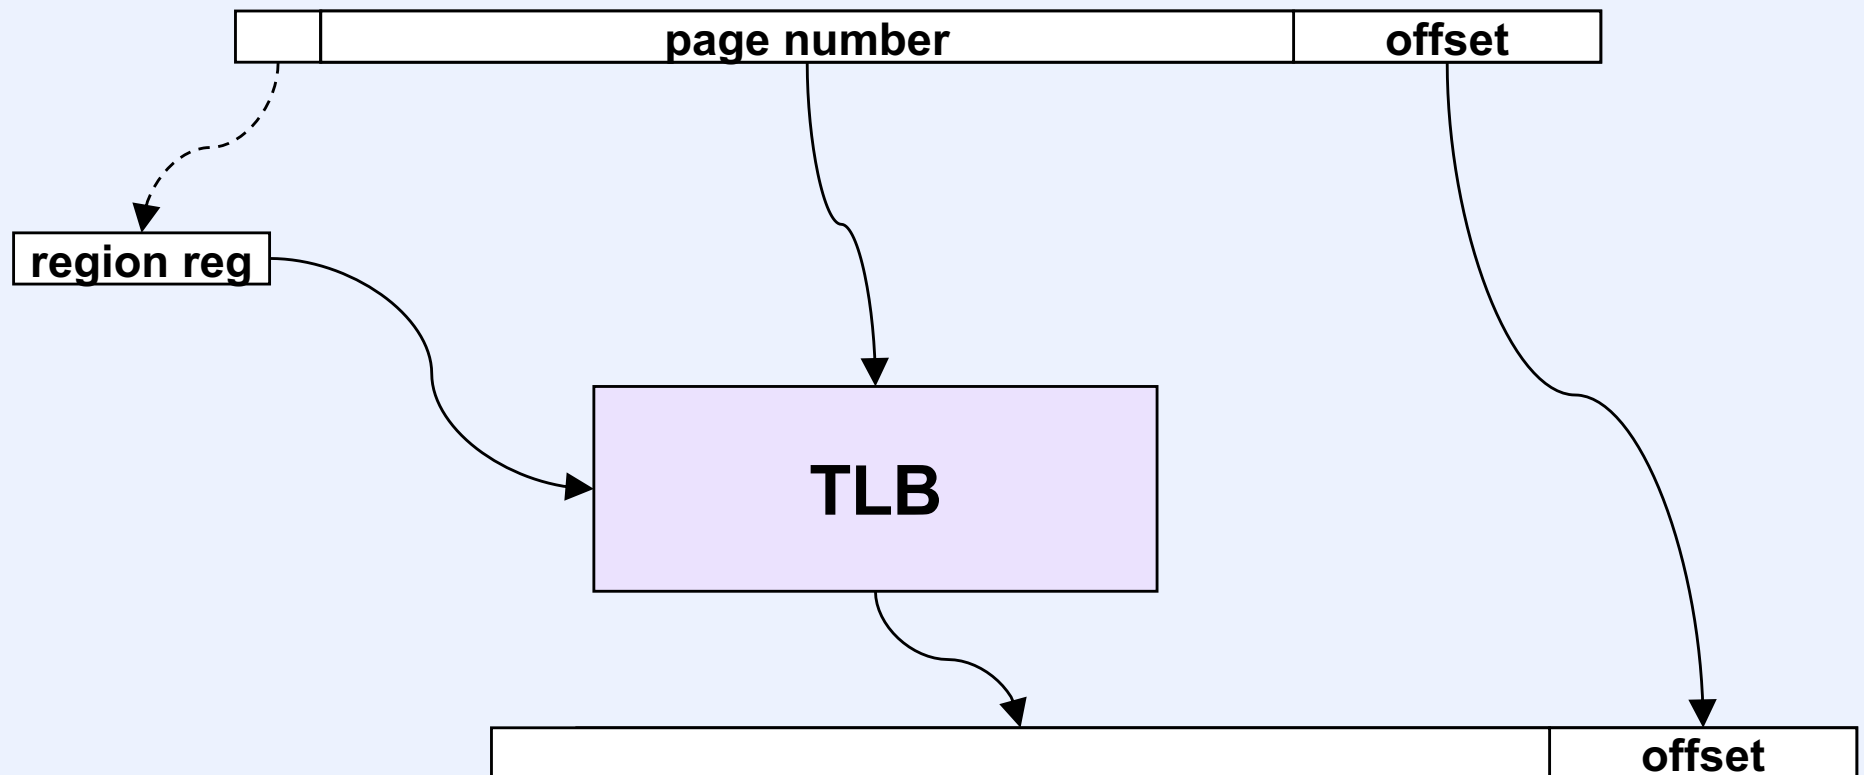
What did the IA-64 designers have in mind as a use for the region registers?

- a) Region registers facilitate increasing the sizes of regions**
- b) They felt that 2^{64} bytes might be too small a limit for address-space size; the region registers provide a means for making it larger**
- c) Region registers are a means for identifying and mapping in shared objects, such as libraries or datasets**
- d) Region-register contents are used in the TLB to uniquely identify regions; TLB flushes are minimized on process switches because it's clear which process an address belongs to**

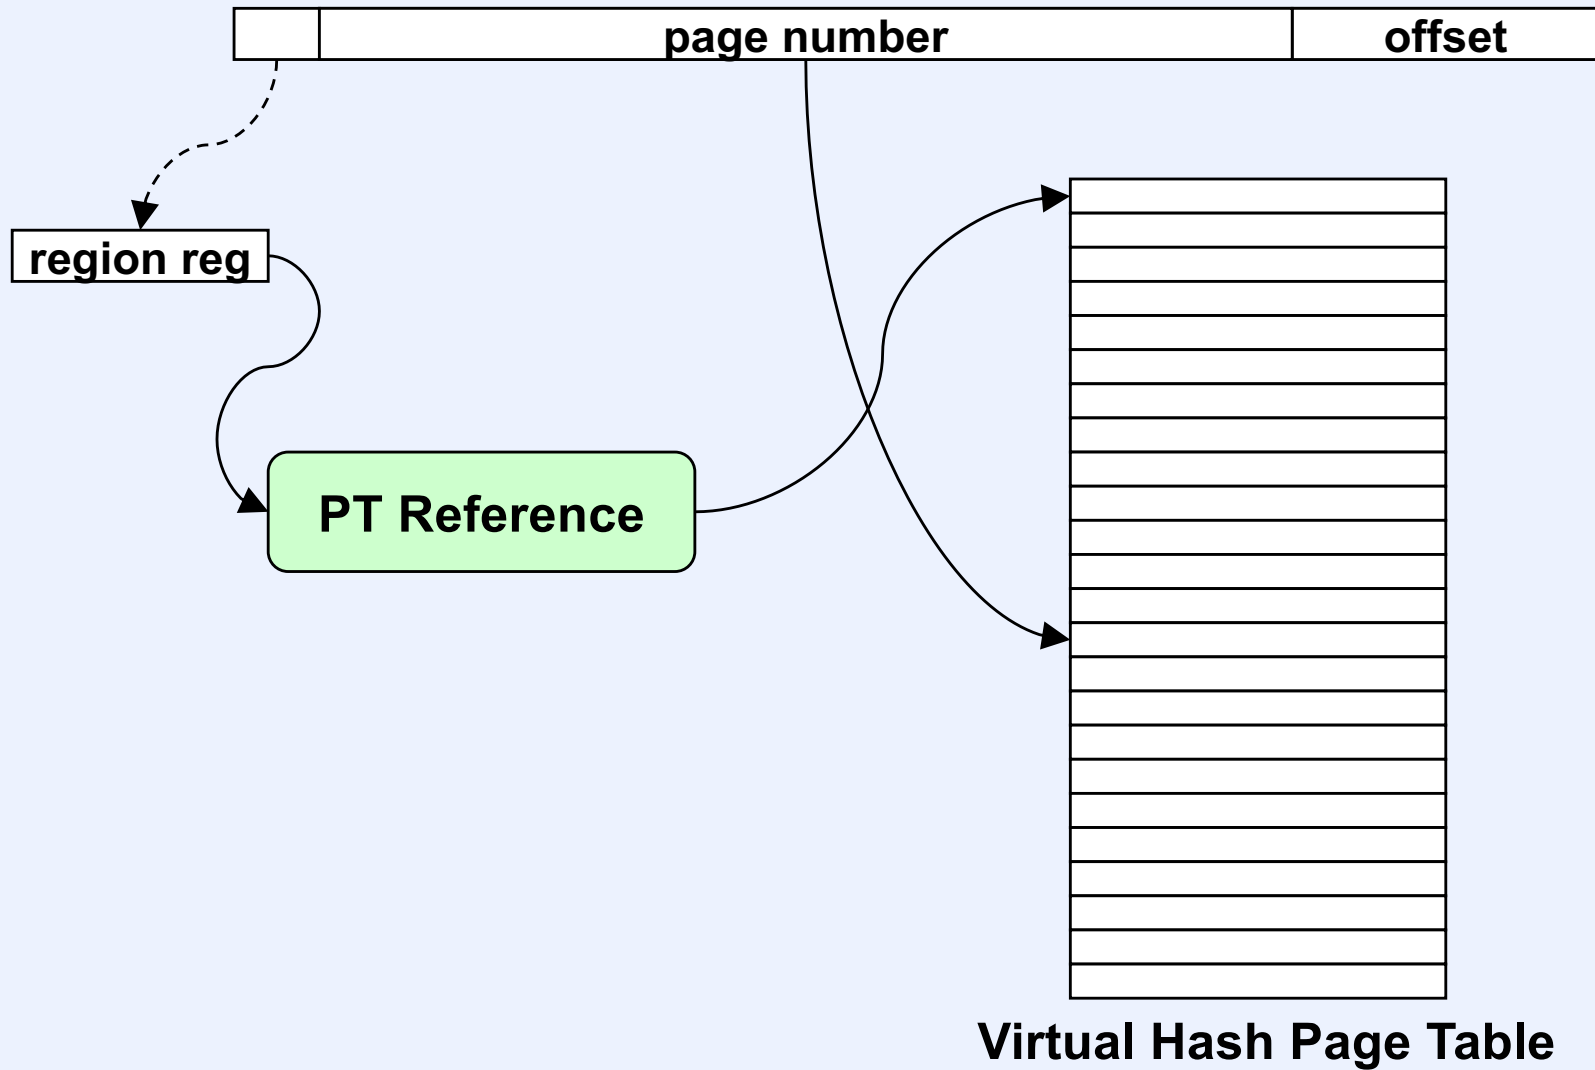
IA-64 Address Translation

- **TLB**
 - software-managed
- **Virtual Hash Page Table (VHPT)**
 - per-region linear page table
 - single large hashed page table

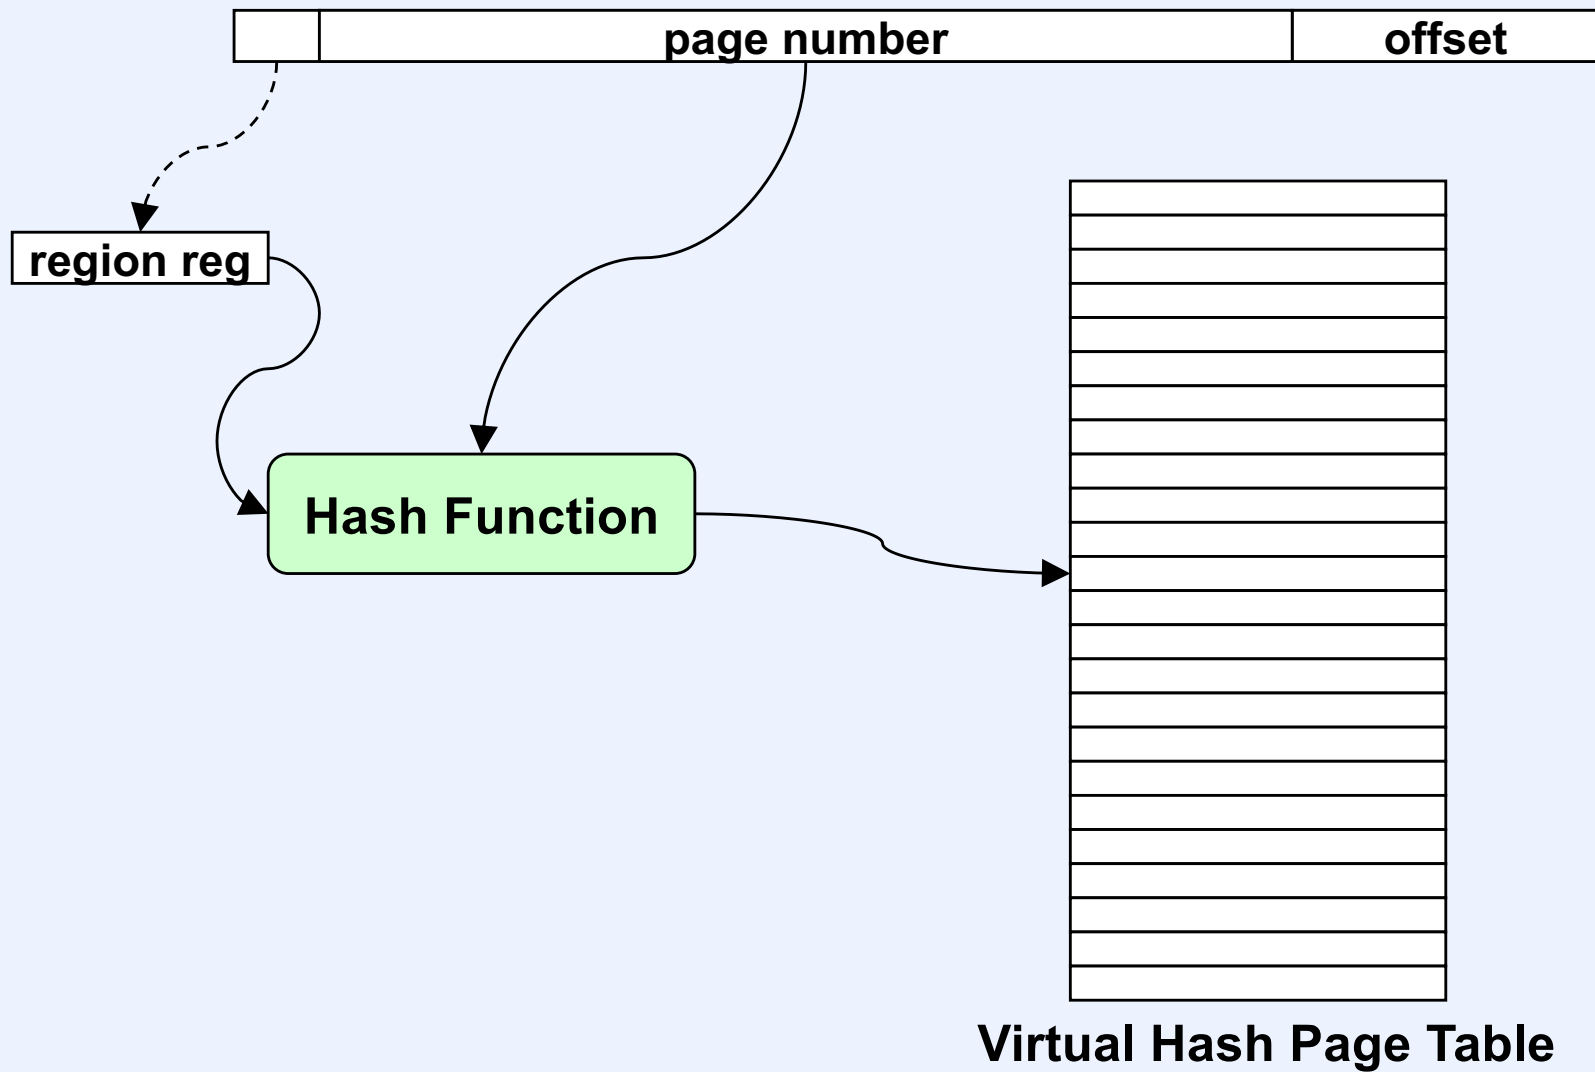
Translation: TLB



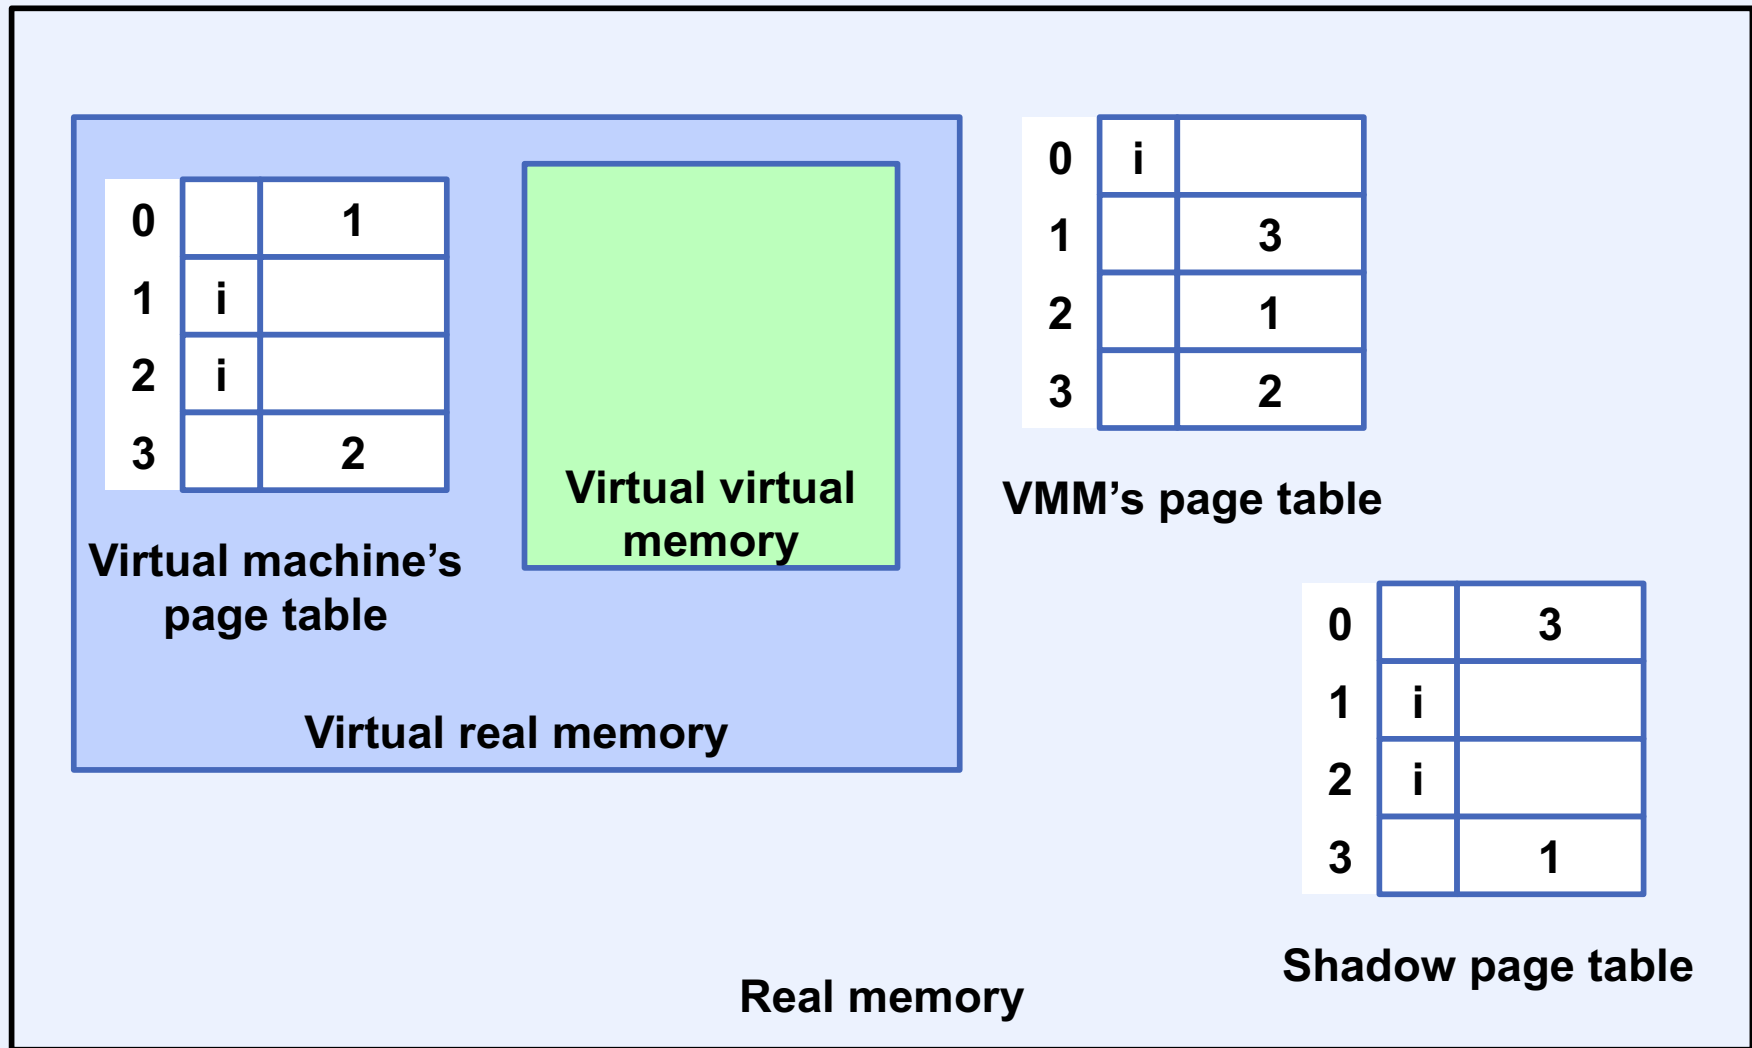
Translation: TLB Miss (LPT)



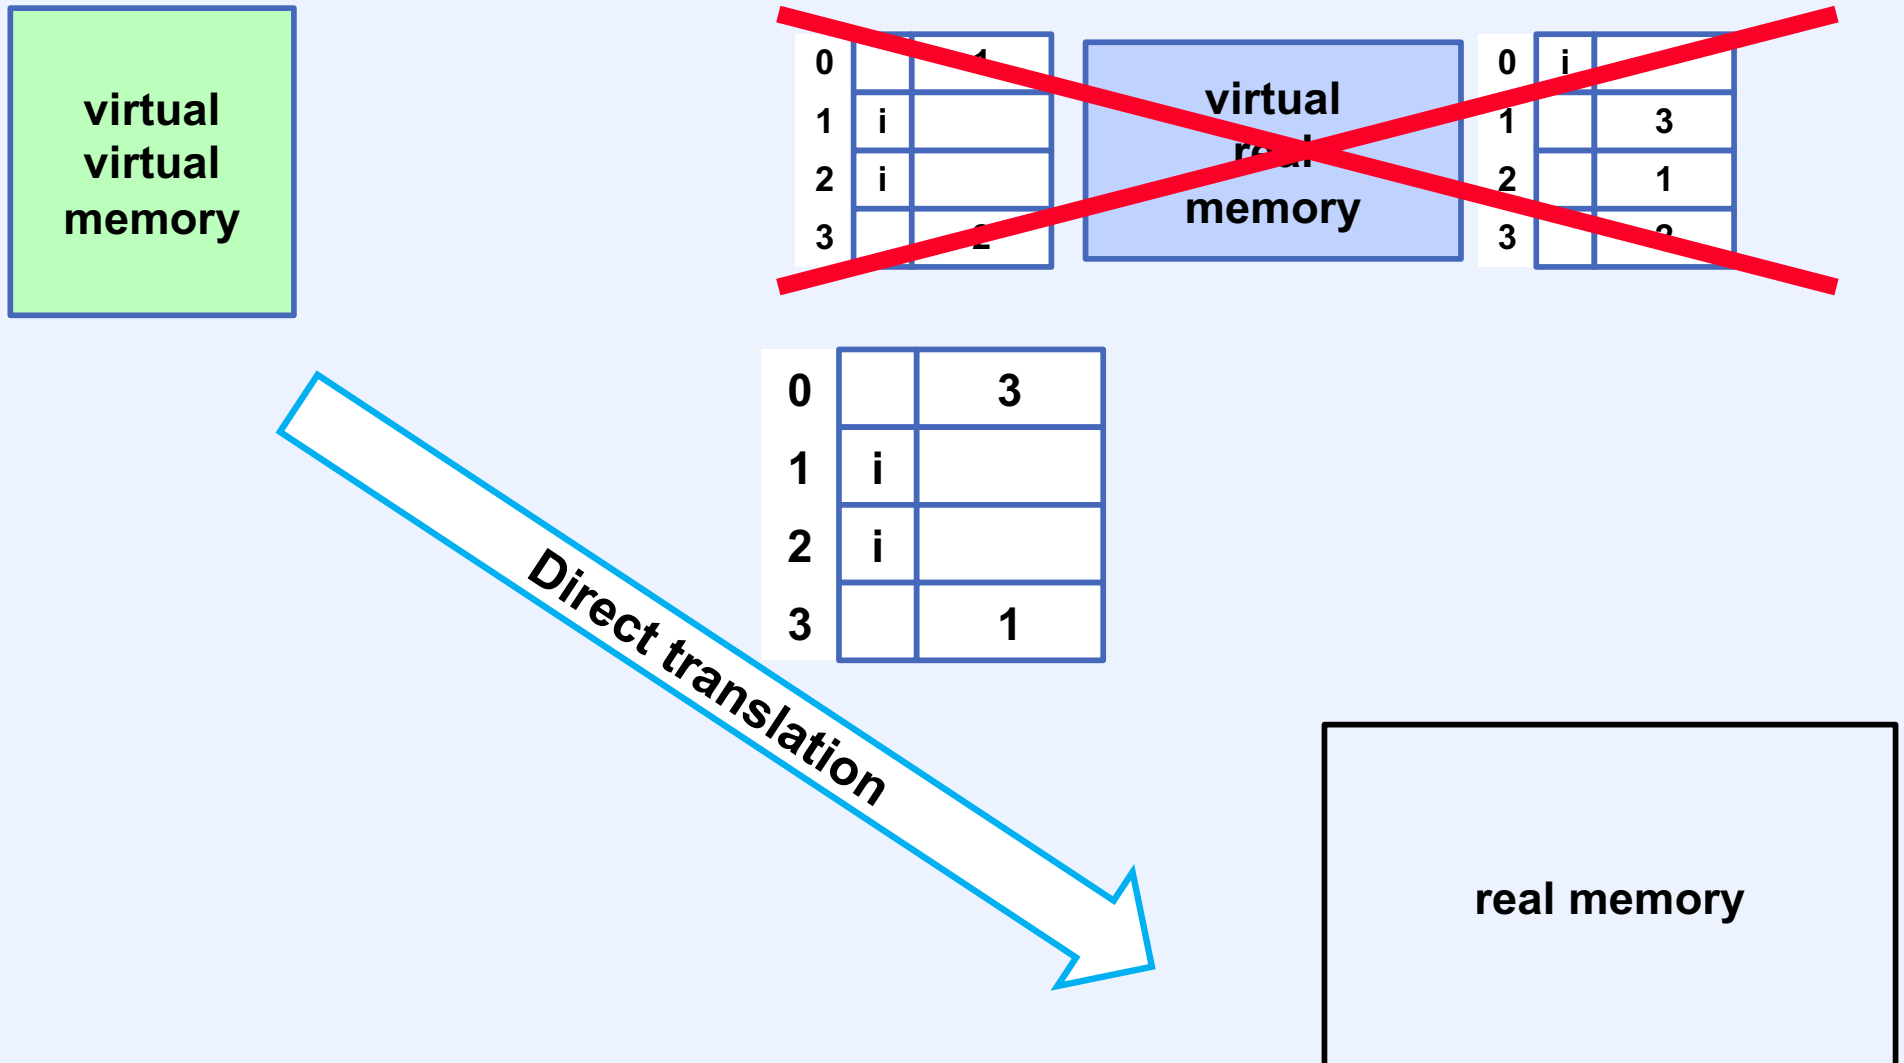
Translation: TLB Miss (HPT)



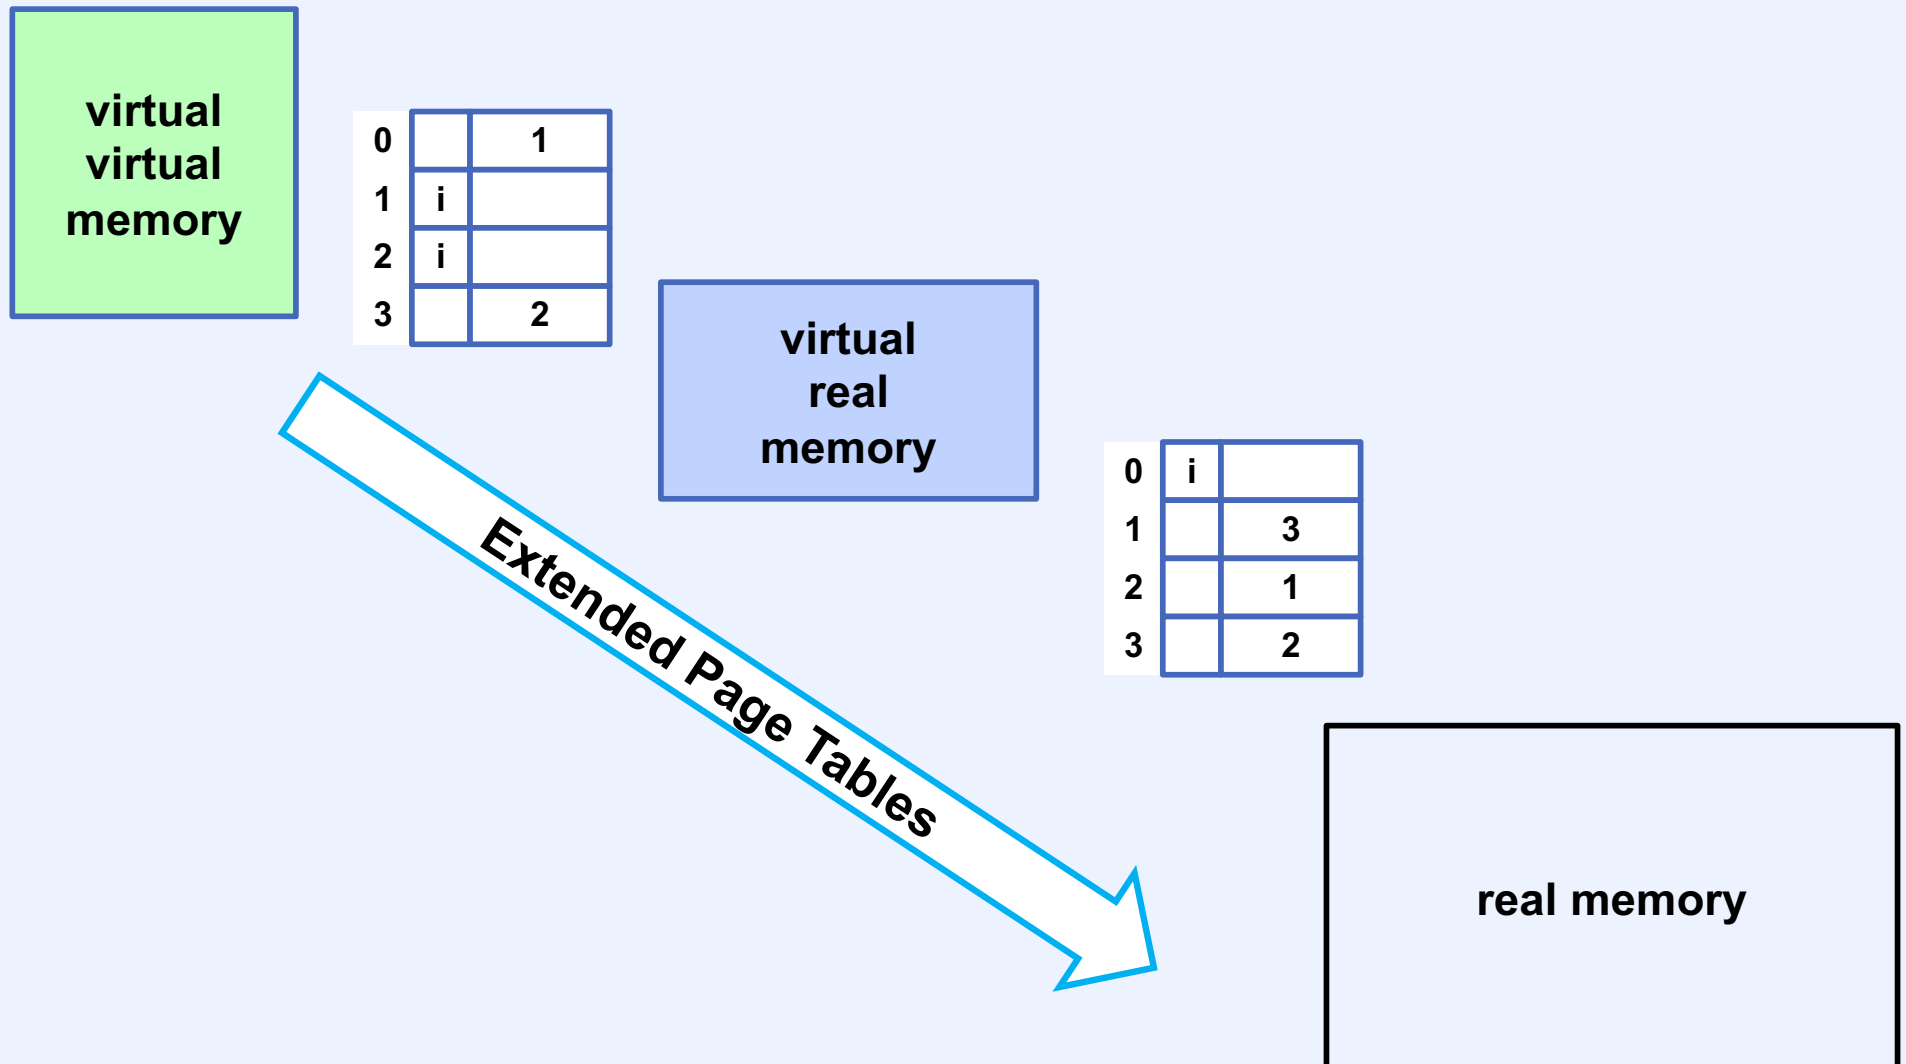
Virtual Machines Meet Virtual Memory



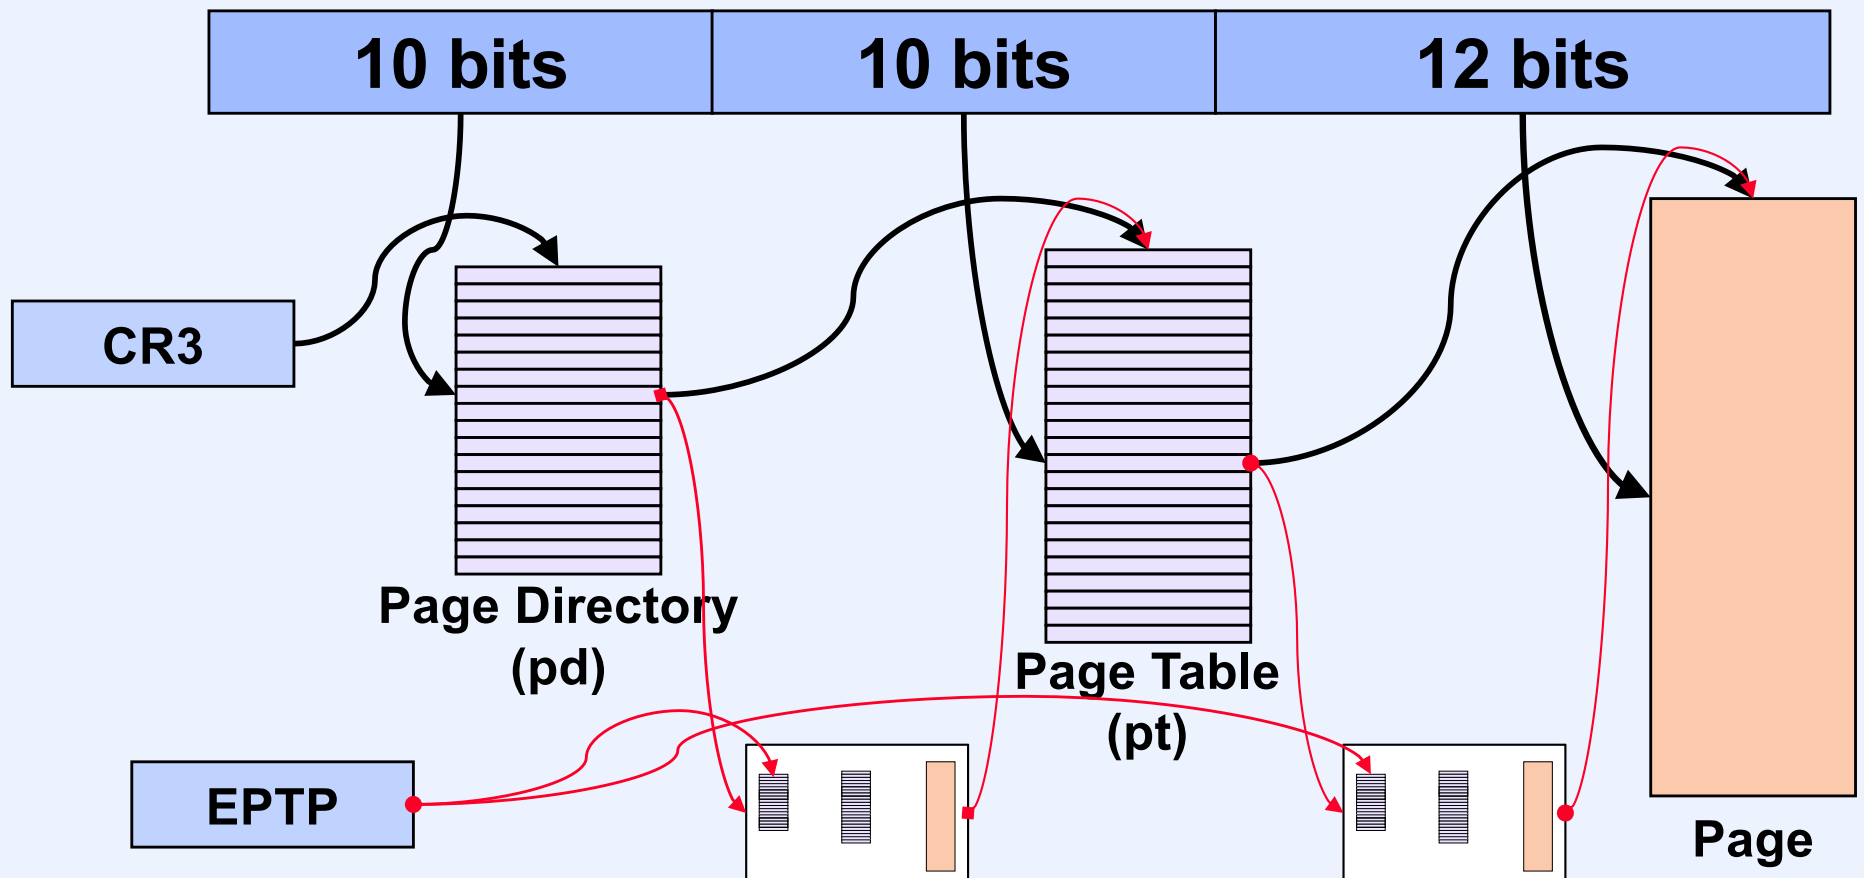
Paravirtualization to the Rescue



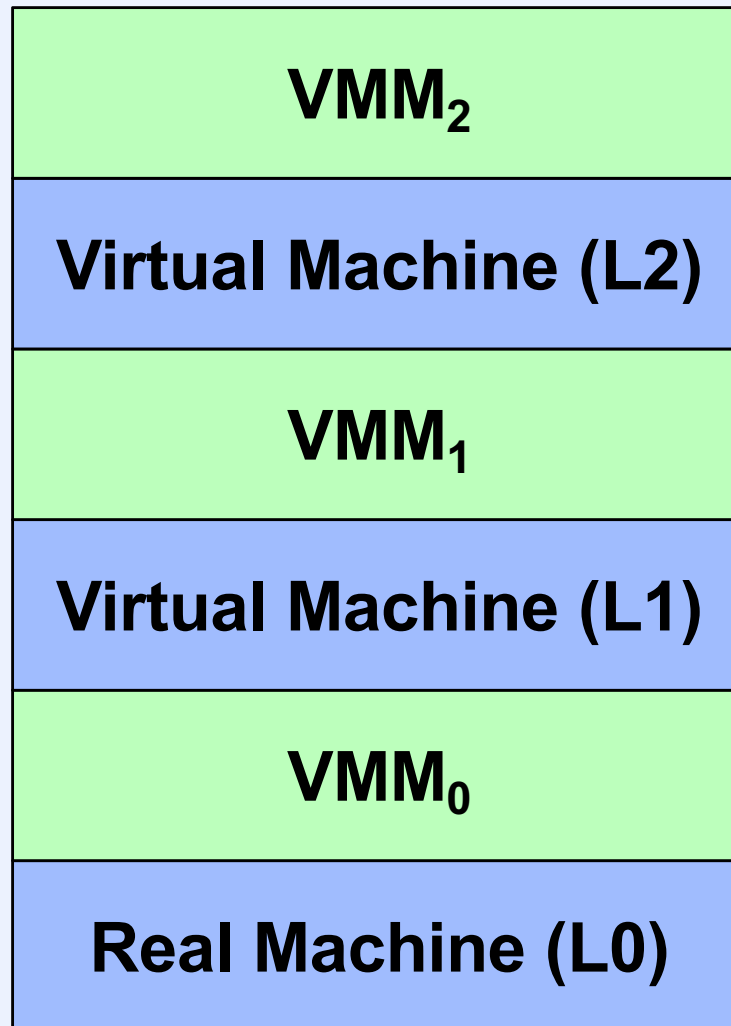
Hardware to the Rescue



x86 Paging with EPT



Nested Virtualization



Quiz 3

We'd like to virtualize EPT. Assume that setting EPTP causes a VMexit if done on a VMM that's not running in real ring -1. What does the VMM running at level 0 (in ring -1) do when it receives such a VMexit from a VMM running at level 1?

- a) it sets EPTP to point to the composition of the page tables mapping VMM_0 's address space to real memory and the page tables pointed to by the value being attempted to be put in EPT**
- b) nothing: the EPT mechanism is virtualized by the hardware**
- c) something else**

VMX

- **New processor mode: root**
 - ring -1: root mode
 - rings 0-3: non-root mode
- **Certain actions cause processor in non-root mode to switch to root mode**
 - VMexit
- **When in root mode, processor can switch back to non-root mode**
 - VMenter

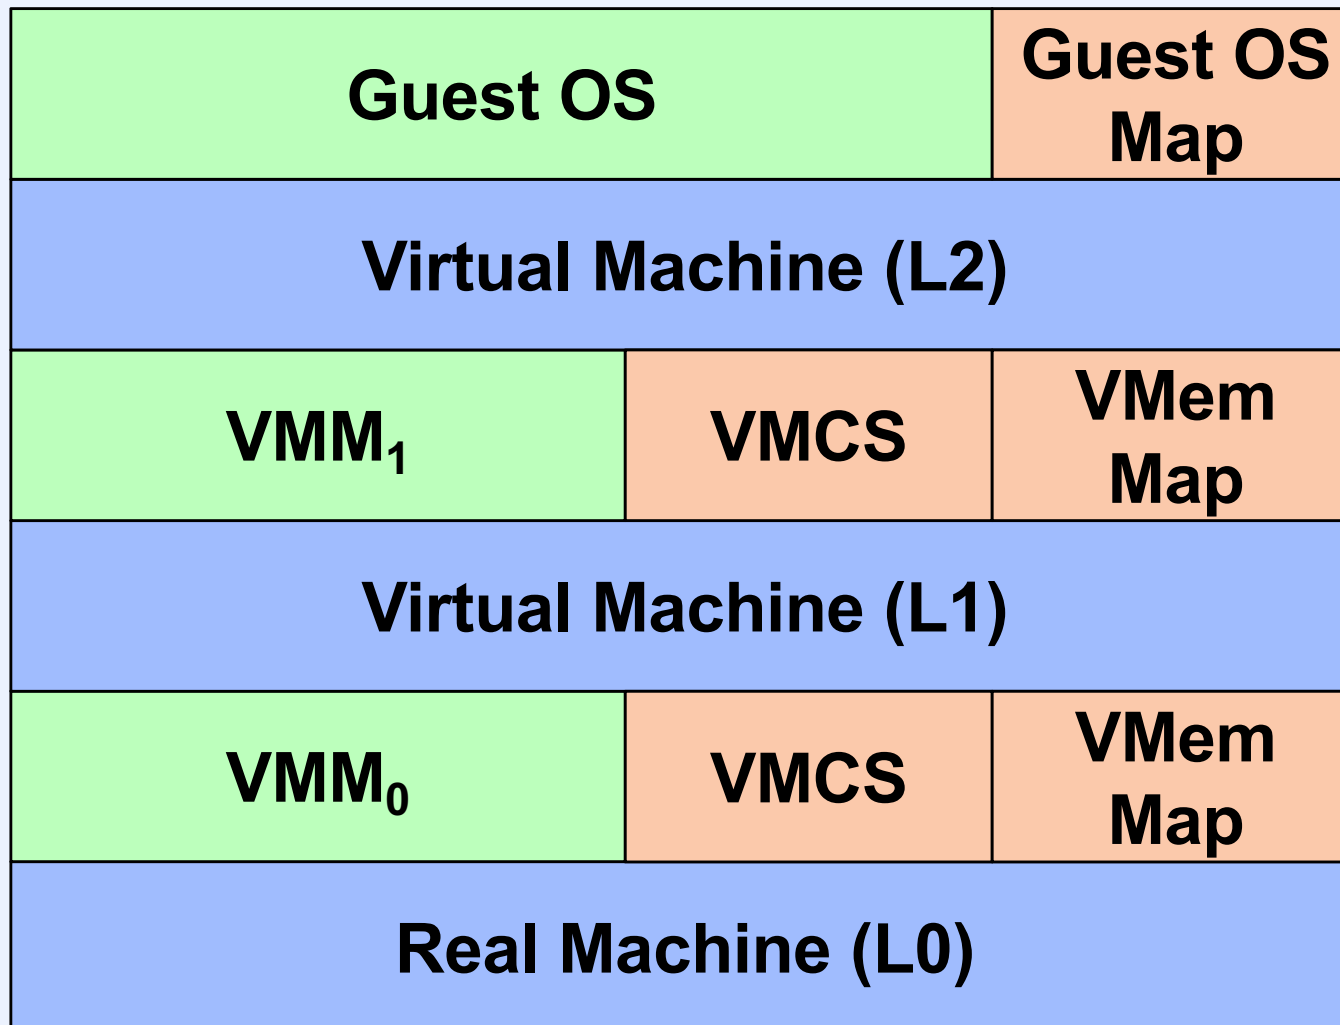
VMCS

- **Virtual machine control structures**
 - **guest state**
 - **virtualized CPU registers (non-root mode)**
 - **host state**
 - **registers to be restored when switching to root mode (VMexit)**
 - **control data**
 - **which events in non-root mode cause VMexits**

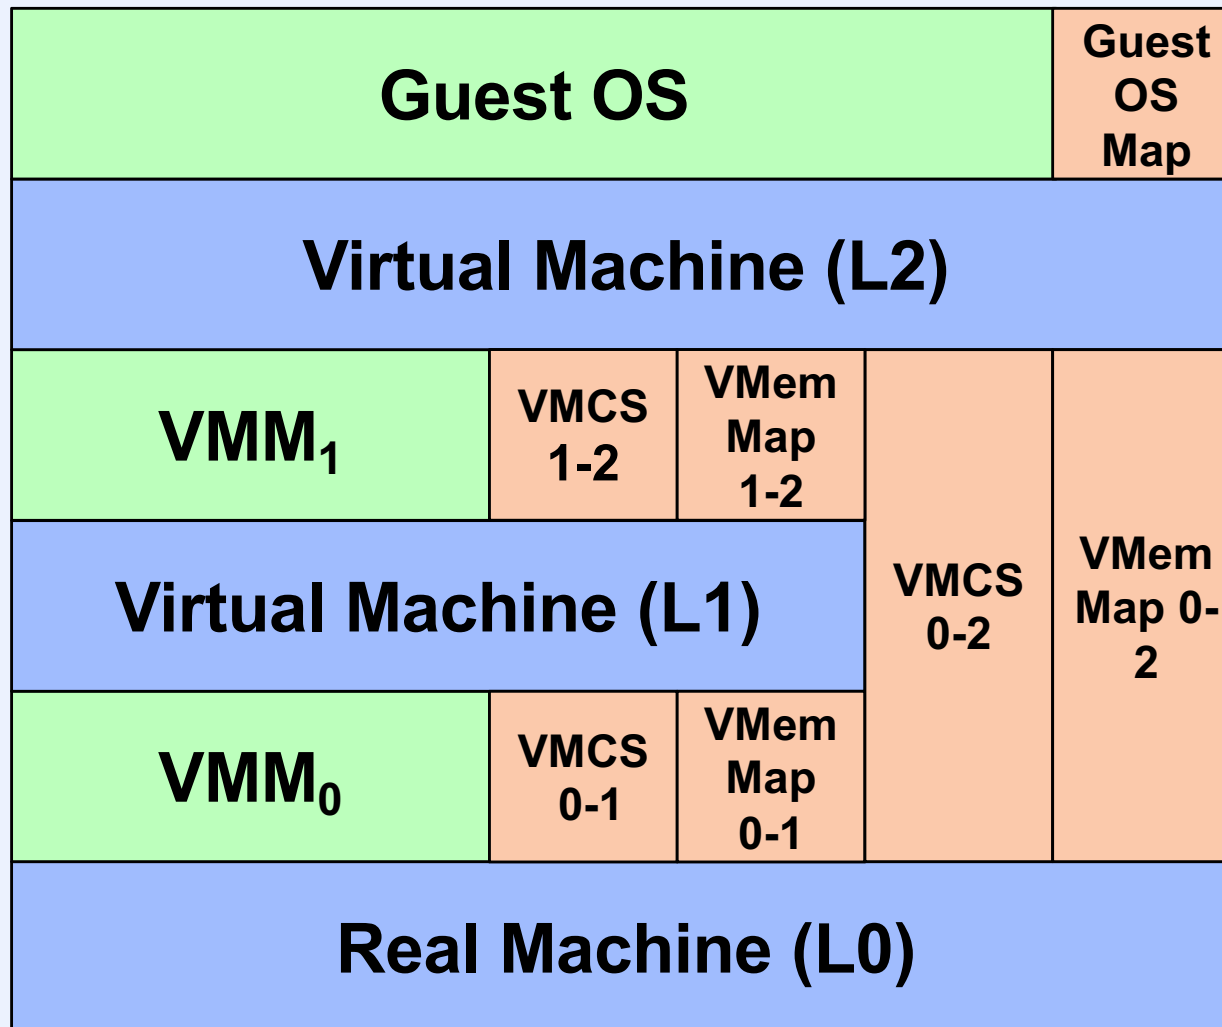
Nested Virtualization on VMX

- A VMM is designed to use VMX extensions (including EPT)
- It supports VMs that appear to be real x86's (but without VMX extensions)
- Can the VMM run in a VM of the level-0 VMM?

Nested Virtualization with VMX



Composed Virtualization



Traditional OS Paging Issues

- **Fetch policy**
- **Placement policy**
- **Replacement policy**

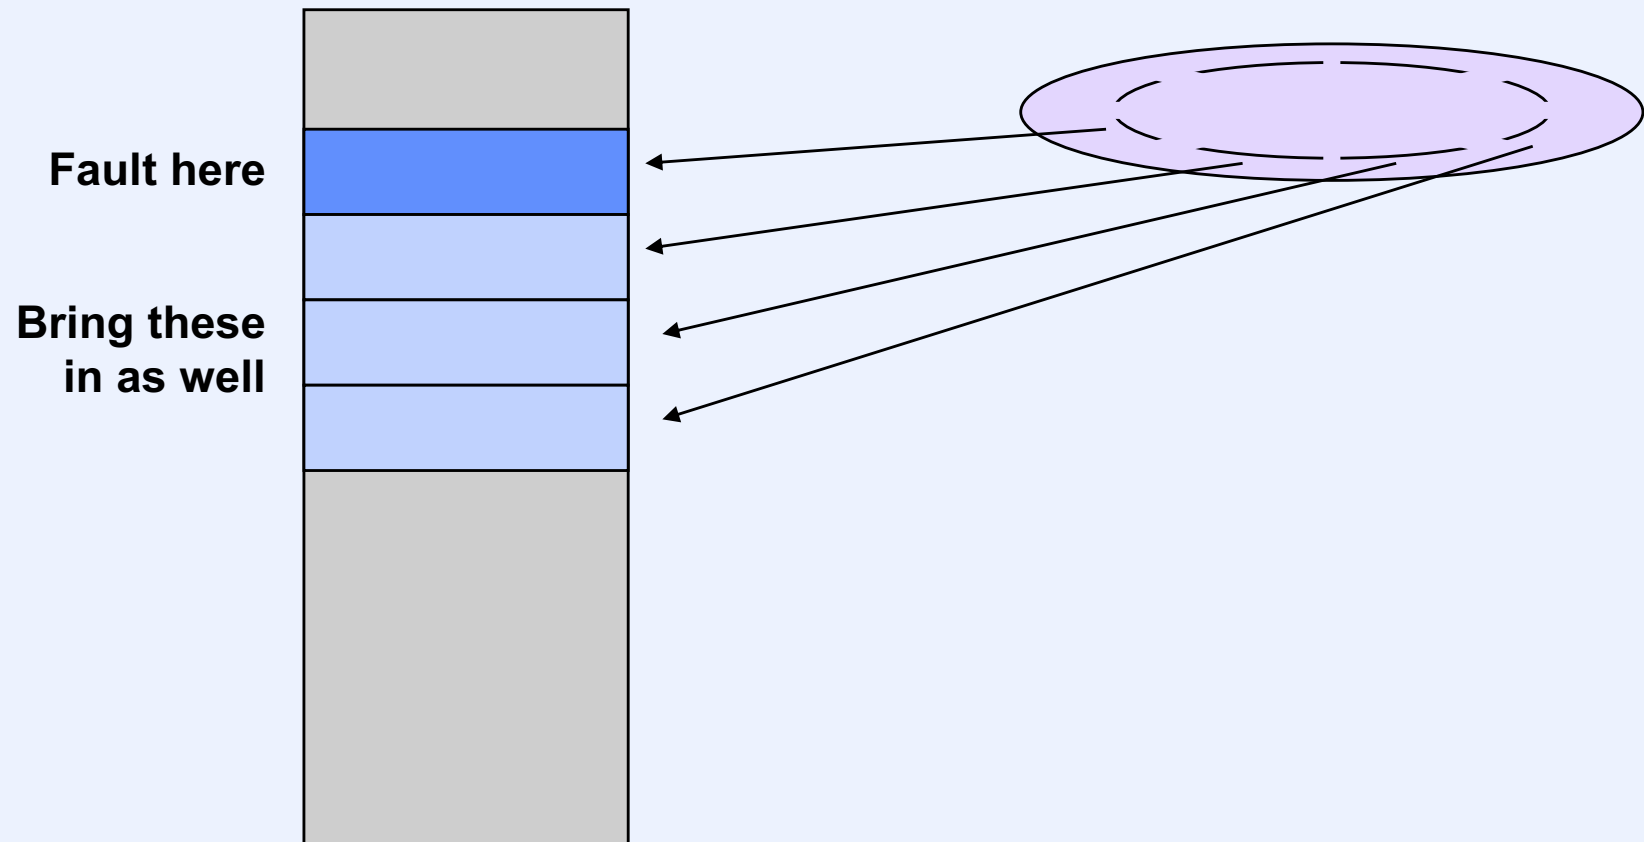
A Simple Paging Scheme

- **Fetch policy**
 - start process off with no pages in primary storage
 - bring in pages on demand (and only on demand — this is known as demand paging)
- **Placement policy**
 - it usually doesn't matter — put the incoming page in the first available page frame
- **Replacement policy**
 - replace the page that has been in primary storage the longest (FIFO policy)

Performance

- 1) Trap occurs (page fault)
- 2) Find free page frame
- 3) Write page out if no free page frame
- 4) Fetch page
- 5) Return from trap

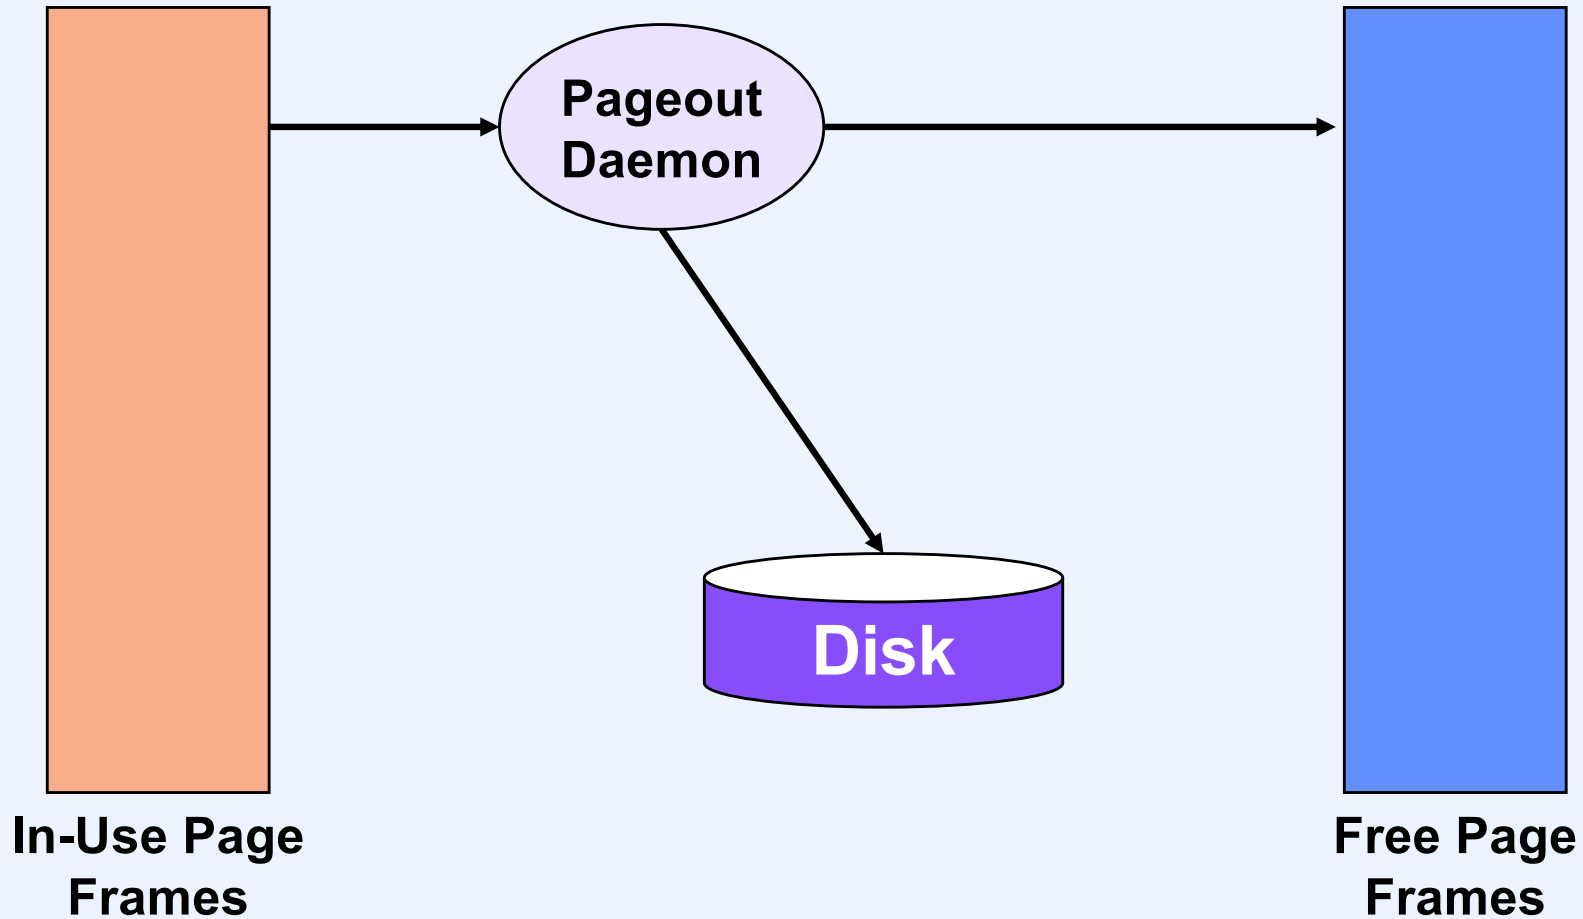
Improving the Fetch Policy



Improving the Replacement Policy

- **When is replacement done?**
 - doing it “on demand” causes excessive delays
 - should be performed as a separate, concurrent activity
- **Which pages are replaced?**
 - FIFO policy is not good
 - want to replace those pages least likely to be referenced soon

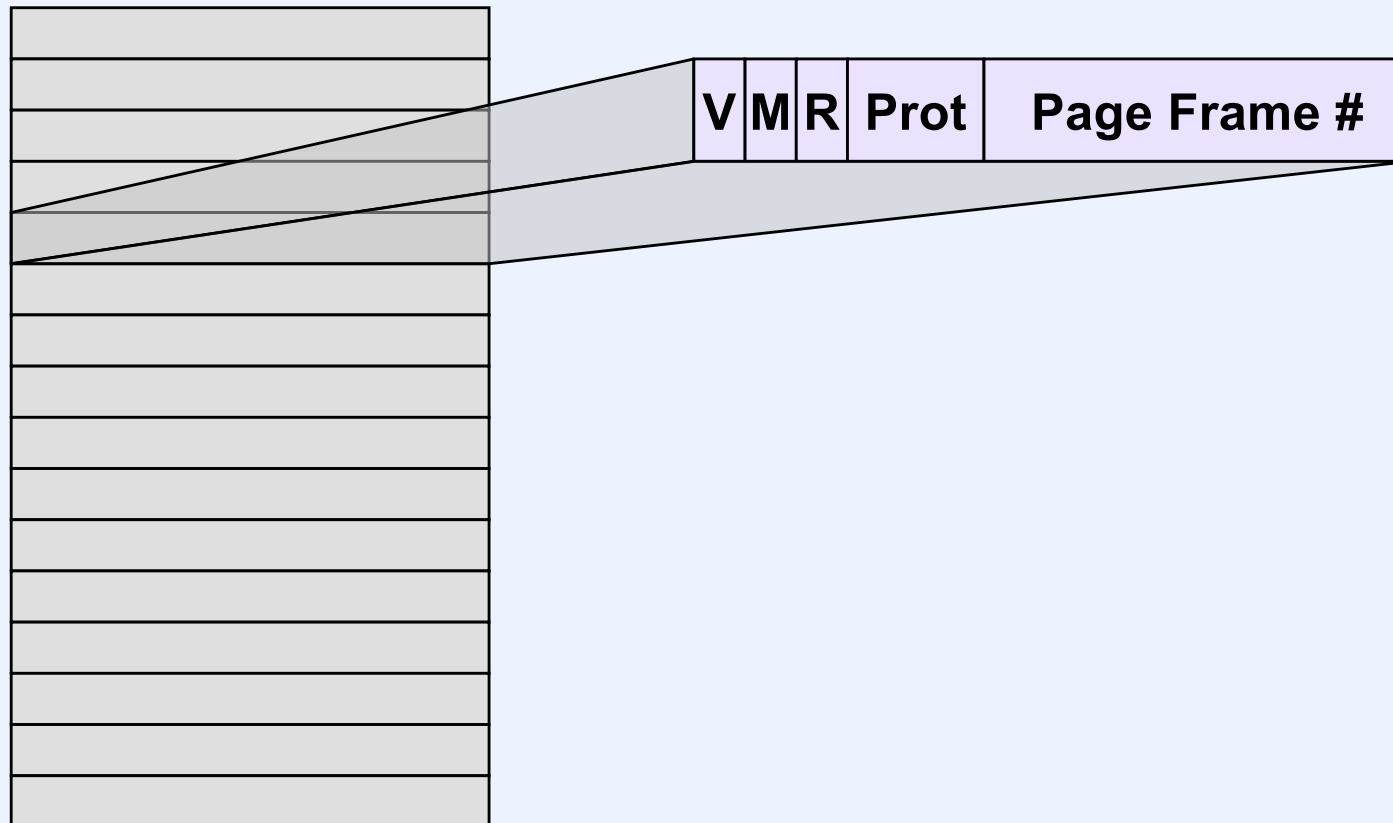
The “Pageout Daemon”



Choosing the Page to Remove

- **Idealized policies:**
 - FIFO (First-In-First-Out)
 - LRU (Least-Recently-Used)
 - LFU (Least-Frequently-Used)
- **Optimal**
 - replace page so as to minimize number of page faults
 - replace page whose next reference is furthest in the future

Implementing LRU



Quiz 4

Your computer is running one process. Pretty much all available real memory is being actively used and processor utilization is around 90%. You now add another process that's similar to the first in terms of both memory and processor utilization (though it's running a different program). Assume the LRU page replacement policy is used.

- a) Processor utilization will rise to nearly 100%**
- b) Processor utilization will stay at around 90%**
- c) Processor utilization will drop precipitously**

Global vs. Local Allocation

- **Global allocation**
 - all processes compete for page frames from a single pool
- **Local allocation**
 - each process has its own private pool of page frames

Thrashing

- **Consider a system that has exactly two page frames:**
 - process A has a page in frame 1
 - process B has a page in frame 2
- **Process A causes a page fault**
- **The page in frame 2 is removed**
- **Process B faults; the page in frame 1 is removed**
- **Process A resumes execution and faults again; the page in frame 2 is removed**
- **...**

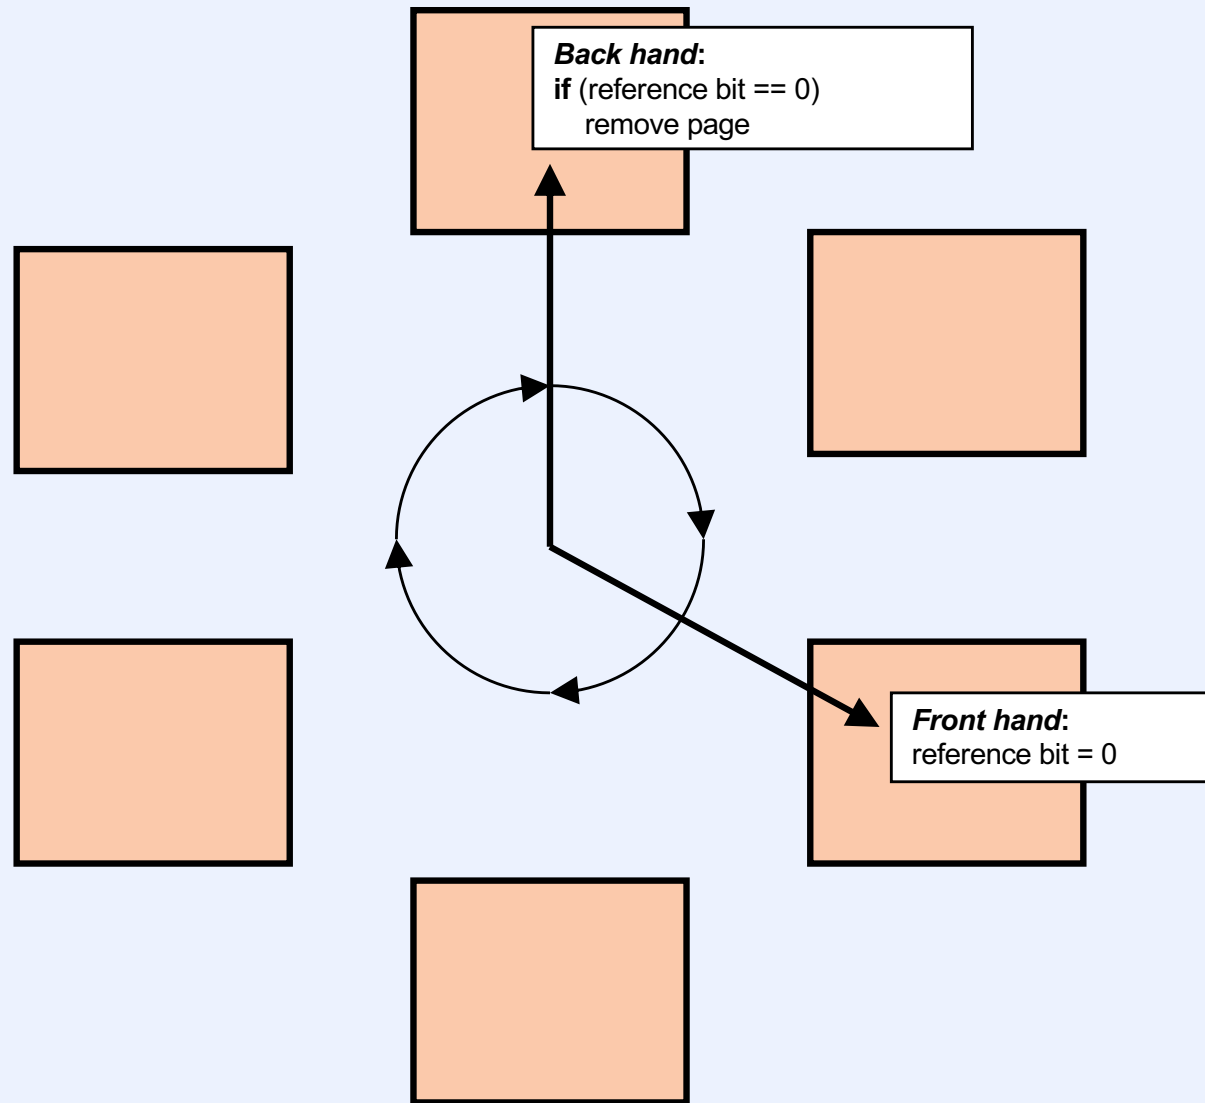
The Working-Set Principle

- The set of pages being used by a program (the working set) is relatively small and changes slowly with time
 - $WS(P,T)$ is the set of pages used by process P over time period T
- Over time period T , P should be given $|WS(P,T)|$ page frames
 - if space isn't available, then P should not run and should be swapped out

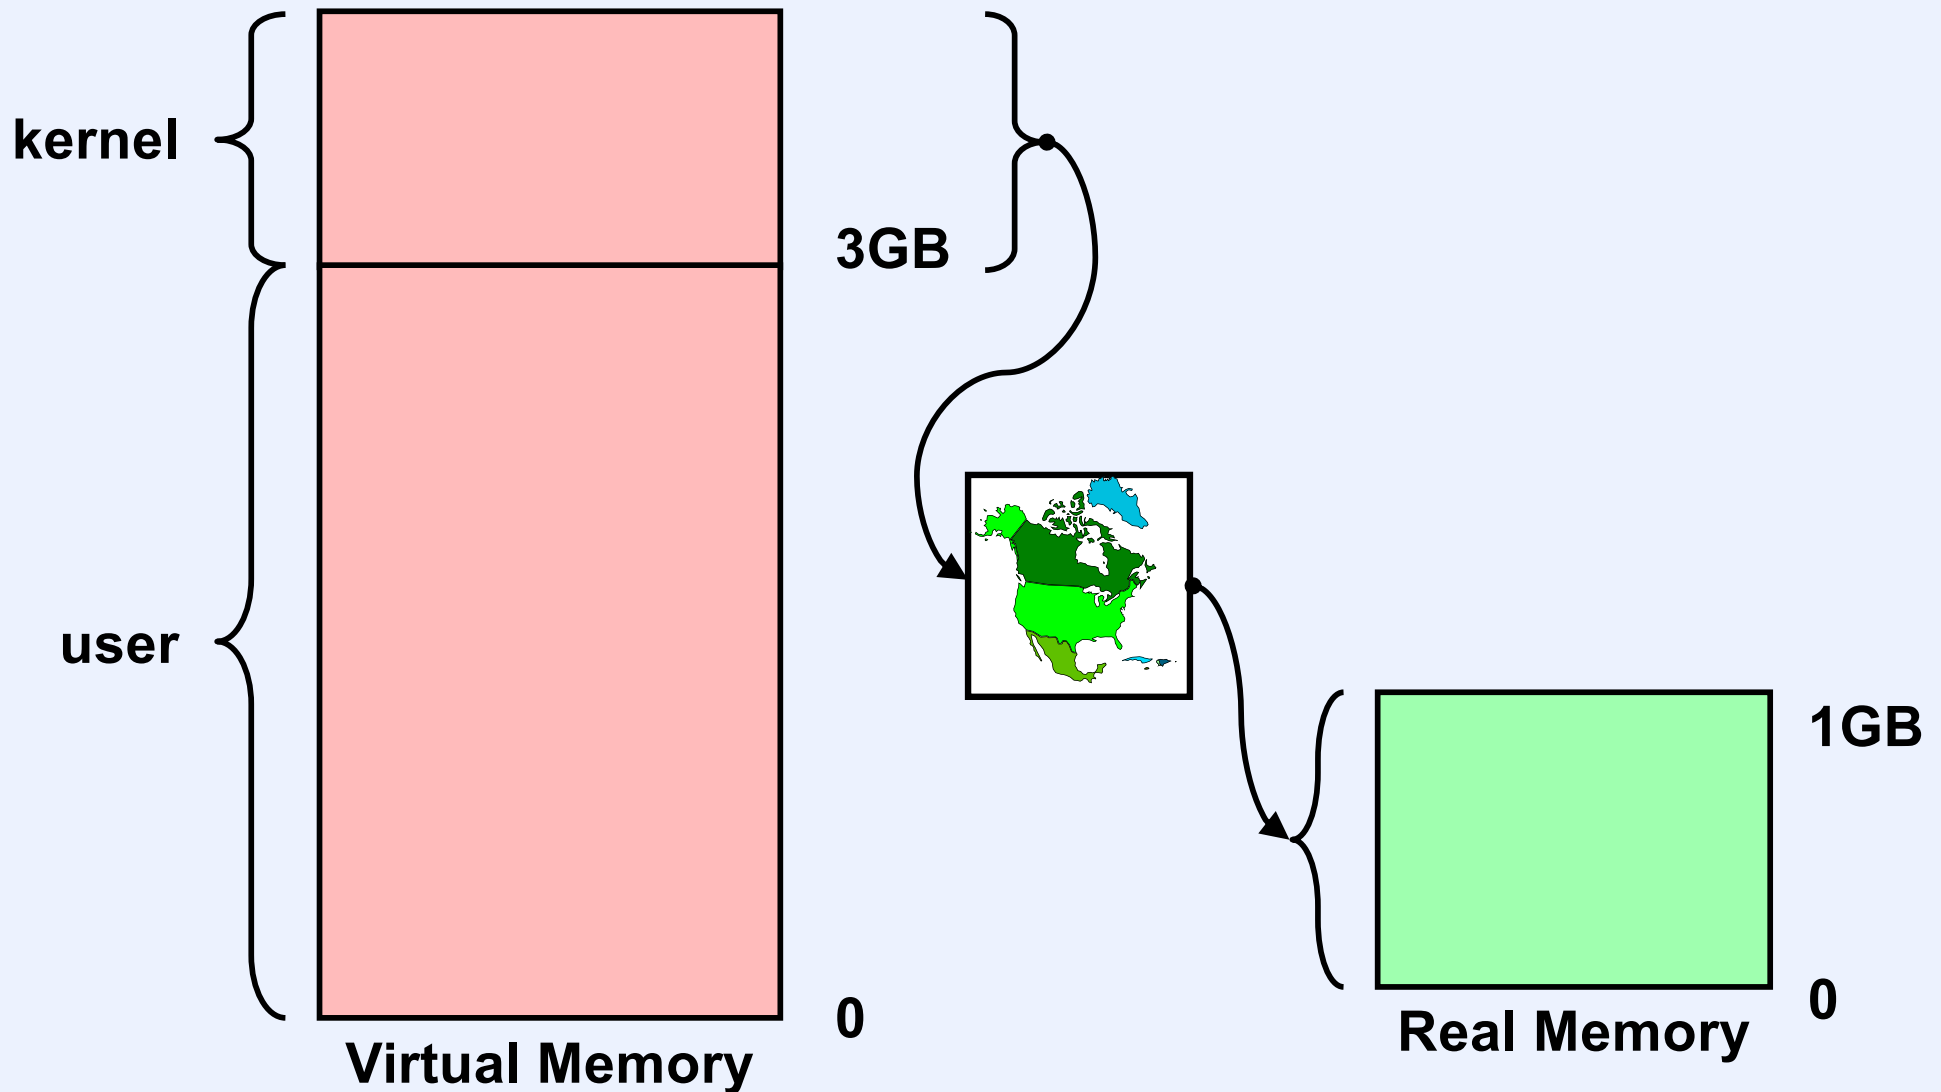
Two Issues

- If a process is active, which of its pages should be in real memory?
- If there is too much of a demand for memory, which processes should run (and which should not run)?

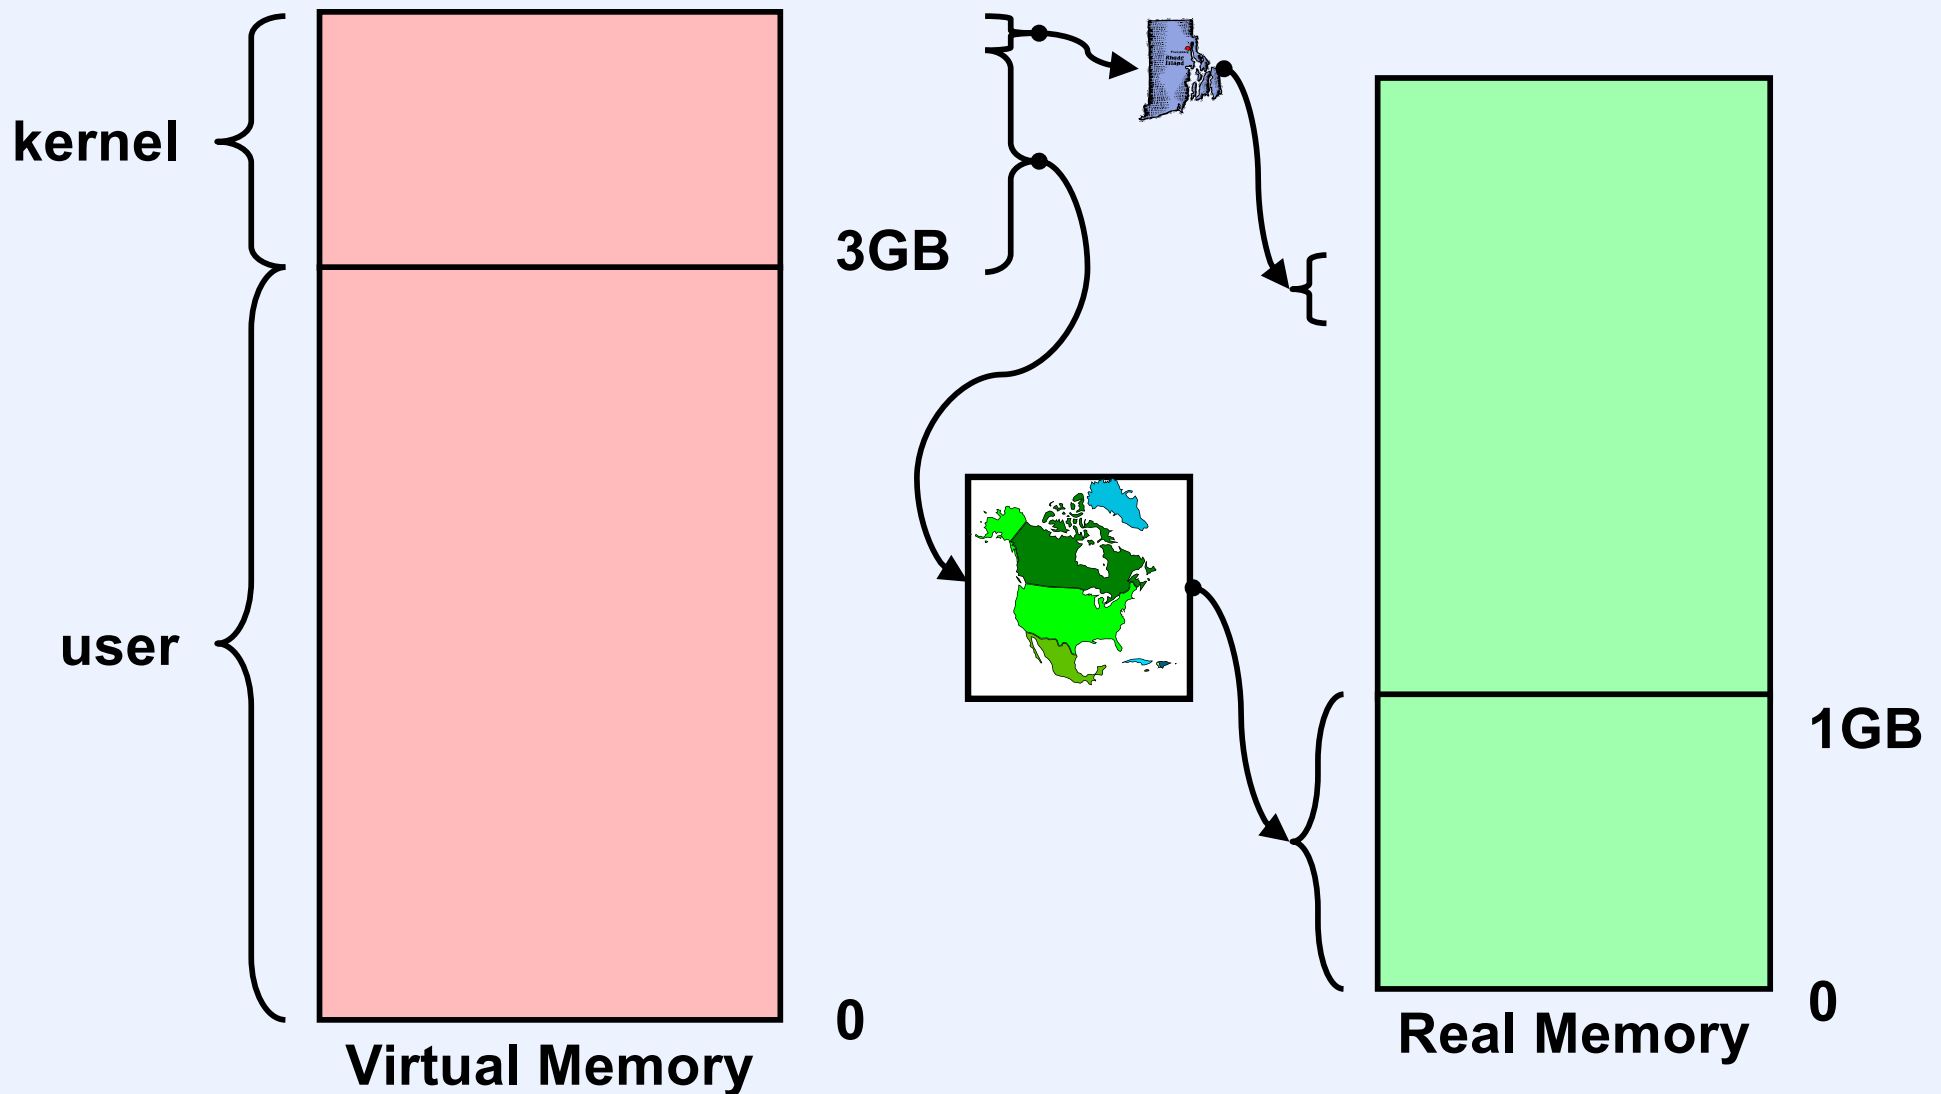
Clock Algorithm



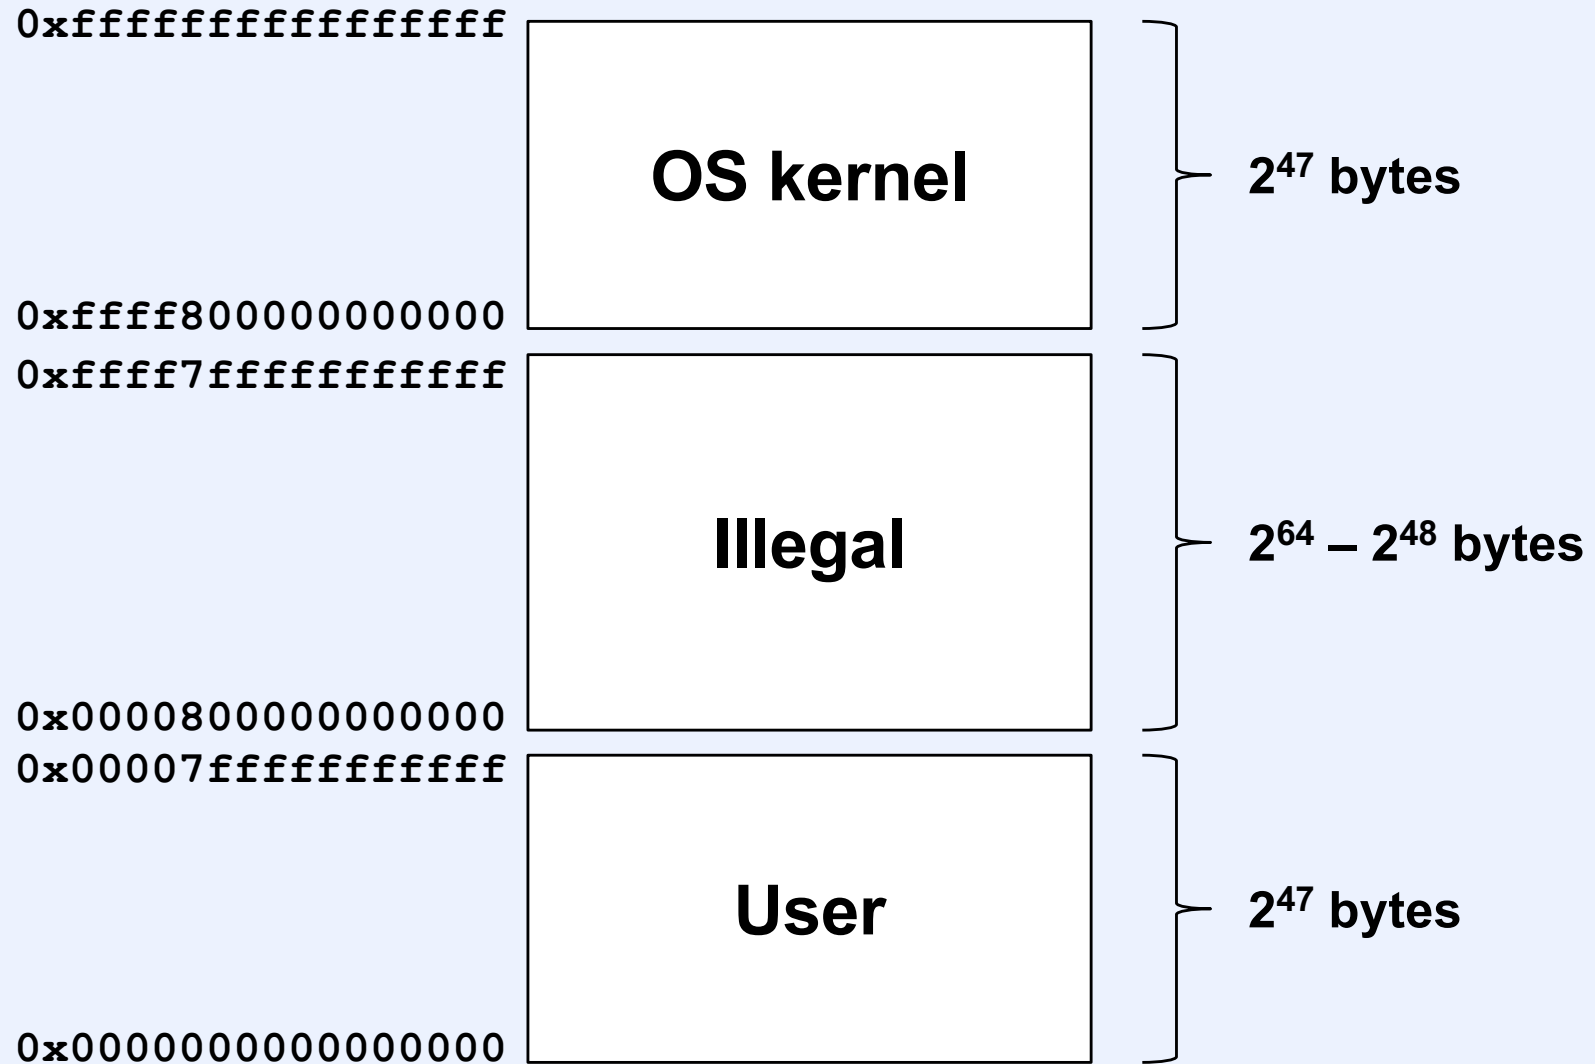
Linux and Real Memory



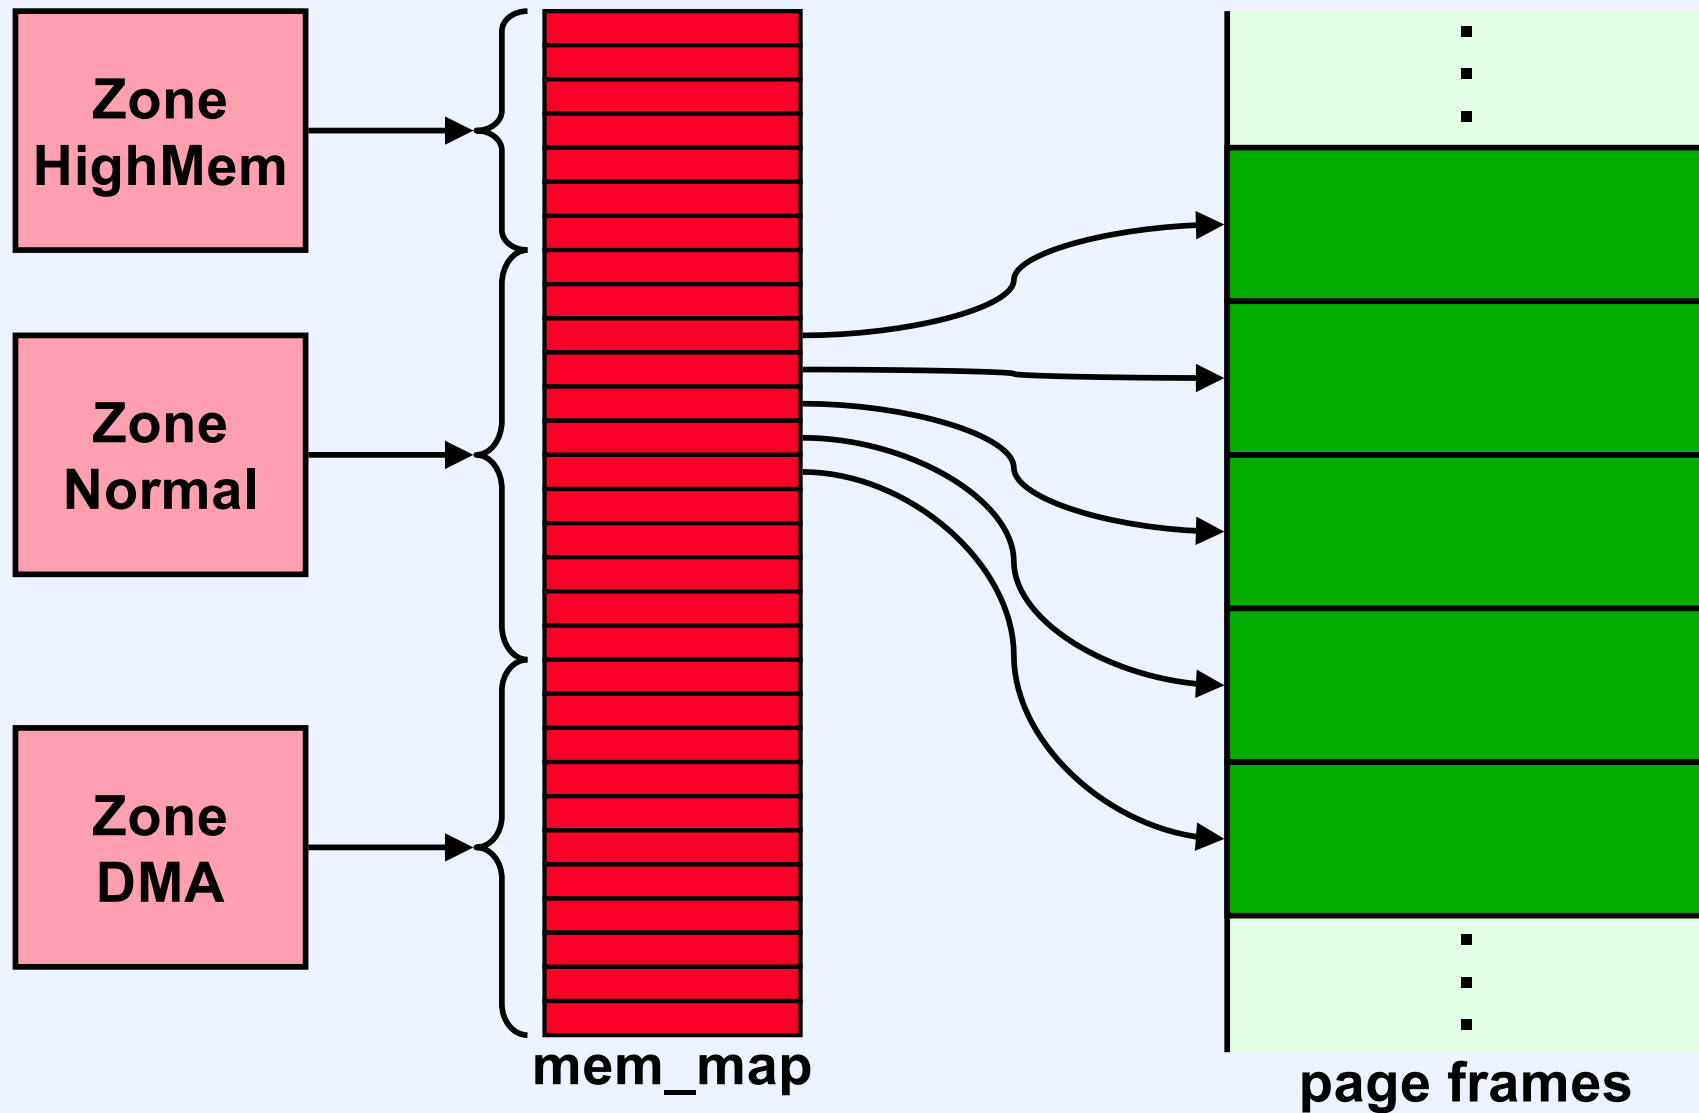
Lots of Real Memory



Address Space



Mem_map and Zones



Quiz 5

We have a disk whose controller uses 64-bit memory addresses. We'd like to set it up for a read that will transfer 32K bytes. Thus the block will be read into a buffer that occupies eight 4KB pages.

- a) The buffer must occupy eight page frames of contiguous real memory**
- b) Since our system uses virtual memory, the buffer occupies contiguous pages of virtual memory, which may be mapped into non-contiguous page frames of real memory**
- c) As in b, except that the contiguous pages of virtual memory need not have valid mappings (and thus page faults are generated and handled)**