

Security Part 5

Live Anonymous Q&A:

<https://tinyurl.com/cs1670feedback>

Recap: TOCTTOU vulnerability

```
/* handin: a setuid-twd  
   program */  
  
if (access(argv[1],  
          R_OK) == 0) {  
    // ... fail  
}  
fd = open(argv[1],  
          O_RDONLY);  
/* copy argv[1] to course  
   directory */
```

```
% handin my_assgn  
...
```

Hidden Code

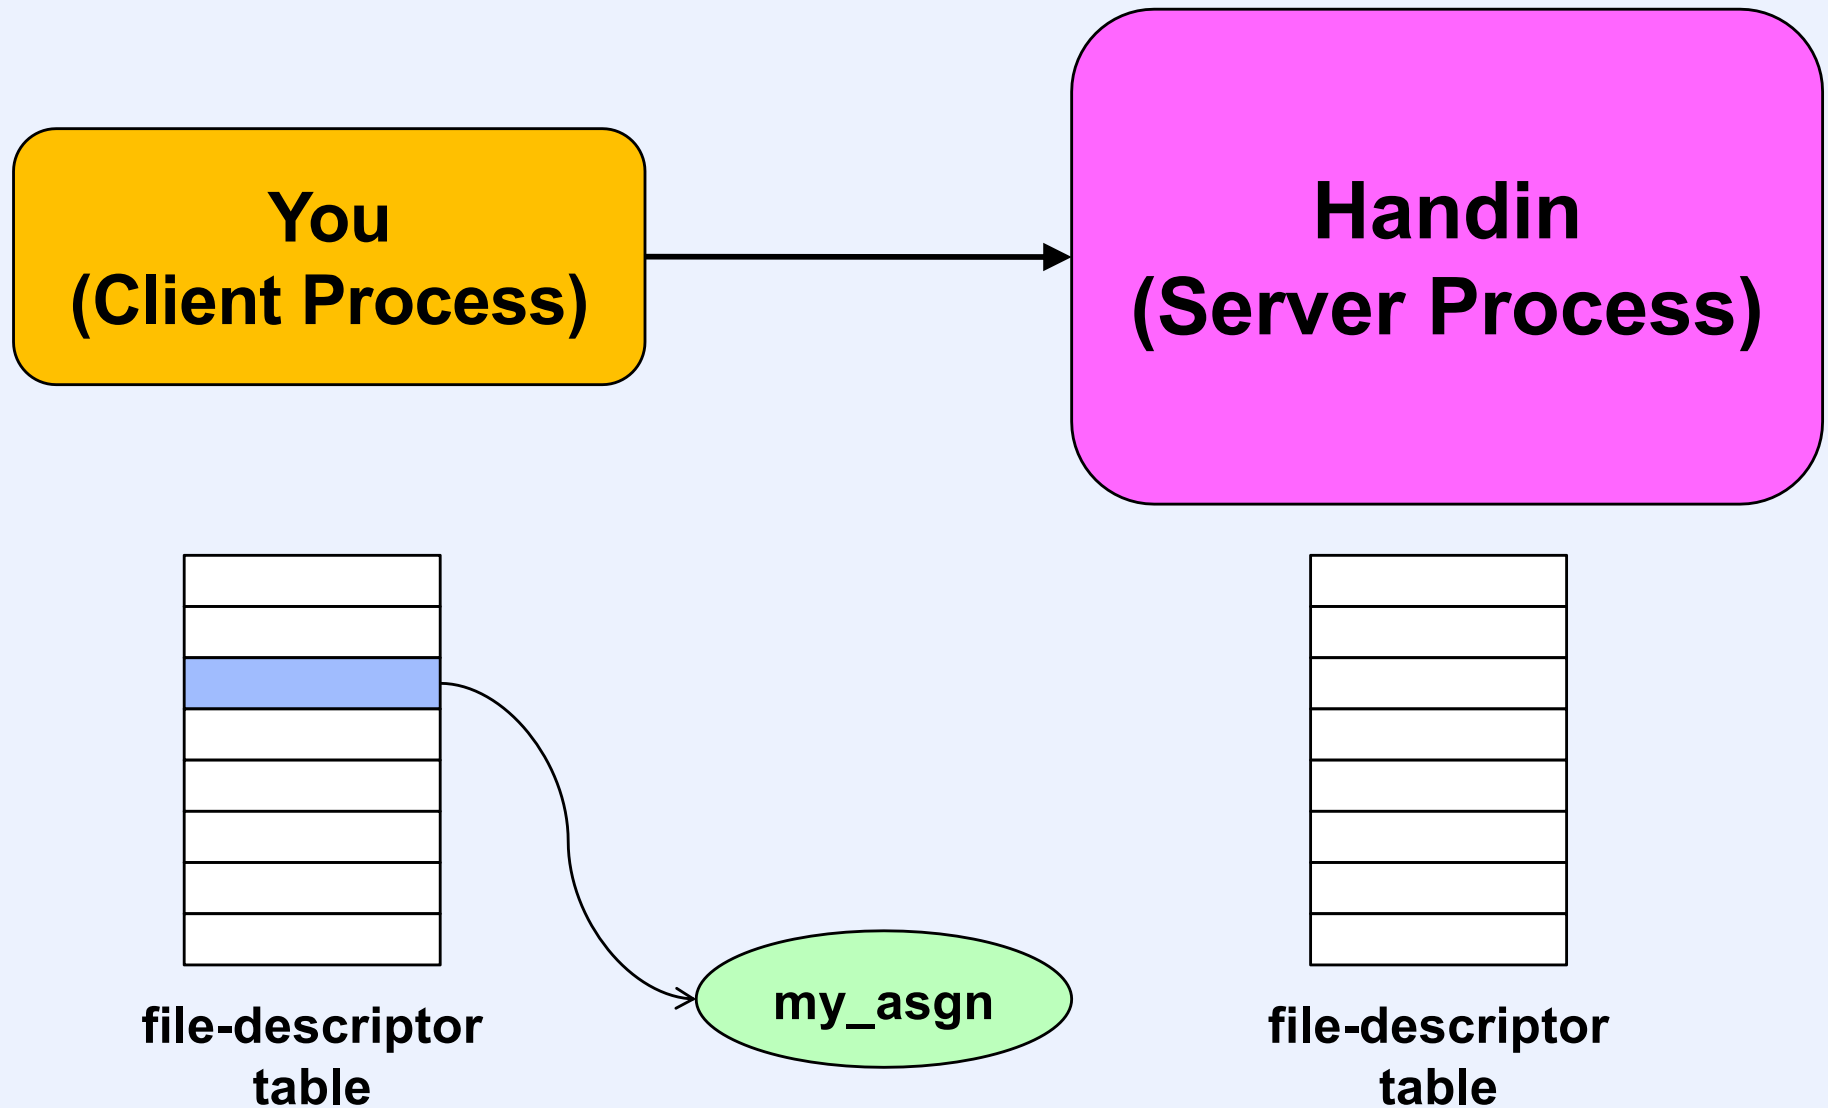
How to Solve?

- **Could use previous solution**
or

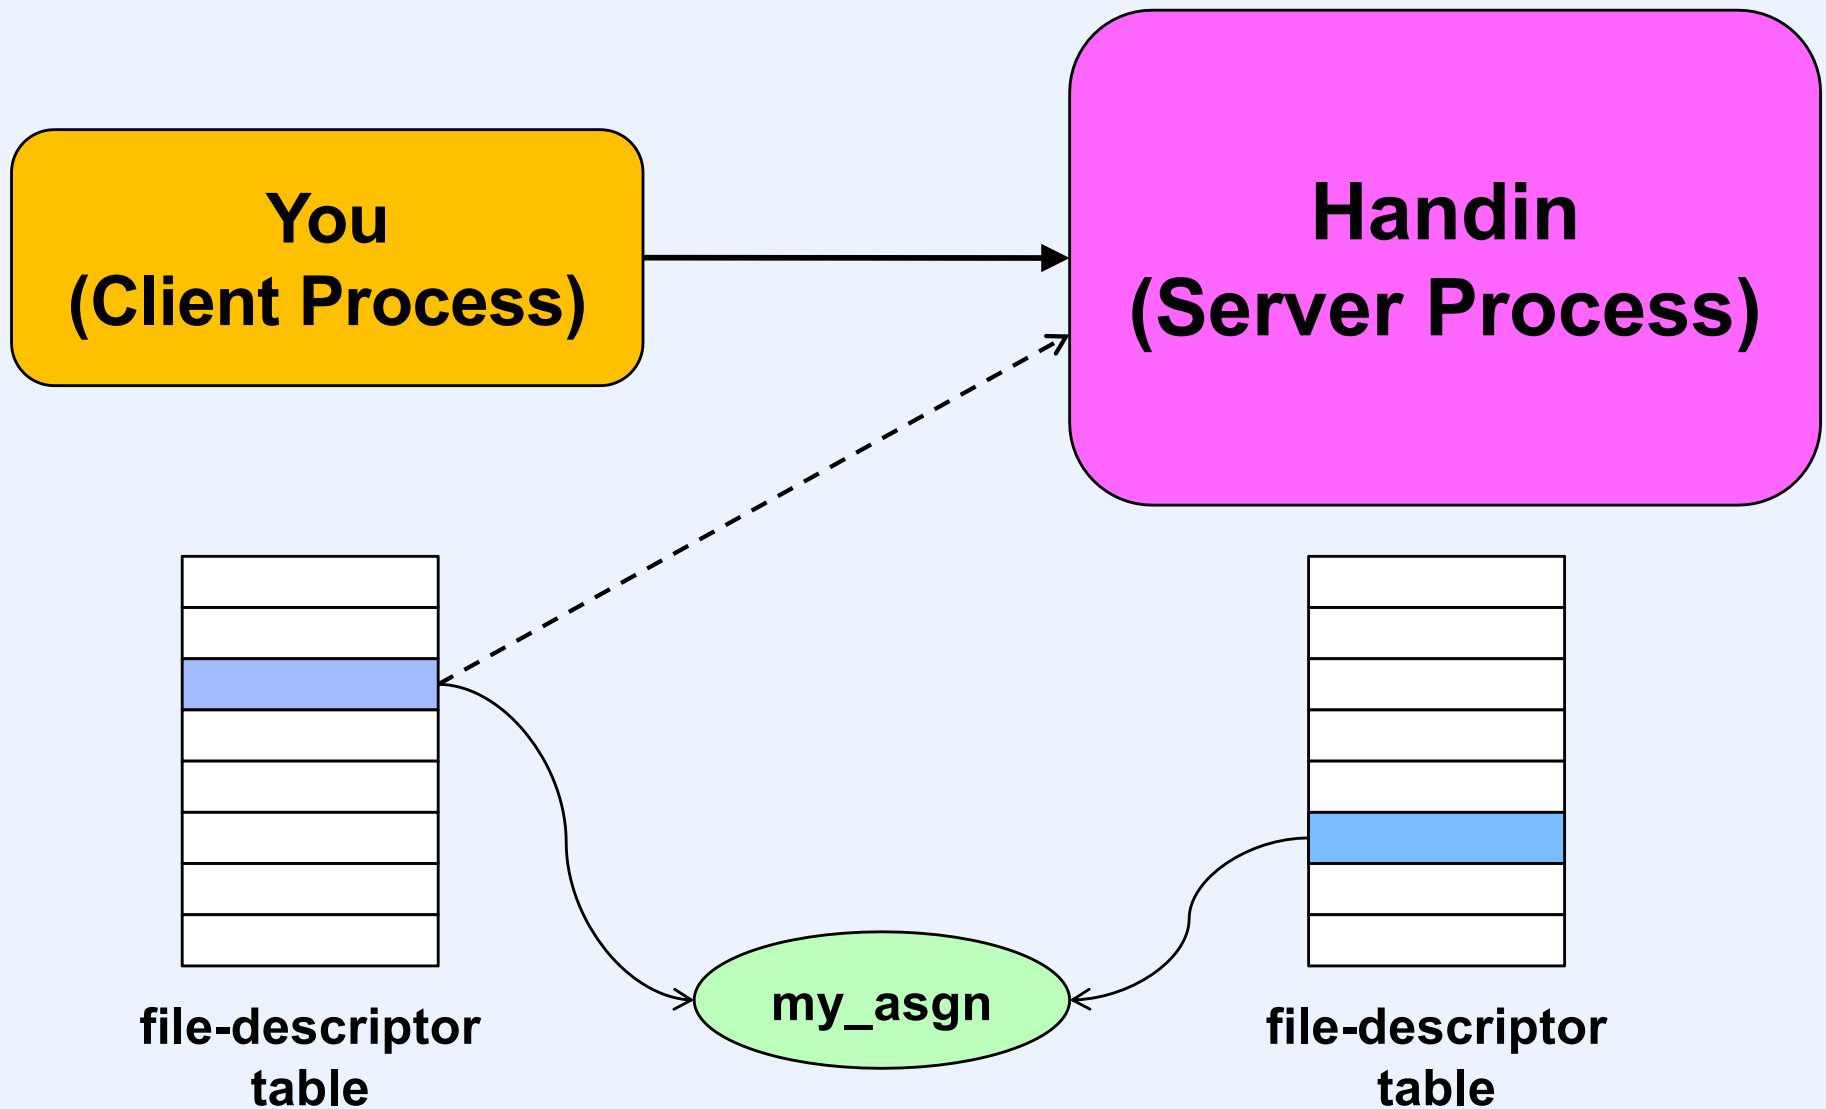
```
afd = open("my_assgn",  
           O_RDONLY);  
close(0);  
dup(afd);  
close(afd);  
execl("handin", 0);
```

```
/* handin */  
int main() {  
    int user = getuid();  
    char fname[256];  
    sprintf(fname,  
            "CourseDirectory/%d", user);  
    int ofd = open(fname,  
                   O_CREAT|O_WRONLY, 0666);  
    while(1) {  
        if ((c = read(0, buf, 256)) == 0)  
            break;  
        write(ofd, buf, c);  
    }  
    return 0  
}
```

Same But Different



File Descriptor as *Capability*



Changing Security Context

Shell Commands

- **su [user_name]**
 - run a new shell with real and effective user IDs being those of user_name
 - if no user_name, then root (super user)
 - must supply correct password for user_name
- **sudo program**
 - run program with appropriate identity and privileges
 - checks to see if caller has permission
 - protected file lists who is allowed to do what
 - must supply your password

Programming Securely

- **It's hard!**
- **Some examples ...**

Truncated Paths

```
int GetFile(char *dirpath, char *name) {
    char FullyQualifiedName[1024];
    if (CheckName(dirpath) == BAD) {
        ...
    }
    strncpy(FullyQualifiedName, dirpath, 512);
    strncat(FullyQualifiedName, name, 512);
    return(open(FullyQualifiedName, O_RDWR));
}

GetFile("////////////////////////...//tmp", vmlinuz);
```

Defense

- **It's not enough to avoid buffer overflow ...**
- **Check for truncation!**

Carelessness

```
char buf[100];
```

```
int len; ← Should be size_t
```

```
read(fd, &len, sizeof(len));
```

} Read data size
into len

```
if (len > 100) {
```

```
    fprintf(stderr, "bad length\n");
```

```
    exit(1);
```

```
}
```

} Intention:
check code
doesn't read
too much

```
read(fd, buf, len);
```

} Read actual data
(len bytes)

A Real-Life Exploit ...

- **sendmail -d6,50**
 - means: set flag 6 to value 50
 - debug option, so why check for min and max?
 - (shouldn't have been turned on for production version ...)
 - (but it was ...)
- **sendmail -d4294967269,117 -d4294967270,110 -d4294967271,113 changed etc to *tmp***
 - /etc/sendmail.cf identifies file containing mailer program, which is executed as root
 - /tmp/sendmail.cf supplied by attacker
 - identifies /bin/sh as mailer program
 - attacker gets root shell

What You Don't Know ...

```
int TrustedServer(int argc, char *argv[]) {  
    ...  
    printf(argv[1]);  
    ...  
}
```

```
% TrustedServer "xyz%n"
```

from the printf man page:

`%n` The number of characters written so far is stored into the integer indicated by the `int *` (or variant) pointer argument. No argument is converted.

→ printf changes arbitrary (?) memory location

Does This Work?

```
% setenv LD_PRELOAD myversions/libcrypt.so.1
```

```
% su
```

```
Password:
```

Isolating Security Contexts



Principle of Least Privilege

- **Perhaps:**
 - run process with a minimal security context
 - special account, etc.
 - send it the capabilities it needs

Complete Isolation

- **Would like to run multiple applications in complete isolation from one another**
 - run them on separate computers with no common file system
 - run them on separate virtual machines
 - run them in separate *containers* on one OS instance

VMs versus Containers



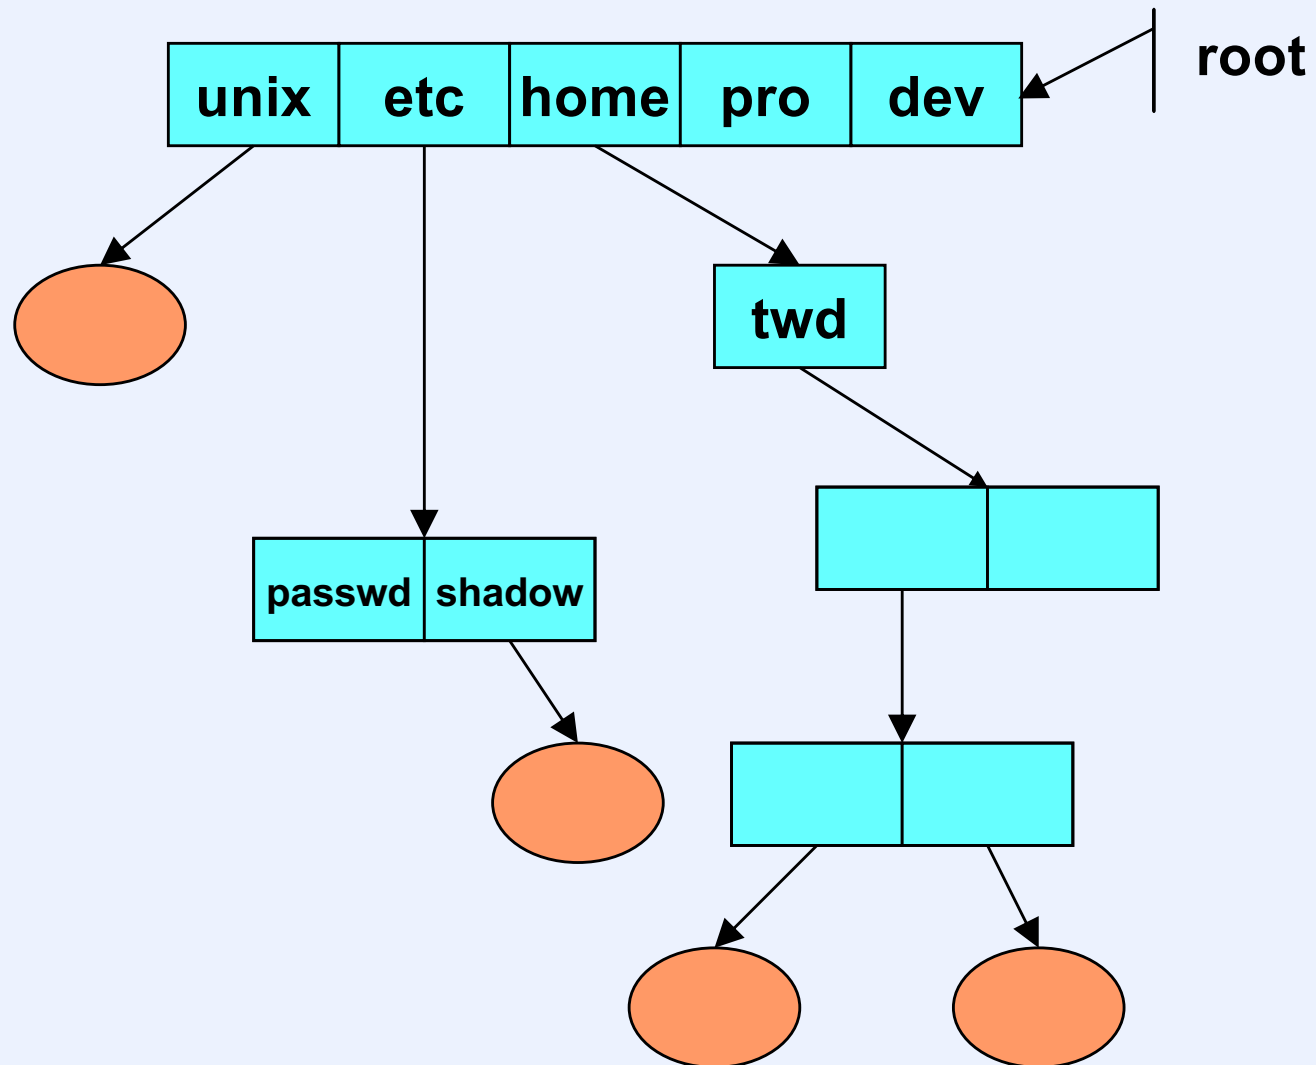
Containers

- **Isolated**
 - processes in a container can't access what's not in the container
 - processes in a container shouldn't even be aware of what's not in the container

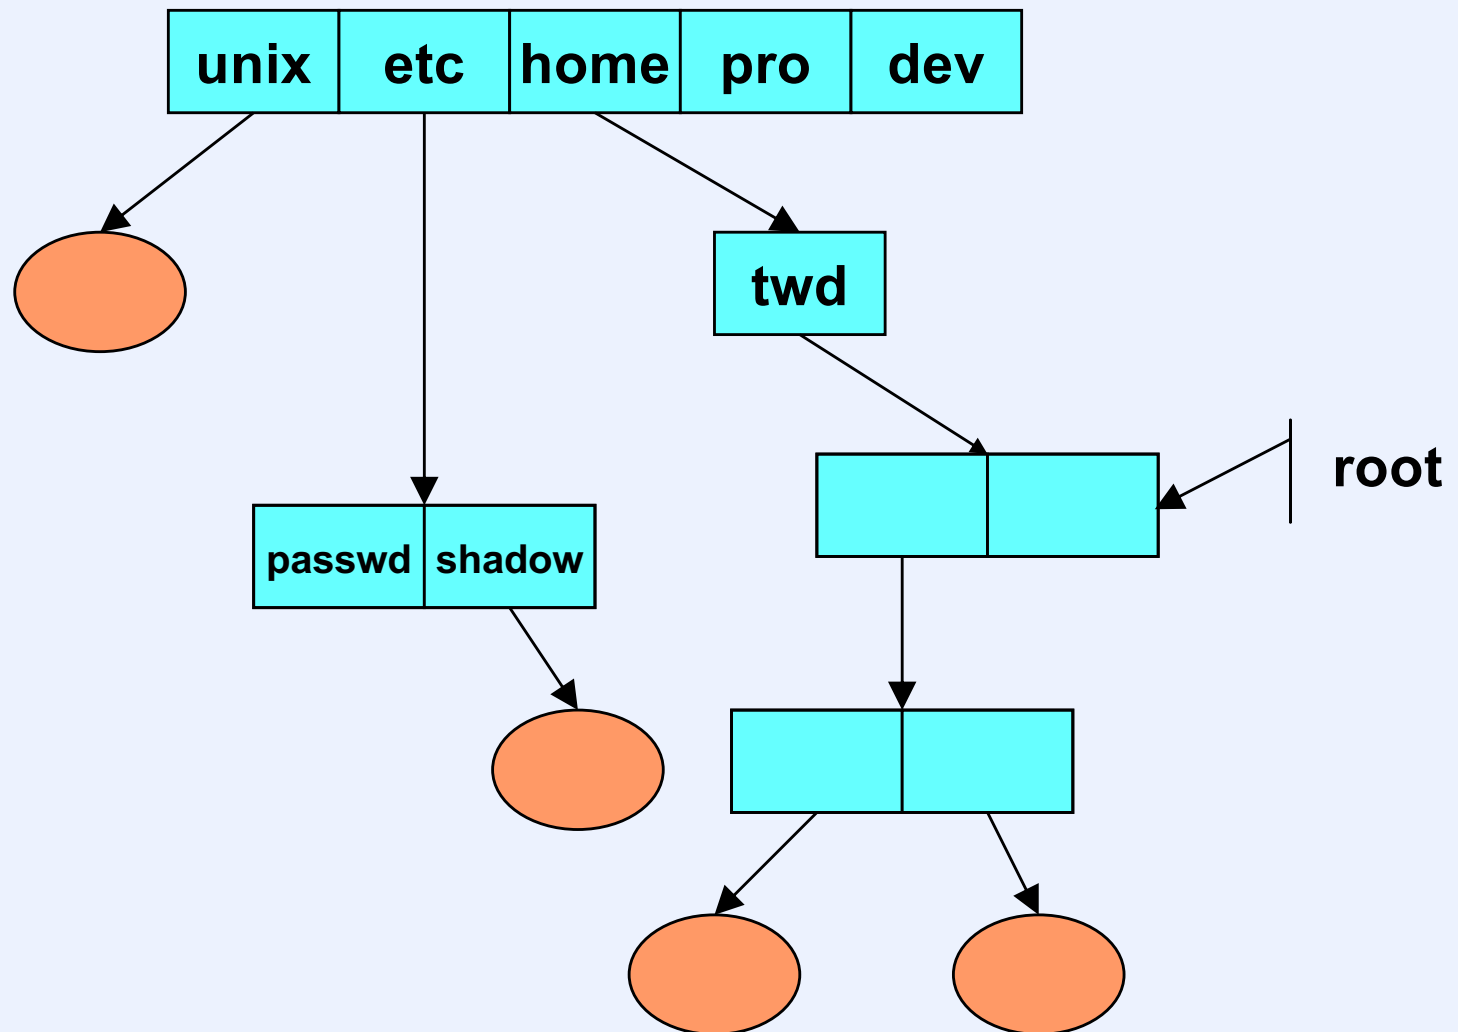
Container Building Blocks

Part 1: File system (chroot)

chroot (before)



chroot (after)




Not a Quiz

Restricting a process to a particular subtree

- a) improves security by effectively running the process in a smaller protection domain**
- b) has little effect on security**
- c) potentially makes security worse**

chroot (after)



Relevant System Calls

- `chroot(path_name)`
- `chdir(path_name)`
- `fchdir(file_descriptor)`

Not a Quiz

After executing *chroot*, “/” refers to the process’s new root directory. Thus “..” is the same as “.” at the process’s root, and the process cannot *cd* directly to the “parent” of its root. Also, recall that hard links may not refer to directories.

- a) *chroot* does effectively limit a process to a subtree
- b) *chroot* does not effectively limit a process to a subtree

Escape!

```
chdir("/");  
pfd = open(".", O_RDONLY);  
mkdir("Houdini", 0700);  
chroot("Houdini");  
fchdir(pfd);  
for (i=0; i<100; i++)  
    chdir("..");  
chroot(".");
```

Namespace Isolation

- **Isolate process by restricting it to a subtree**
 - chroot isn't foolproof
- **Fix chroot**
 - make it superuser only
 - make sure processes don't have file descriptors referring to directories above their roots

Fixed in BSD

- jail
 - can't *cd* above root
 - all necessary files for standard environment present below root
 - *ps* doesn't see processes in other jails



Container Building Blocks

Part 2: Resources & Namespaces

Linux Responds ...

- **cgroups**
 - group together processes for
 - resource limiting
 - prioritization
 - accounting
 - control
- **name-space isolation**
 - isolate processes in different name spaces
 - mount points
 - PIDs
 - UIDs
 - etc.

Linux Containers

- **Reside in isolated subtrees**
 - (fixed) chroot restricts processes in a container to the subtree
 - file systems are mounted in container namespaces, so that other containers can't see them
- **Separate UID and PID spaces**
 - PIDs start at 1 for each container
 - container UIDs mapped to OS UIDs
 - UID 0 has privileges in container, but not outside of container
- **Limits placed on CPU, I/O and other usages**

Docker

- **Runs in Linux containers (also runs on Windows)**
 - **container contains all software and files needed for execution**
 - **provides standard API for applications**
 - **even if on Windows**
- **On macOS, actually runs a hypervisor and runs containers inside (Linux) VM**

Windows Security

Back to Windows

- **Security history**
 - **DOS and early Windows**
 - no concept of logging in
 - no authorization
 - all programs could do everything
 - **later Windows**
 - good authentication
 - good authorization with ACLs
 - default ACLs are important
 - few understand how ACLs work ...
 - many users ran with admin privileges
 - all programs can do everything ...

Privileges in Windows

- **Properties of accounts**
 - administrator \approx superuser
 - finer breakdown for service applications
- **User account control (starting with Vista)**
 - accounts with administrator privileges have two access tokens
 - one for normal usage
 - another with elevated rights

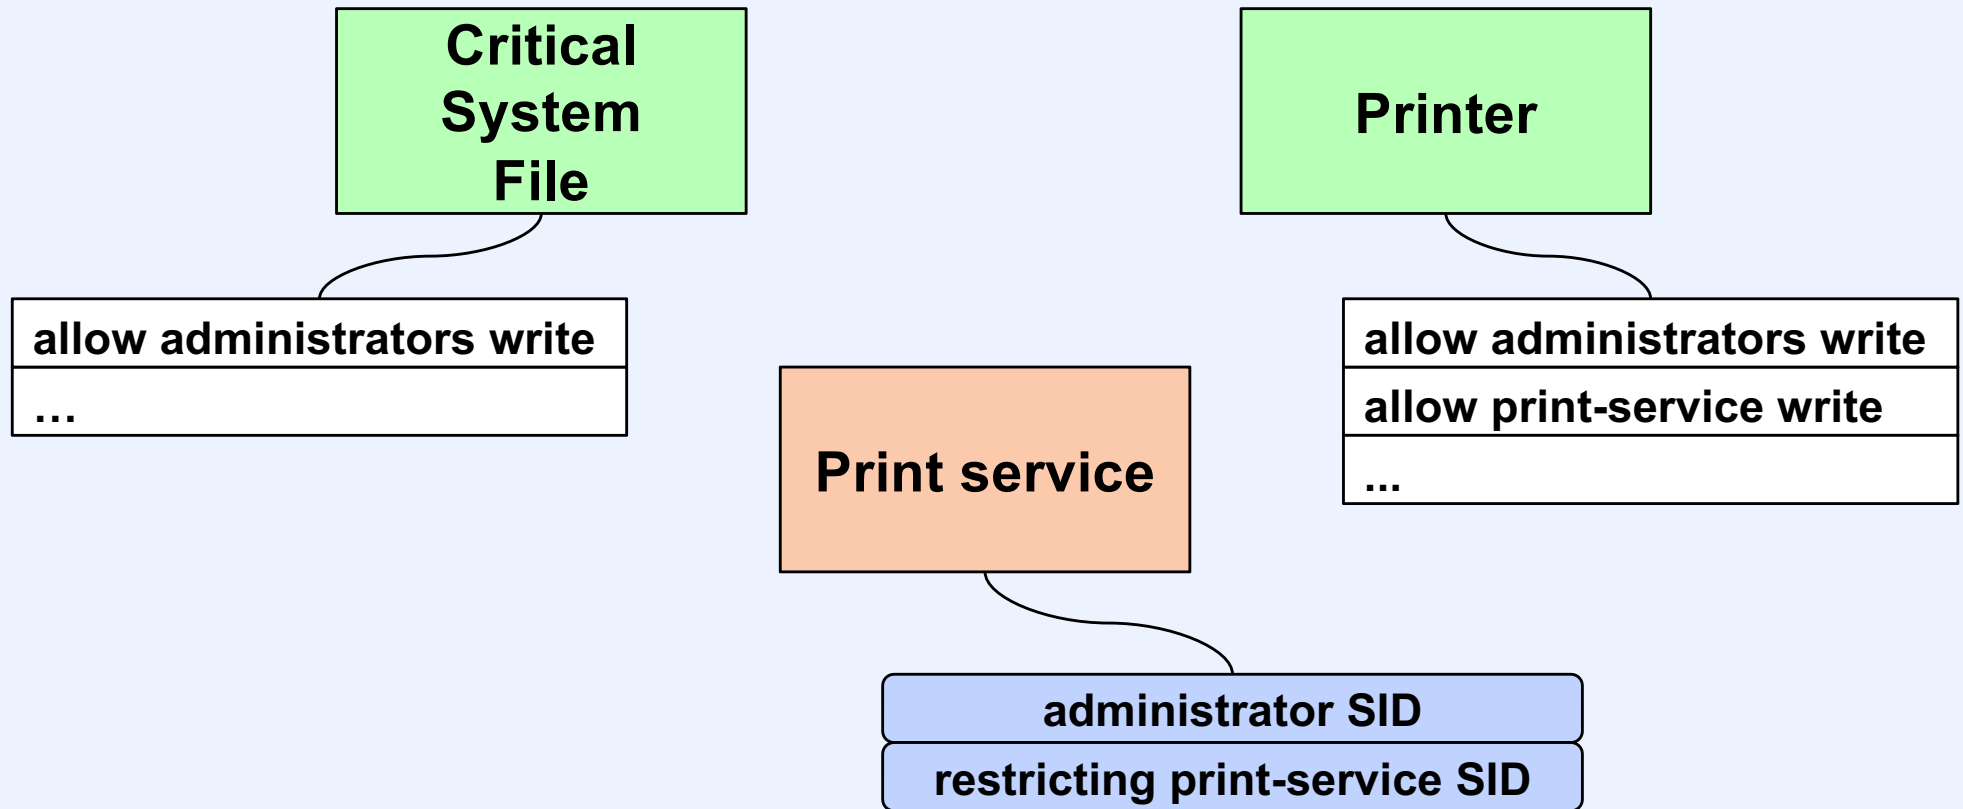
Least Privilege

- **Easy answer**
 - disable privileges
 - works only if the process has any ...
- **Another answer**
 - restricting SIDs
 - limit what a server can do
 - two passes over ACL for access check
 - first: as previously specified
 - second: using only restricting SIDs

Least Privilege for Servers

- **Pre-Vista:**
 - **services ran in local system account**
 - all possible privileges
 - successful attackers “owned” system
 - too complicated to give special account to each service
- **Vista and beyond**
 - **services still run in system account**
 - **per-service SIDs created**
 - used in DACLs to indicate just what service needs
 - marked *restricting* in service token

Example



Not a Quiz

- **Why are there two passes made over the ACL?**
- **Answer: a restricting SID is not an additional access right, but it diminishes what can be done with existing rights**
 - one must first show that one has an access right, then check if it has been diminished

Least Privilege for Clients

- **Pre Vista**
 - no
- **Vista and beyond**
 - windows integrity mechanism
 - a form of MAC

Print Server

- **Client sends request to server**
 - print contents of file X
- **Server acts on request**
 - does client have read permission?
 - server may have (on its own) read access, but client does not
 - server might not have read access, but client does

Unix Solution

- **Client execs print-server, passing it file name**
 - **set-uid-root program**
 - **it (without races!) checks that client has access to file, then prints it**

Windows Solution

- **Server process started when system is booted**
- **Clients send it print requests**
 - **how does client prove to server it has access?**
 - **how does server prove to OS that client has said ok?**

Impersonation

- **Client sends server *impersonation token***
 - subset of its access token
- **Server temporarily uses it in place of its own access token**

Quiz 1


I've written a print server. You would like to use it to print a file. However, you don't trust me — you're concerned that my print server software might read some of your files that you don't want me to read. My print server uses either the Unix approach (setuid-to-twd) or the windows approach (you send it an impersonation token) to deal with access control.

- a) You have nothing to worry about**
- b) You have nothing to worry about if it uses the Unix approach**
- c) You have nothing to worry about if it uses the Windows approach**
- d) You have a lot to worry about with both**

Security Models

Serious Security


- National defense
- Proprietary information
- Personal privacy



Mandatory vs. Discretionary Access Control

- **Discretionary**
 - ACLs, capabilities, etc.
 - access is at the discretion of the owner
- **Mandatory**
 - government/corporate security, etc.
 - access is governed by strict policies

Mandatory Access Control (1)



Mandatory Access Control (2)

- **Privacy/confidentiality policies**
 - compartmentalization

**student
records**

registrar

**faculty
salaries**

**dean of the
faculty**

**medical
records**

**University-
affiliated
hospitals**

Mandatory Access Control (3)

- **Local computer policy**
 - **web-server**
 - **may access only designated web-server data**
 - **administrators**
 - **may execute only administrative programs**
 - **(may not execute code supplied by ordinary users)**

Bell-LaPadula Model

1) Simple security property

no-read-up


- no subject may read from an object whose classification is higher than the subject's clearance

2) *-property

no-write-down

- no subject may write to an object whose classification is lower than the subject's clearance

Information Black Hole



Attack!




**Not
cleared
for top-
secret
orders**


Managing Confidentiality

- **Black-hole avoidance**
 - trusted vs. untrusted subjects
 - trusted subjects may write down

Espionage



**agent X learns of
invasion plans**




**communication
not possible**



**agent Y can send
email to spymaster
(but doesn't know
what to send)**

Covert Channels



**agent X runs
resource-intensive
program**



**sneaky
communication
possible**

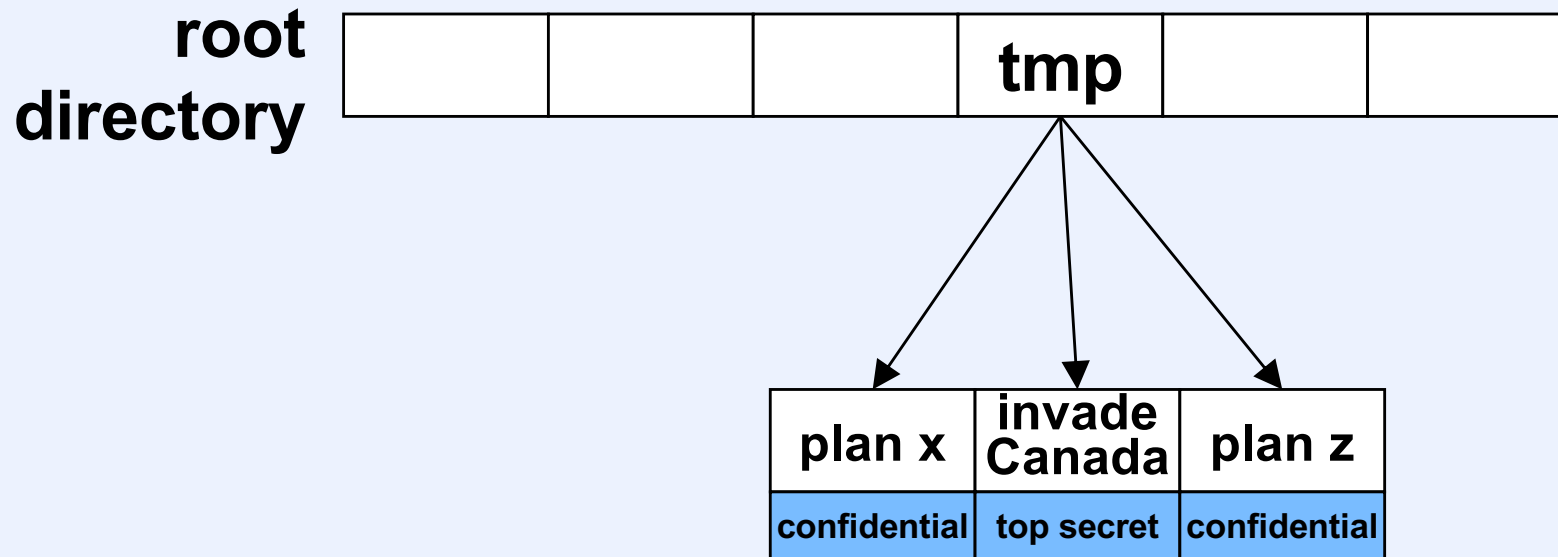


**agent Y monitors load
sends email to
spymaster**

Defense

- **Identify all covert channels**
 - (good luck ...)
- **Eliminate them**
 - find a suitable scheduler
 - eliminates just one channel

Multi-Level Directories (1)



Multi-Level Directories (2)

