

Interrupts, Etc.

Disk I/O

```
int disk_write(...) {  
    ...  
    startIO(); // start disk operation  
    ...  
    enqueue(disk_waitq, CurrentThread);  
    thread_switch();  
    // wait for disk operation to  
    // complete  
    ...  
}
```

```
void disk_intr(...) {  
    thread_t *thread;  
    ...  
    // handle disk interrupt  
    ...  
    thread = dequeue(disk_waitq);  
    if (thread != 0) {  
        enqueue(RunQueue, thread);  
        // wakeup waiting thread  
    }  
    ...  
}
```

Improved Disk I/O

```
int disk_write(...) {  
    ...  
    oldIPL = setIPL(diskIPL);  
    startIO();          // start disk operation  
    ...  
    enqueue(disk_waitq, CurrentThread);  
    thread_switch();  
    // wait for disk operation to complete  
    setIPL(oldIPL);  
    ...  
}
```

Modified *thread_switch*

```
void thread_switch() {
    thread_t *OldThread;
    int oldIPL;
    oldIPL = setIPL(HIGH_IPL);
    // protect access to RunQueue by masking all interrupts
    while(queue_empty(RunQueue)) {
        // repeatedly allow interrupts, then check RunQueue
        setIPL(0); // IPL == 0 means no interrupts are masked
        setIPL(HIGH_IPL);
    }
    // We found a runnable thread
    OldThread = CurrentThread;
    CurrentThread = dequeue(RunQueue);
    swapcontext(OldThread->context, CurrentThread->context);
    setIPL(oldIPL);
}
```

Preemptive Kernels on MP

- What's different?
- A thread accesses a shared data structure:
 1. it might be *interrupted* by an interrupt handler (running on its processor) that accesses the same data structure
 2. *another thread* running on another processor might access the same data structure
 3. it might be forced to *give up its processor* to another thread, either because its time slice has expired or it has been preempted by a higher-priority thread
 4. an *interrupt handler* running on *another processor* might access the same data structure

Solution?

```
int X = 0;
SpinLock_t L = UNLOCKED;
```

```
void AccessXThread() {
    SpinLock(&L);
    X = X+1;
    SpinUnlock(&L);
}
```

```
void AccessXInterrupt() {
    ...
    SpinLock(&L);
    X = X+1;
    SpinUnlock(&L);
    ...
}
```

Solution ...

```
int X = 0;  
SpinLock_t L = UNLOCKED;
```

```
void AccessXThread() {  
    MaskInterrupts();  
    SpinLock(&L);  
    X = X+1;  
    SpinUnlock(&L);  
    UnMaskInterrupts();  
}
```

```
void AccessXInterrupt() {  
    ...  
    SpinLock(&L);  
    X = X+1;  
    SpinUnlock(&L);  
    ...  
}
```

Quiz 1

We have a **single-core** system with a preemptible kernel. We're concerned about data structure *X*, which is accessed by kernel threads as well as by the interrupt handler for dev.

- a) It's sufficient for threads to mask dev interrupts while accessing *X*
- b) In addition, threads must lock (blocking) mutexes before masking interrupts and accessing *X*
- c) b doesn't work. Instead, threads must lock spinlocks before accessing *X*
- d) In addition to c, the dev interrupt handler must lock a spinlock before accessing *X*
- e) Something else is needed

Deferred Work

- **Interrupt handlers run with interrupts masked**
 - may interfere with handling of other interrupts, particularly if they do a lot of computation
- **Solution**
 - do minimal work now
 - do rest later without interrupts masked

Deferred Processing

```
void TopLevelInterruptHandler(int dev) {  
    InterruptVector[dev](); // call appropriate handler  
    if (PreviousContext == ThreadContext) {  
        UnMaskInterrupts();  
        while (!Empty(WorkQueue)) {  
            Work = DeQueue(WorkQueue);  
            Work();  
        }  
    }  
}  
  
void NetworkInterruptHandler() {  
    // deal with interrupt  
    ...  
    EnQueue(WorkQueue, MoreWork);  
}
```

Windows Interrupt Priority Levels

hardware	31	High
	30	Power fail
	29	Inter-processor
	28	Clock
	.	
software	4	Device 2
	3	Device 1
	2	DPC
	1	APC
	0	Thread

Deferred Procedure Calls

```
void InterruptHandler( ) {  
    // deal with interrupt  
    ...  
    QueueDPC(MoreWork, arg);  
    /* enqueues MoreWork on  
       the DPC queue and  
       requests a DPC  
       interrupt  
    */  
}
```

```
void DPCHandler( ... ) {  
    while (!Empty(DPCQueue)) {  
        Work = DeQueue(DPCQueue);  
        Work();  
    }  
}
```

Software Interrupt Threads

```
void InterruptHandler() {  
    // deal with interrupt  
    ...  
    EnQueue (WorkQueue,  
             MoreWork);  
    SetEvent (Work);  
}
```

```
void SoftwareInterruptThread() {  
    while (TRUE) {  
        WaitEvent (Work)  
        while (!Empty (WorkQueue)) {  
            Work = DeQueue (  
                WorkQueue);  
            Work ();  
        }  
    }  
}
```

Preemption: User-Level Only

```
void ClockHandler() {  
    // deal with clock interrupt  
    ...  
    if (TimeSliceOver())  
        ShouldReschedule = 1;  
}
```

```
void TopLevelInterruptHandler(int dev) {  
    InterruptVector[dev]();  
    if (PreviousMode == UserMode) {  
        // the clock interrupted user-mode code  
        if (ShouldReschedule)  
            Reschedule();  
    }  
    ...  
}
```

```
void TopLevelTrapHandler(...) {  
    SpecificTrapHandler();  
    ...  
    if (ShouldReschedule) {  
        /* the time slice expired while the thread  
           was in kernel mode */  
        Reschedule();  
    }  
}
```

Preemption: Full

```
void ClockInterruptHandler( ) {  
    // deal with clock interrupt  
  
    ...  
    if (TimeSliceOver)  
        QueueDPC (Reschedule) ;  
}
```

Directed Processing

- **Signals: Unix**
 - perform given action in context of a particular thread in user mode
- **APC: Windows asynchronous procedure calls**
 - roughly same thing, but also may be done in kernel mode

Asynchronous Procedure Calls

- **Two uses**
 - **kernel APC: release of kernel resources**
 - **user APC: notifying a thread of an external event**

Kernel APC

- **Release of kernel resources**
 - interrupt handler has information that must be copied to user process
 - can't be done unless in context of process
 - otherwise address space not mapped in
 - interrupt handler requests kernel APC to have user thread, running in kernel mode, copy information to user space, and then free data in the kernel

User APC

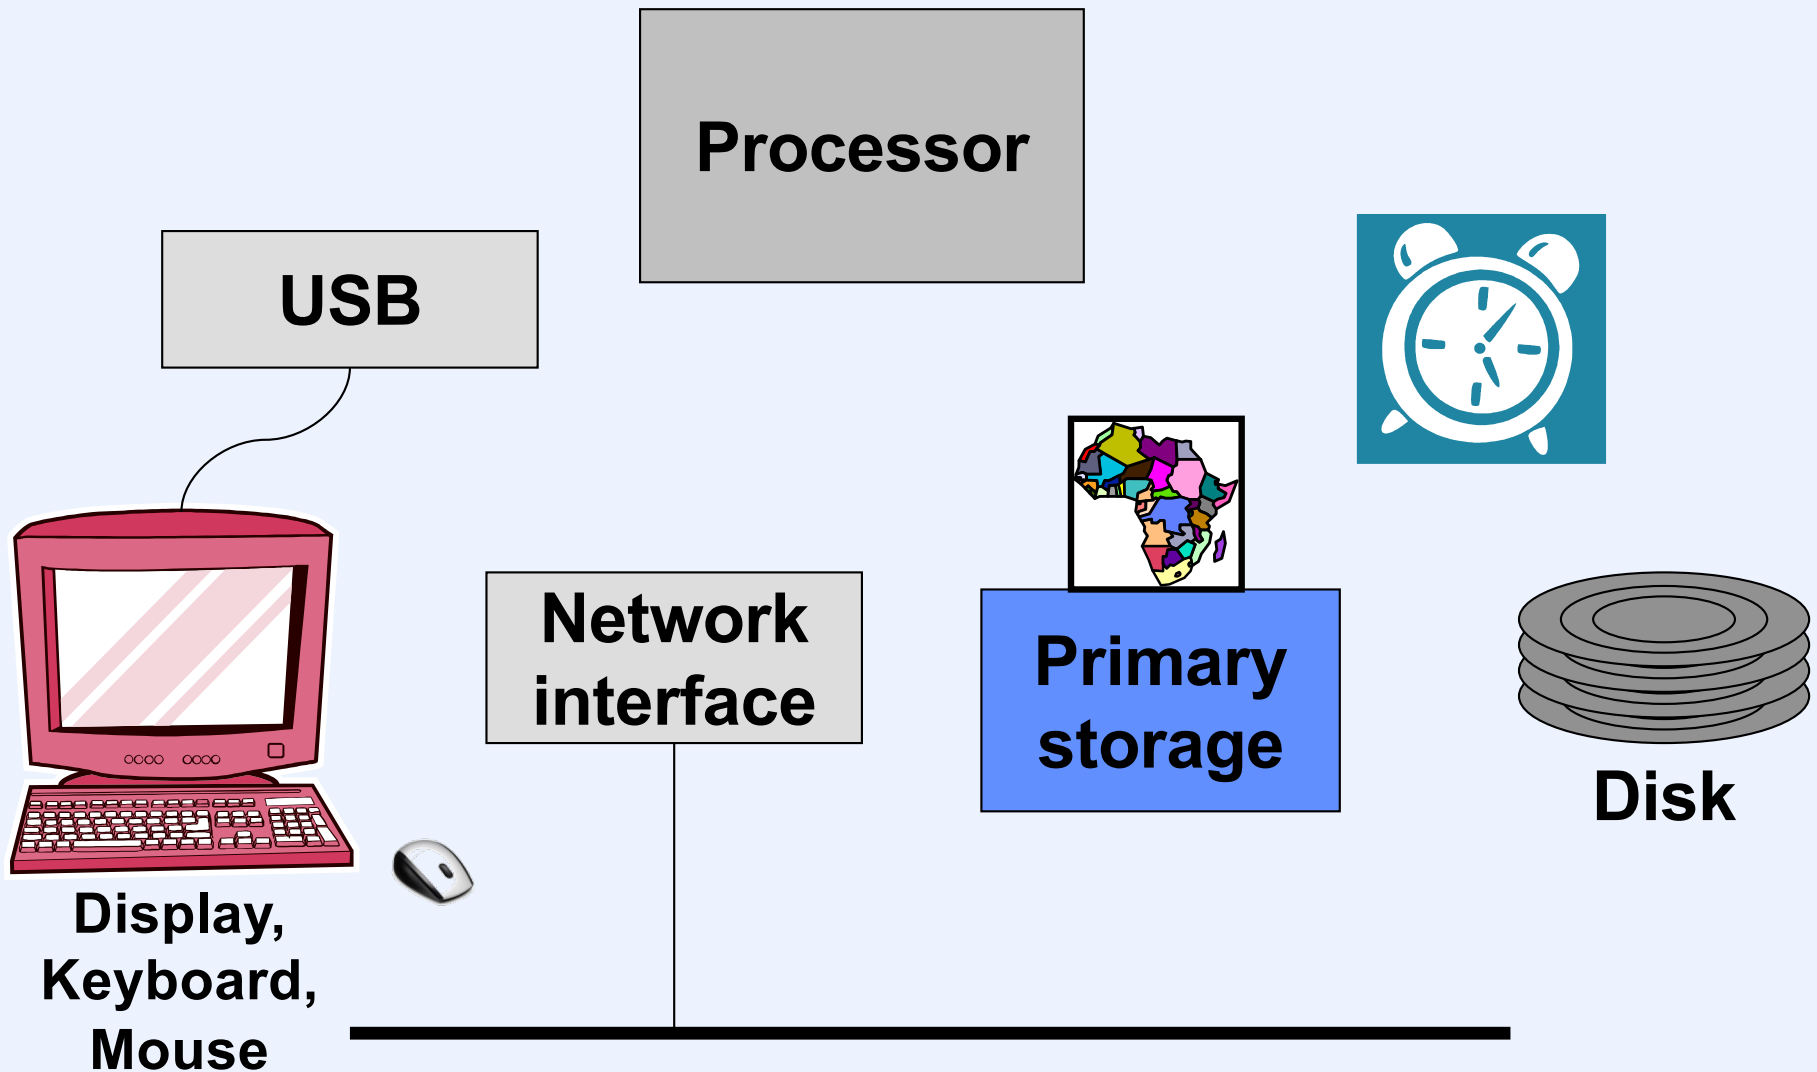
- **Notifying thread of external event**
 - **example: asynchronous I/O**
 - thread supplies *completion routine* when starting asynchronous I/O request
 - called in thread's context when I/O completes
 - similar to a Unix signal
 - called only when thread is in *alertable wait state*
 - an option in certain blocking system calls

APC Implementation

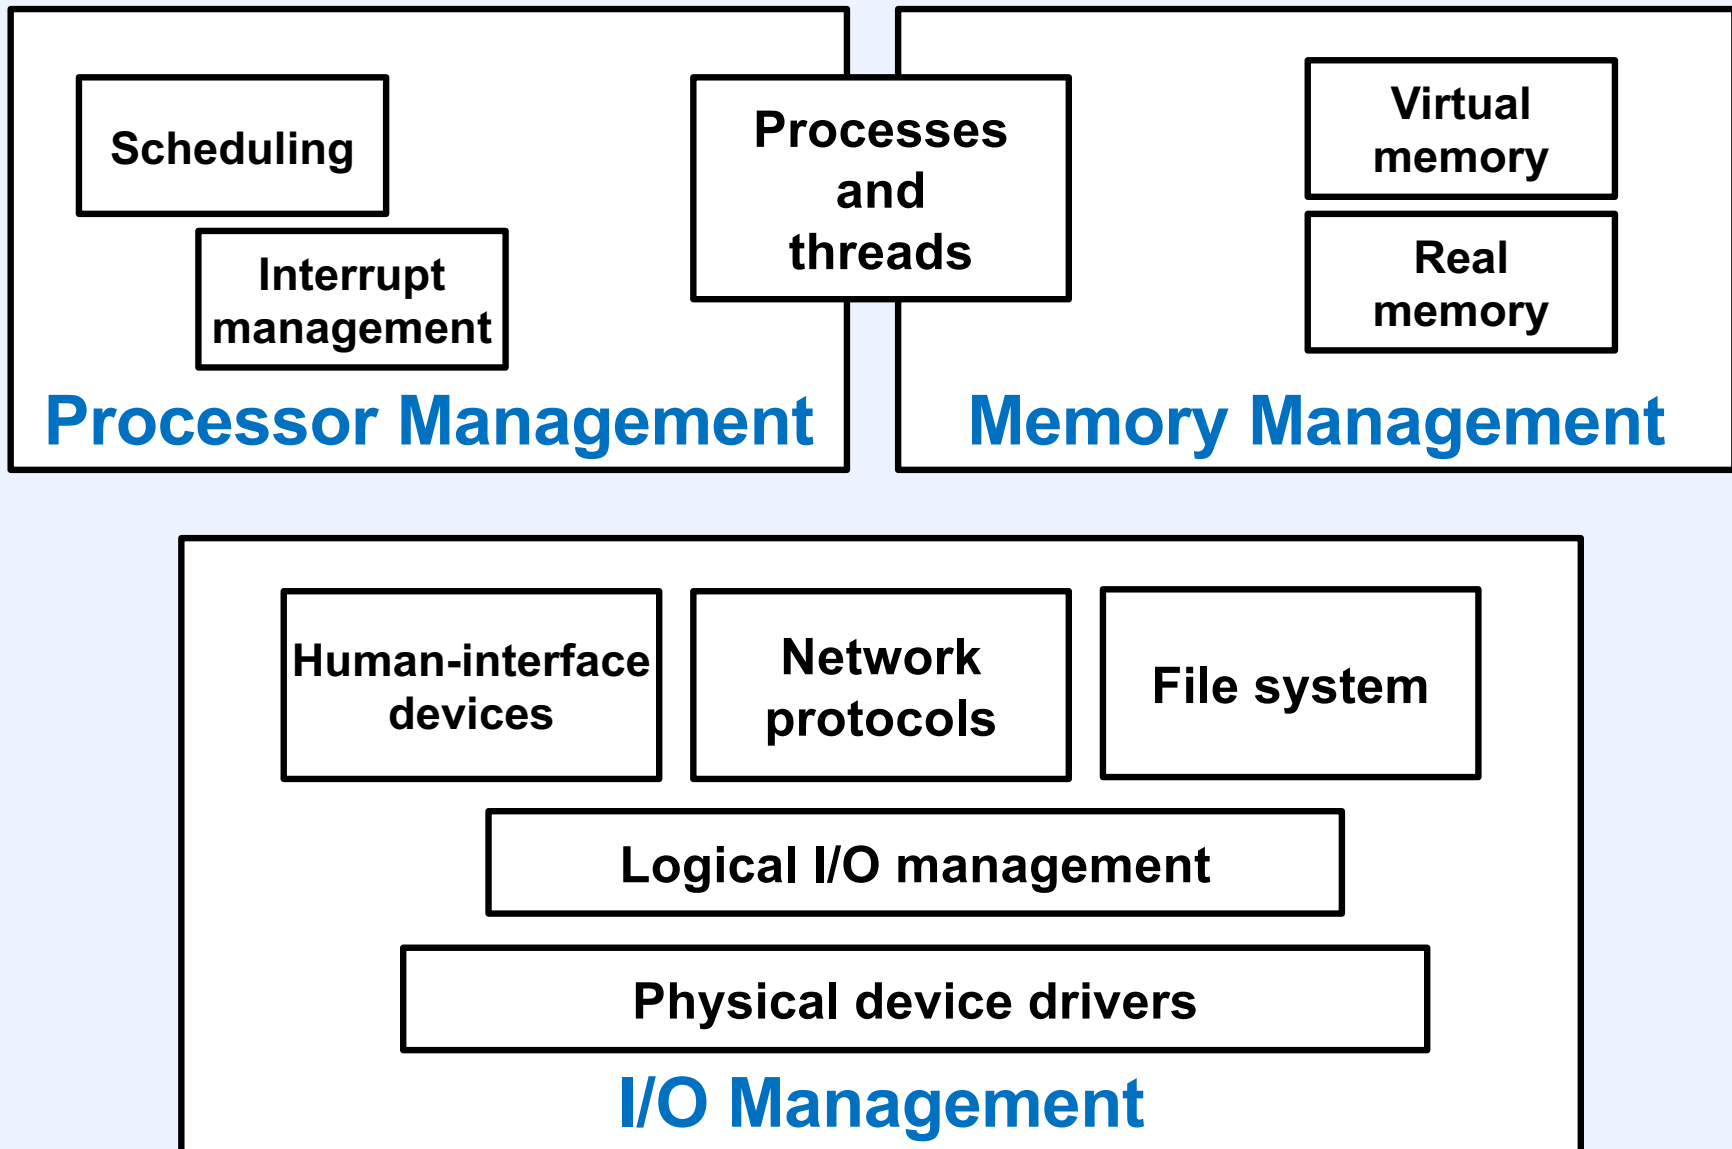
- **Per-thread list of pending APCs**
 - on notification, thread executes them
- **User APC**
 - thread in alertable state is woken up and executes pending APCs when it returns to user mode
- **Kernel APC**
 - running thread interrupted by APC interrupt (lowest-priority interrupt)
 - waiting thread is “unwaited”
 - execute pending kernel APCs

I/O


Simple Configuration




OS Components: Functional



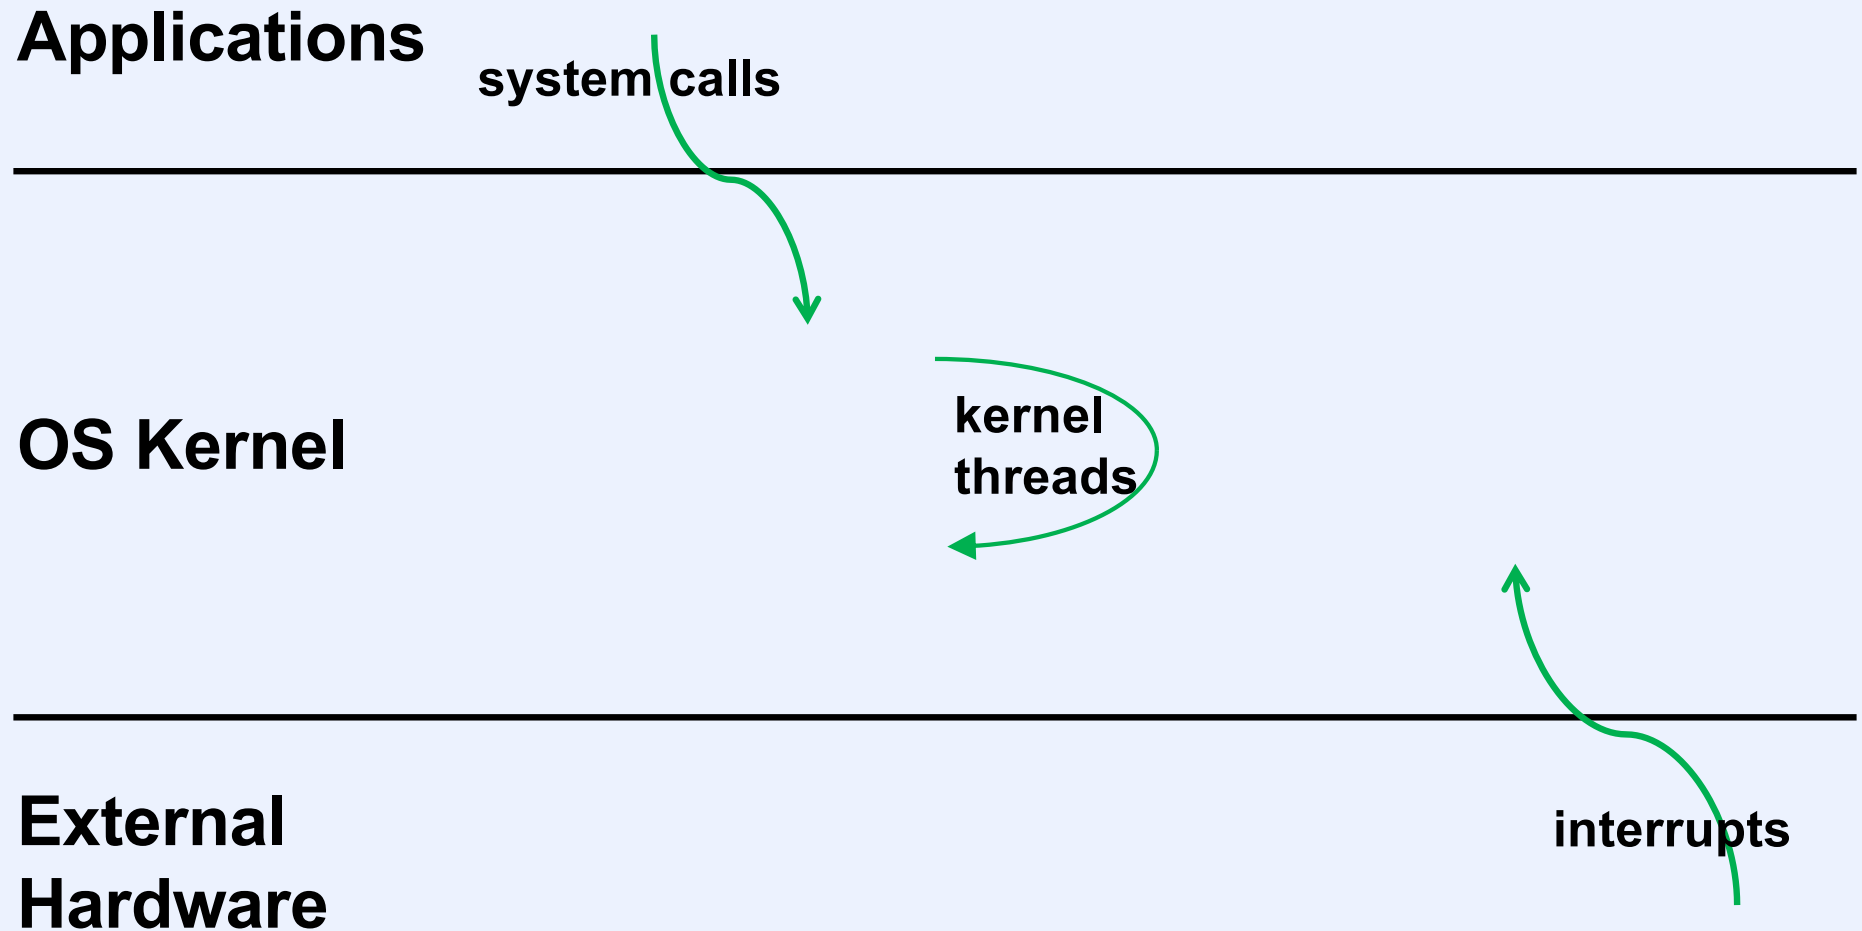
OS Components: Portability



OS Components: Importance



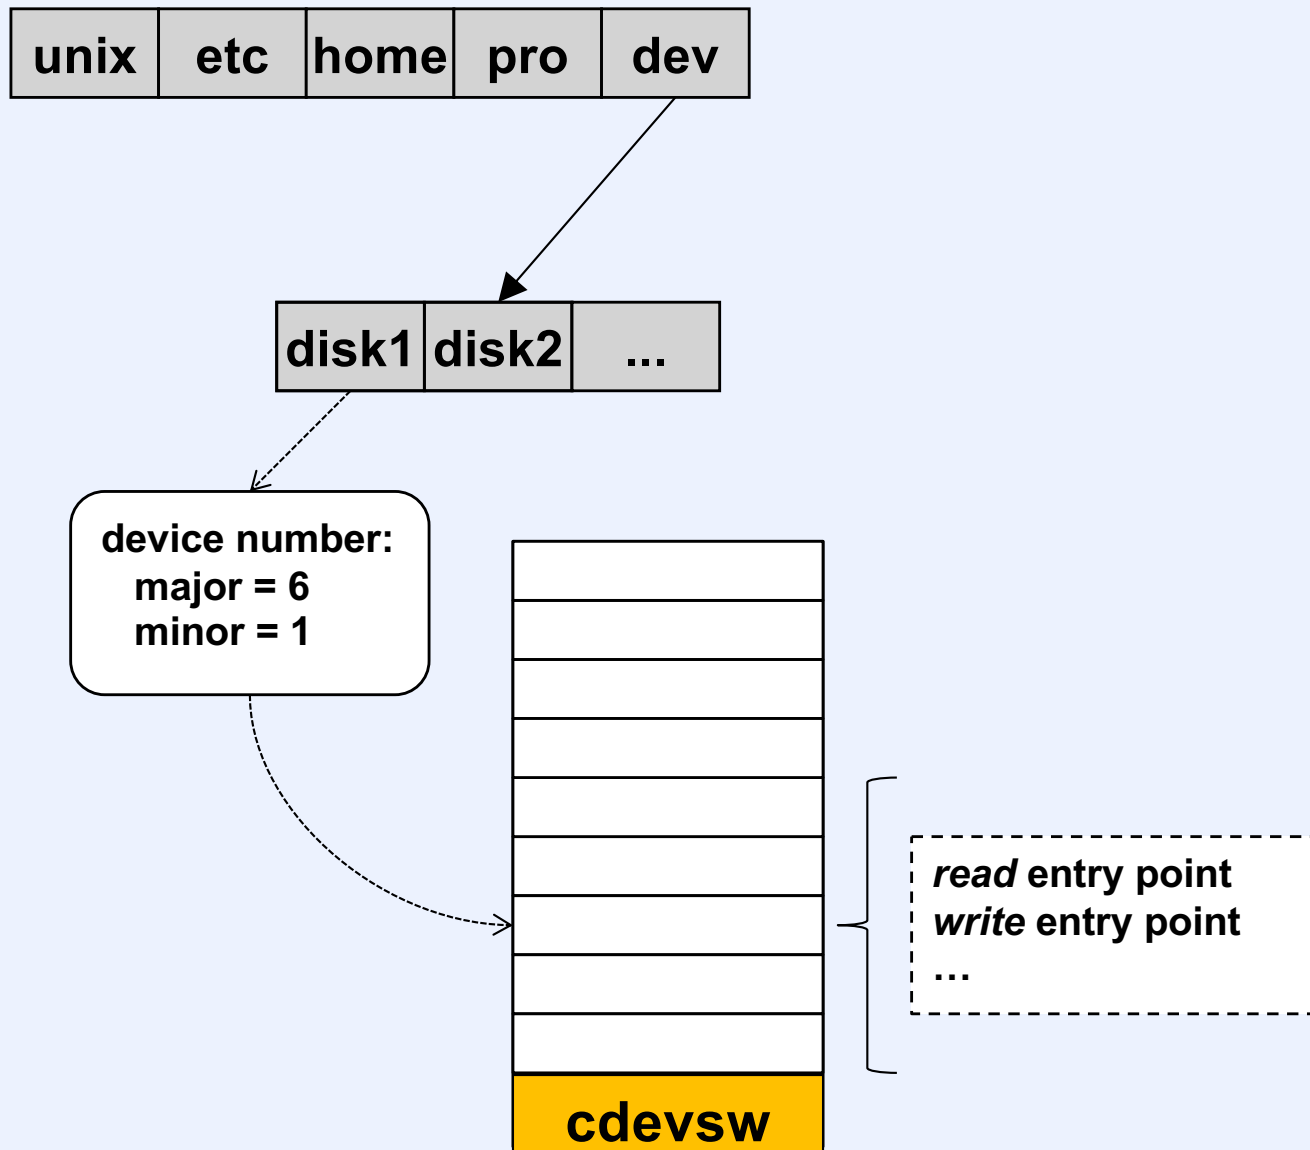
OS Components: Flow of Control



To Be Discussed

- **What is the functionality of the components?**
- **What are the key data structures?**
- **How is the system broken up into modules?**
- **To what extent is the system extensible?**
- **What parts run in the OS kernel in privileged mode? What parts run as library code in user applications? What parts run as separate applications?**
- **In which execution contexts do the various activities take place?**

Finding Devices



Discovering Devices

- **You plug in a new device to your computer ...**
 - **OS must notice**
 - **must find a device driver**
 - **what kind of device is it?**
 - **where is the driver?**
 - **must assign a name**
 - **how chosen?**
 - **multiple similar devices**
 - **how does application choose?**

Computer Terminal



A “tty”



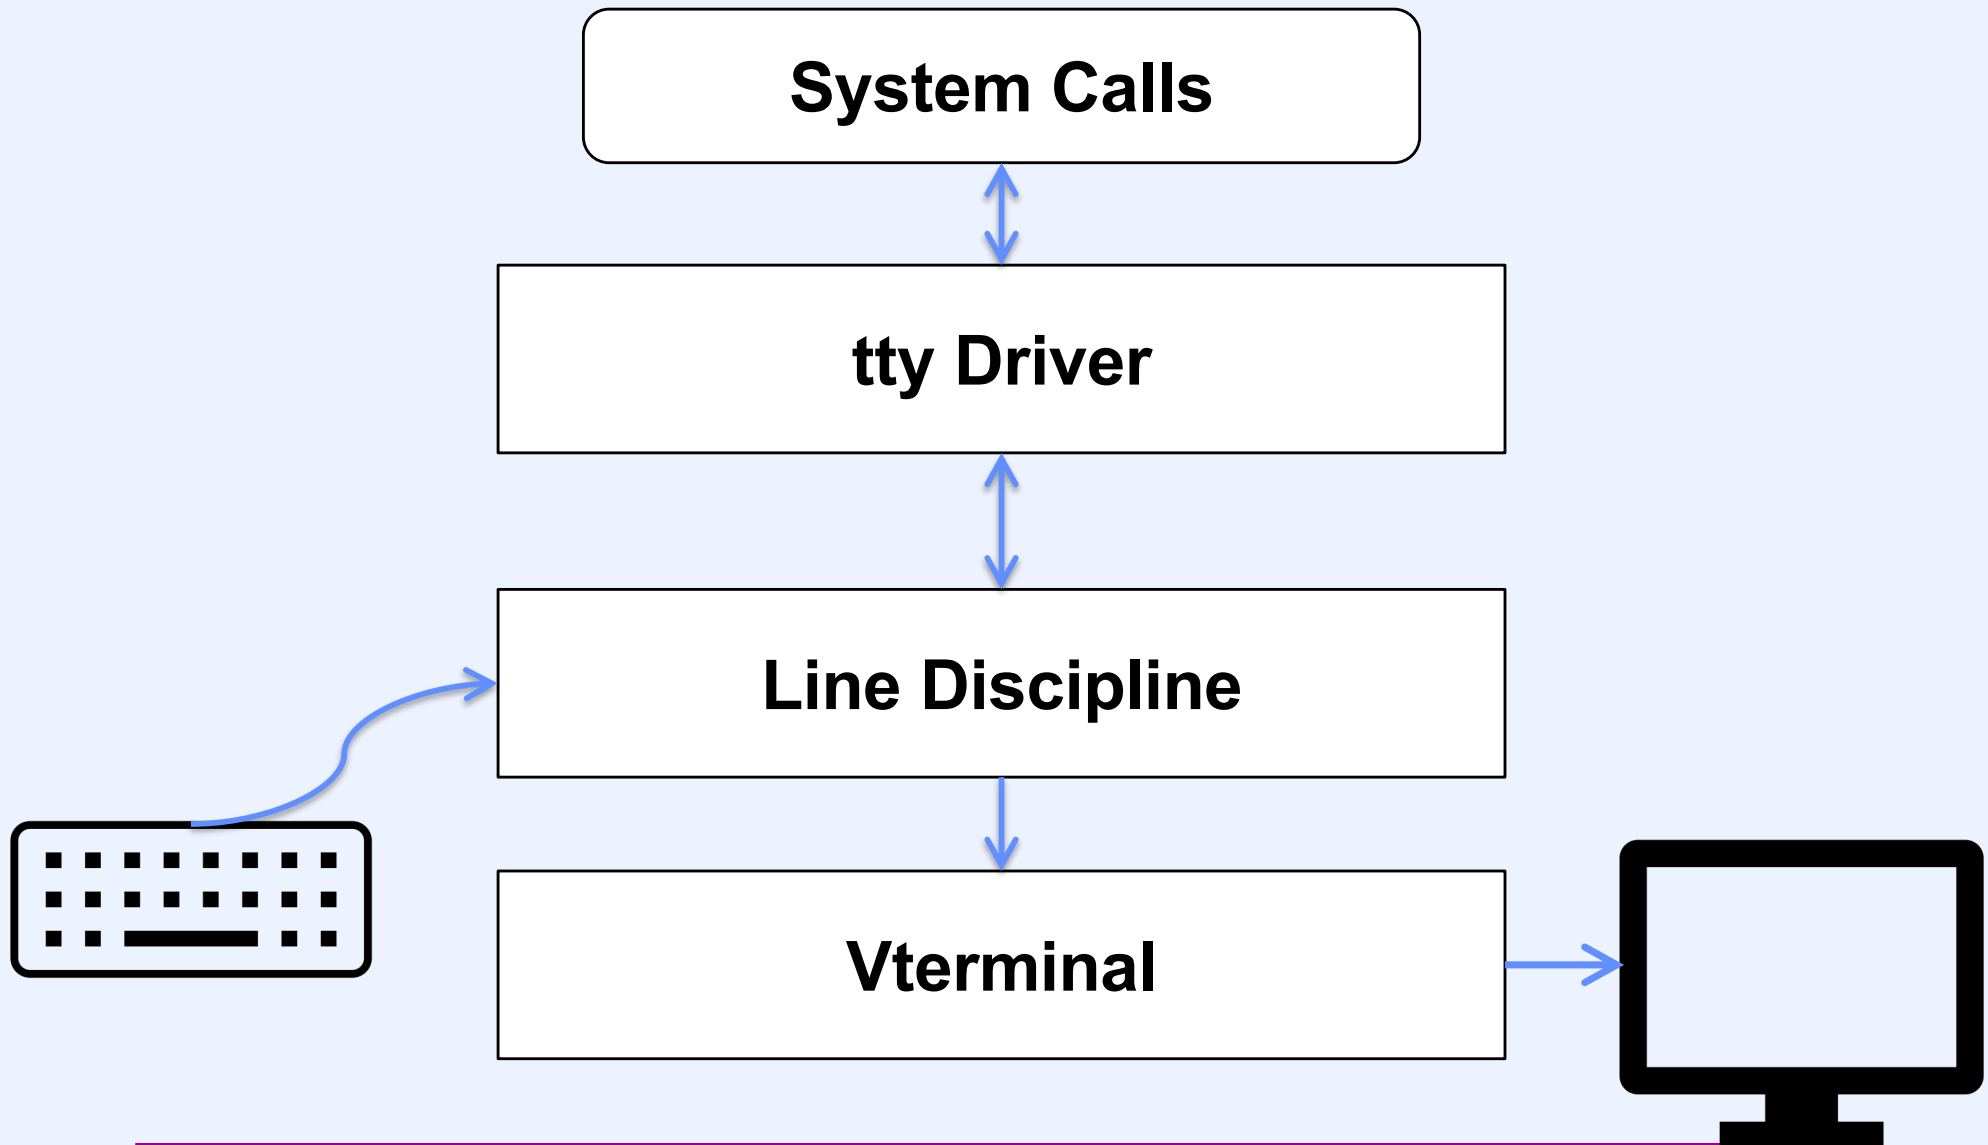
A Typewriter



Terminals

- **Long obsolete, but still relevant**
- **Issues**
 - 1) **characters are generated by the application faster than they can be sent to the terminal**
 - 2) **characters arrive from the keyboard even though there isn't a waiting read request from an application**
 - 3) **input characters may need to be processed in some way before they reach the application**

Terminals



Quiz 2


In which context are characters transformed from raw into cooked?

- a) In the interrupt context (i.e., on a “borrowed” stack)**
- b) In the context of the thread performing the *read* system call**
- c) Some other context**

Input/Output

- **Architectural concerns**
 - memory-mapped I/O
 - programmed I/O (PIO)
 - direct memory access (DMA)
 - I/O processors (channels)
- **Software concerns**
 - device drivers
 - concurrency of I/O and computation

Simple I/O Architecture



PIO Registers

GoR	GoW	IER	IEW				
-----	-----	-----	-----	--	--	--	--

Control register

RdyR	RdyW						
------	------	--	--	--	--	--	--

Status register

--	--	--	--	--	--	--	--

Read register

--	--	--	--	--	--	--	--


Write register

Legend:	GoR	Go read (start a read operation)
	GoW	Go write (start a write operation)
	IER	Enable read-completion interrupts
	IEW	Enable write-completion interrupts
	RdyR	Ready to read
	RdyW	Ready to write

Programmed I/O

- E.g.: Terminal controller
- Procedure (write)
 - write a byte into the *write register*
 - set the WGO bit in the *control register*
 - wait for WREADY bit (in *status register*) to be set (if interrupts have been enabled, an interrupt occurs when this happens)

DMA Registers



Legend:	Go	Start an operation
	Op Code	Operation code (identifies the operation)
	IE	Enable interrupts
	Rdy	Controller is ready

Direct Memory Access

- E.g.: Disk controller
- Procedure
 - set the *disk address* in the *device address register* (only relevant for a seek request)
 - set the *buffer address* in the *memory address register*
 - set the *op code* (SEEK, READ or WRITE), the GO bit and, if desired, the interrupt ENABLE bit in the *control register*
 - wait for interrupt or for READY bit to be set

Device Drivers

