


Unix Structure (2)

Representing the Address Space

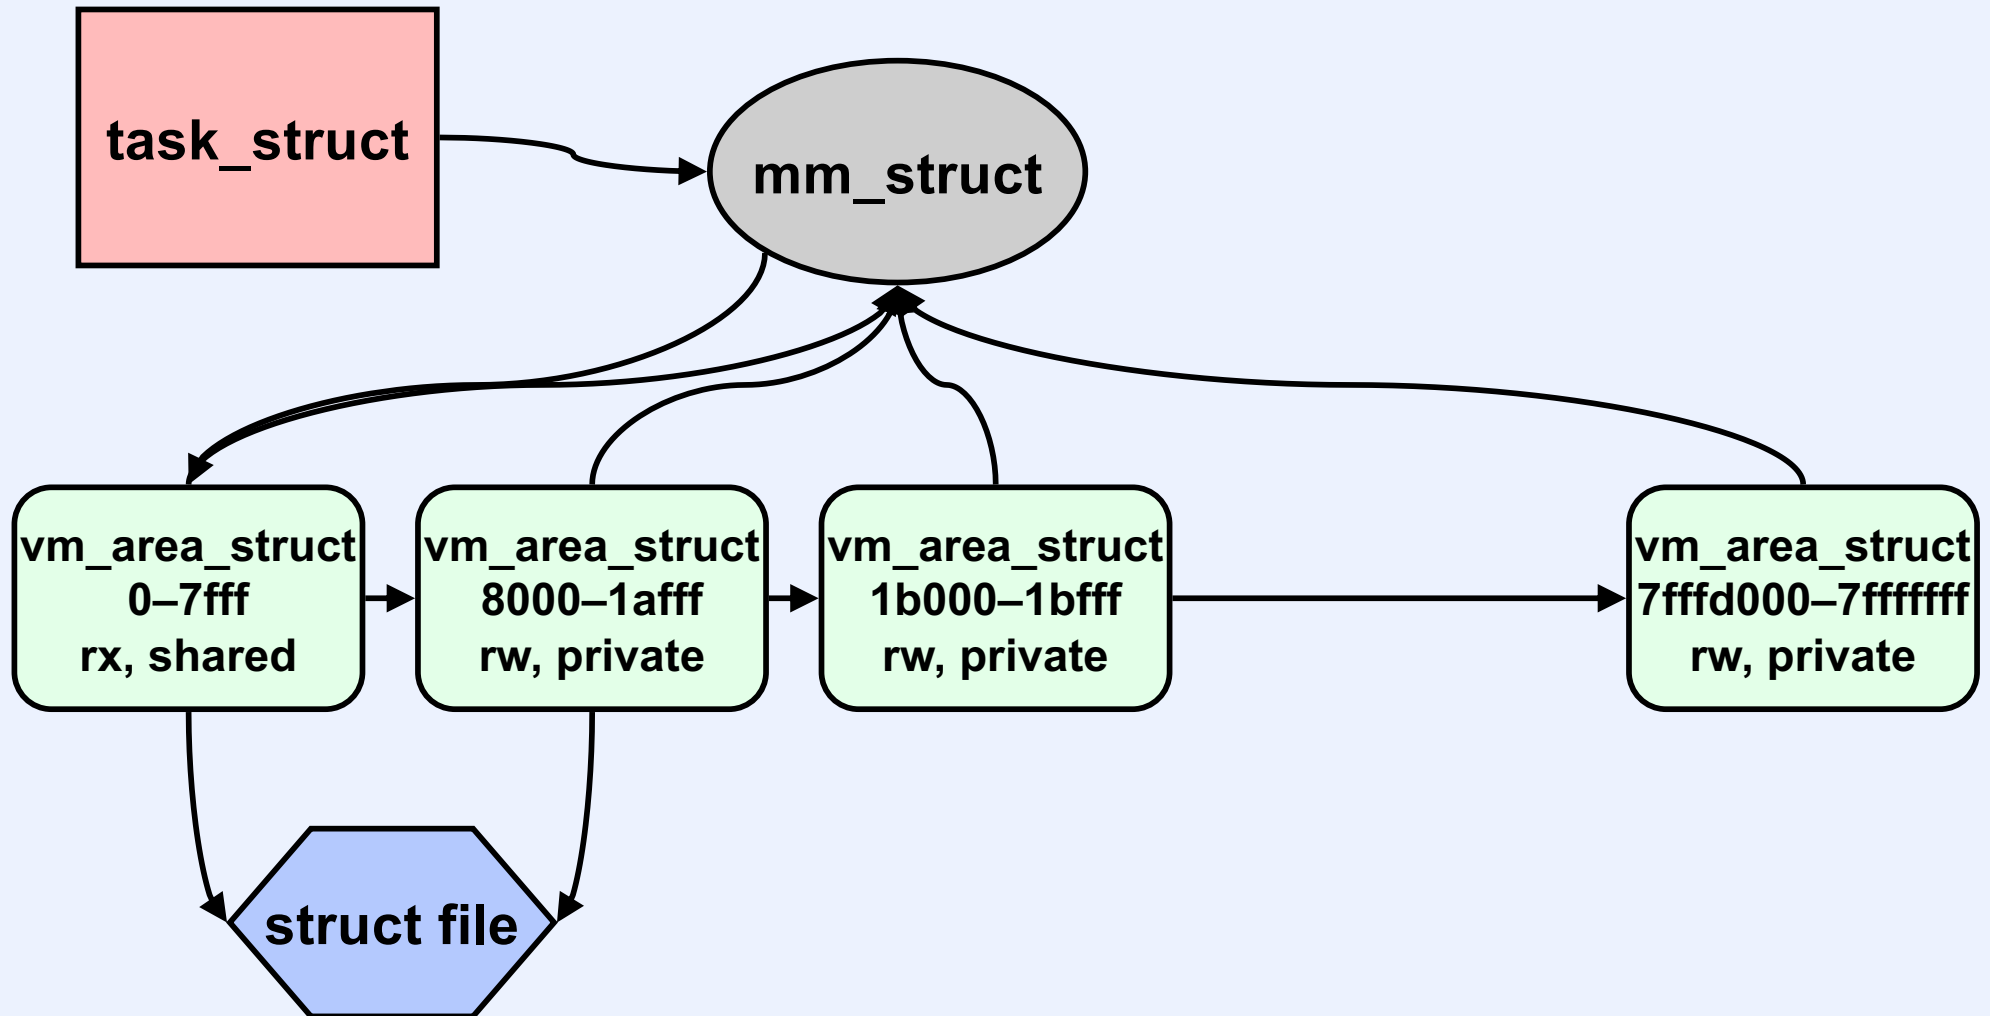
- Important component of a process is its address space
 - how is it represented?
- Can page tables represent a process's address space?

Simple User Address Space




Address-Space Representation

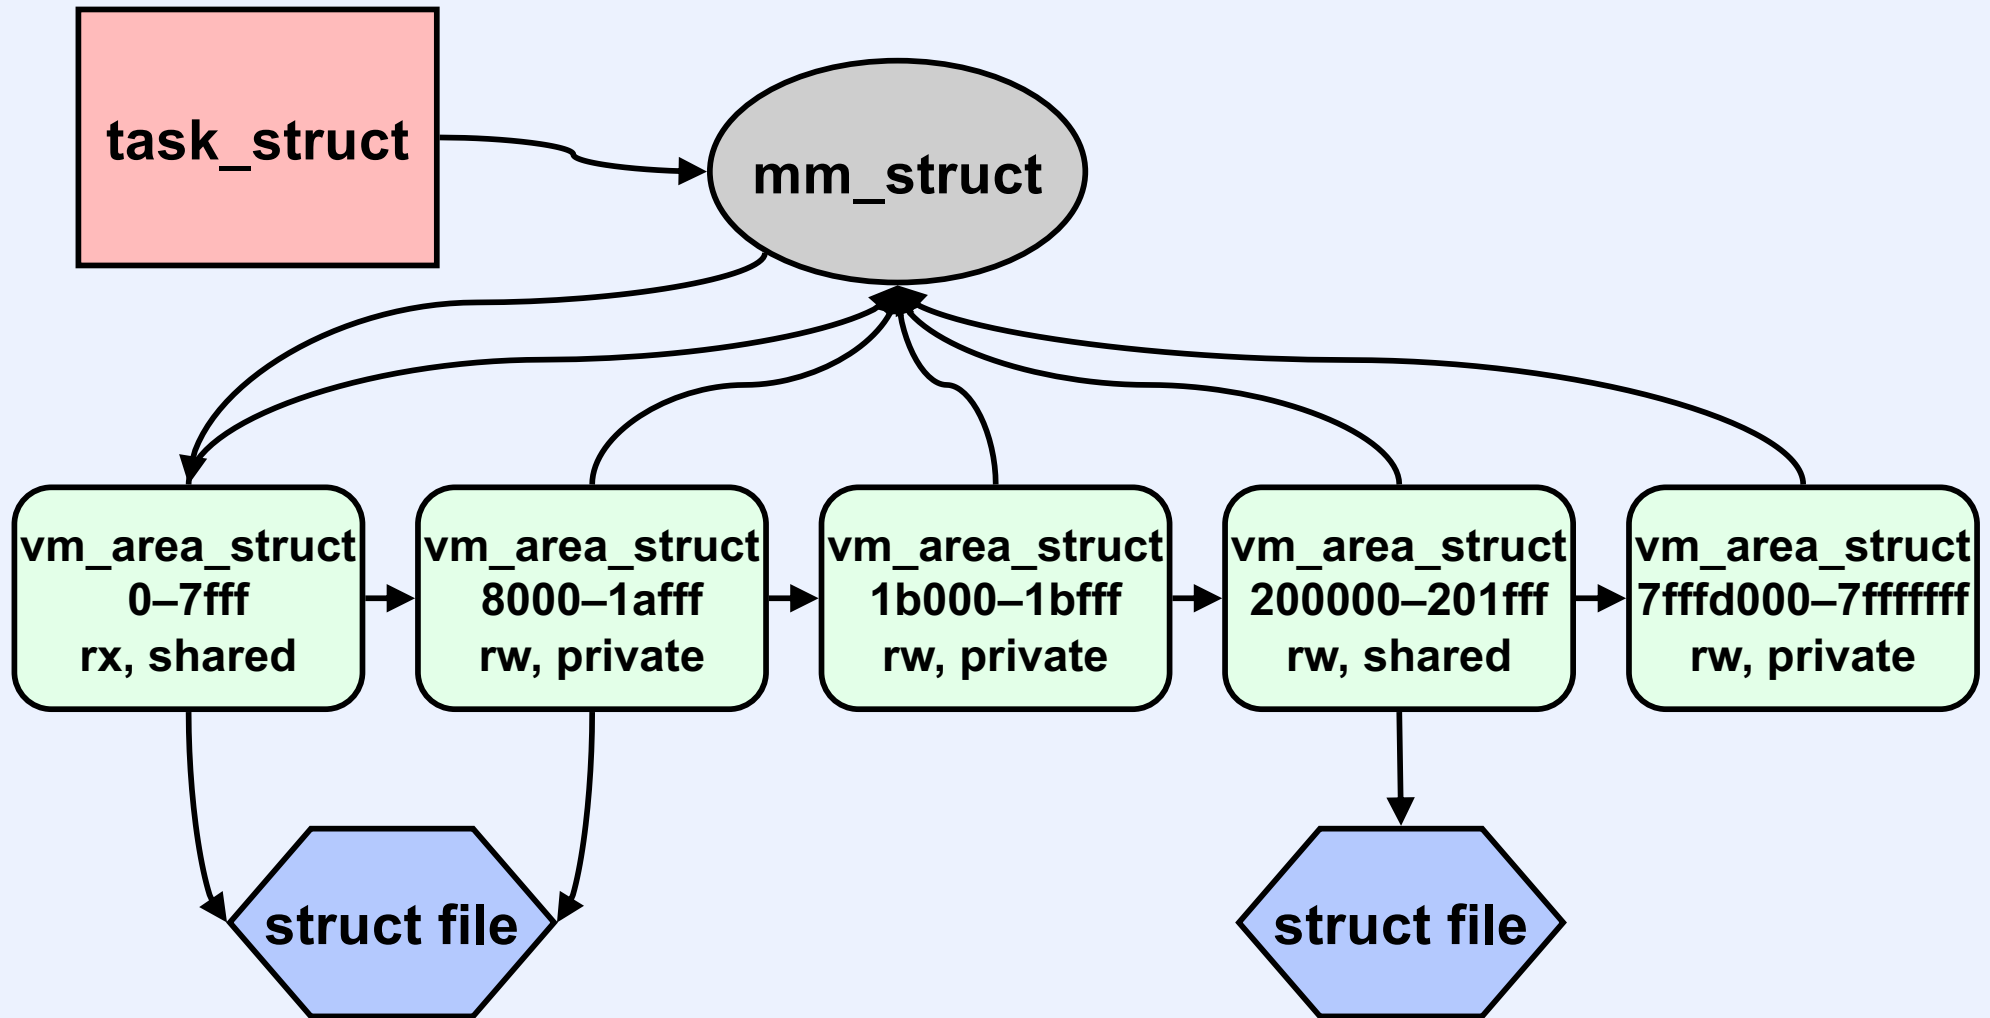
Somewhat Simplified



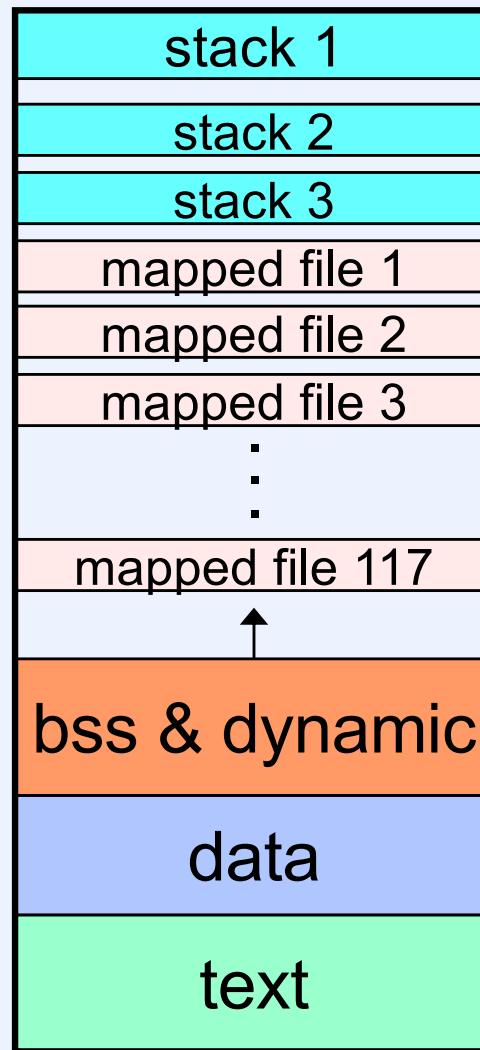
Adding a Mapped File



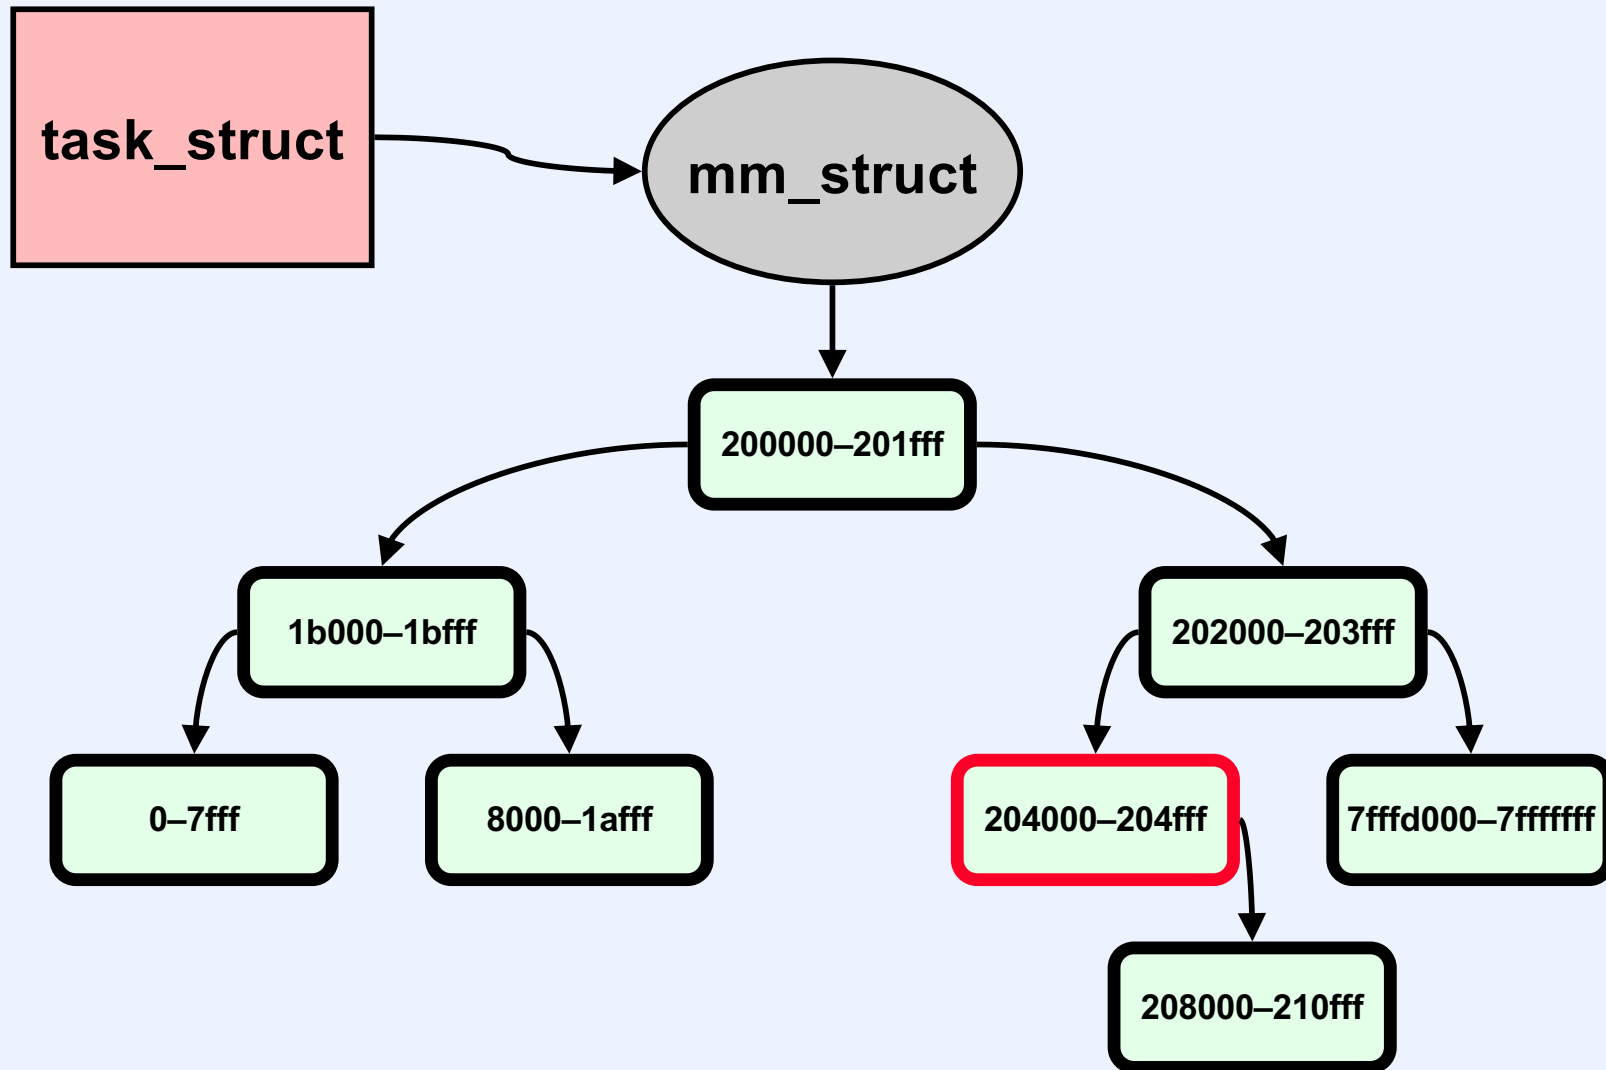
Address-Space Representation: More Areas



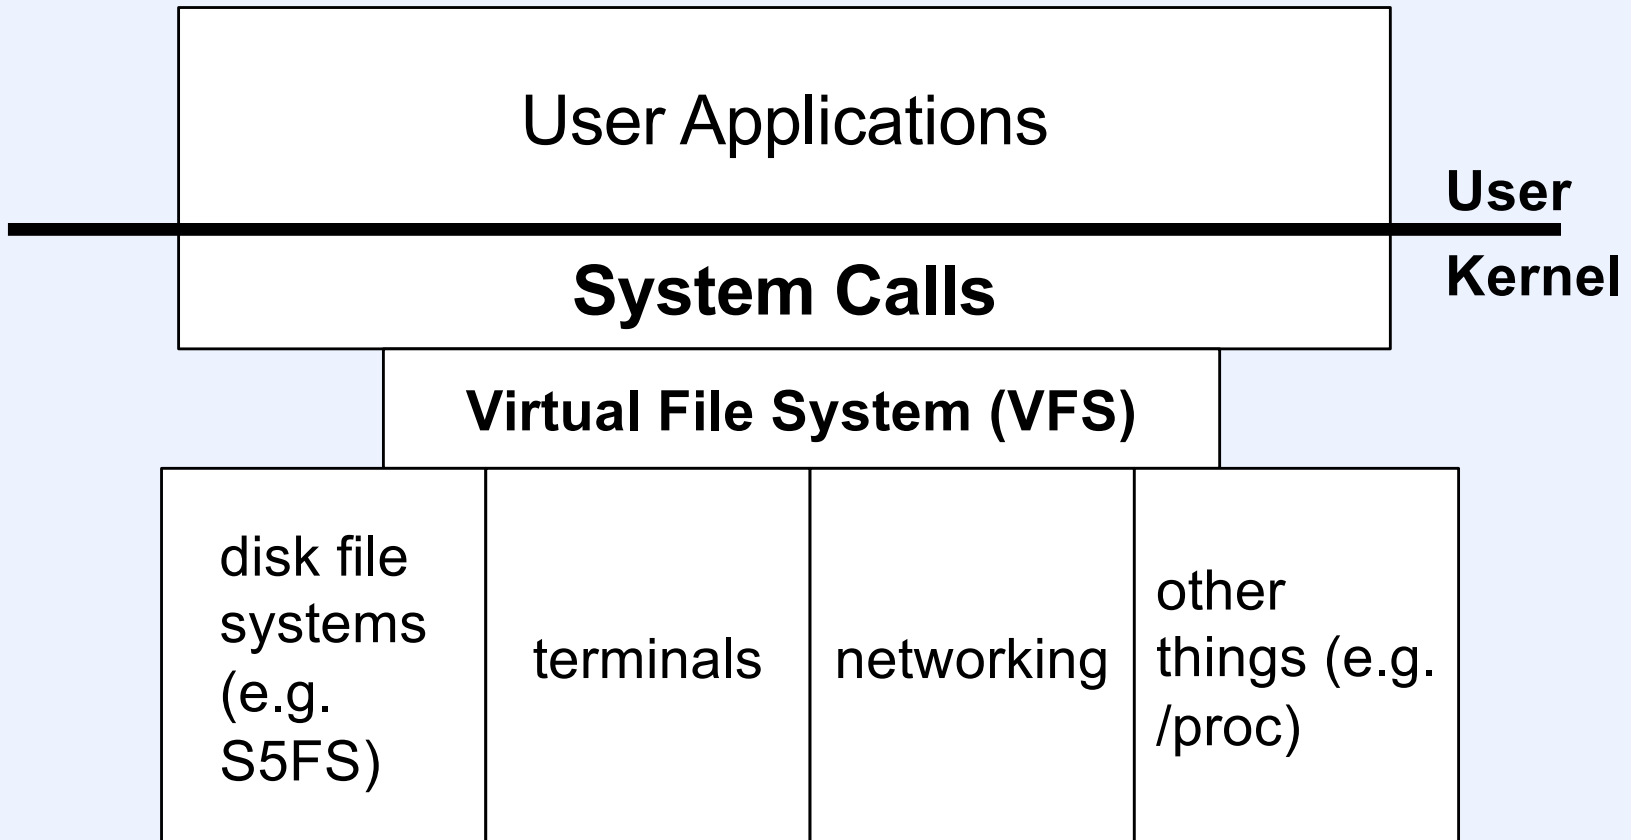
Adding More Stuff



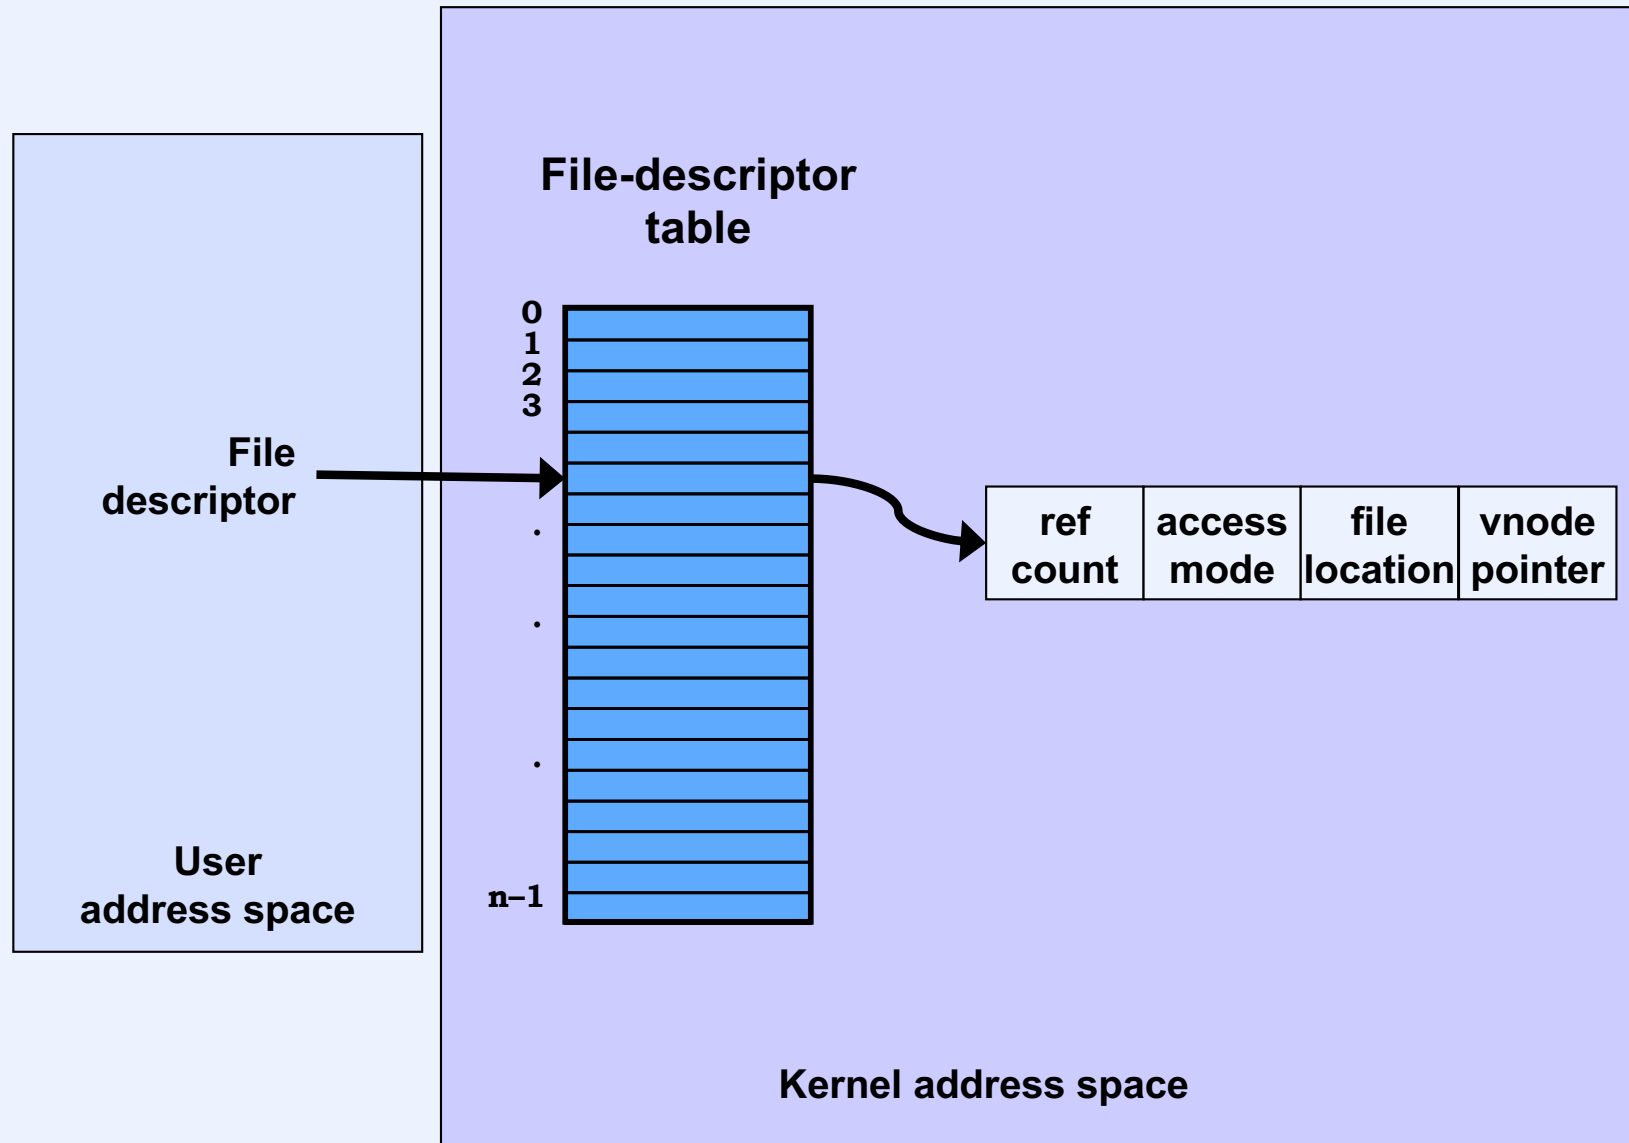
Address-Space Representation: Reality



Layering



File-Descriptor Table



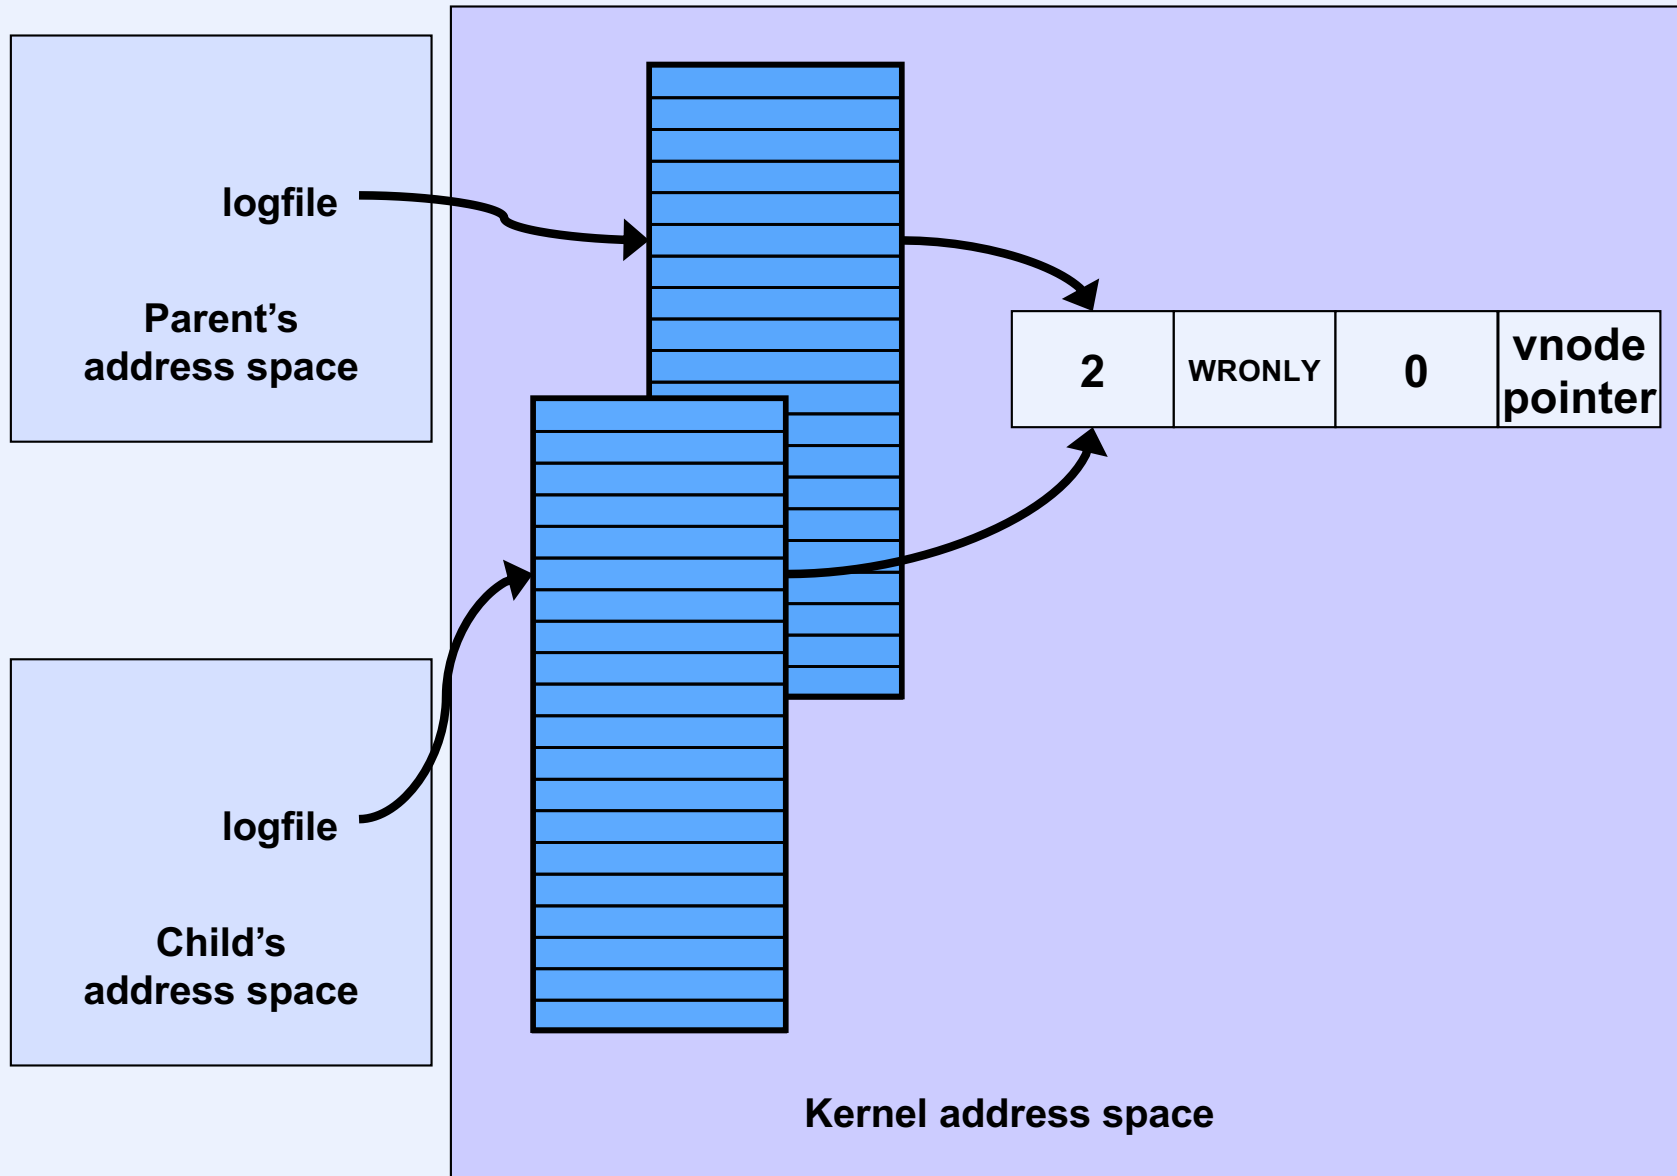
Fork and File Descriptors

```
int logfile = open("log", O_WRONLY);
if (fork() == 0) {
    /* child process computes something, then does: */
    write(logfile, LogEntry, strlen(LogEntry));
    ...
    exit(0);
}

/* parent process computes something, then does: */

write(logfile, LogEntry, strlen(LogEntry));
...
```

File Descriptors After Fork



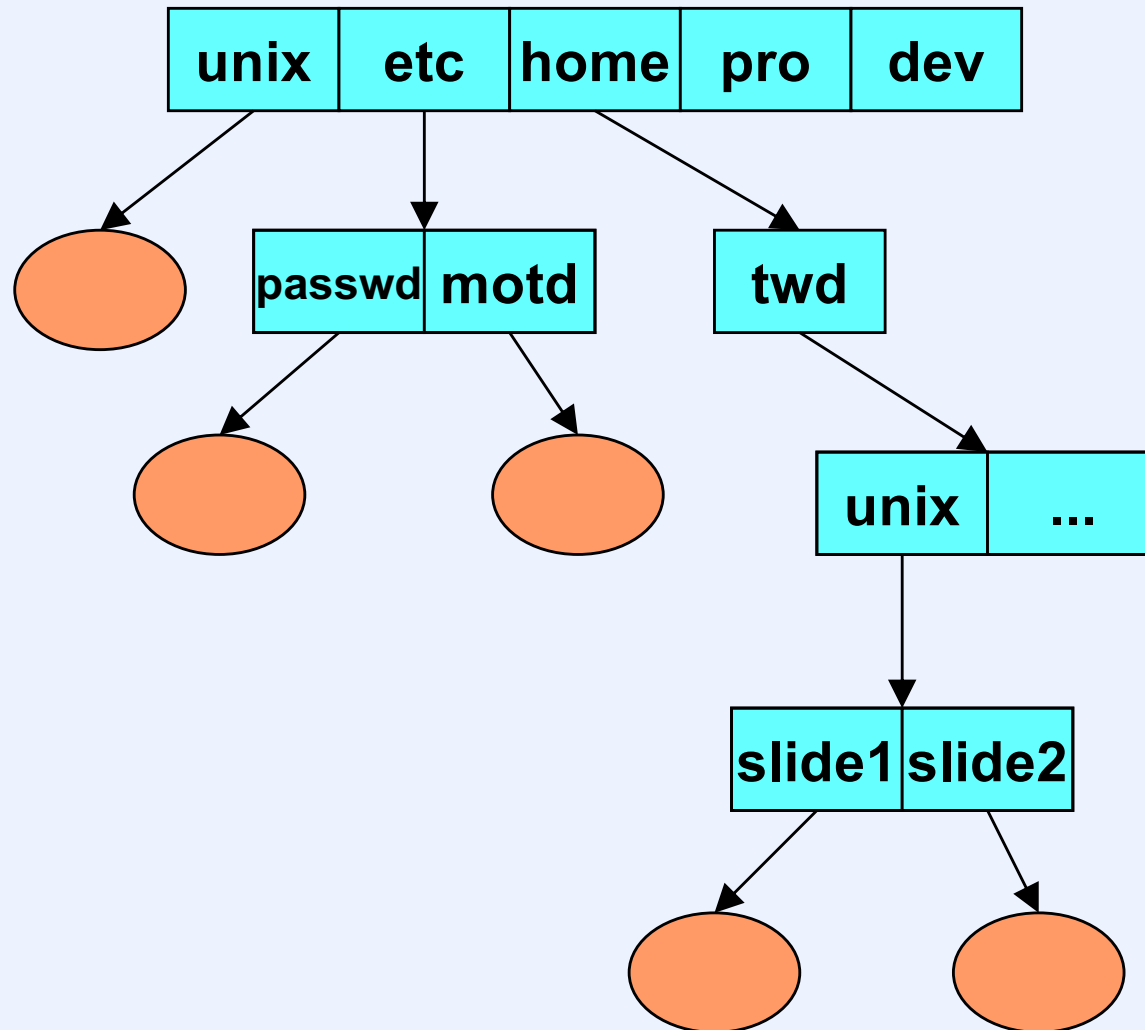
Quiz 1

```
int fd1 = open("file", O_CREAT|O_RDWR, 0666);  
unlink("file");  
write(fd1, "123", 3);  
int fd2 = open("file", O_CREAT|O_RDWR, 0666);  
write(fd2, "4", 1);  
if (fork() == 0) {  
    write(fd1, "5", 1);  
}  
exit(0);
```

The final contents of file are:

- a) 4
- b) 45
- c) 453
- d) 12345

Directories




Directory Representation

Component Name	Inode Number
----------------	--------------

directory entry

.	1
..	1
unix	117
etc	4
home	18
pro	36
dev	93

Hard Links




% ln /unix /etc/image
link system call

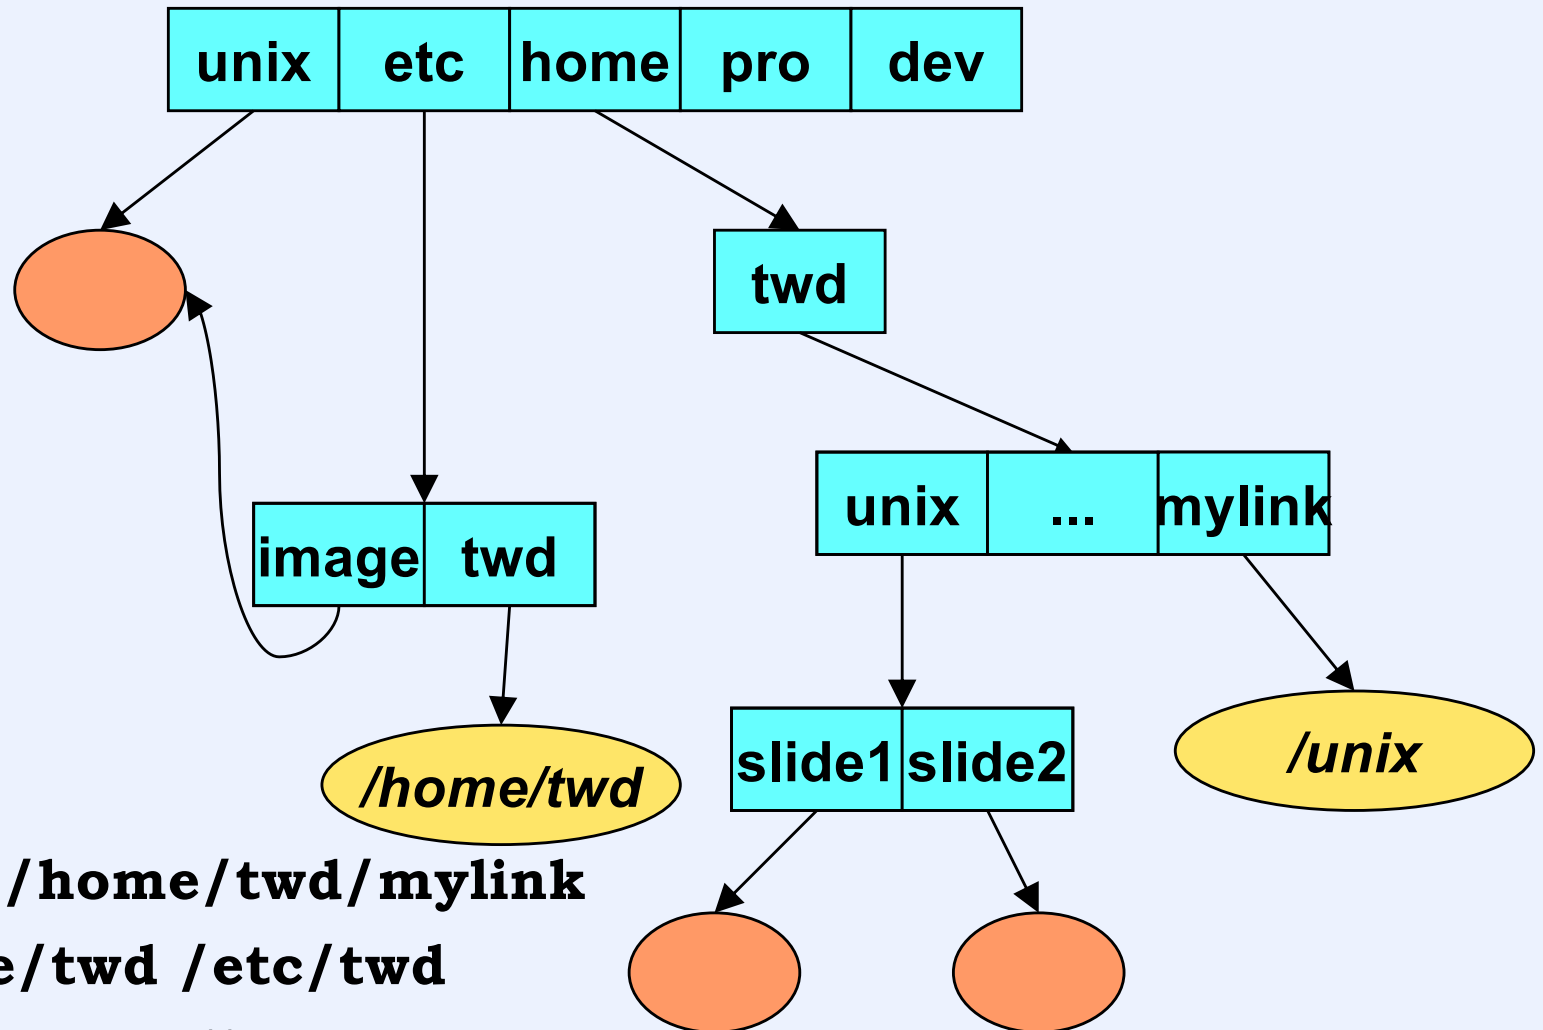
Directory Representation

.	1
..	1
unix	117
etc	4
home	18
pro	36
dev	93

.	4
..	1
image	117
motd	33



Soft Links



```
% ln -s /unix /home/twd/mylink
```

```
% ln -s /home/twd /etc/twd
```

```
# symlink system call
```



Working Directory

- **Maintained in kernel for each process**
 - paths not starting with “/” start with the working directory
 - changed by use of the *chdir* system call
 - displayed (via shell) using “pwd”
 - how is this done?


pwd

- **Print Working Directory**
 - suppose the current working directory is **/A/B/C**
 - how can a program determine it?


root



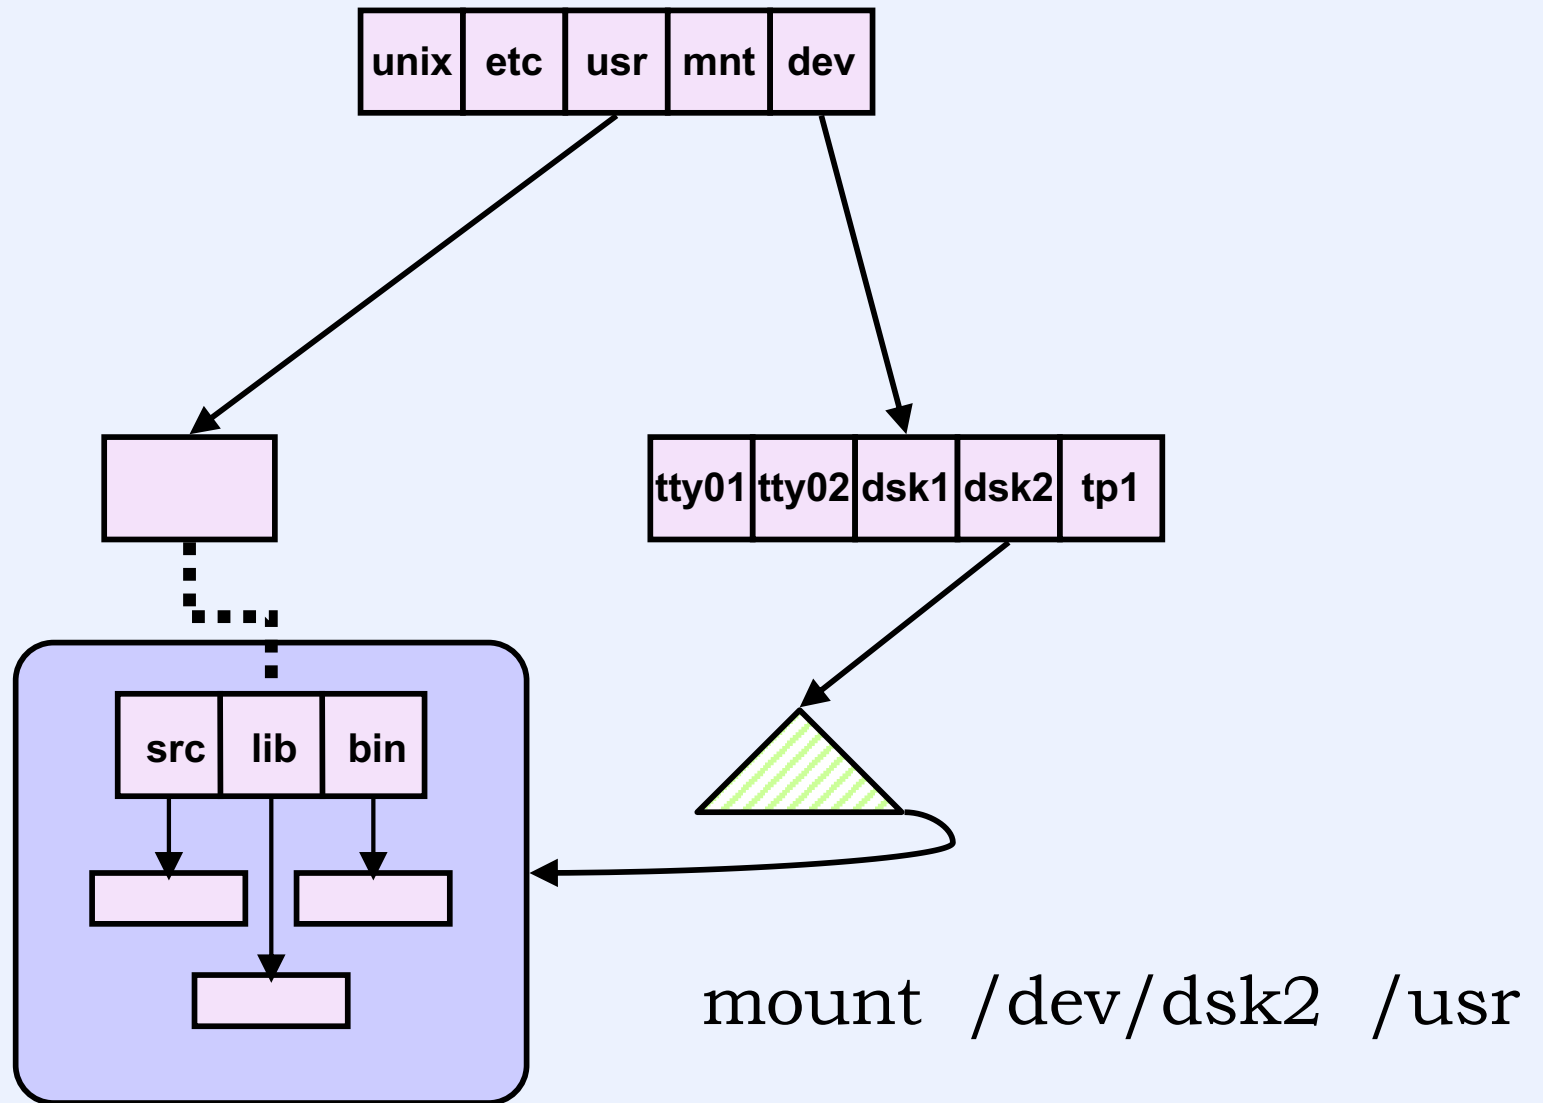
**current
directory**



Mount Points (1)



Mount Points (2)



Representing File Systems

```
class fs {  
    char dev[STR_MAX];           // device containing the f.s.  
    char mountpt[STR_MAX];       // where the f.s. is mounted  
    vnode *vnodecovered;         // file on which f.s. is mounted  
    vnode *root;                // root of the f.s.  
    virtual void read_vnode(vnode *);  
    virtual void delete_vnode(vnode *);  
};
```

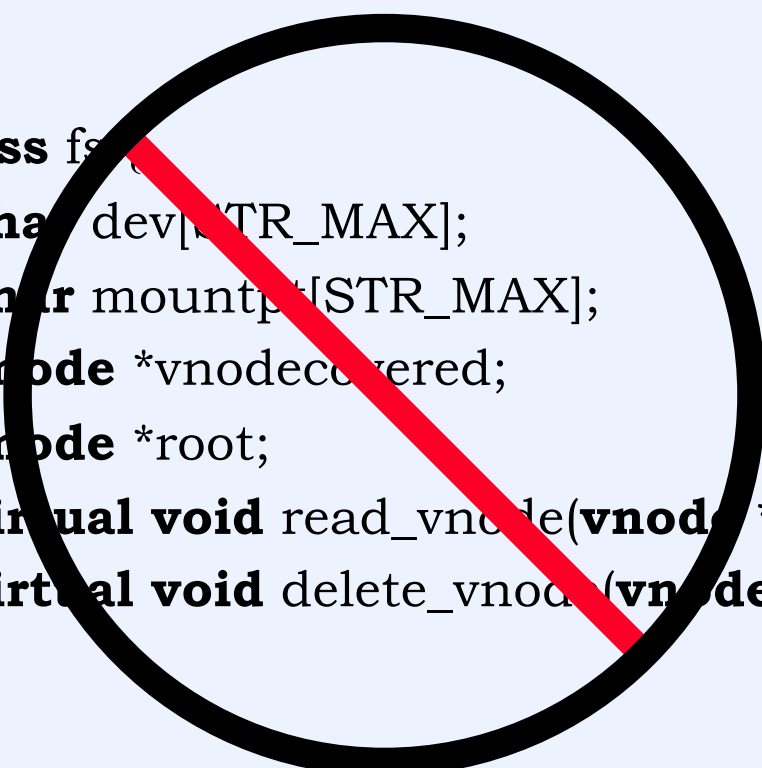

Representing Files

```
class vnode {  
    unsigned short refcount;  
    fs *vfsmounted;  
    fs *vfs;  
    unsigned long vno;  
    int mode;  
    int len;  
    link_list_t link;  
    kmutex_t mutex;  
    virtual int create(const char *, int, vnode **);  
    virtual int read(int, void *, int);  
    virtual int write(int, const void *, int);  
    ...  
    class mobj mobj;  
};
```

But Wait ...

- **What's this about C++?**
 - **real operating systems are written in C ...**

fs



```
class fs {  
    char dev[STR_MAX];  
    char mountpt[STR_MAX];  
    vnnode *vnodecovered;  
    vnnode *root;  
    virtual void read_vnode(vnnode *);  
    virtual void delete_vnode(vnnode *);  
};
```

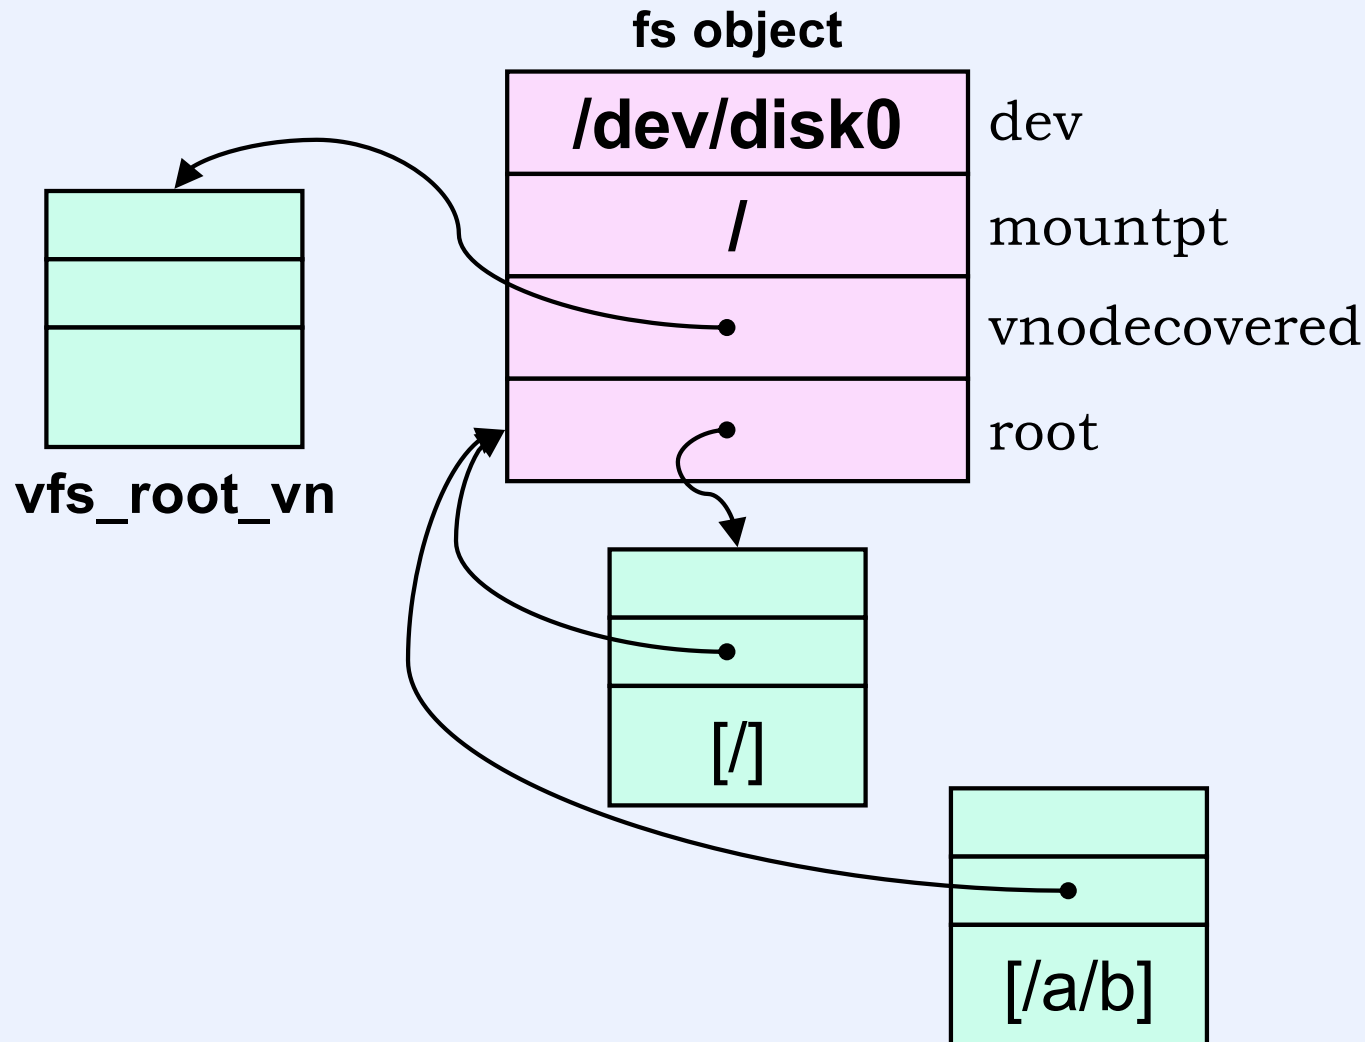
```
typedef struct fs {  
    char fs_dev[STR_MAX];  
    char fs_mountpt[STR_MAX];  
    struct vnnode *fs_vnodecovered;  
    struct vnnode *fs_root;  
    fs_ops_t *fs_op;  
    /* function pointers */  
    void *fs_i;  
    /* extra stuff in subclasses */  
} fs_t;
```

vnode

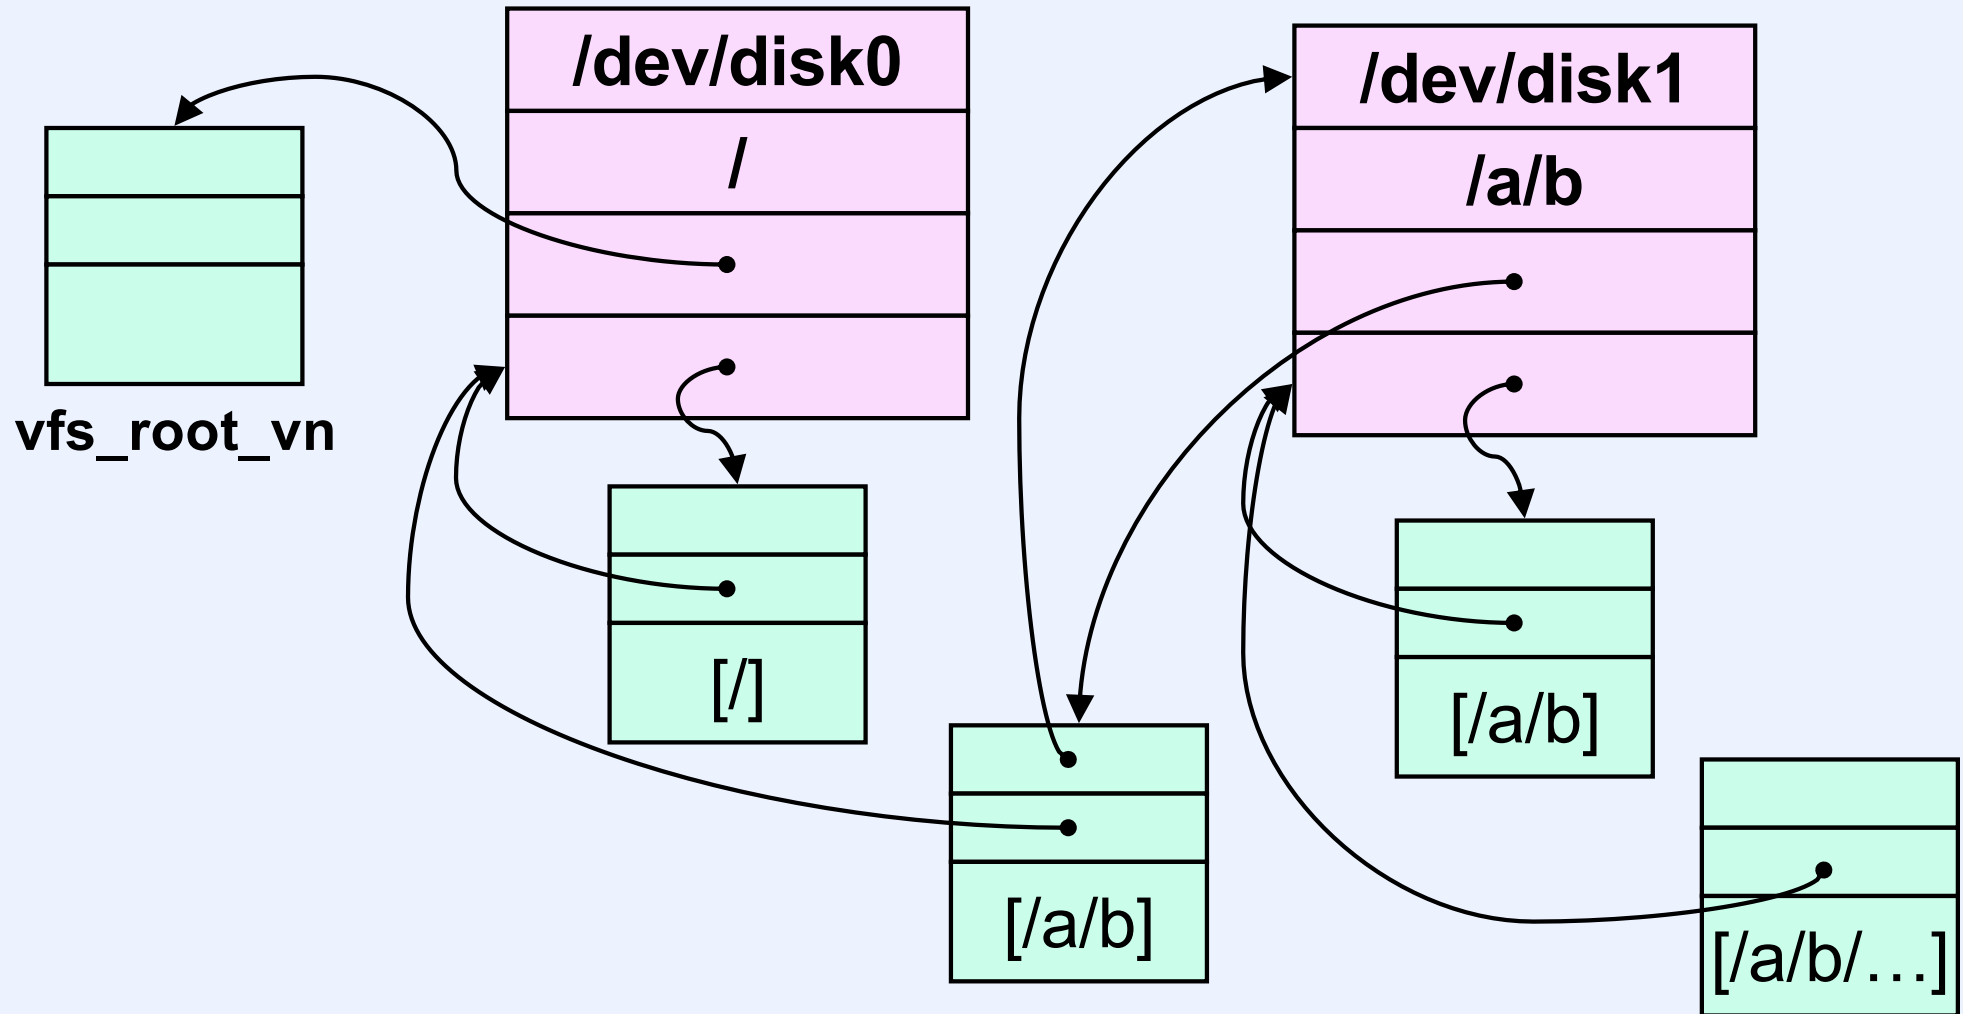
```
class vnode {  
    unsigned short refcount;  
    fs *vfsmounted;  
    fs *vfs;  
    unsigned long vno;  
    int mode;  
    int len;  
    link_list_t link;  
    kmutex_t mutex;  
    virtual int create(const char *, int,  
        vnode **);  
    virtual int read(int, void *, int);  
    virtual int write(int, const void *,  
        int);  
    ...  
    class mobj mobj;  
};
```

```
typedef struct vnode {  
    unsigned short vn_refcount;  
    struct fs *vn_vfsmounted;  
    struct fs *vn_vfs;  
    unsigned long vn_vno;  
    int vn_mode;  
    int vn_len;  
    link_list_t vn_link;  
    kmutex_t vn_mutex;  
    struct vnode_ops *vn_op;  
    /* function pointers */  
    struct mobj mobj;  
    void *vn_i;  
    /* extra stuff in subclasses */  
} vnode_t;
```

Mounting a File System (1)



Mounting a File System (2)




Quiz 2

- **Suppose the current working directory is /A/B/M. /A/B is a path in file system 1. /M is a directory in file system 2. File system 2 is mounted on file system 1 at /A/B. How does the pwd command handle this case?**
 - a) **it's easy: the “..” entry in /A/B/M refers to the directory /A/B and one simply continues backwards towards the root as if there were no mount point**
 - b) **a special system call is required to determine the path**
 - c) **one must start at the root and continue forwards towards the working directory**


Digression (Sort of)

- **Weenix has two file systems**
 - **ramfs**
 - trivial in-memory file system for testing purposes
 - caching of blocks isn't necessary
 - **s5fs**
 - non-trivial file system that you implement
 - caching of blocks is important
 - caching uses virtual-memory subsystem so that *mmap* works right



Traditional I/O



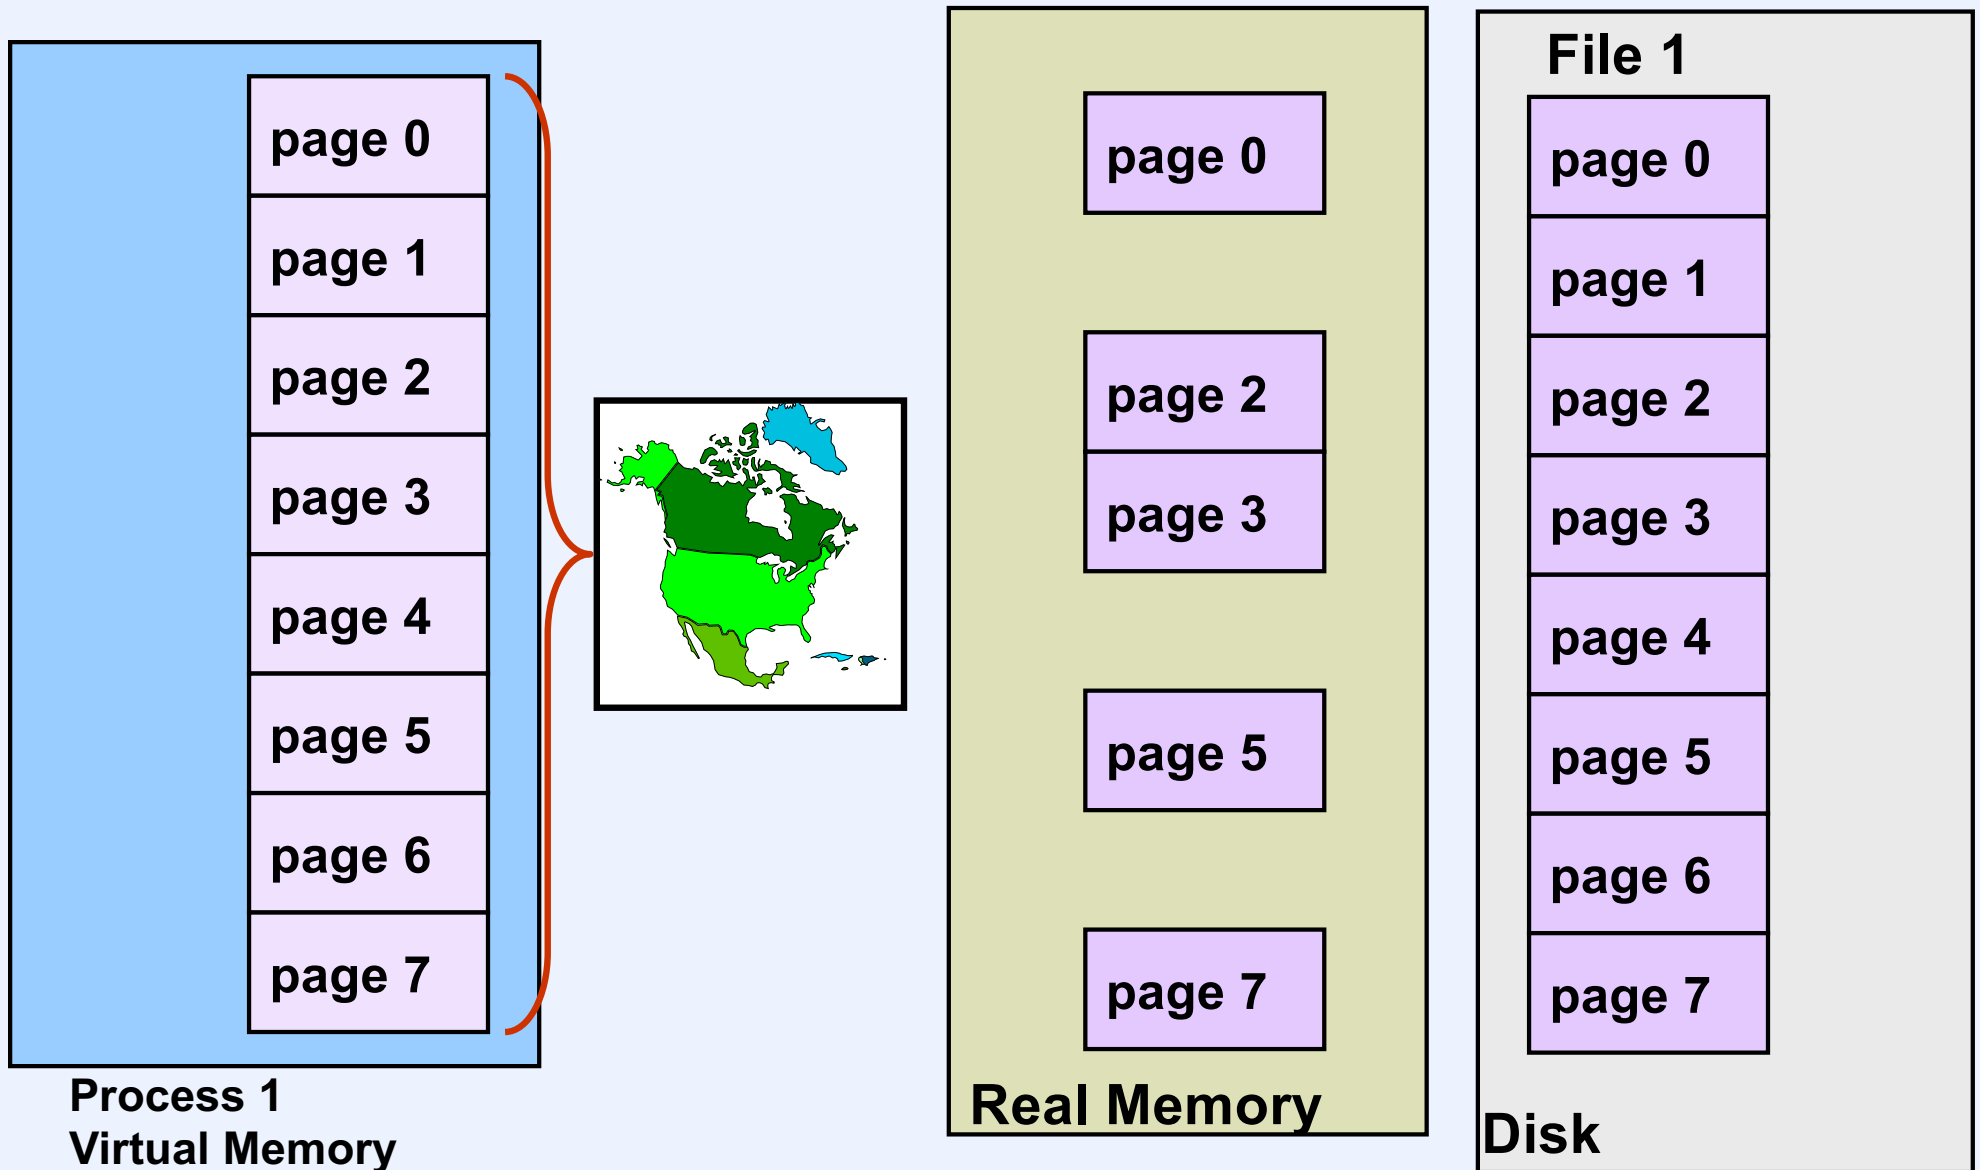
User Process 1



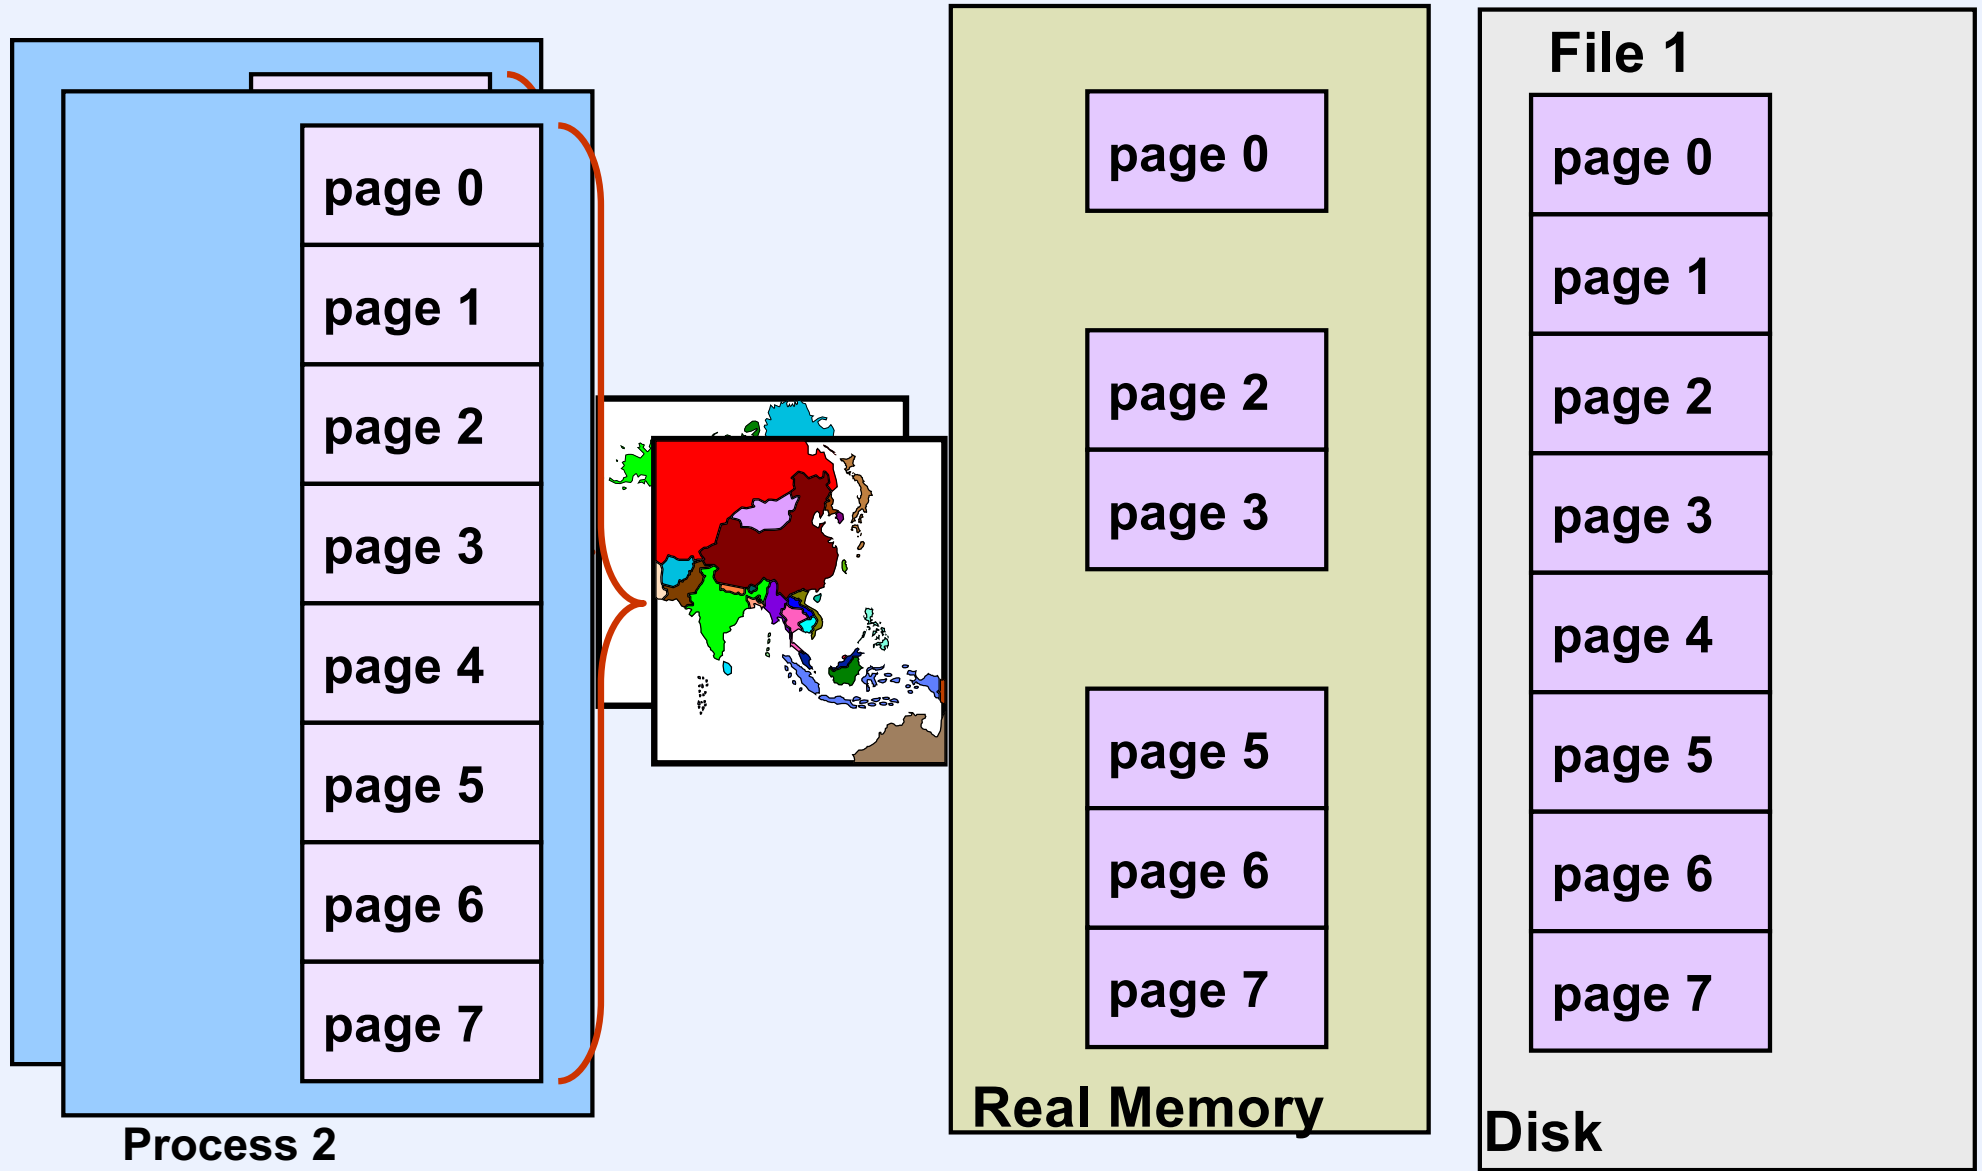
User Process 2



Mapped File I/O



Multi-Process Mapped File I/O



Memory Objects

