

Implementing Threads 2

Time Slicing

- **Periodically**
 - current thread forced to do a thread yield



```
void ClockInterrupt(int sig) {  
    thread_yield();  
}
```

- **Implement ClockInterrupt with VTALRM signal**

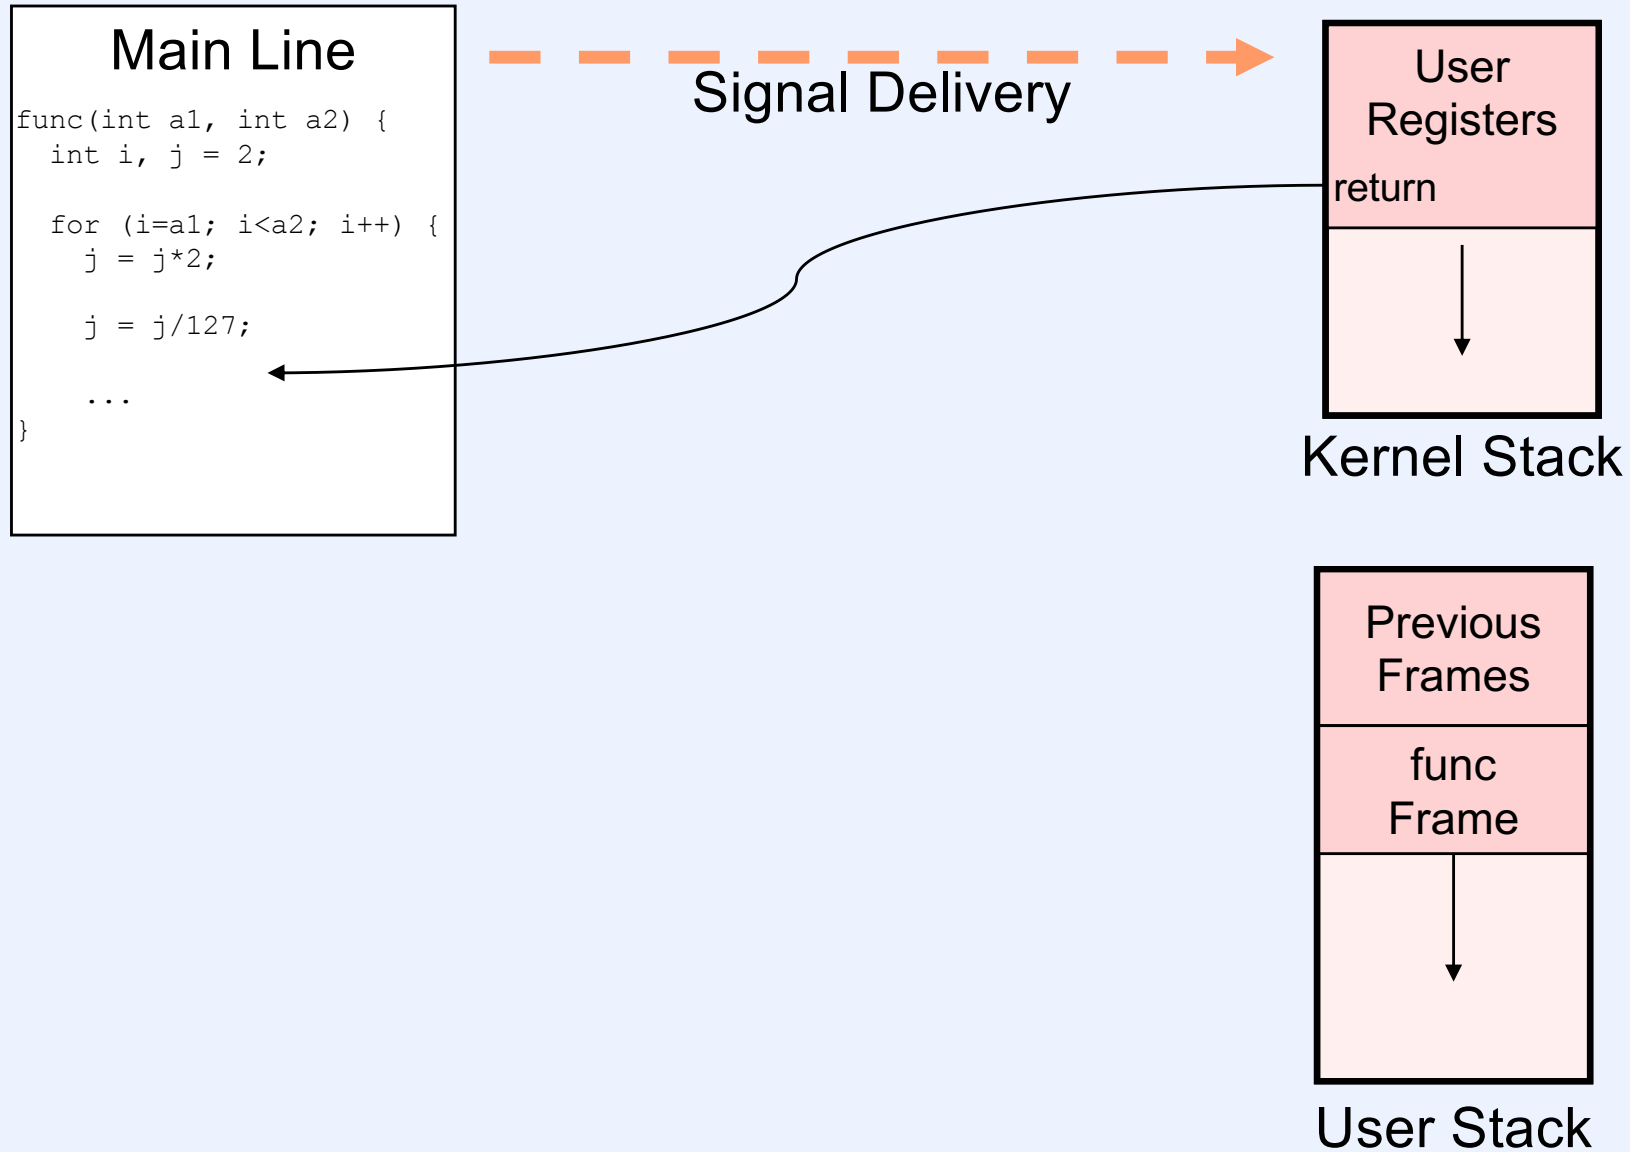
Invoking the Signal Handler

- **Basic idea is to set up the user stack so that the handler is called as a subroutine and so that when it returns, normal execution of the thread may continue**
- **Complications:**
 - saving and restoring registers
 - signal mask

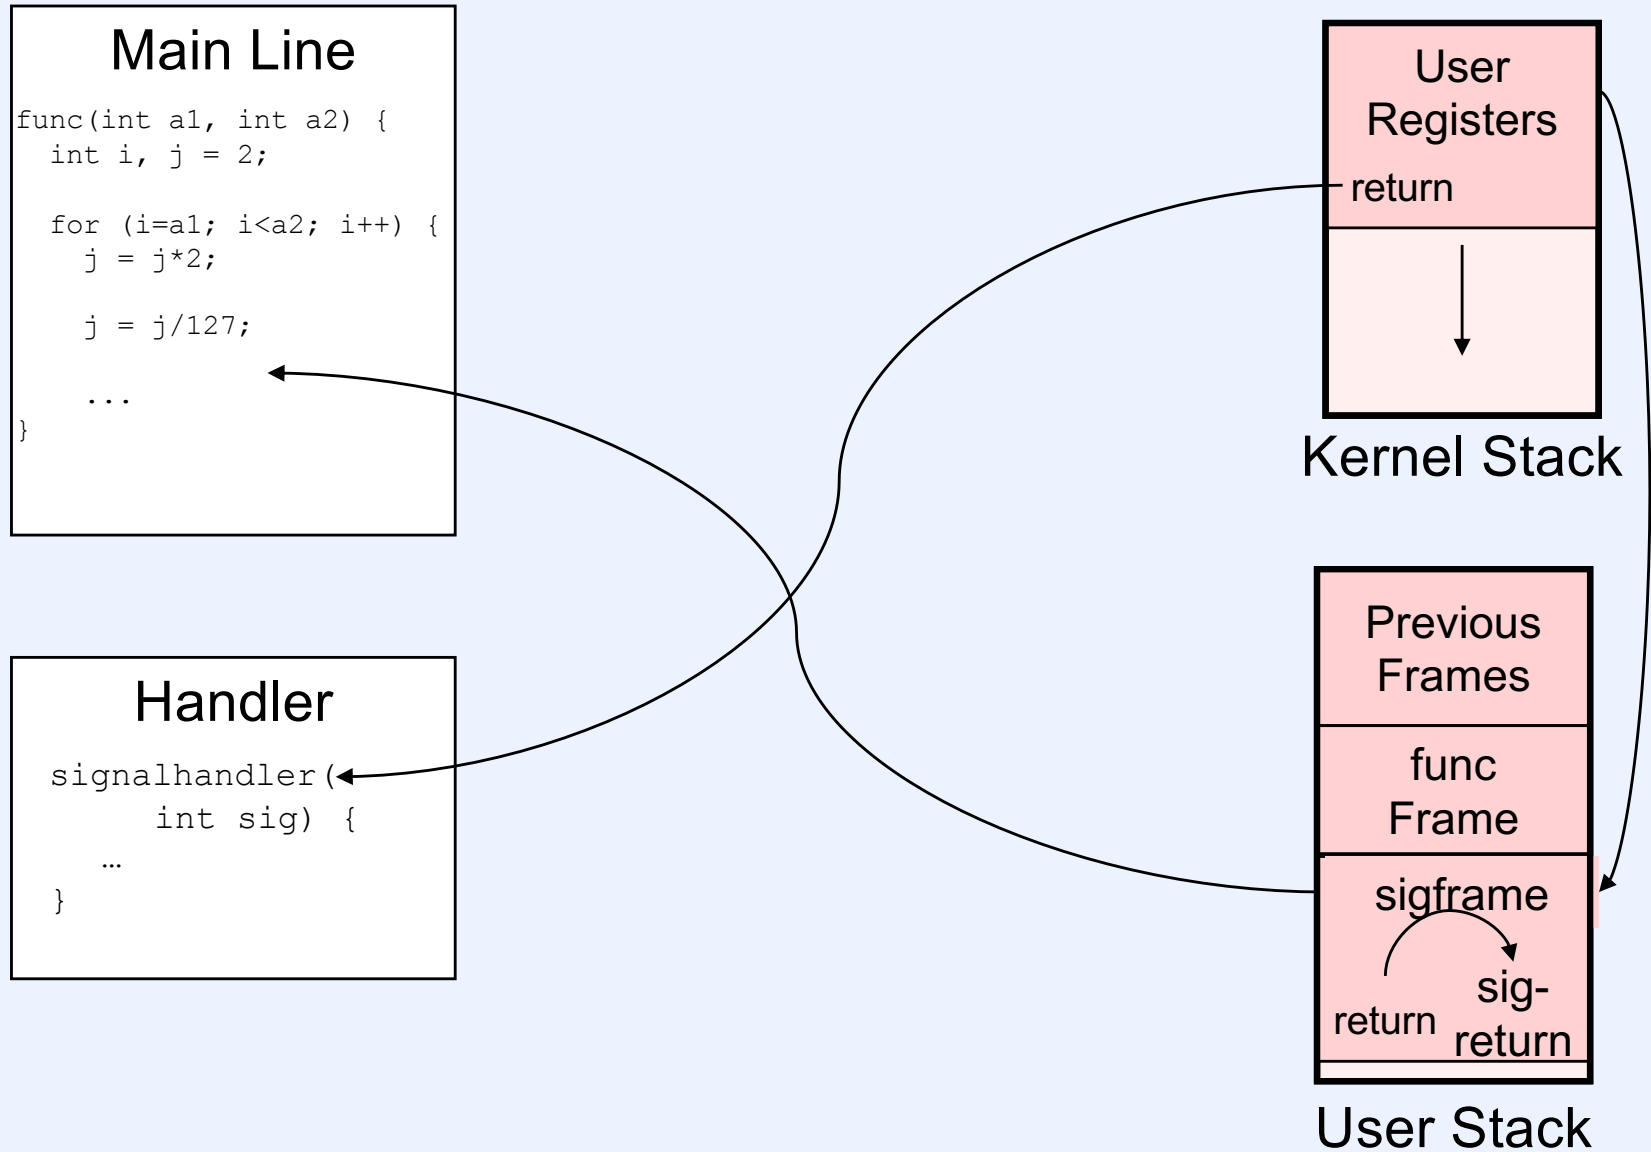
Invoking the Signal Handler (1)




Invoking the Signal Handler (2)



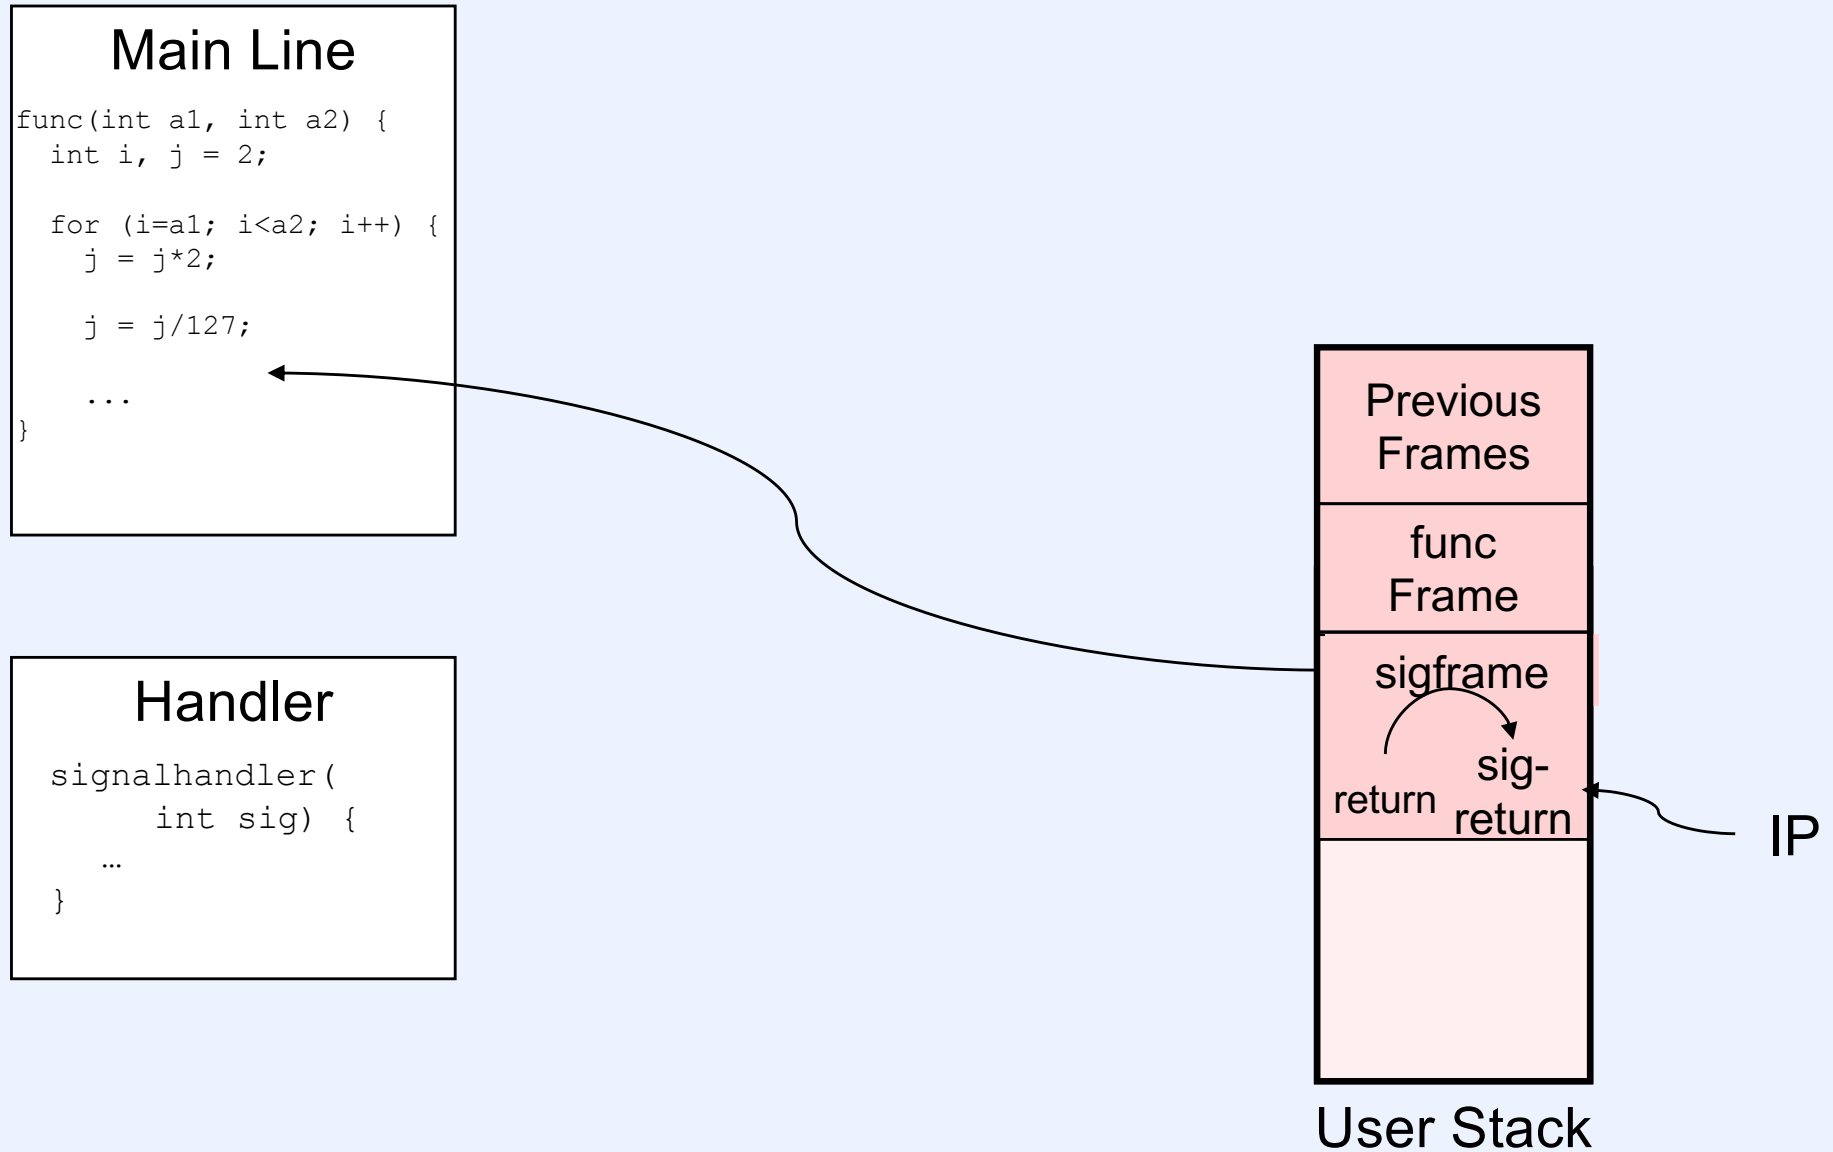
Invoking the Signal Handler (3)



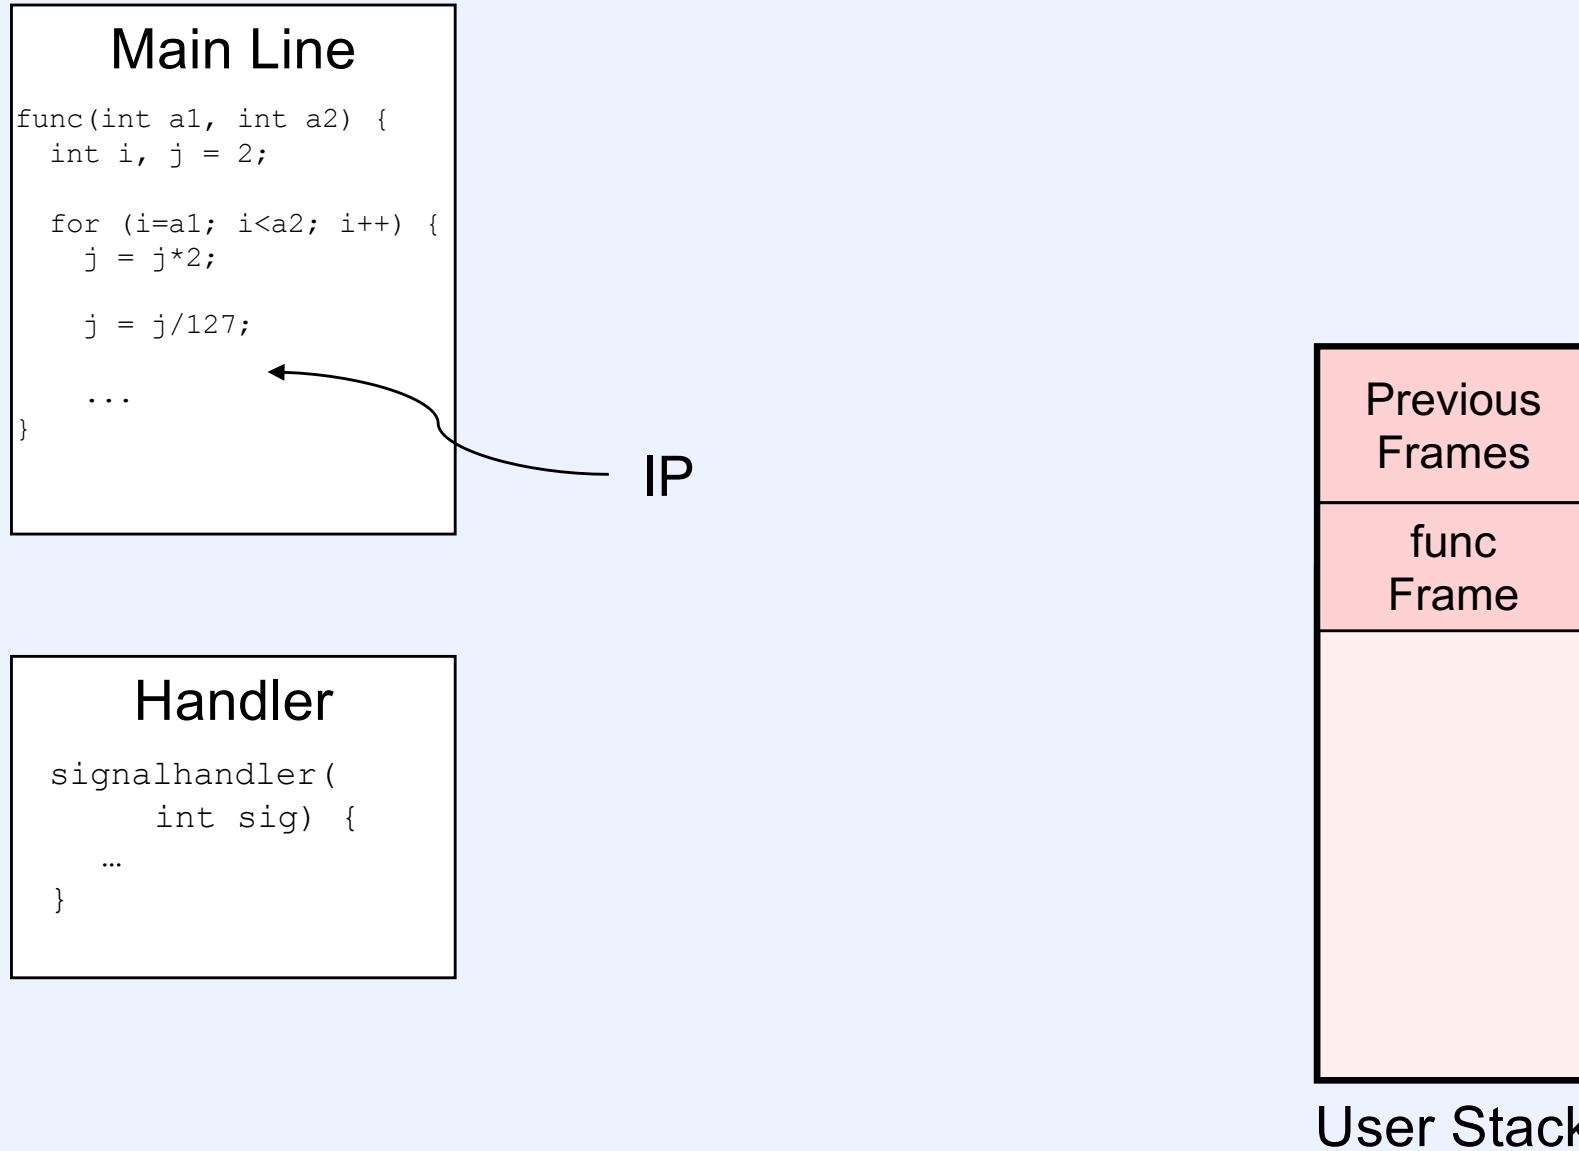
Invoking the Signal Handler (4)



Invoking the Signal Handler (5)



Invoking the Signal Handler (6)



Quiz 1

The description of invoking the signal handler:

- a) works fine**
- b) is susceptible to buffer-overflow attacks**
- c) is rendered useless because of an approach for making such attacks more difficult (by turning off execute permission on some memory)**
- d) doesn't work, period**

Time Slicing

- **Periodically**
 - **current thread forced to do a thread yield**

```
void ClockInterrupt(int sig) {  
    // SIGVTALRM is now masked  
    sigprocmask(SIG_UNBLOCK, &VTALRMmask, 0);  
    // SIGVTALRM is now unmasked  
    utthread_yield();  
    // thread resumes here  
}
```

- **Implement ClockInterrupt with VTALRM signal**

Setting Up Time Slicing

```
struct sigaction timesliceact;  
timesliceact.sa_handler = ClockInterrupt;  
timesliceact.sa_mask = VTALRMmask;  
timesliceact.sa_flags = SA_RESTART; // avoid EINTR  
struct timeval interval = {0, 1}; // every microsecond  
struct itimerval timerval;  
timerval.it_value = interval;  
timerval.it_interval = interval;  
sigaction(SIGVTALRM, &timesliceact, 0);  
setitimer(ITIMER_VIRTUAL, &timerval, 0);  
    // time slicing is started!
```

Async-Signal Safety

- A function is asynchronous-signal safe if it may be used in the handler for an asynchronous signal (such as SIGVTALRM)
 - malloc and free
 - no
 - mutex_lock
 - no
 - read and write
 - yes

Achieving Async-Safety

- **The problem: an action in the signal handler interferes with an action in the main-line code**
 - while in malloc/free, a signal occurs and the handler calls malloc/free
 - while holding the lock on a mutex, a thread is interrupted and the handler attempts to lock the mutex
- **The solution: mask signals while in malloc/free and when holding locks**

Caution!

- ***uthread_switch*** is not async-signal safe
 - it's called from *uthread_yield*, which is called from the signal handler for SIGVTALRM
 - must mask signals before calling it (and unmask afterwards)

Masking/Unmasking Signals

```
sigset_t VTALRMmask;
```

...

```
sigemptyset (&VTALRMmask);
```

```
sigaddset (&VTALRMmask, SIGVTALRM);
```

...

```
sigprocmask (SIG_BLOCK, &VTALRMmask, 0);
```

...

```
sigprocmask (SIG_UNBLOCK, &VTALRMmask, 0);
```


Doing It Cheaply

```
void uthread_nopreempt_on() {  
    uthread_no_preempt = 1;  
}
```

```
uthread_no_preempt_on();  
  
uthread_switch();
```

```
void uthread_nopreempt_off() {  
    uthread_no_preempt = 0;  
}
```

```
uthread_no_preempt_off();
```

```
void ClockInterrupt(int sig) {  
    if (uthread_no_preempt)  
        return;  
    ...  
}
```

Nested Calls

```
void uthread_nopreempt_on() {  
    uthread_no_preempt_count++;  
}
```

```
void uthread_nopreempt_off() {  
    uthread_no_preempt_count--;  
}
```

```
void ClockInterrupt(int sig) {  
    if (uthread_no_preempt_count > 0)  
        return;  
    ...  
}
```

Corrected Nested Calls

```
void uthread_nopreempt_on() {  
    ut_curthr->uthread_no_preempt_count++;  
}  
  
void uthread_nopreempt_off() {  
    ut_curthr->uthread_no_preempt_count--;  
}  
  
void ClockInterrupt(int sig) {  
    if (ut_curthr->uthread_no_preempt_count > 0)  
        return;  
    ...  
}
```

Limitations of User Threads

- **Threads are implemented strictly at user level**
 - the OS kernel is unaware of their existence
- **What happens if a user thread makes a blocking system call, e.g., *read*?**

Quiz 2

```
void uthread_switch( ) {  
    volatile int first = 1;  
    getcontext(&CurrentThread->ctx);  
    if (first) {  
        first = 0;  
        CurrentThread = dequeue(RunQueue);  
        setcontext(&CurrentThread->ctx);  
    }  
    return;  
}
```

Given the discussion so far, will RunQueue ever be empty (in a program that has no deadlocks)?

- a) yes**
- b) no**

Multiple Processors

```
void uthread_switch( ) {  
    volatile int first = 1;  
    getcontext(&CurrentThread->ctx);  
    if (first) {  
        first = 0;  
        CurrentThread = dequeue(RunQueue);  
        setcontext(&CurrentThread->ctx);  
    }  
    return;  
}
```

- **How do we employ multiple processors?**
 - code merely switches the caller's processor to another thread
 - **What if the RunQueue is empty?**
 - it could be empty, particularly if we have multiple processors
-

Solution Sketch

- Introduce “idle threads”, one for each processor
- Thread calling *uthread_switch* switches to idle thread for its current processor
- Idle thread then switches to first thread on *RunQueue*, if any
- If *RunQueue* is empty, idle thread repeatedly checks *RunQueue* until it's not empty, then switches to first thread

Solution Details

```
1 void uthread_switch() {
2     volatile int first = 1;
3     getcontext(&CurrentThread[processor_ID]->context);
4     if (!first)
5         return;
6     first = 0;
7     setcontext(&IdleThread[processor_ID]->context);
8 }

9 void IdleThread_switch() {
10    getcontext(&IdleThread[processor_ID]->context);
11    while (1) {
12        if (queue_empty(RunQueue))
13            continue;
14        CurrentThread[processor_ID] = dequeue(RunQueue);
15        setcontext(&CurrentThread[processor_ID]->context);
16    }
17 }
```


MThreads

- **Idle threads are pthreads**
 - called LWPs (lightweight processes), following standard usage
 - each LWP represents a processor (core)
 - a “virtual processor”
 - *lwp_switch* rather than *IdleThread_switch*
 - rather than “busy-wait” for a uthread to run, it calls *lwp_park* to wait for one (using a POSIX condition variable)

MP Mutual Exclusion

- **Two sorts**
 - **spin locks**
 - threads wait by repeatedly testing the lock
 - **blocking locks**
 - threads wait by sleeping, depending on some other thread to wake them up

Hardware Support for Spin Locks

- Compare and swap instruction

```
int CAS(int *ptr, int old, int new) {  
    int tmp = *ptr;  
    if (*ptr == old)  
        *ptr = new;  
    return tmp;  
}
```

Naive Spin Lock

```
void spin_lock(int *spin) {  
    while(CAS(spin, 0, 1))  
        ;  
}
```

```
void spin_unlock(int *spin) {  
    *spin = 0;  
}
```

Better Spin Lock

```
void spin_lock(int *spin) {  
    while (1) {  
        if (*spin== 0) {  
            // the mutex was at least momentarily unlocked  
            if (!CAS(spin, 0, 1)  
                break; // we have locked the mutex  
            // some other thread beat us to it, so try again  
        }  
    }  
}
```

Spin Locks in MThreads

- **Since LWPs represent virtual processors, we don't want busy waiting**
- **Use POSIX mutexes instead**

Blocking Locks

```
void blocking_lock(mutex_t *mut) {  
    if (mut->holder != 0) {  
        enqueue(mut->wait_queue,  
                CurrentThread);  
        uthread_switch();  
    } else  
        mut->holder = CurrentThread;  
}
```

```
void blocking_unlock(mutex_t *mut) {  
    if (queue_empty(mut->wait_queue))  
        mut->holder = 0;  
    else {  
        mut->holder =  
            dequeue(mut->wait_queue);  
        enqueue(RunQueue, mut->holder);  
    }  
}
```

Does it work?

Working Blocking Locks (?)

```
void blocking_lock(mutex_t *mut) {  
    spin_lock(&mut->spinlock);  
    if (mut->holder != 0) {  
        enqueue(mut->wait_queue,  
                CurrentThread);  
        spin_unlock(&mut->spinlock);  
        uthread_switch();  
    } else {  
        mut->holder = CurrentThread;  
        spin_unlock(&mut->spinlock);  
    }  
}
```

```
void blocking_unlock(mutex_t *mut) {  
    spin_lock(&mut->spinlock);  
    if (queue_empty(  
        mut->wait_queue)) {  
        mut->holder = 0;  
    } else {  
        mut->holder =  
            dequeue(mut->wait_queue);  
        enqueue(RunQueue,  
                mut->holder);  
    }  
    spin_unlock(&mut->spinlock);  
}
```

Quiz 3

This

- a) always works
- b) sometimes doesn't work
- c) never works

Futexes

- **Safe, *efficient* kernel conditional queueing in Linux**
- **All operations performed atomically**
 - `futex_wait(futex_t *futex, int val)`
 - **if `futex->val` is equal to `val`, then sleep**
 - **otherwise return**
 - `futex_wake(futex_t *futex)`
 - **wake up one thread from `futex`'s wait queue, if there are any waiting threads**

Ancillary Functions

- `int atomic_inc(int *val)`
 - **add 1 to `*val`, return its original value**
- `int atomic_dec(int *val)`
 - **subtract 1 from `*val`, return its original value**

Attempt 1

```
void lock(futex_t *futex) {  
    int c;  
    while ((c = atomic_inc(&futex->val)) != 0)  
        futex_wait(futex, c+1);  
}
```

```
void unlock(futex_t *futex) {  
    futex->val = 0;  
    futex_wake(futex);  
}
```

Attempt 2

```
void lock(futex_t *futex) {
    int c;
    if ((c = CAS(&futex->val, 0, 1) != 0)
        do {
            if (c == 2 || (CAS(&futex->val, 1, 2) != 0))
                futex_wait(futex, 2);
            while ((c = CAS(&futex->val, 0, 2)) != 0))
        }

void unlock(futex_t *futex) {
    if (atomic_dec(&futex->val) != 1) {
        futex->val = 0;
        futex_wake(futex);
    }
}
```

Quiz 4
Does it work?

- a) Yes
- b) No

Blocking Locks in MThreads

- We could use futexes, but don't
- *uthread_switch* gets an additional argument
 - a POSIX mutex (representing a spin lock)
 - unlock it after getting out of the context of the calling thread


Actual Code

```
uthread_mtx_lock(uthread_mtx_t *mtx) {  
    uthread_nopreempt_on();  
    pthread_mutex_lock(&mtx->m_pmut);  
    if (mtx->m_owner == NULL) {  
        mtx->m_owner = ut_curthr;  
        pthread_mutex_unlock(&mtx->m_pmut);  
        uthread_nopreempt_off();  
    } else {  
        ut_curthr->ut_state = UT_WAIT;  
        uthread_switch(&mtx->m_waiters, 0, &mtx->m_pmut);  
        uthread_nopreempt_off();  
    }  
}
```


MP Memory Issues

- **Naive view is that all processors in MP system see same memory contents at all times**
 - they don't

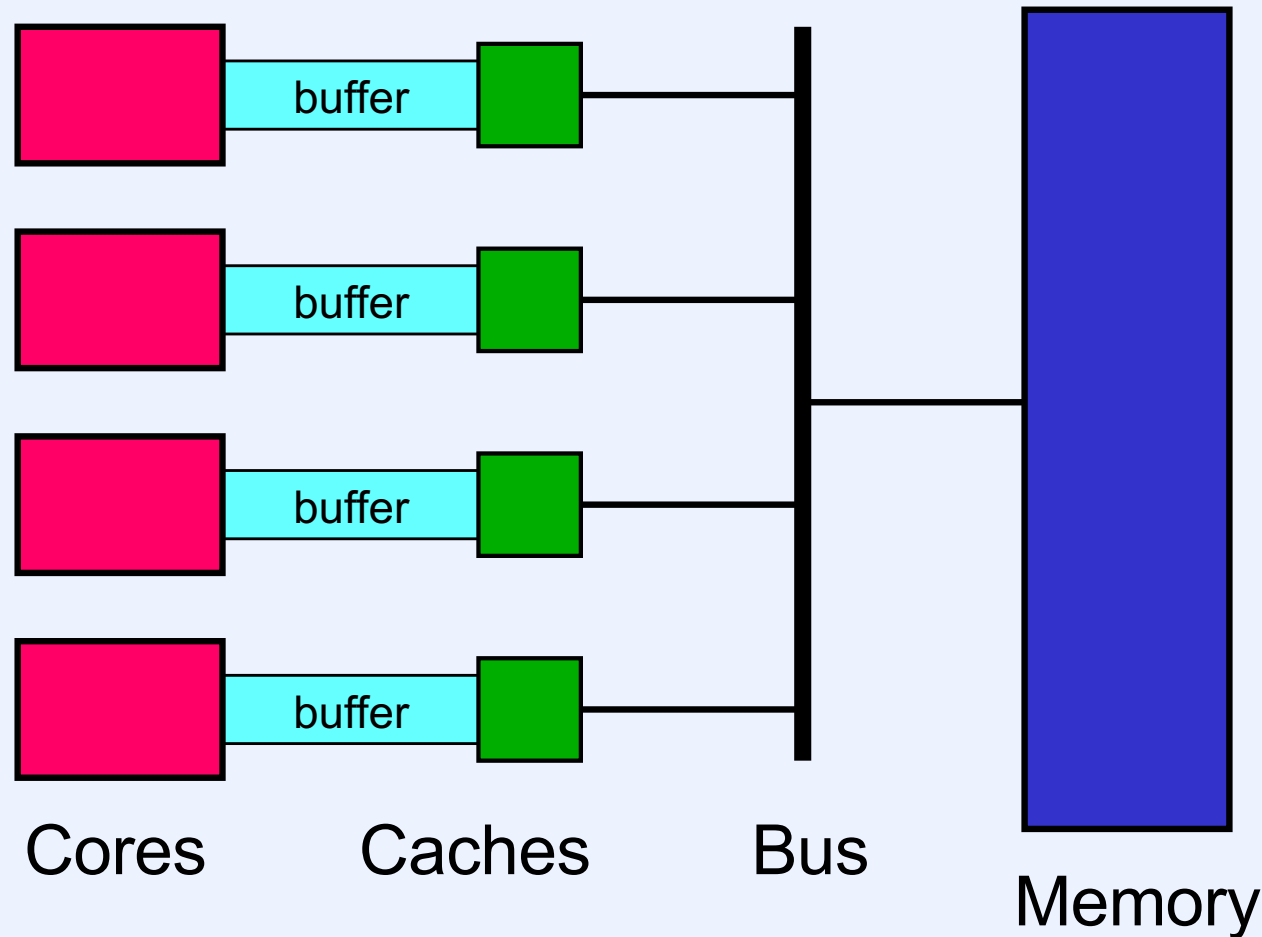
Multi-Core Processor: Simple View



Multi-Core Processor: More Realistic View



Multi-Core Processor: Even More Realistic



Concurrent Reading and Writing

Thread 1:

```
i = shared_counter;
```

Thread 2:

```
shared_counter++;
```

Mutual Exclusion w/o Mutexes

```
void peterson(long me) {  
    static long loser;           // shared  
    static long active[2] = {0, 0}; // shared  
    long other = 1 - me;        // private  
    active[me] = 1;  
    loser = me;  
    while (loser == me && active[other])  
        ;  
    // critical section  
    active[me] = 0;  
}
```

Busy-Waiting Producer/Consumer

```
void producer(char item) {  
  
    while(in - out == BSIZE)  
        ;  
  
    buf[in%BSIZE] = item;  
  
    in++;  
}
```

```
char consumer( ) {  
    char item;  
    while(in - out == 0)  
        ;  
  
    item = buf[out%BSIZE];  
  
    out++;  
  
    return(item);  
}
```

Coping

- **Use what's available in the architecture to make sure all cores have the same view of memory (when necessary)**
 - lock prefix on x86
 - mfence x86 instruction
- **Use the synchronization primitives**
 - presumably the implementers knew what they were doing