



Implementing Threads 4

Two-Level Model

One Kernel Thread



Two-Level Model: Multiple Kernel Threads




Quiz 1

One kernel thread for each user thread is clearly a sufficient number of kernel threads in the two-level model. Is it necessary for maximum concurrency?

- a) there are no situations in which that number of threads is necessary, as long as there are at least as many kernel threads as processors**
- b) there must always be that number of kernel threads for the two-level model to work well**
- c) there are situations in which that number is necessary, but they occur rarely**


Deadlock



MThreads

- **Two-level threads implementation of Uthreads**
 - kernel-supported threads are POSIX threads
 - user threads based on your implementation of Uthreads
- **Effectively a multiprocessor implementation**
 - use POSIX mutexes rather than spin locks
 - use POSIX condition variables rather than the idle loop

Two-Level Model: MThreads



Synchronizing LWPs

```
uthread_switch(...) {  
    uthread_mtx_lock(&runq_mtx)  
    volatile int first = 1;  
    getcontext(&ut_curthr->ut_ctx);  
    if (!first) {  
        ...  
    }  
    setcontext(&curlwp->lwp_ctx);  
}  
lwp_switch() {  
    ...  
    ut_curthr = top_priority_thread(&runq);  
    uthread_mtx_unlock(&runq_mtx);  
    setcontext(&ut_curthr->ut_ctx);  
    ...  
}
```

Synchronizing LWPs (2)

```
uthread_switch(...) {  
    spin_lock(&runq_mtx)  
    volatile int first = 1;  
    getcontext(&ut_curthr->ut_ctx);  
    if (!first) {  
        ...  
    }  
    setcontext(&curlwp->lwp_ctx);  
}  
lwp_switch() {  
    ...  
    ut_curthr = top_priority_thread(&runq);  
    spin_unlock(&runq_mtx);  
    setcontext(&ut_curthr->ut_ctx);  
    ...  
}
```

Synchronizing LWPs (3)

```
uthread_switch(...) {  
    pthread_mutex_lock(&runq_mtx)  
    volatile int first = 1;  
    getcontext(&ut_curthr->ut_ctx);  
    if (!first) {  
        ...  
    }  
    setcontext(&curlwp->lwp_ctx);  
}  
lwp_switch() {  
    ...  
    ut_curthr = top_priority_thread(&runq);  
    pthread_mutex_unlock(&runq_mtx);  
    setcontext(&ut_curthr->ut_ctx);  
    ...  
}
```

POSIX Mutexes and MThreads

- **POSIX mutexes used to synchronize activity among LWPs**
- **Problem case**
 - **uthread (running on LWP) locks mutex**
 - **clock interrupt occurs, uthread yields LWP to another uthread**
 - **that uthread (running on same LWP) locks same mutex**
 - **deadlock: LWP attempting to lock mutex it currently has locked**
- **Solution**
 - **mask interrupts while thread has mutex locked**

Example

```
void uthread_wake(uthread_t *uthr) {  
  
    pthread_mutex_lock(&runq_mtx);  
  
    ...  
  
    // wake up thread, put it on runq  
  
    ...  
  
    pthread_mutex_unlock(&runq_mtx);  
  
}
```

Example: Fixed


```
void uthread_wake(uthread_t *uthr) {  
    uthread_nopreempt_on();  
    pthread_mutex_lock(&runq_mtx);  
  
    ...  
  
    // wake up thread, put it on runq  
  
    ...  
  
    pthread_mutex_unlock(&runq_mtx);  
    uthread_nopreempt_off();  
}
```

Thread-Local Storage in Mthreads


- `__thread thread_t *ut_curthr;`
 - **reference to the current uthread**
- `__thread lwp_t *curlwp`
 - **reference to the current LWP (POSIX thread)**
- **Thread-Local Storage accesses are not async-signal safe!**
- **Must turn off preemption while using TLS**
 - **otherwise thread could be preempted and later resumed on another LWP**
 - **TLS pointer refers to the wrong item!**

Interrupts, Etc.

Interrupts

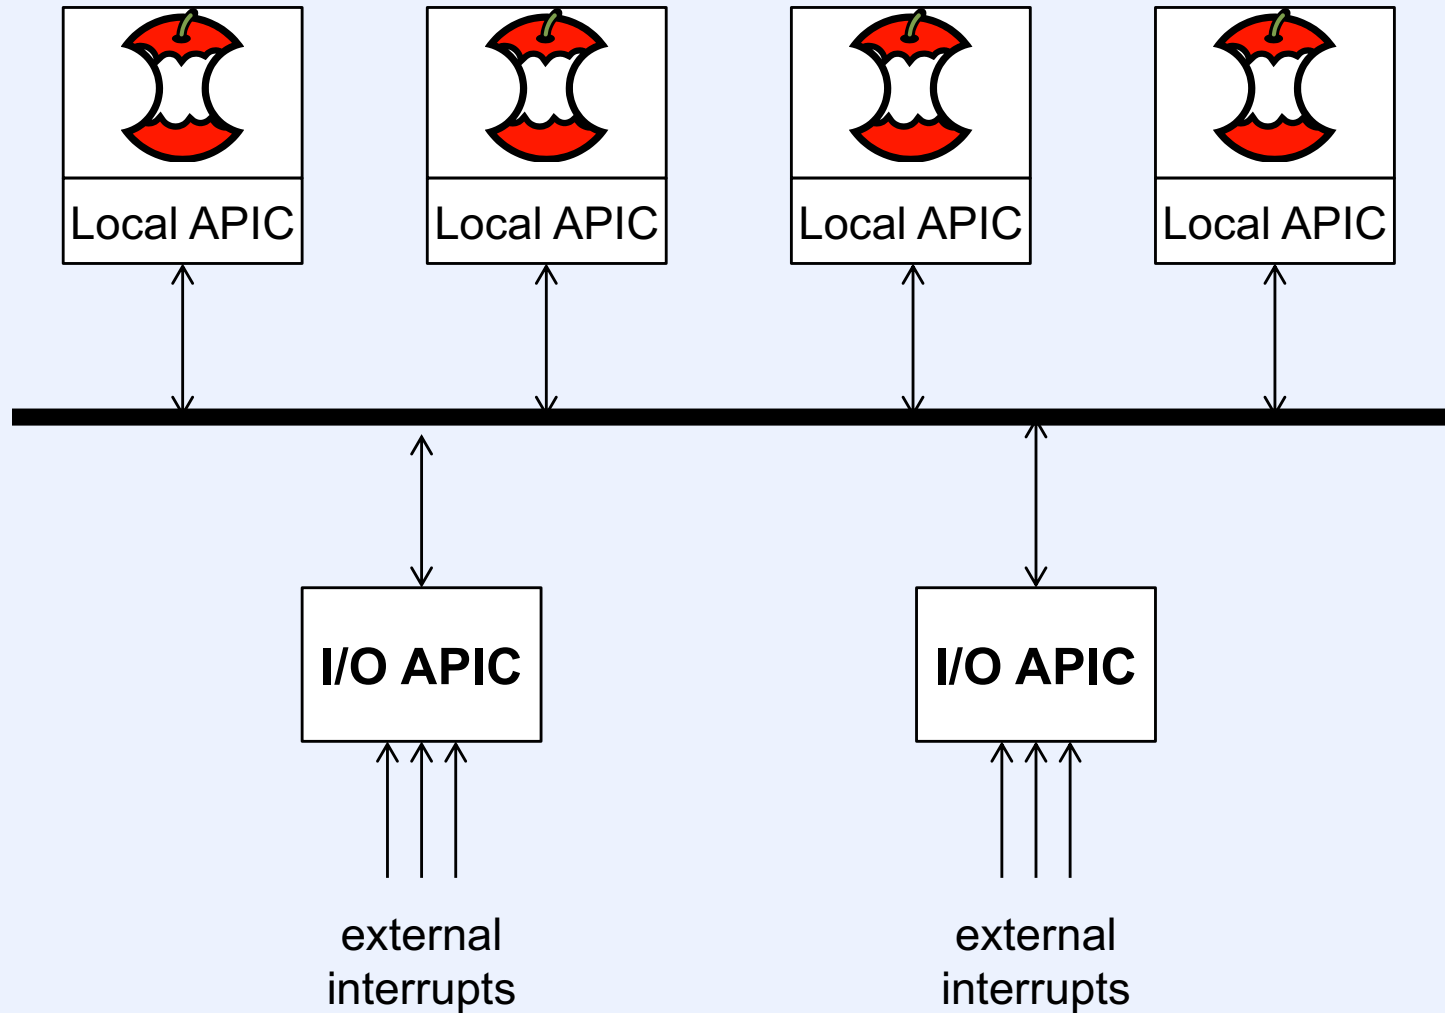


Current thread's
user stack



Current thread's
kernel stack

x86 Interrupt Architecture



Dealing with Interrupts

- **Interrupt comes from external source**
- **Must execute code to handle it**
- **Which stack?**
 - **use current (kernel) stack**
 - or**
 - **use separate stack**

Use the Current (Kernel) Stack

- **Borrowed from the current thread**
 - requires all threads to have sufficiently large kernel stacks
- **Interrupted thread may not concurrently execute**
 - would corrupt interrupt handler's stack frames
- **Interrupt handler should not block**
 - may be waiting on interrupted thread
 - deadlock
 - threads must mask interrupts to protect data structures shared with interrupt handlers

Hierarchical Interrupt Masking

- **Interrupts assigned priorities 1 through n**
 - **current interrupt priority level i**
 - **interrupts with priorities 1 through i masked**
 - **IPL set to i**
 - **when interrupt of priority i is handled**
 - **when IPL set explicitly to i**
- **Raising the IPL**
 - **protects data structures**
 - **good!**
 - **delays response to lower priority interrupts**
 - **bad!**

Separate Interrupt Stack

- **Single stack used strictly by interrupt handlers**
 - hardware saves current register state on interrupt stack
 - no interrupt-handling space required for kernel stacks
 - all can be smaller
 - in principle, interrupted thread may continue to execute
 - won't corrupt interrupt handler's frames

Quiz 2

Is interrupt-masking still required?

- a) yes
- b) no

(Non-Quiz) Questions

- 1) Is interrupt masking still required? *Done***
- 2) May interrupt handler block?**
- 3) Does it work on multiprocessors?**

Answers

1) Masking is still required

- if nothing else, to protect data structures shared by multiple interrupt handlers
- to prevent deadlock if spin locks are used

2) Interrupt handlers should not block

- would have to block with raised IPL
- results in lengthy time period with interrupts masked
 - delayed response

3) Yes, but each processor has its own interrupt stack (if separate interrupt stacks are part of the architecture)

Interrupt Threads

- **Give each interrupt instance its own stack**
 - handlers effectively execute as separate threads
 - interrupted thread continues to run
 - but interrupts remain masked while interrupt thread is processing interrupt

Effect of Interrupts

- **Normally don't directly affect current thread**
 - thread is interrupted
 - interrupt is dealt with
 - thread is resumed
- **I/O-completion interrupts**
 - may result in waking up higher-priority thread (that was waiting for the I/O completion)
- **Clock interrupts**
 - may trigger end of time slice

Synchronization and Interrupts

- **Non-preemptive kernels**
 - threads running in privileged mode yield the processor only voluntarily
 - involuntary thread switches happen only to threads in user mode
 - end of time slice
 - higher-priority thread is made runnable
 - inter-thread in-kernel synchronization is easy
- **Preemptive kernels**
 - threads running in privileged mode may be forced to yield the processor
 - inter-thread sync in kernel is not easy

Non-Preemptive Kernel Sync.

```
int X = 0;
```

```
void AccessXThread() {  
    int oldIPL;  
    oldIPL = setIPL(IXLevel);  
    X = X+1;  
    setIPL(oldIPL);  
}
```

```
void AccessXInterrupt() {  
    ...  
    X = X+1;  
    ...  
}
```

Quiz 2

```
int X = 0;
```

```
void AccessXThread() {  
    int oldIPL;  
    oldIPL = setIPL(IXLevel);  
    X = X+1;  
    setIPL(oldIPL);  
}
```

```
void AccessXInterrupt() {  
    ...  
    X = X+1;  
    ...  
}
```

This works

- a) only on a single-core system with a non-preemptive kernel
- b) also on a single-core system with a preemptive kernel
- c) also on multi-core systems

Disk I/O

```
int disk_write(...) {  
    ...  
    startIO(); // start disk operation  
    ...  
    enqueue(disk_waitq, CurrentThread);  
    thread_switch();  
    // wait for disk operation to  
    // complete  
    ...  
}
```

```
void disk_intr(...) {  
    thread_t *thread;  
    ...  
    // handle disk interrupt  
    ...  
    thread = dequeue(disk_waitq);  
    if (thread != 0) {  
        enqueue(RunQueue, thread);  
        // wakeup waiting thread  
    }  
    ...  
}
```

Improved Disk I/O

```
int disk_write(...) {  
    ...  
    oldIPL = setIPL(diskIPL);  
    startIO();          // start disk operation  
    ...  
    enqueue(disk_waitq, CurrentThread);  
    thread_switch();  
    // wait for disk operation to complete  
    setIPL(oldIPL);  
    ...  
}
```

Modified *thread_switch*

```
void thread_switch() {
    thread_t *OldThread;
    int oldIPL;
    oldIPL = setIPL(HIGH_IPL);
    // protect access to RunQueue by masking all interrupts
    while(queue_empty(RunQueue)) {
        // repeatedly allow interrupts, then check RunQueue
        setIPL(0); // IPL == 0 means no interrupts are masked
        setIPL(HIGH_IPL);
    }
    // We found a runnable thread
    OldThread = CurrentThread;
    CurrentThread = dequeue(RunQueue);
    swapcontext(OldThread->context, CurrentThread->context);
    setIPL(oldIPL);
}
```

Preemptive Kernels on MP

- What's different?
- A thread accesses a shared data structure:
 1. it might be *interrupted* by an interrupt handler (running on its processor) that accesses the same data structure
 2. *another thread* running on another processor might access the same data structure
 3. it might be forced to *give up its processor* to another thread, either because its time slice has expired or it has been preempted by a higher-priority thread
 4. an *interrupt handler* running on *another processor* might access the same data structure

Solution?

```
int X = 0;
SpinLock_t L = UNLOCKED;
```

```
void AccessXThread() {
    SpinLock(&L);
    X = X+1;
    SpinUnlock(&L);
}
```

```
void AccessXInterrupt() {
    ...
    SpinLock(&L);
    X = X+1;
    SpinUnlock(&L);
    ...
}
```

Solution ...

```
int X = 0;  
SpinLock_t L = UNLOCKED;
```

```
void AccessXThread() {  
    MaskInterrupts();  
    SpinLock(&L);  
    X = X+1;  
    SpinUnlock(&L);  
    UnMaskInterrupts();  
}
```

```
void AccessXInterrupt() {  
    ...  
    SpinLock(&L);  
    X = X+1;  
    SpinUnlock(&L);  
    ...  
}
```

Quiz 3

We have a single-core system with a preemptible kernel. We're concerned about data structure X , which is accessed by kernel threads as well as by the interrupt handler for dev.

- a) It's sufficient for threads to mask dev interrupts while accessing X**
- b) It's sufficient for threads to mask all interrupts while accessing X**
- c) In addition to a, threads must lock (blocking) mutexes before accessing X**
- d) c doesn't work. Instead, threads must lock spinlocks before accessing X**