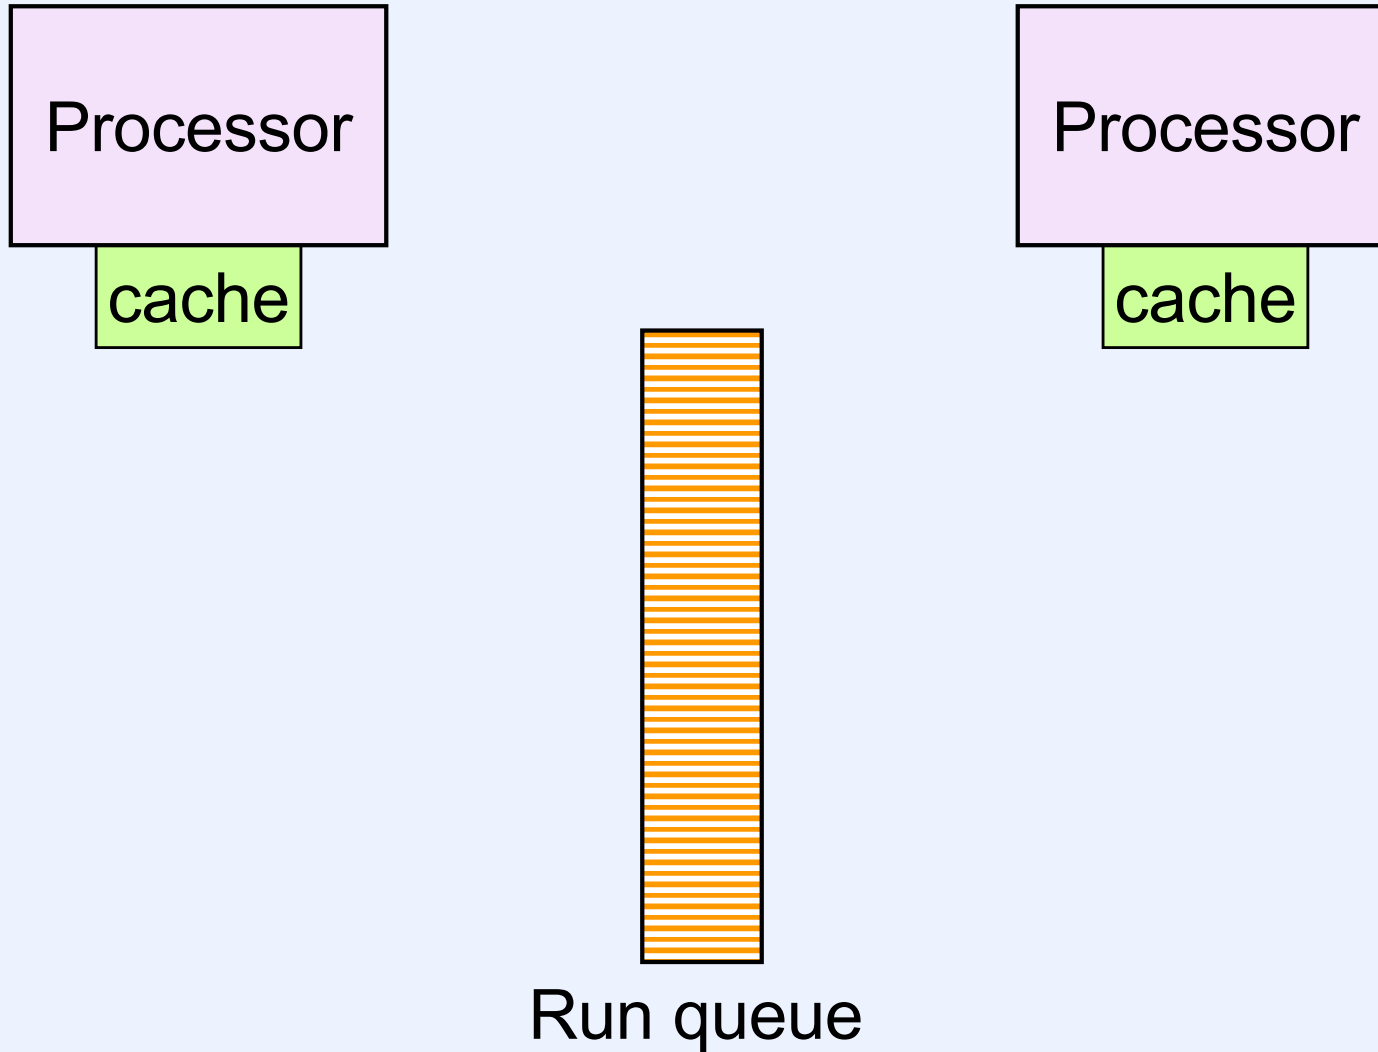


Scheduling Part 3

Diagram



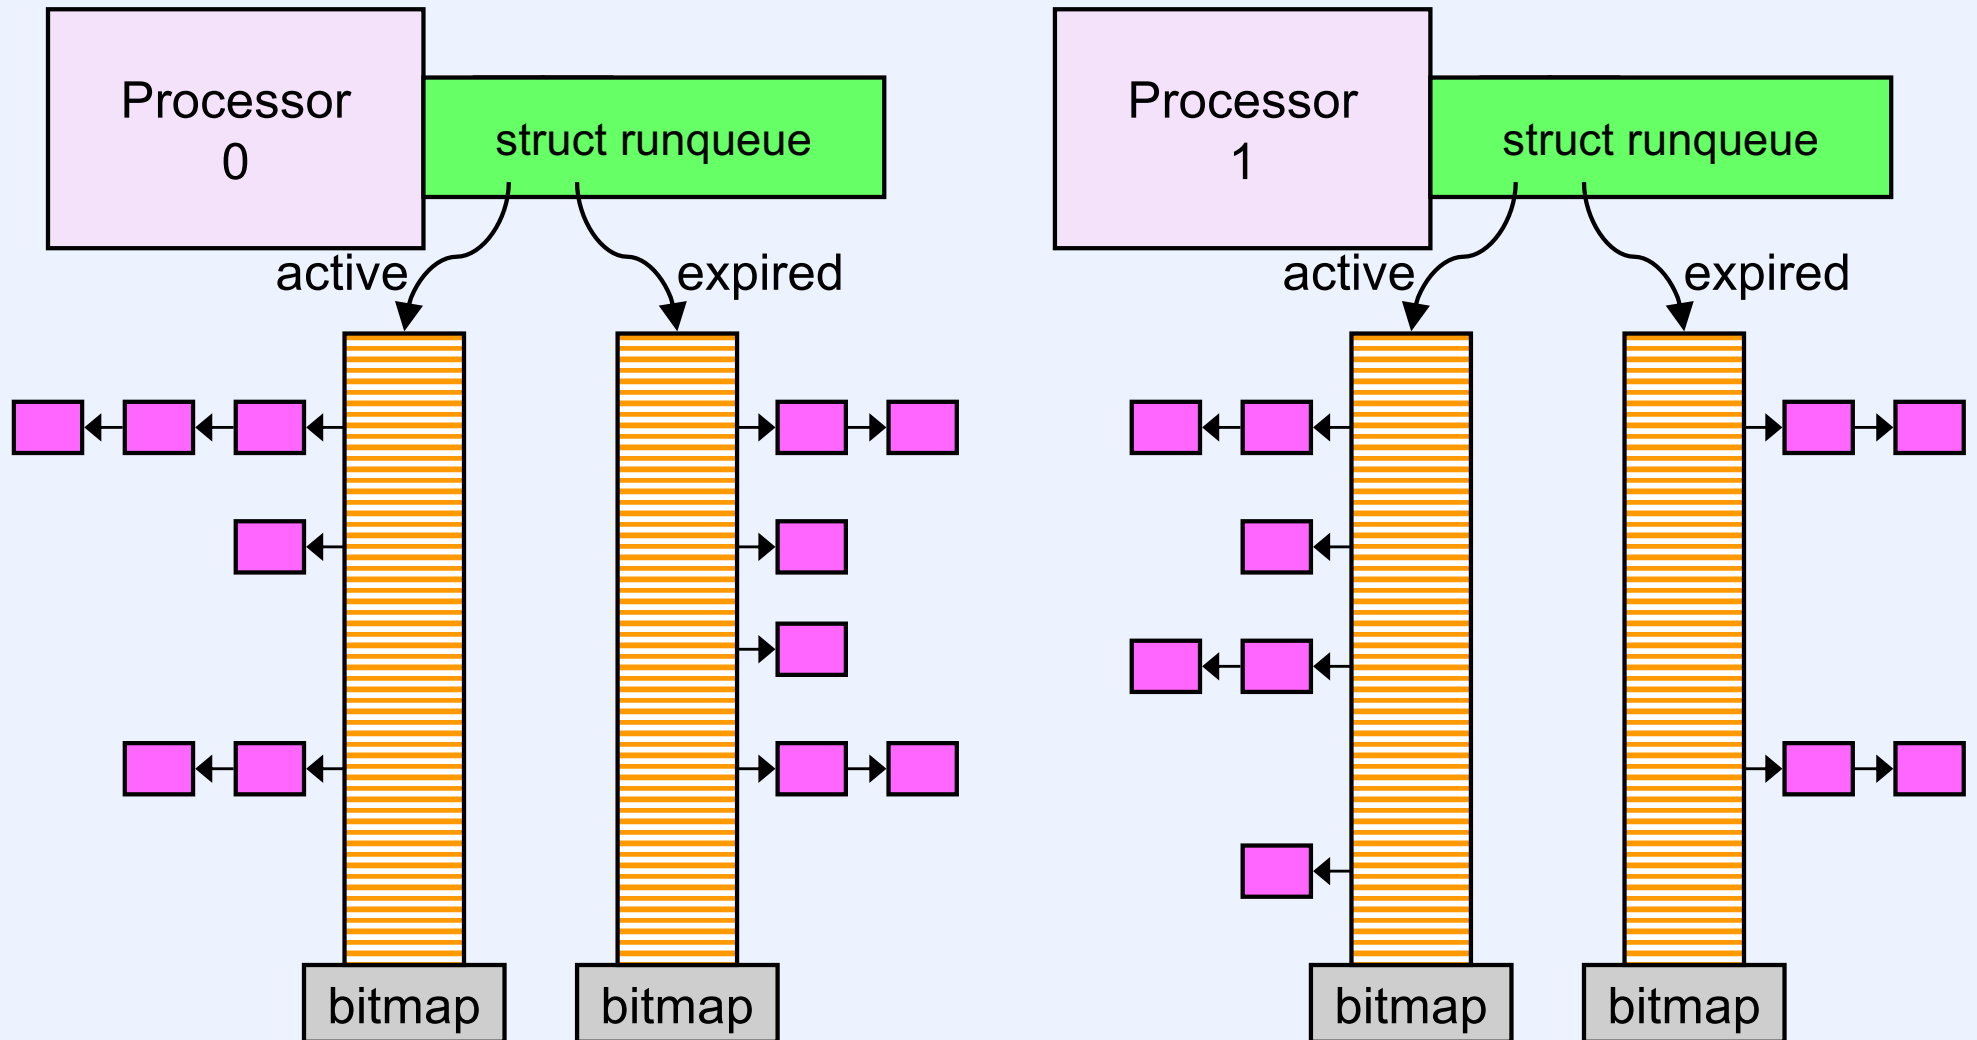
Old Scheduler: Problems

- **$O(n)$ execution**
- **Poor interactive performance with heavy loads**
- **SMP contention for run-queue lock**
- **SMP affinity**
 - **cache “footprint”**

O(1) Scheduler

- **All concerns of old scheduler plus:**
 - efficient, scalable execution
 - identify and favor interactive processes
 - good SMP performance
 - minimal lock overhead
 - processor affinity

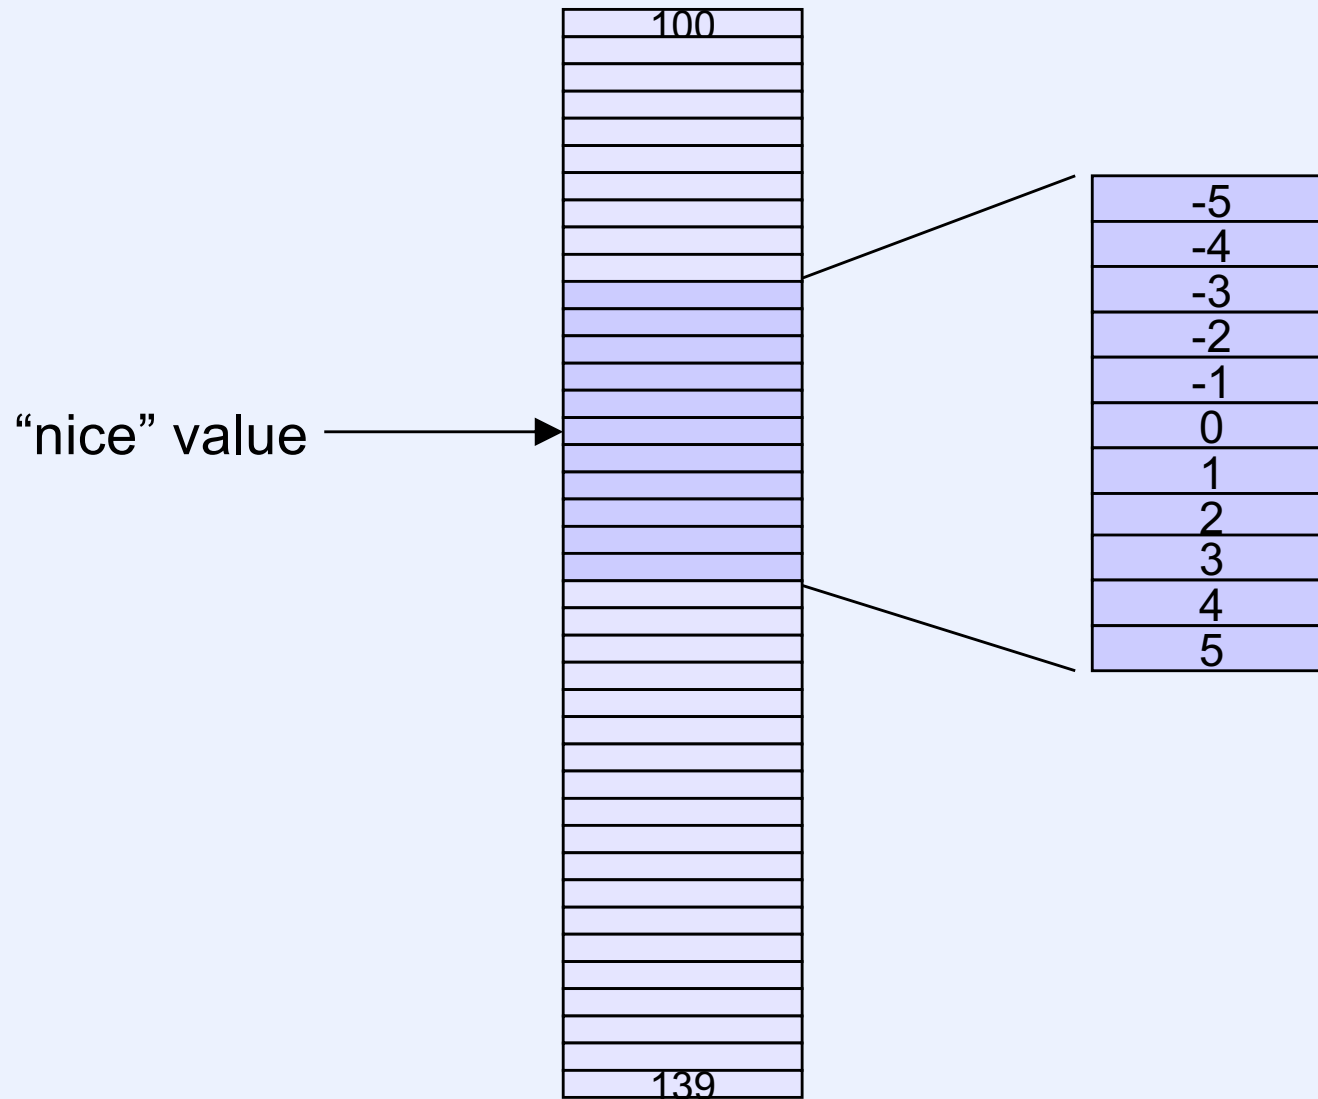
O(1) Scheduler: Data Structures



O(1) Scheduler: Queues

- **Two queues per processor**
 - **active: processes with remaining time slice**
 - **expired: processes whose time slices expired**
 - **each queue is an array of lists of processes of the same priority**
 - **bitmap indicates which priorities have processes**
 - **processors scheduled from private queues**
 - **infrequent lock contention**
 - **good affinity**

O(1) Scheduler: Priorities



O(1) Scheduler: Actions

- **Process switch**
 - pick best priority from active queue
 - if empty, switch active and expired
 - new process's time slice is function of its priority
- **Wake up**
 - priority is boosted or dropped depending on sleep time
 - interactive processes are defined as those whose priority is above a certain threshold
- **Time-slice expiration**
 - normal processes join expired queue
 - real-time join active queue


O(1) Scheduler: Load Balancing

- **Processors with empty queues steal from busiest processor**
 - checked every millisecond
- **Processors with relatively small queues also steal from busiest processor**
 - checked every 250 milliseconds

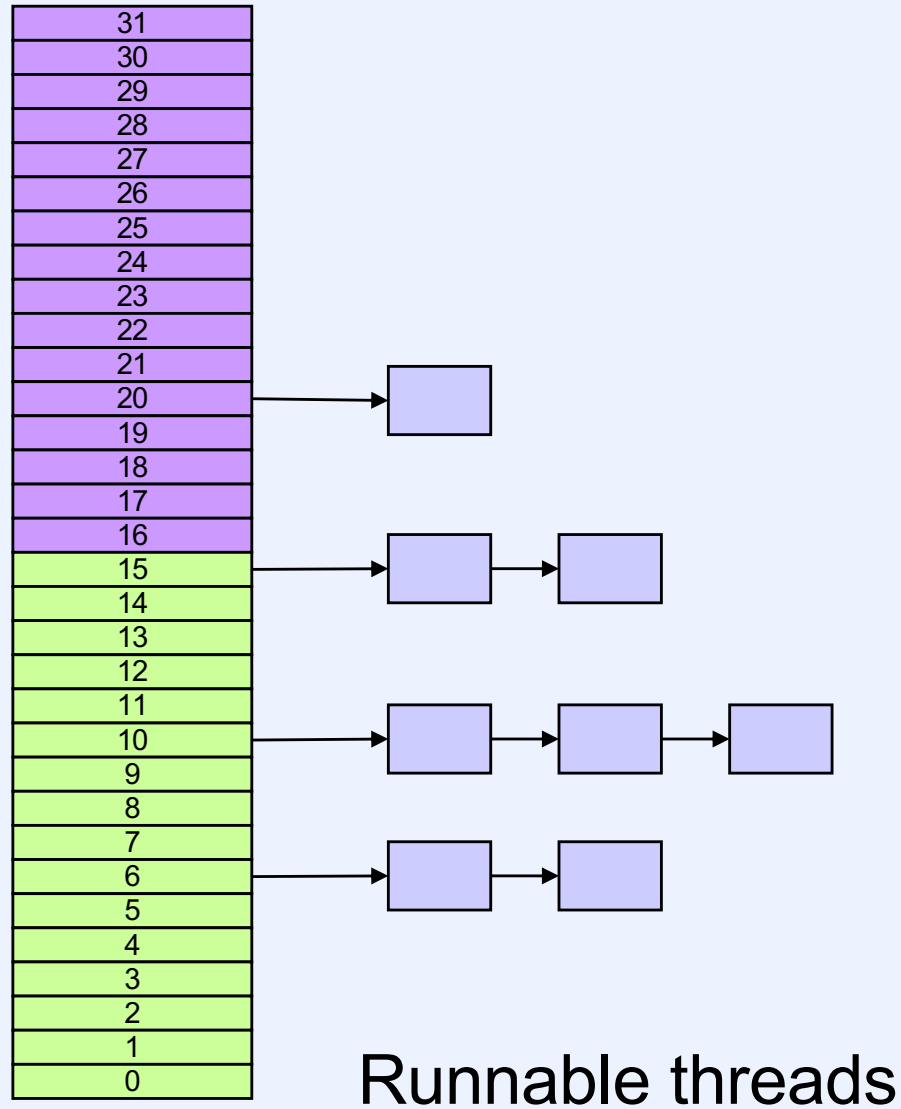
Scheduling in Windows

- Handling “normal” interactive and compute-bound threads
- Real-time threads
- Multiple processors

Priorities



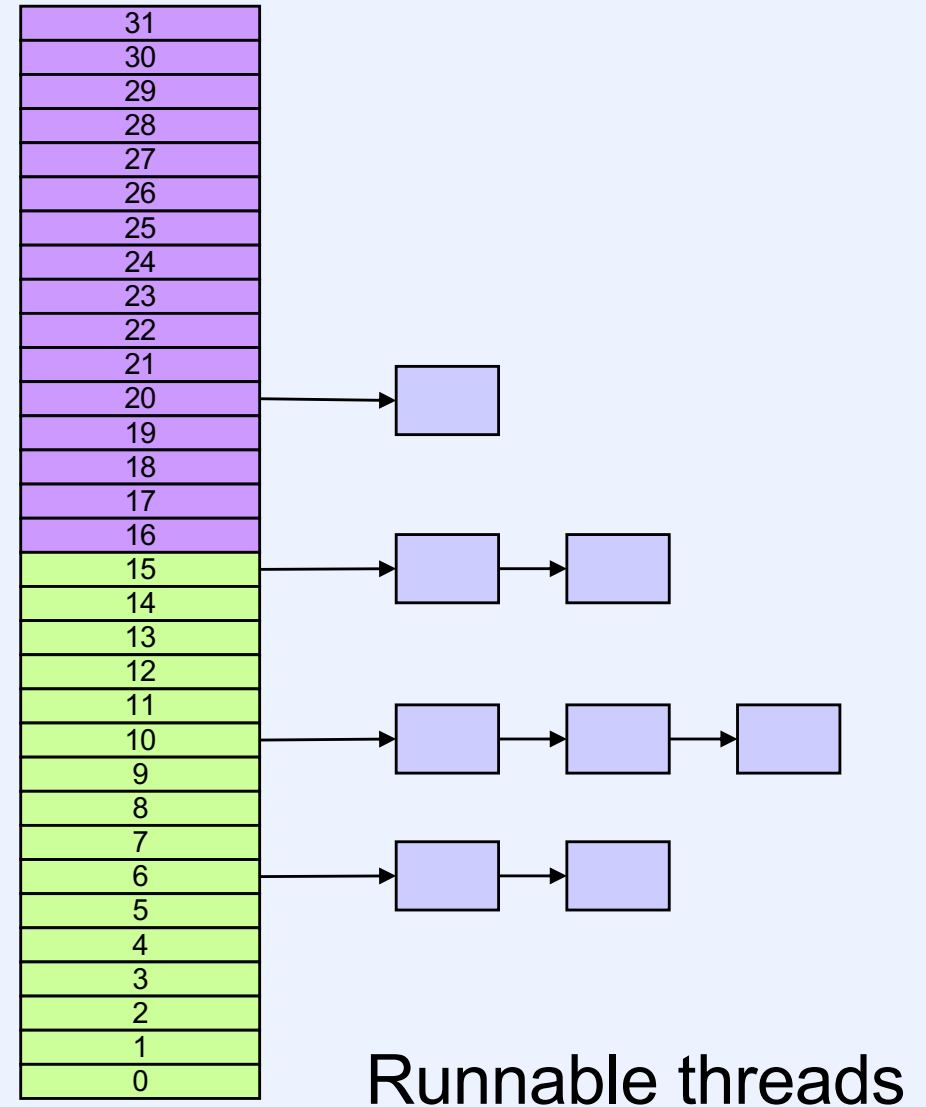
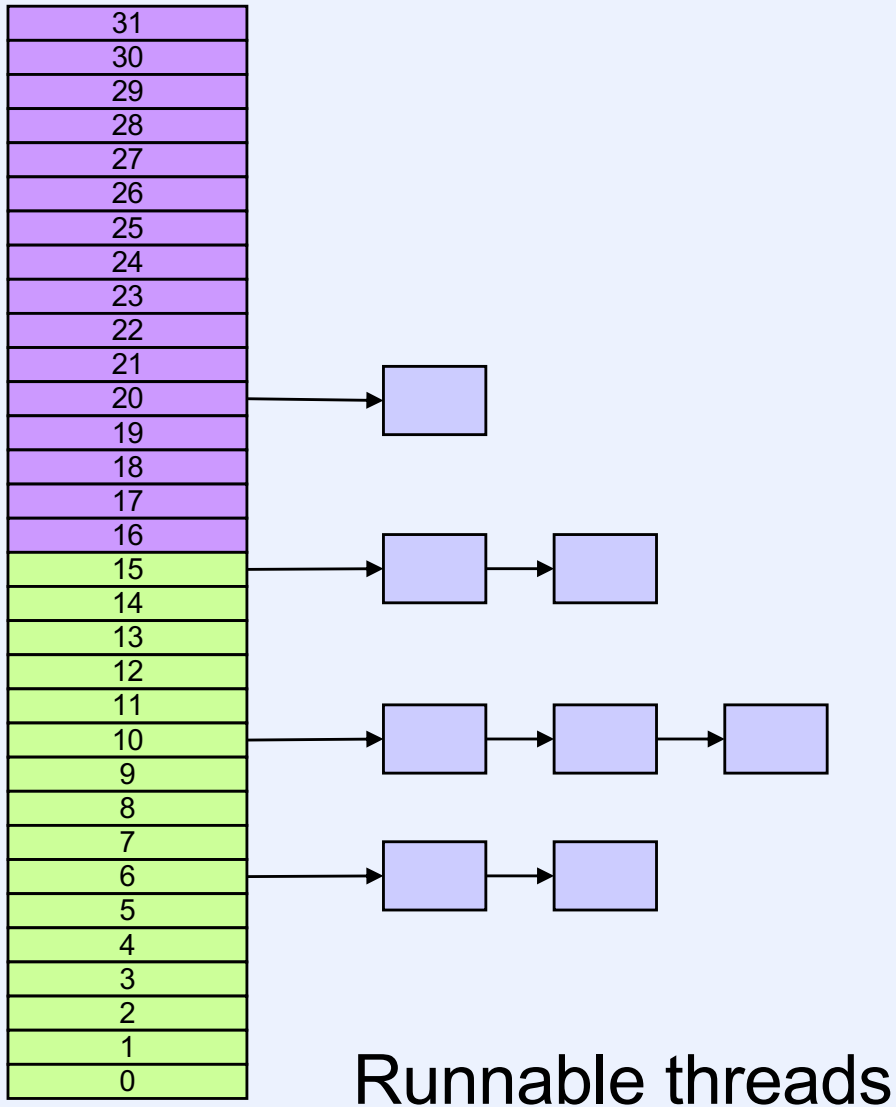
Uniprocessor Windows



Improving Real Time

- **Multimedia applications need 80% of processor time**
- **Make sure normal applications get at least 20%**
- **How?**
- **Windows solution: MMCSS**
 - **multimedia class scheduler service**
 - **dynamically manage multimedia threads**
 - **run at real-time priority 80% of time**
 - **run at normal priority 20% of time**

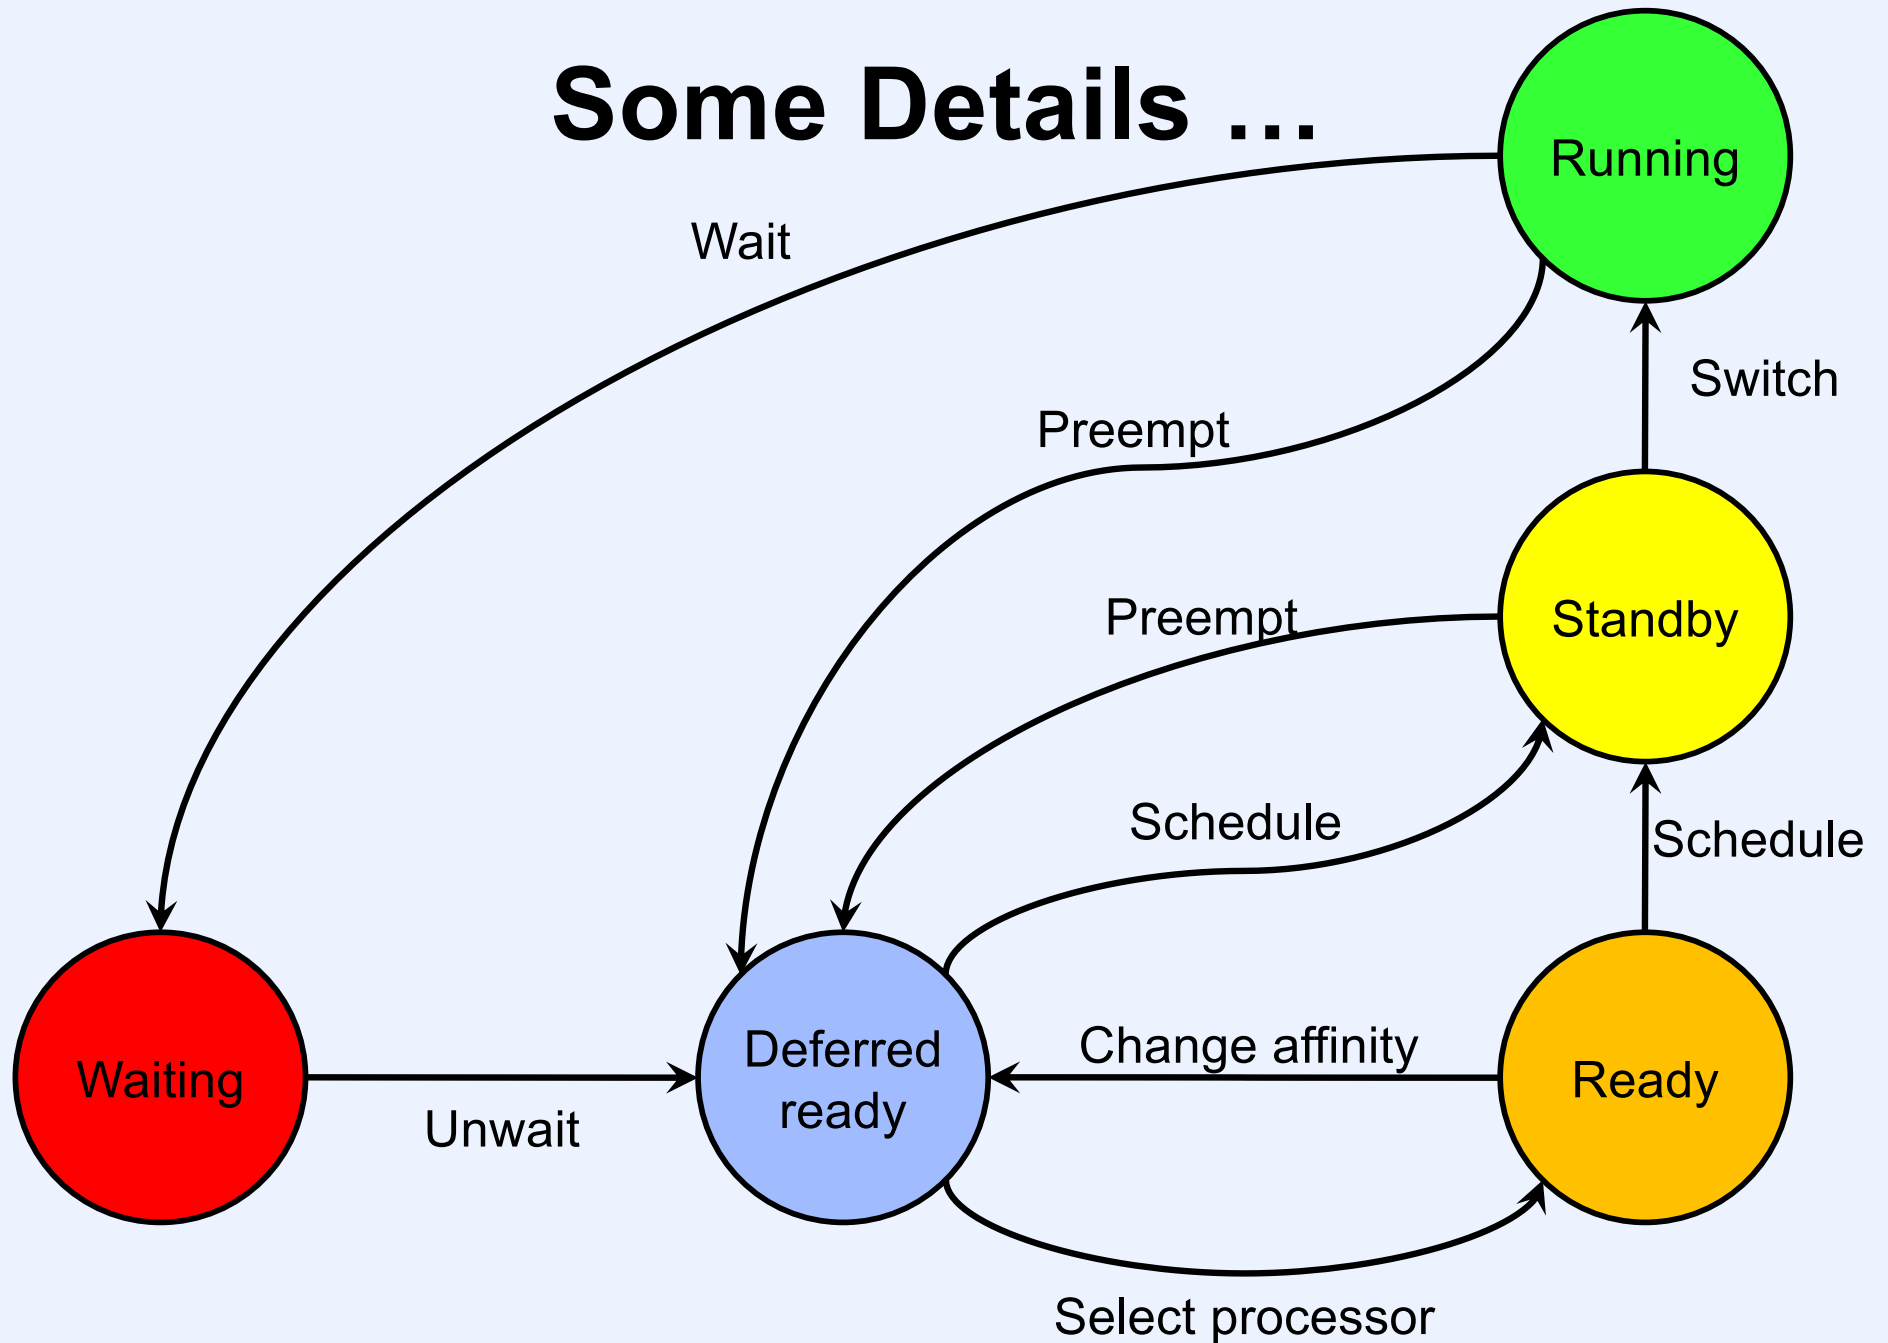
Multiprocessor Windows



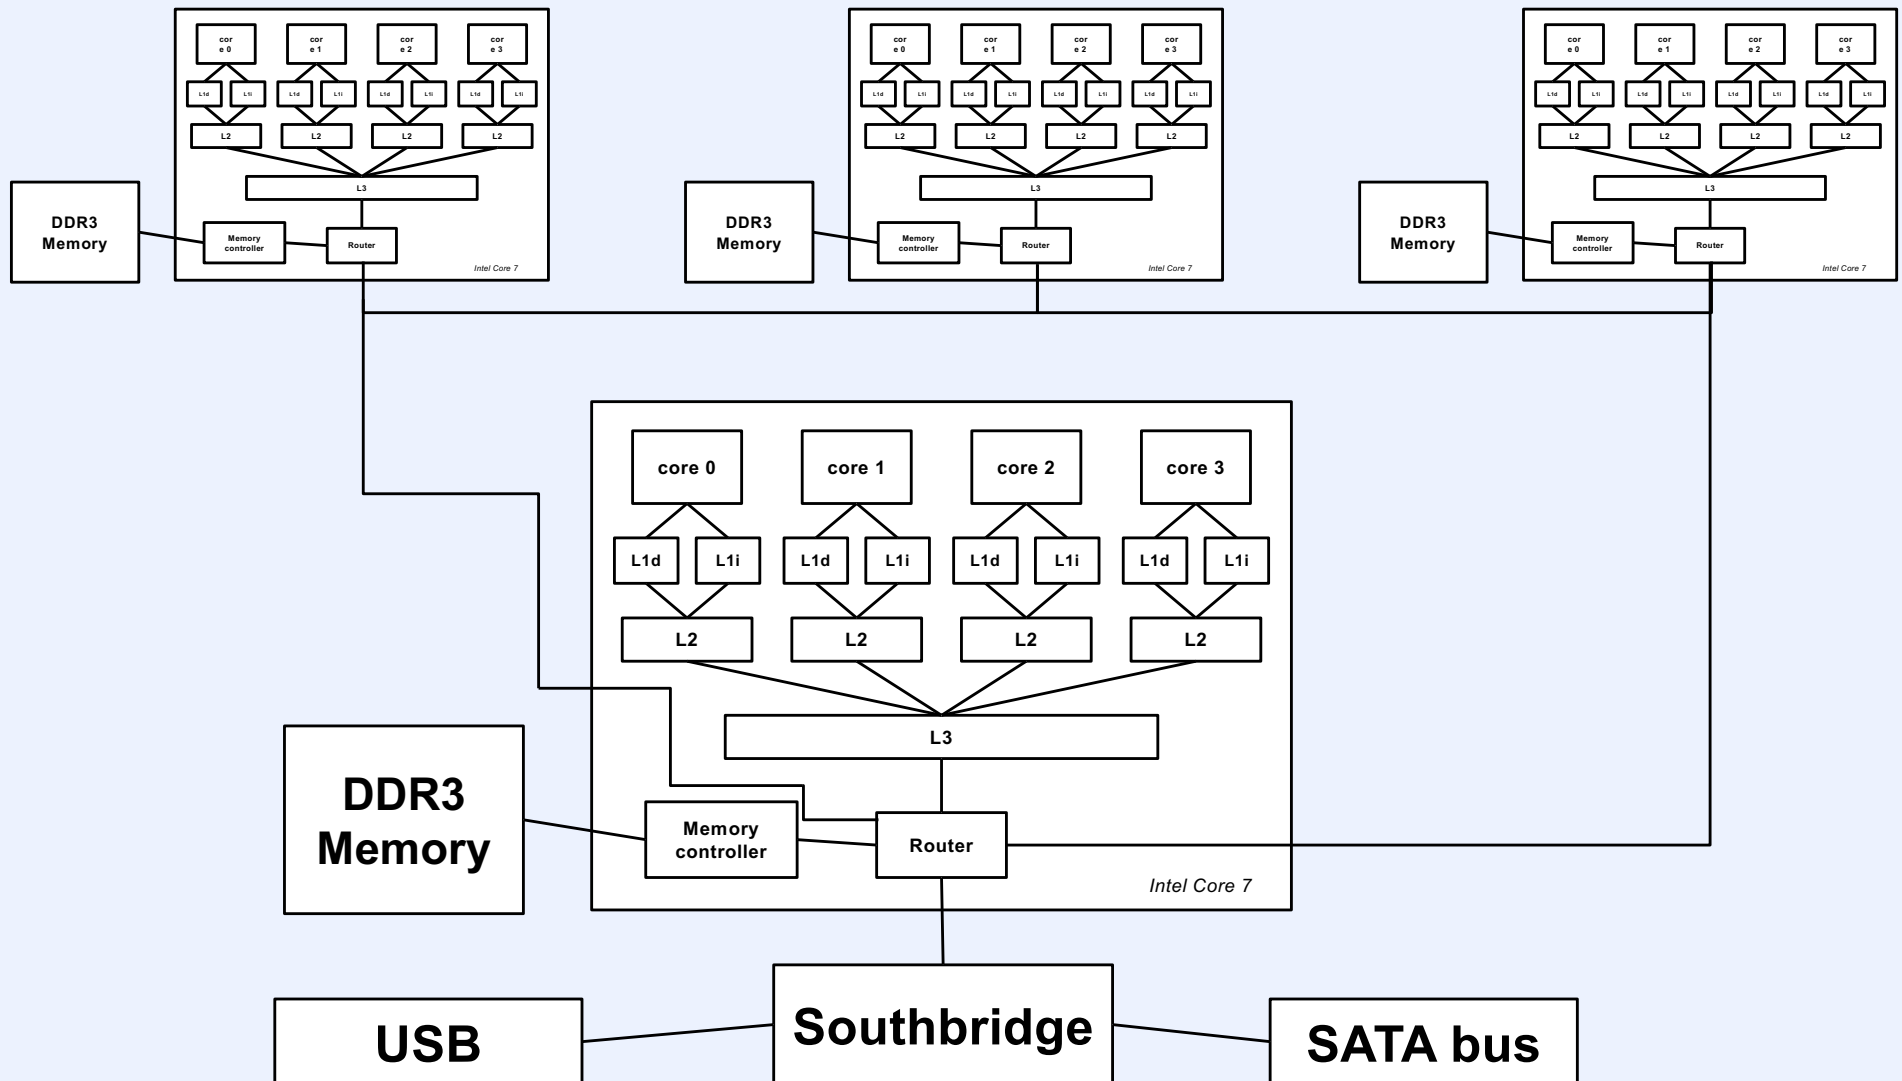
Which Processor?

- Newly created thread assigned *ideal processor*
 - randomly chosen
- May also set *affinity mask*
 - may be scheduled only on processors in mask
- Scheduling decision:
 - if idle processors available
 - first preference: ideal processor (if idle)
 - second preference: most recent processor (if idle)
 - otherwise
 - joins run queue of ideal processor

Some Details ...



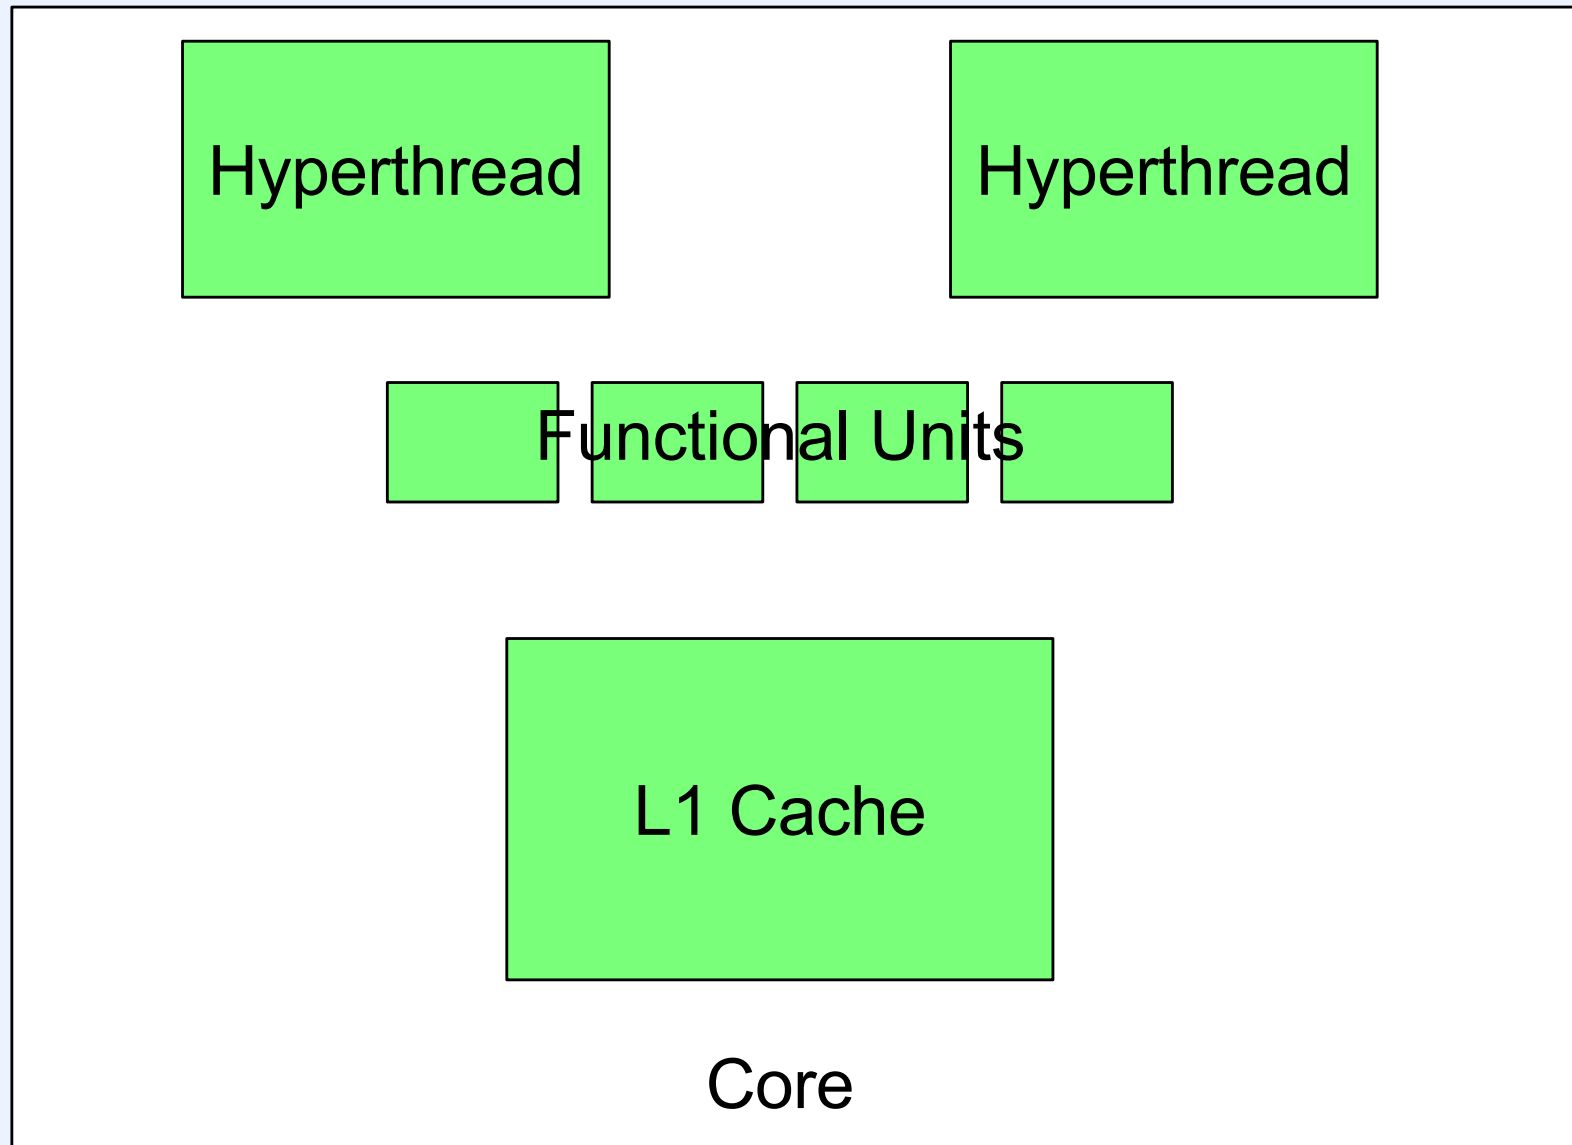
NUMA



Scheduling Concerns

- **Hyperthreads**
 - two instruction streams sharing same functional units and same L1 cache
- **How long does cache footprint matter?**
 - what cache parameters are important?
- **When is it a good idea to put a thread on:**
 - a different core?
 - a different NUMA node?

Hyperthreads


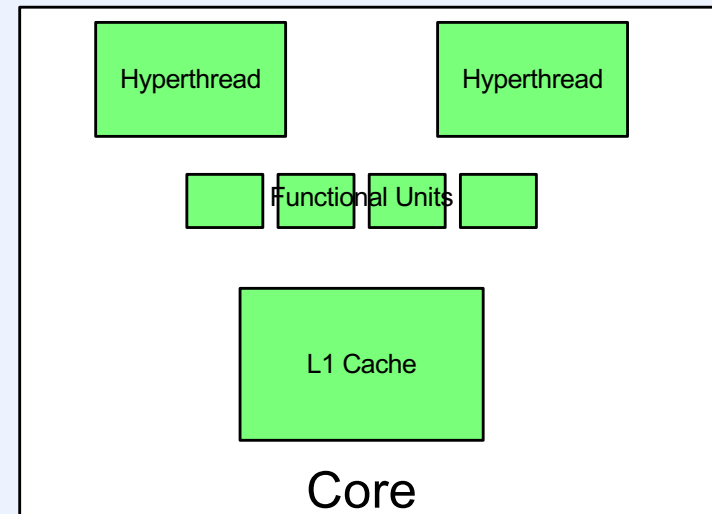



Quiz 1

We have a two-core processor with two hyperthreads per core. We have exactly two runnable threads.

- a) for best performance, each thread should run on a separate core**
- b) it doesn't matter which hyperthreads are used to run the two threads, as long as two hyperthreads are used**

Cores

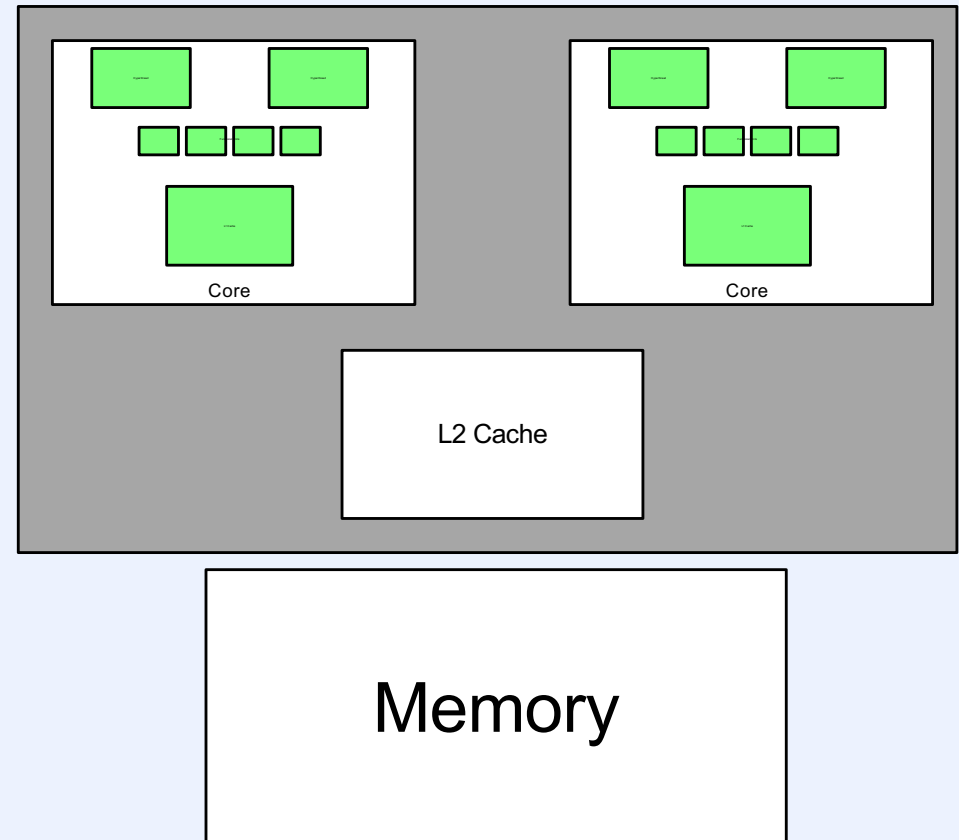



Quiz 2

We have a two-core processor with four hyperthreads per core. We have four runnable threads.

- a) it doesn't matter which hyperthreads are used to run the four threads, as long as four hyperthreads are used**
- b) it does matter: we should make sure that two hyperthreads of each core are used**
- c) it does matter: not only should we make sure that two hyperthreads of each core are used, but we also need to consider whether any of the threads are in the same process**

NUMA Nodes






Quiz 3

We have a system with two NUMA nodes. Most of the hyperthreads are busy on one of the nodes, the other node is completely idle. What operation performed by a thread would make the thread (or the thread it creates) a good candidate to move to the other node?

- a) pthread_create**
- b) fork**
- c) execv**
- d) waitpid**

File Systems Part 1

Files



Requirements

- **Permanent storage**
 - resides on disk (or alternatives)
 - survives software and hardware crashes
 - (including loss of disk?)
- **Quick, easy, and efficient**
 - satisfies needs of most applications
 - how do applications use permanent storage?

Applications

- **Software development**
 - text editors
 - linkers and loaders
 - source-code control
 - **Document processing**
 - editing
 - browsing
 - **Web stuff**
 - serving
 - browsing
 - **Program execution**
 - paging
-

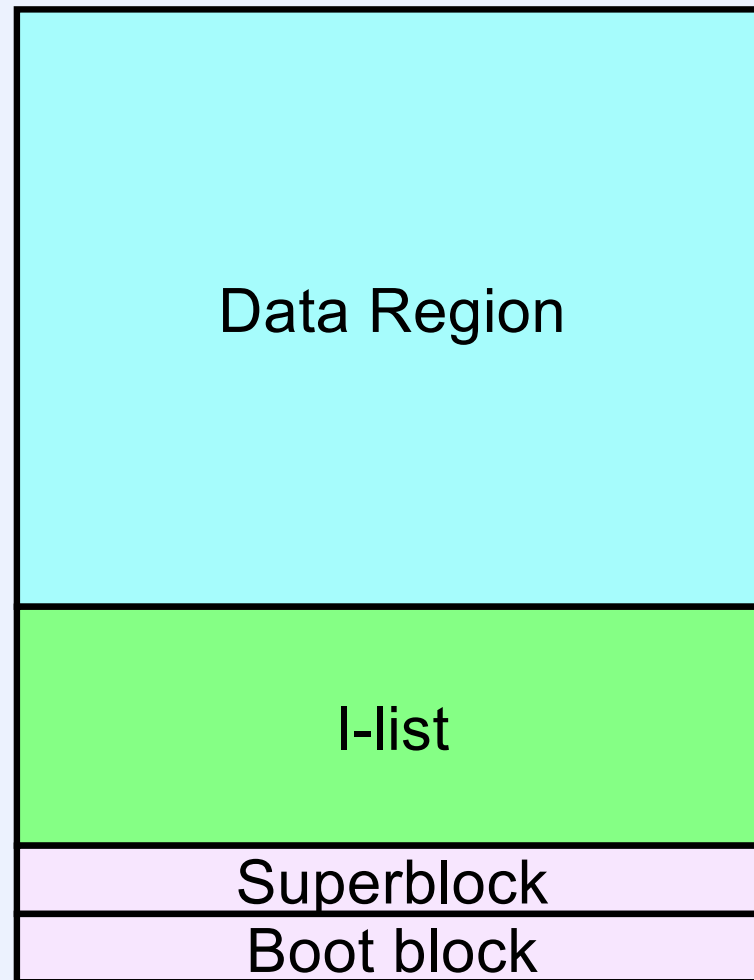
Needs

- **Directories**
 - convenient naming
 - fast lookup
- **File access**
 - sequential is very common!
 - “random access” is relatively rare

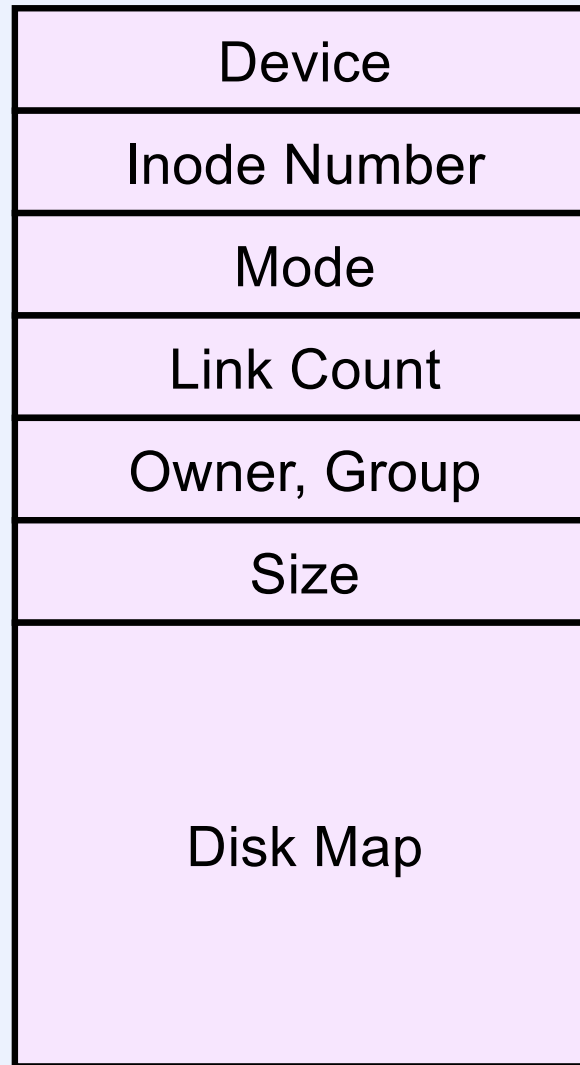
S5FS

- **A simple file system**
 - slow
 - not terribly tolerant to crashes
 - reasonably efficient in space
 - no compression
- **Concerns**
 - on-disk data structures
 - file representation
 - free space
- **It's the Weenix file system!**


S5FS Layout



S5FS: Inode



Disk Map



Quiz 4

Suppose a new file is created. (At this point it occupies zero blocks.) Then one byte is written to it at byte offset $(266 \times 2^{10}) + 1$. Assume the block size is 2^{10} and block addresses occupy four bytes. How many blocks are required to represent the file, not counting its inode?

- a) 1**
- b) 3**
- c) 270**
- d) more than 270**

Quiz 5

Suppose one now writes to all locations in the file, from its beginning up to the location written to in the previous slide (byte offset $(266 \times 2^{10}) + 1$). How many blocks are required to represent the file, not counting its inode?

- a) 1**
- b) 3**
- c) 270**
- d) more than 270**