

Next edition of OS @ Brown

Fall 2026, Instructor: Malte Schwarzkopf (me)

If you enjoyed the material and want to help make the next version of the course, **apply to TA in Fall 2026.**

Either way, I want to hear from you:

- 1. what is great about the course, and*
- 2. what you'd change!*

Reach out to meet or share your thoughts:




malte@brown.edu

Security Part 7

Live Anonymous Q&A:

<https://tinyurl.com/cs1670feedback>

Integrity



Biba Model

- Integrity is what's important
 - no-write-up
 - no-read-down

Quiz 1

You're concerned about downloading malware to your computer and very much want to prevent it from affecting your computer. Which would be the most appropriate policy to use?

- a) no write up**
- b) no read up**
- c) no write down**
- d) no read down**

Windows and MAC

- **Concerns**
 - viruses
 - spyware
 - etc.
- **Installation is an integrity concern**
- **Solution**
 - adapt Biba model

Windows Integrity Control

- **No-write-up**
- **All subjects and objects assigned a level**
 - **untrusted**
 - **low integrity**
 - **Internet Explorer/Edge**
 - **medium integrity**
 - **default**
 - **high integrity**
 - **system integrity**
- **Object owners may lower integrity levels**
- **May set *no-read-up* on an object**

Industrial-Strength Security

- **Target:**
 - embezzlers



Clark-Wilson Model

- Integrity and confidentiality aren't enough
 - there must be control over how data is produced and modified
 - well formed transactions



Cash account

withdrawals here



**Accounts-payable
account**

**must be matched
by entries here**

- Separation of duty
 - steps of transaction must involve multiple people

Mandatory Access Control (MAC)

Implementing MAC

- Label subjects and objects
- Security policy makes decisions based on labels and context

**registrar
person**

**d.o.f.
person**

**CS
person**

**web-server
process**

**student
record**

**salary
record**

**password
file**

**public
database**


Quiz 2

I have a file that I accidentally set as having rw permission for everyone (0666). You have a process that has opened my file rw. I discover this and immediately change the permissions to 0600 (access only by me). Can your process still read and write the file?

- a) It can read and write**
- b) It can read, but not write**
- c) It can write, but not read**
- d) It can do neither**

SELinux

- **Security-Enhanced Linux**
 - MAC-based security
 - labels on all subjects and objects
 - policy-specification language
- **Use in Android (since v4.3)**
- **Deny by default**



SELinux Examples (1)

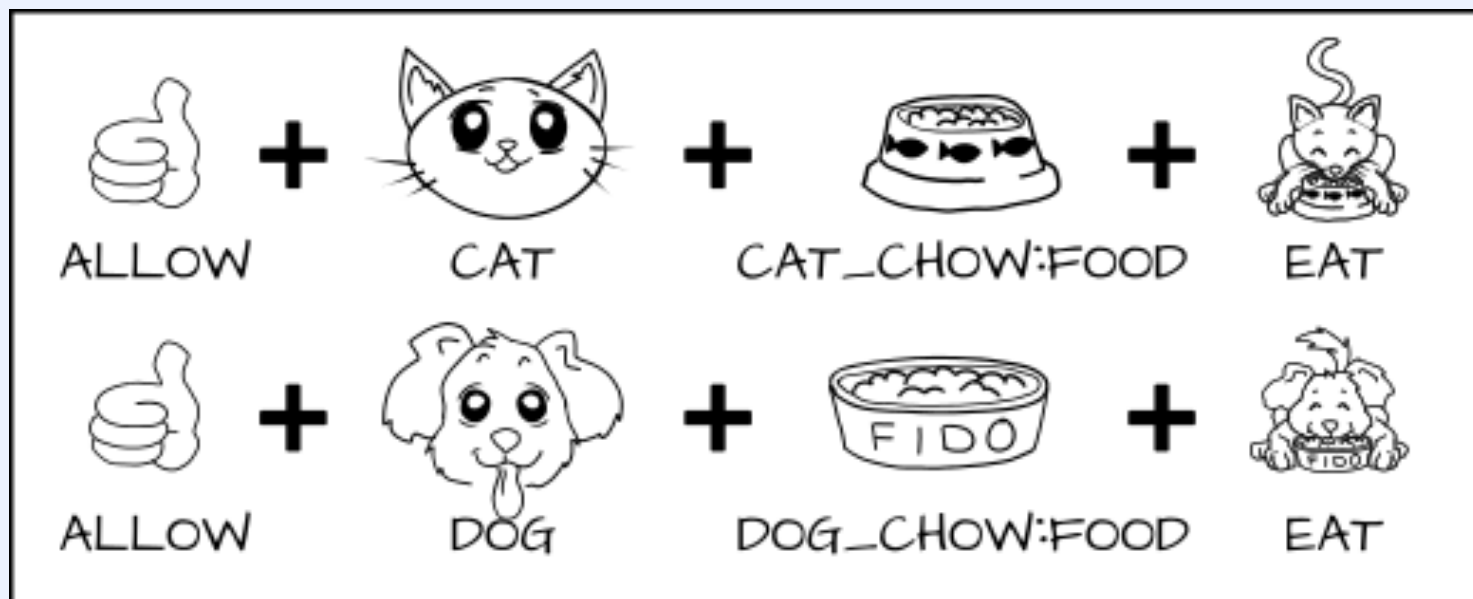


- Publicly readable files assigned type *public_t*
- Subjects of normal users run in domain *user_t*
- */etc/passwd*: viewable, but not writable, by all
- */etc/shadow*: protected
- SELinux rules

```
allow user_t public_t : file read
```

- normal users may read public files


```
allow passwd_t passwd_data_t : file {read write}
```

- */etc/shadow* is of type *passwd_data_t*
- subjects in *passwd_t* domain may read/write */etc/shadow*



Cartoon credit: <https://opensource.com/business/13/11/selinux-policy-guide>

Result



Cartoon credit: <https://opensource.com/business/13/11/selinux-policy-guide>

SELinux Examples (2)

- **How does a program get into the *passwd_t* domain?**
 - **assume passwd program is of type *passwd_exec_t***

```
allow passwd_t passwd_exec_t : file entrypoint
allow user_t passwd_exec_t : file execute
allow user_t passwd_t : process transition
type_transition user_t passwd_exec_t : process
passwd_t
```

Quiz 3

We've seen how the setuid feature in Unix is used to allow normal users to change their passwords in /etc/shadow.

- a) This approach actually isn't secure, which is among the reasons why SELinux exists**
- b) The approach is secure and thus SELinux doesn't really add any additional protection to /etc/shadow**
- c) The approach is secure but there are other potential /etc/shadow-related vulnerabilities that SELinux helps deal with**

Off-the-Shelf SELinux

- **Strict policy**
 - normal users in *user_r* role
 - users allowed to be administrators are in *staff_r* role
 - but may run admin commands only when in *sysadm_r* role
 - policy requires > 20,000 rules
 - tough to live with
- **Targeted policy**
 - targets only “network-facing” applications
 - everything else in *unconfined_t* domain
 - ~11,000 rules

Capability-Based Systems

Confused-Deputy Problem

- **The system has a pay-per-use compiler**
 - keeps billing records in file `/u/sys/comp/usage`
 - puts output in file you provide
 - `/u/you/comp.out`
- **The concept of a pay-per-use compiler annoys you**
 - you send it a program to compile
 - you tell it to put your output in `/u/sys/comp/usage`
 - it does
 - it's confused
 - you win

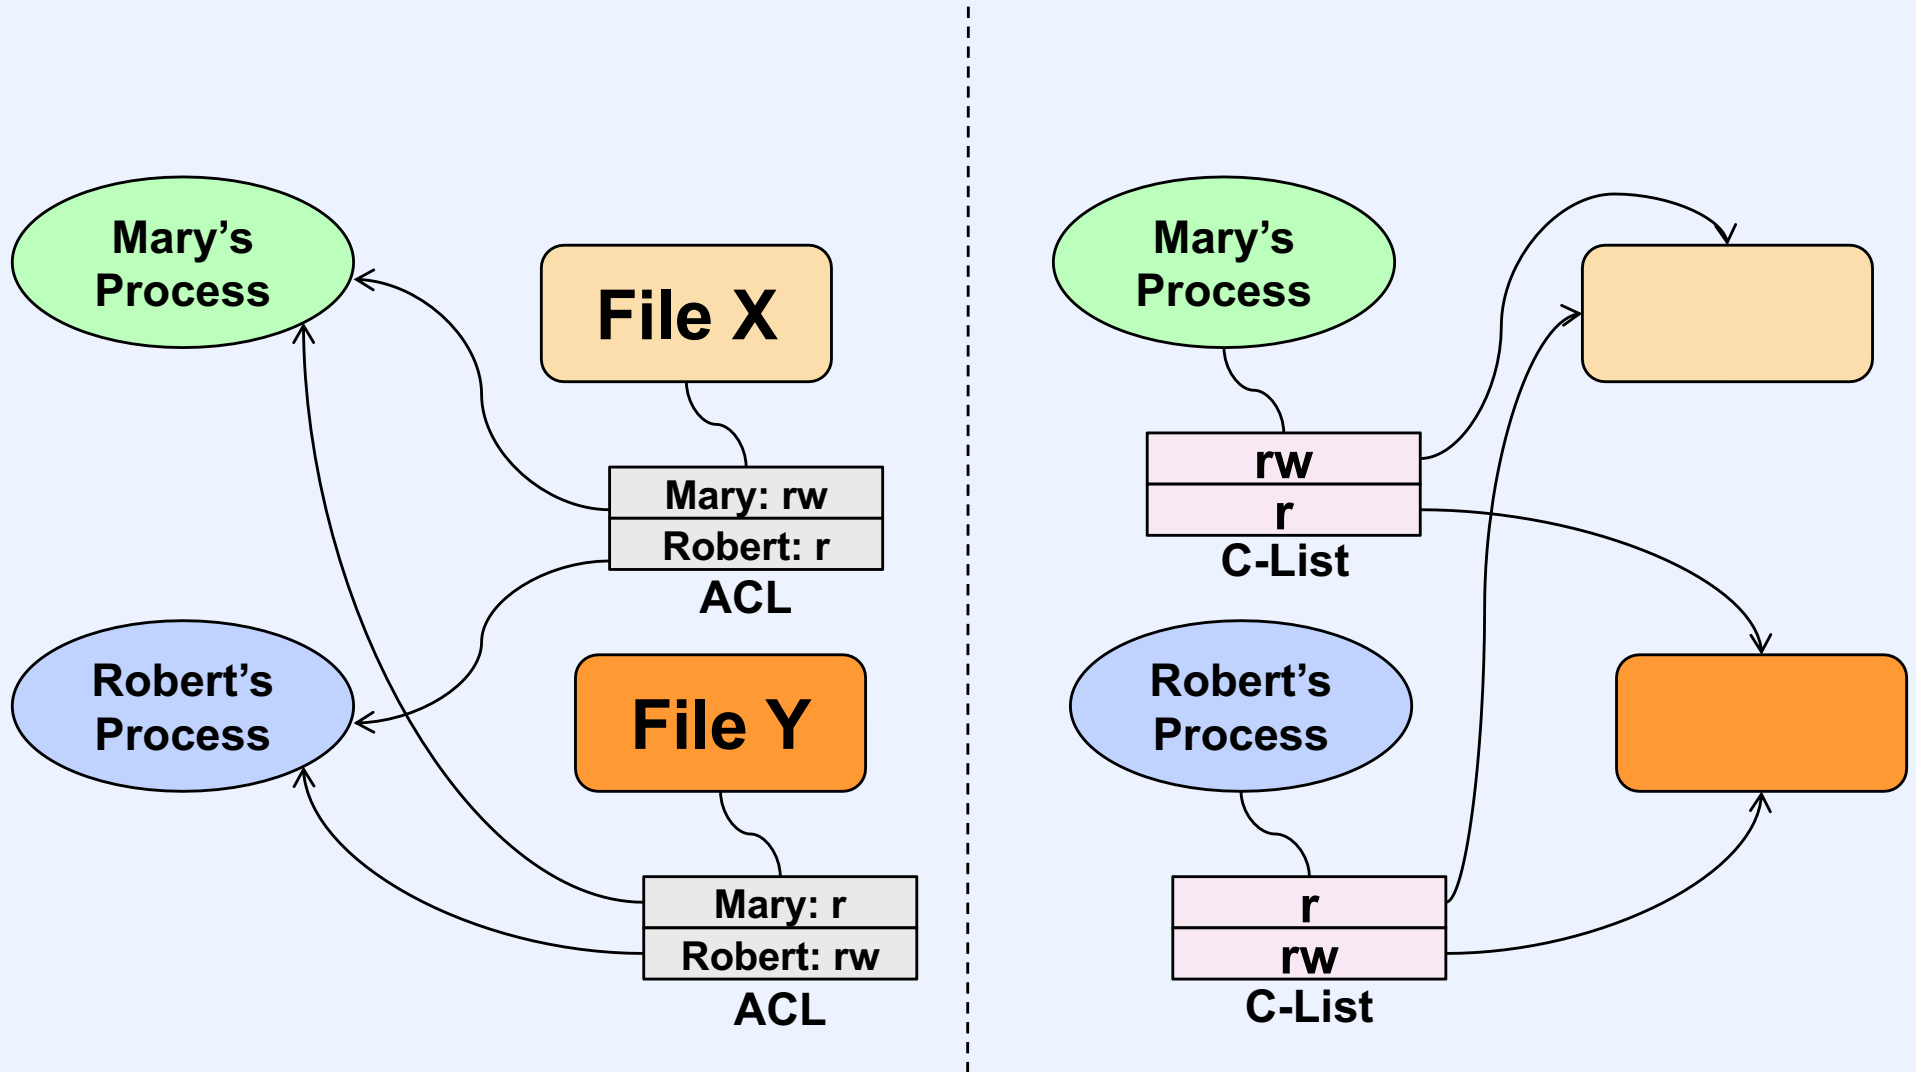
Unix and Windows to the Rescue

- **Unix**
 - compiler is “su-to-compiler-owner”
- **Windows**
 - client sends impersonation token to compiler
- **Result**
 - malicious deputy problem
- **Could be solved by passing file descriptors**
 - not done
 - should be ...

Authority


- **Pure ACL-based systems**
 - authority depends on subject's user and group identities
- **Pure capability-based systems**
 - authority depends upon capabilities possessed by subject

ACLs vs. C-Lists




More General View

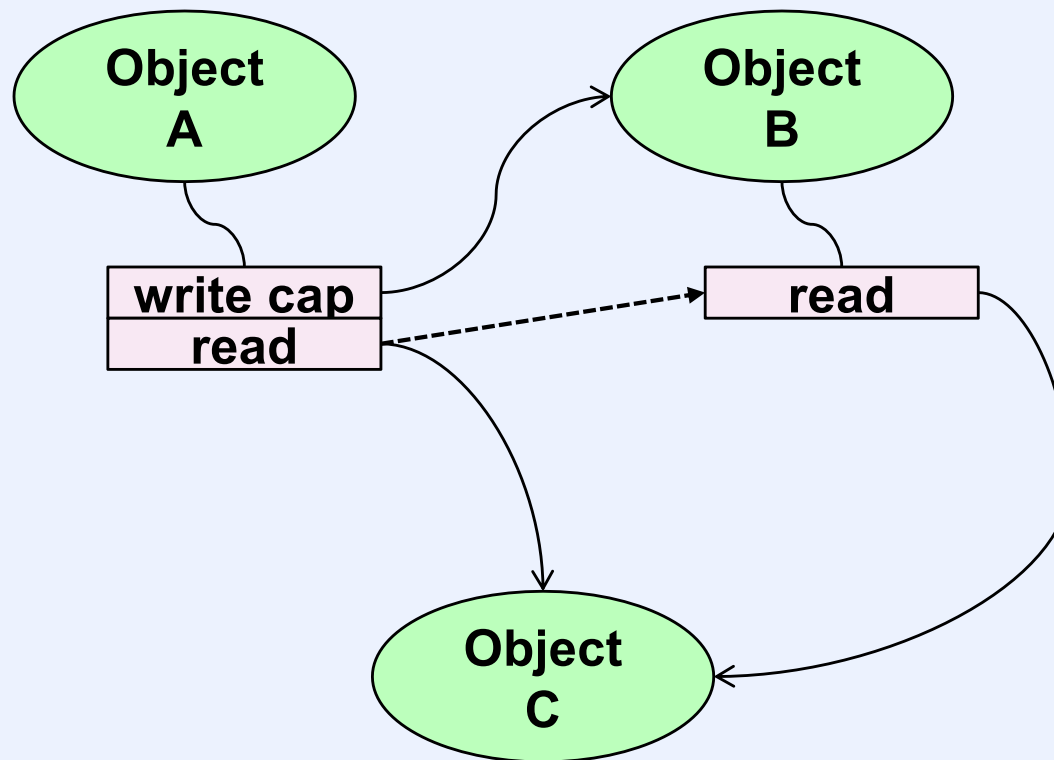
- Subjects and resources are *objects* (in the OO sense)



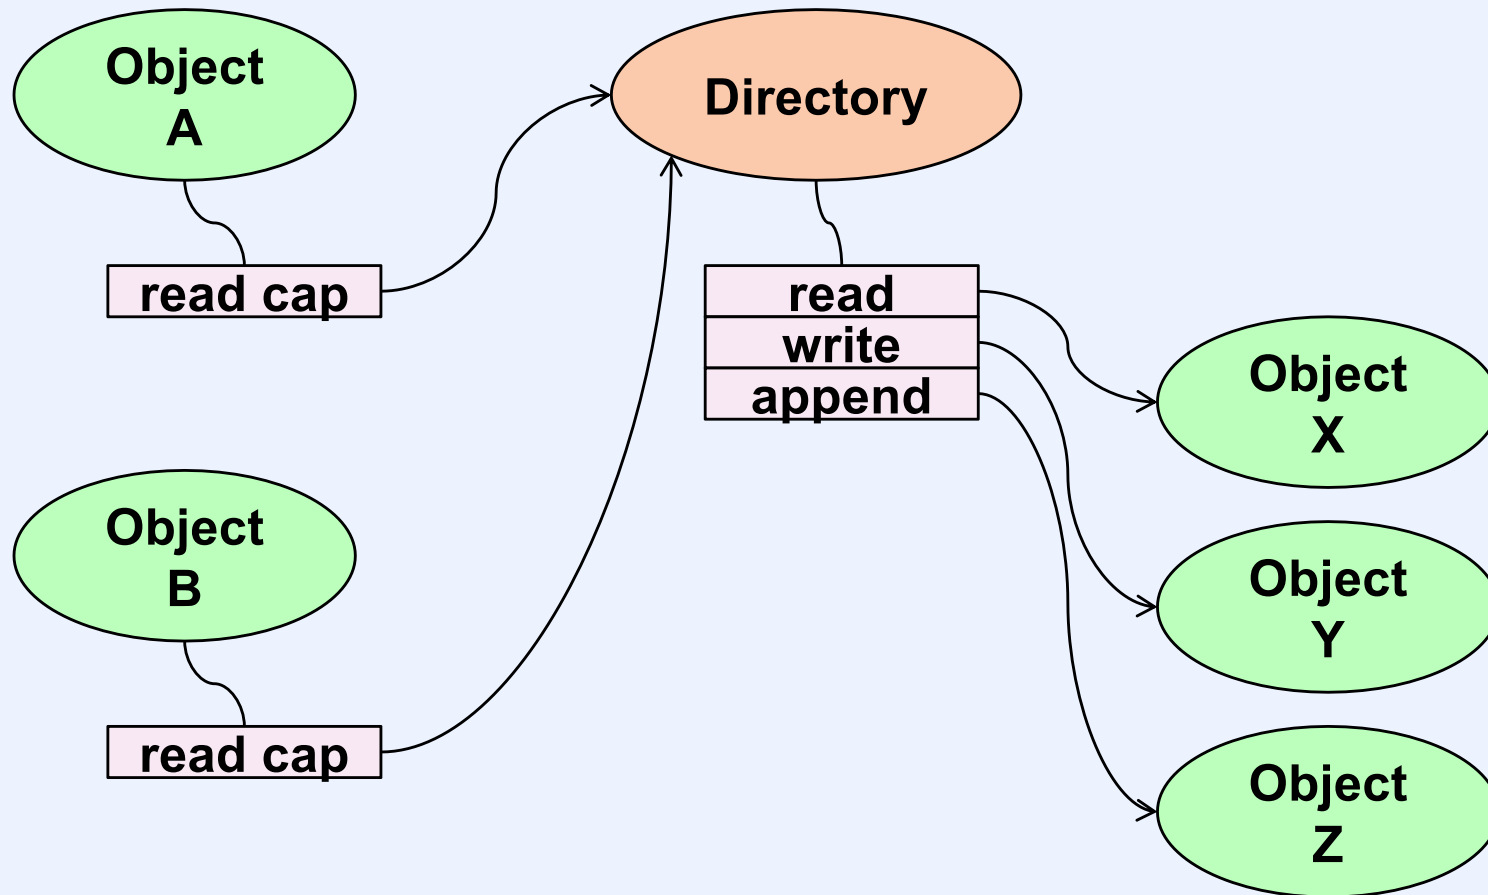
Copying Capabilities (1)



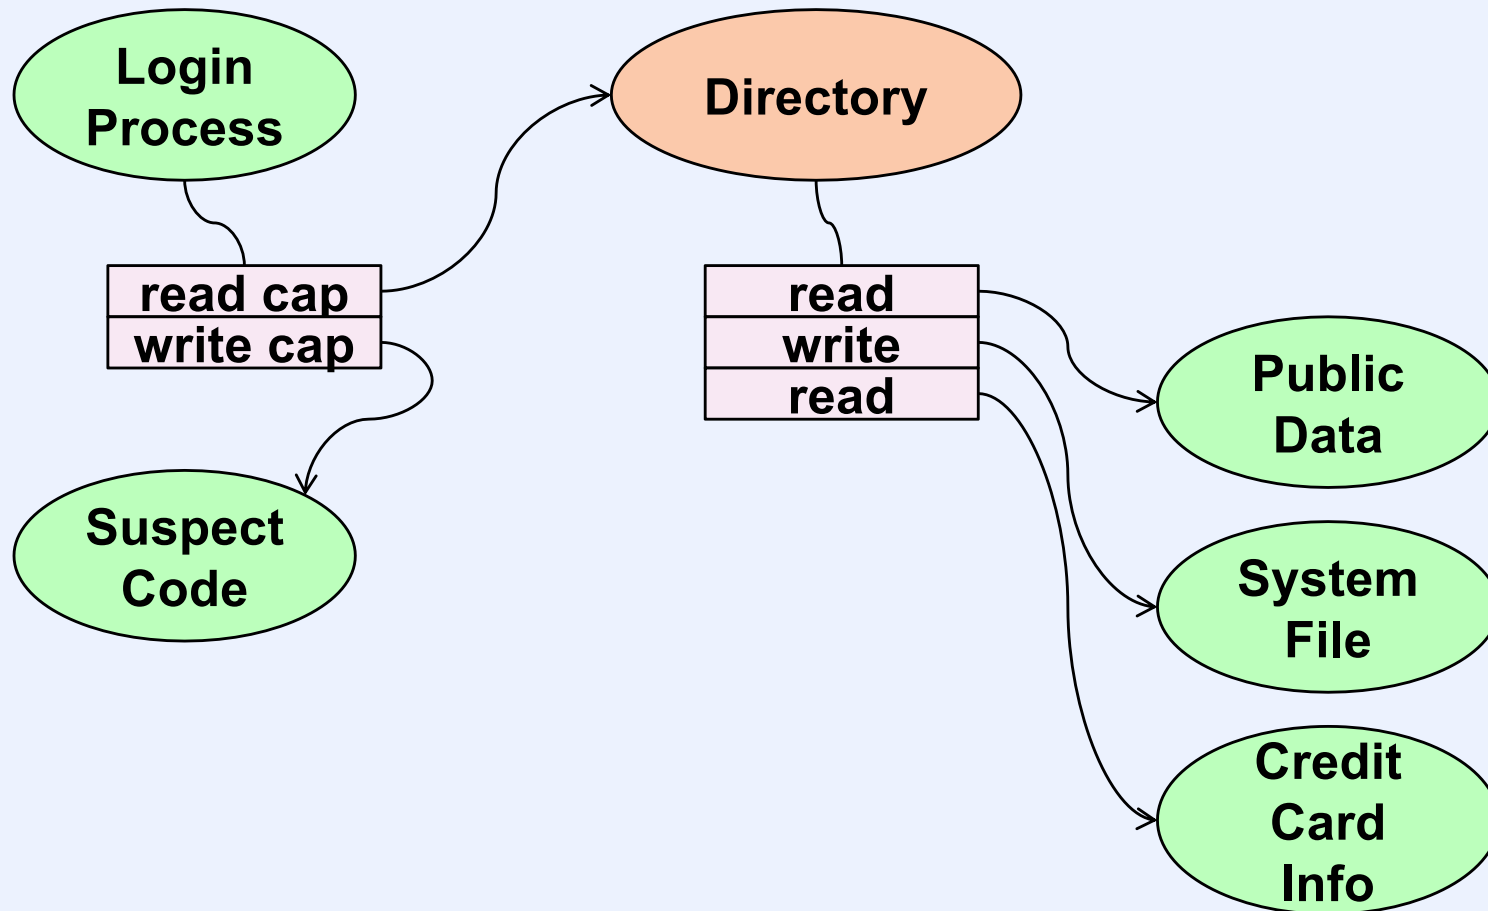
Copying Capabilities (2)



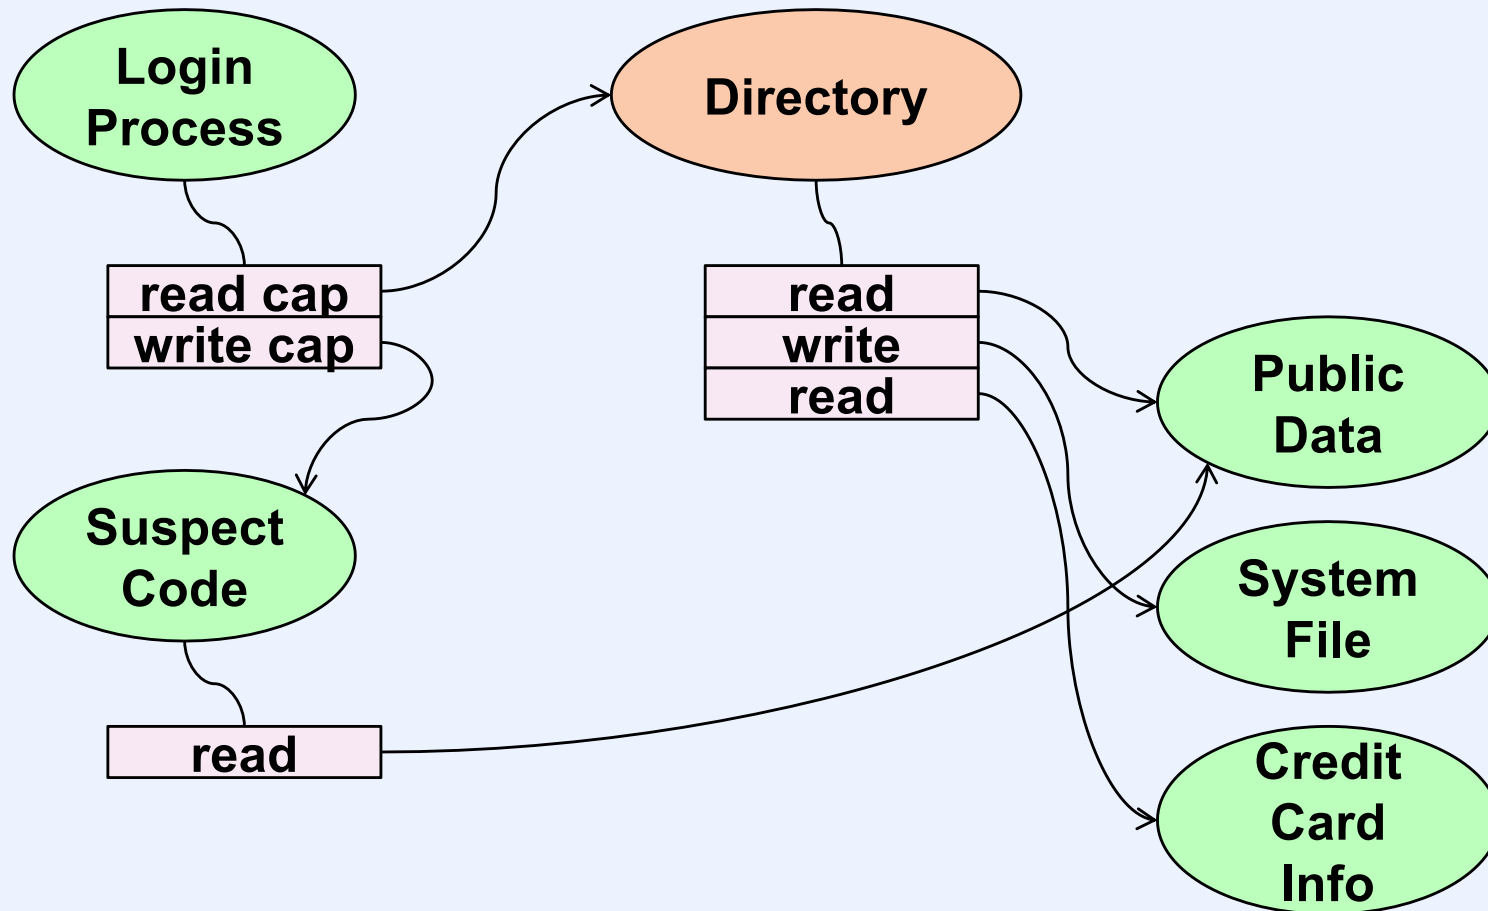
“Directories”



Least Privilege (1)



Least Privilege (2)



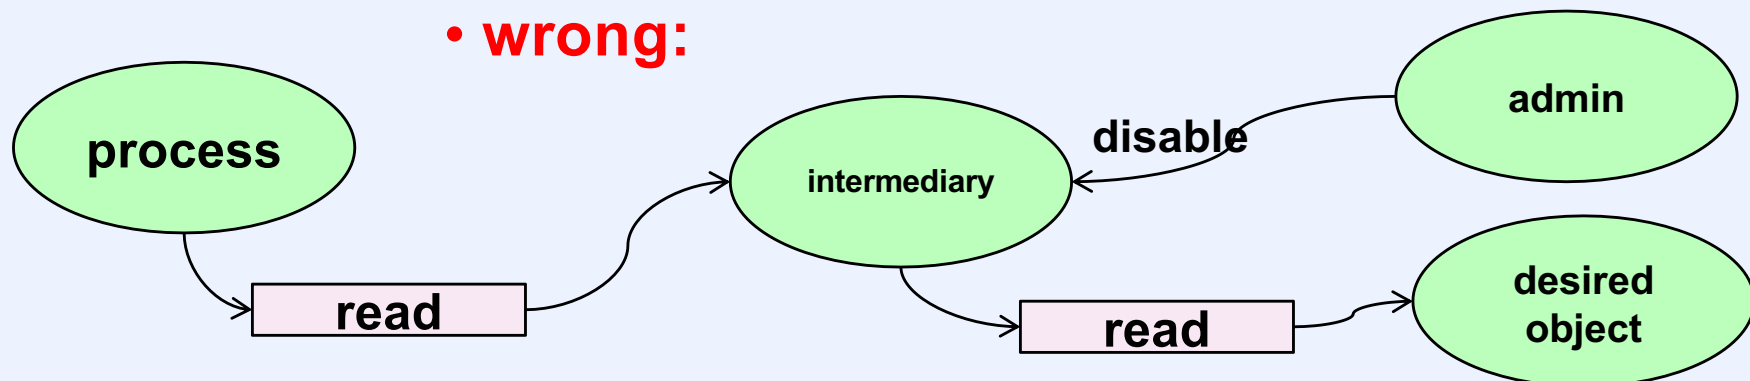
Issues

- **Files aren't referenced by names. How do your processes get capabilities in the first place?**
 - **your “account” is your login process**
 - **created with all capabilities it needs**
 - **persistent: survives log-offs and crashes**

Issues

- Can MAC be implemented on a pure capability system?
 - proven impossible twice
 - capabilities can be transferred to anyone
 - **wrong: doesn't account for write-capability and read-capability capabilities**
 - capabilities can't be retracted once granted

• **wrong:**




Do Pure Capability Systems Exist?

- **Yes!**
 - long history
 - Cambridge CAP System
 - Plessey 250
 - IBM System/38 and AS/400
 - Intel iAPX 432
 - KeyKOS
 - EROS
 - CHERI

A Real Capability System

- **KeyKOS**
 - commercial system
 - capability-based microkernel
 - used to implement Unix
 - (sort of defeating the purpose of a capability system ...)
 - used to implement KeySafe
 - designed to satisfy “high B-level” orange-book requirements
 - probably would have worked
 - company folded before project finished

KeySafe



Speculative Execution Attacks

An Important Assumption

- **Pages cannot be read if read permission is off**
 - **on Intel x86-64**
 - **rings 0-2 can read all mapped pages**
 - **ring 3 can read only those pages for which read permission is explicitly given (in the page-table entries)**

What If the Assumption is Wrong?

- **Code in ring 3 can read everything that's currently mapped**
- **In Linux, all physical memory is mapped into kernel's address space**
 - **kernel address space mapped into every process**
 - **every process can read all physical memory**

Making the Assumption True

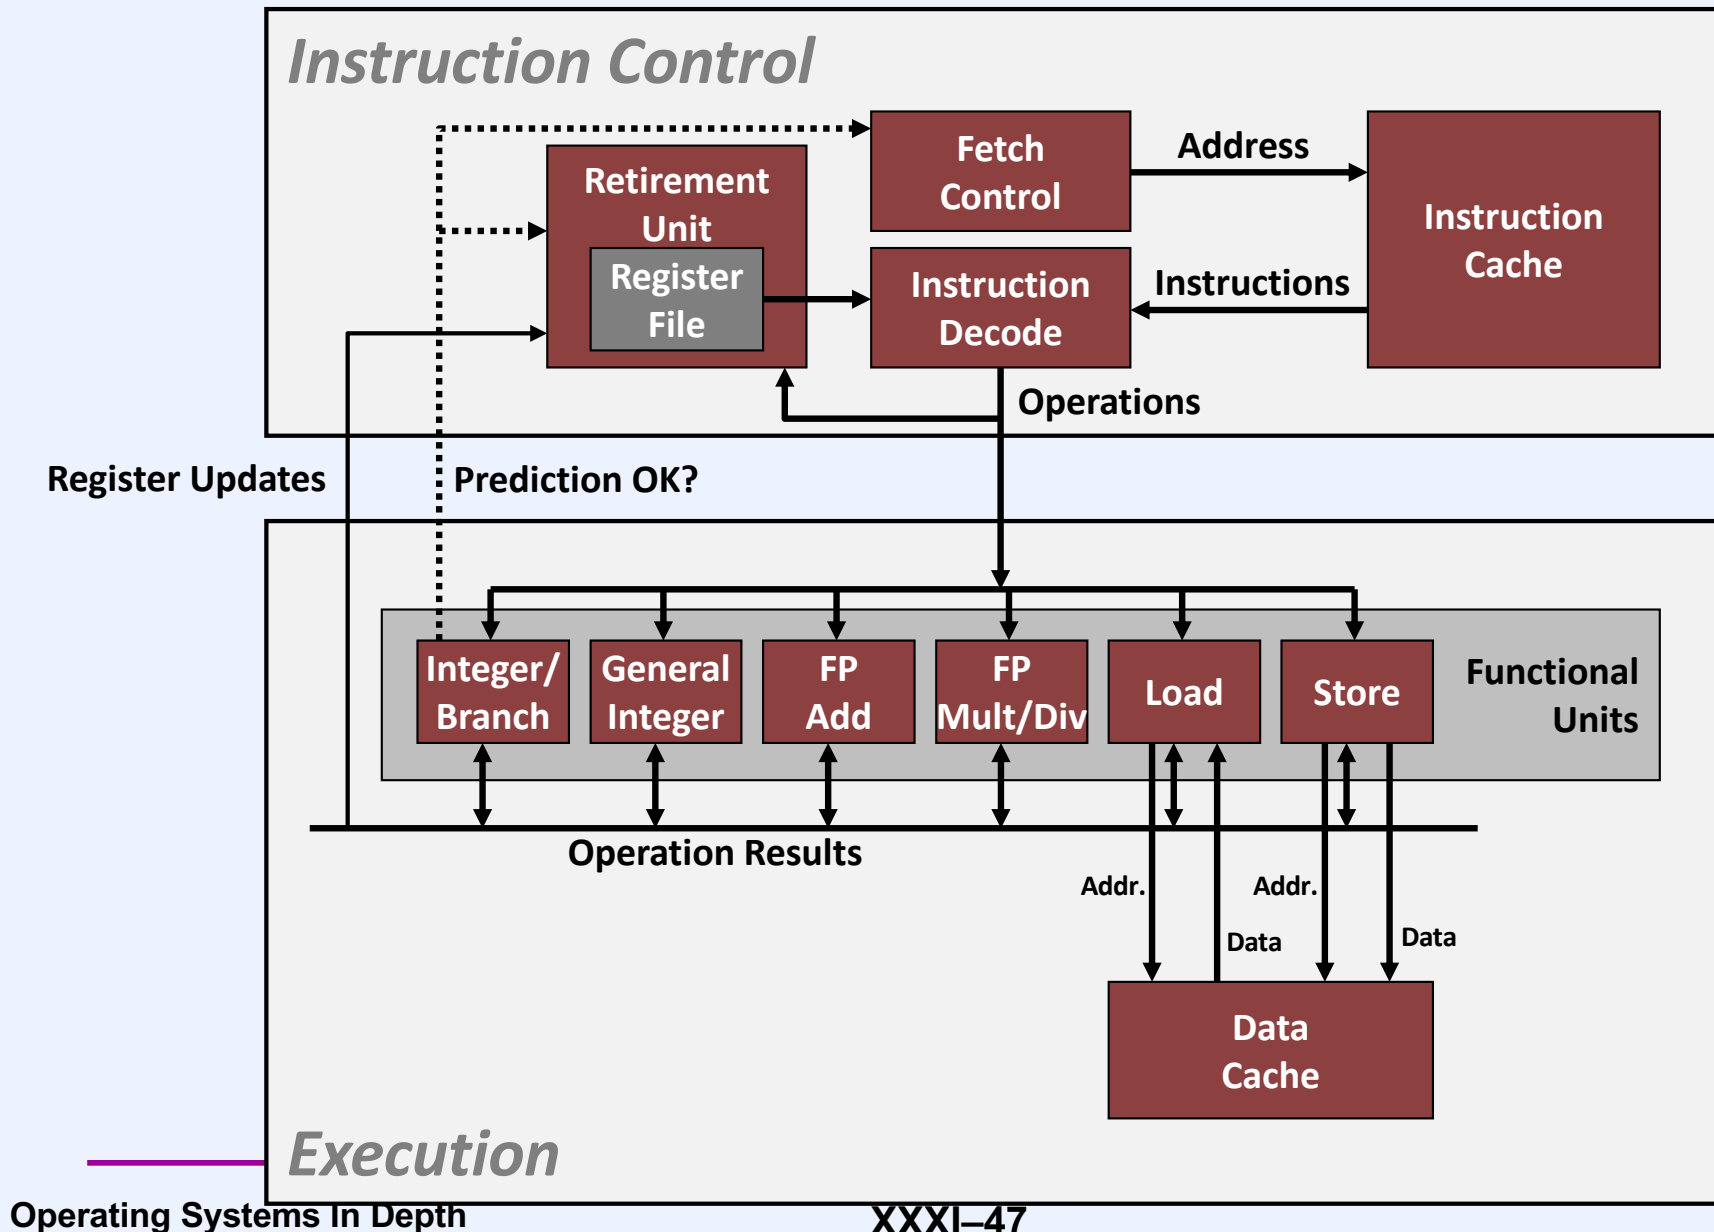
- **Processor checks page protection on each access of memory**
 - a fault occurs and access is not allowed if page is marked not readable

Speculative Execution

```
1:  movl    $0x1,%ecx
2:  xorq    %rdx,%rdx
3:  cmpq    %rsi,%rdx
4:  jne     someplace_else
5:  movl    %esi,%edi
6:  imull    (%rax,%rdx,4),%ecx
```

} perhaps execute
these instructions

Modern CPU Design



Speculative Execution

- If speculatively executed instruction turns out to be used (branch not taken), its results are *retired*
 - e.g., if a load into a register, the register is updated
 - if speculatively executed instruction results in exception, exception is taken
- If not used (branch taken), register is not updated
 - exceptions are ignored

Micro Operations

- Machine instructions are actually composed of micro operations, which can be executed concurrently by the various functional units
- To fetch from memory into register, (at least) two micro operations are used:
 - load value from memory
 - value goes into cache
 - check if access is allowed
 - if not, register not updated
 - exception occurs
 - unless executed speculatively and instruction not used

Speculative Execution

```
1:  movq    kernel_address,%rcx
2:  movq    random_value,%rdx
3:  cmpq    %rcx,%rdx
4:  jne     someplace_else    jump is taken
5:  movb    (%rcx),%al        No exception, %rax not  
                             modified, but cache is  
                             updated
```

Not a security problem, because there's no way to determine what went into the cache.

Correct?!?!?

Speculative Execution

```
1:  movq    kernel_address,%rcx
2:  movq    random_value,%rdx
3:  cmpq    %rcx,%rdx
4:  jne     someplace_else
5:  movb    (%rcx),%al
6:  shl     %rax,$0xc    # multiply by 4096
7:  movq    %rax,0(%rbx,%rax)
    # %rbx contains the address of a
    # 1MB (256 x 4096 byte) array in
    # user-accessible memory
```

The Array

(%rcx)



0:

1:

2:

3:

...

x:

--	--


...

253:

254:

255:

Accessing the Array



Speculative Execution

```
1:  movq    kernel_address,%rcx
2:  movq    random_value,%rdx
3:  cmpq    %rcx,%rdx
4:  jne     someplace_else
5:  movb    (%rcx),%al
6:  shl     %rax,$0xc    # multiply by 4096
7:  movq    %rax,0(%rbx,%rax)

# %rbx contains the address of a
# 1MB (256 x 4096 byte) array in
# user-accessible memory
```

Can we be sure processor won't be able to guess whether jump is taken?

Alternative Approach

```
1:  movq    kernel_address,%rcx
2:  movq    random_value,%rdx
3:  cmpq    %rcx,%rdx
4:  movb     $1,0          # cause a segfault
5:  movb     (%rcx),%al
6:  shl     %rax,$0xc      # multiply by 4096
7:  movq     %rax,0(%rbx,%rax)
   # %rbx contains the address of a
   # 1MB (256 x 4096 byte) array in
   # user-accessible memory
```

Kernel Address Mapping

- **On Linux and OSX**
 - kernel is mapped into every user process
 - all physical memory is mapped into kernel
 - thus all physical memory can be accessed via speculative execution

Meltdown

- **Can read all of kernel memory on a Linux machine at the rate of 503 KB/sec**
 - **it's not dependent on software bugs: works on a bug-free kernel**
 - **achieved by using Intel TSX so that the speculatively executed instructions are not retired**
 - **slows down to 122 KB/sec if segfaults are used**

Countermeasures

- **KASLR**
 - kernel address-space layout randomization
 - kernel starts at a randomly chosen address
 - 40 bits of randomization
 - try all possibilities ...
- **KAISER**
 - “kernel address-space isolation to have side channels efficiently removed”
 - kernel address space is not mapped when in user mode
 - prevents meltdown!

Quiz 4

As stated, KAISER requires that there are no mappings of kernel virtual addresses when the processor is running in ring 3 (user mode).

- a) This is completely innocuous and has no adverse effects**
- b) This may “break” some otherwise correct user code**
- c) This may have noticeable performance effects**
- d) Both b and c are true**

Other Similar Attacks (1)

- **RIDL: Rogue In-Flight Data Load**
 - certain processor buffers hold data coming from memory
 - loaded by victim thread
 - before doing address translations and protection checks, hardware guesses that it might be what's needed by attacker thread
 - attacker puts value in array (and thus in cache)
 - hardware determines addresses are different and does not retire results
 - attacker determines value by access times in array
 - not prevented by Kaiser

Other Similar Attacks (2)

- **Fallout**
 - stores deposited in store buffer
 - allowing thread to continue executing without waiting for memory to be updated
 - subsequent loads check store buffer to see if it contains data to be loaded
 - address/protections checks are optimistic
 - attacker can determine what data is being stored by victim

Summary

- **Attacks rely on extreme optimization done in hardware**
 - depend on detailed knowledge of hardware architecture
 - much of this was learned via reverse engineering
- **Attacks allow one to extract information across virtual machine and container boundaries**