


Introduction to Networking


Distributed File Systems



What We Cover

- Communication protocols
- Remote procedure call protocols
- **Distributed file systems**

Communication



Some Issues

- **Quantity: how many are communicating?**
 - unicast
 - broadcast
 - multicast
- **Quality: how good/reliable is the communication?**
 - best effort
 - fully reliable
 - guaranteed bandwidth and delay


Some More Issues

- Naming and addressing
 - www.cs.brown.edu vs. 128.148.32.110
- Routing
- Congestion

Circuits vs. Packets

- **Circuit switching**
 - reserve a circuit between communicating entities
- **Packet switching**
 - break data into packets
 - transmit each packet separately — each packet could travel a different route

Internetworking with TCP/IP



IP Header

vers	hlen	type of serv	total length	
identification			flags	fragment offset
time-to-live		protocol	header checksum	
source address				
destination address				
options				padding
data				

User Datagram Protocol

vers	hlen	type of serv	total length	
identification			flags	fragment offset
time-to-live		protocol	header checksum	
source address				
destination address				
options				padding
data				
Source Port			Destination Port	
Length			UDP Checksum	

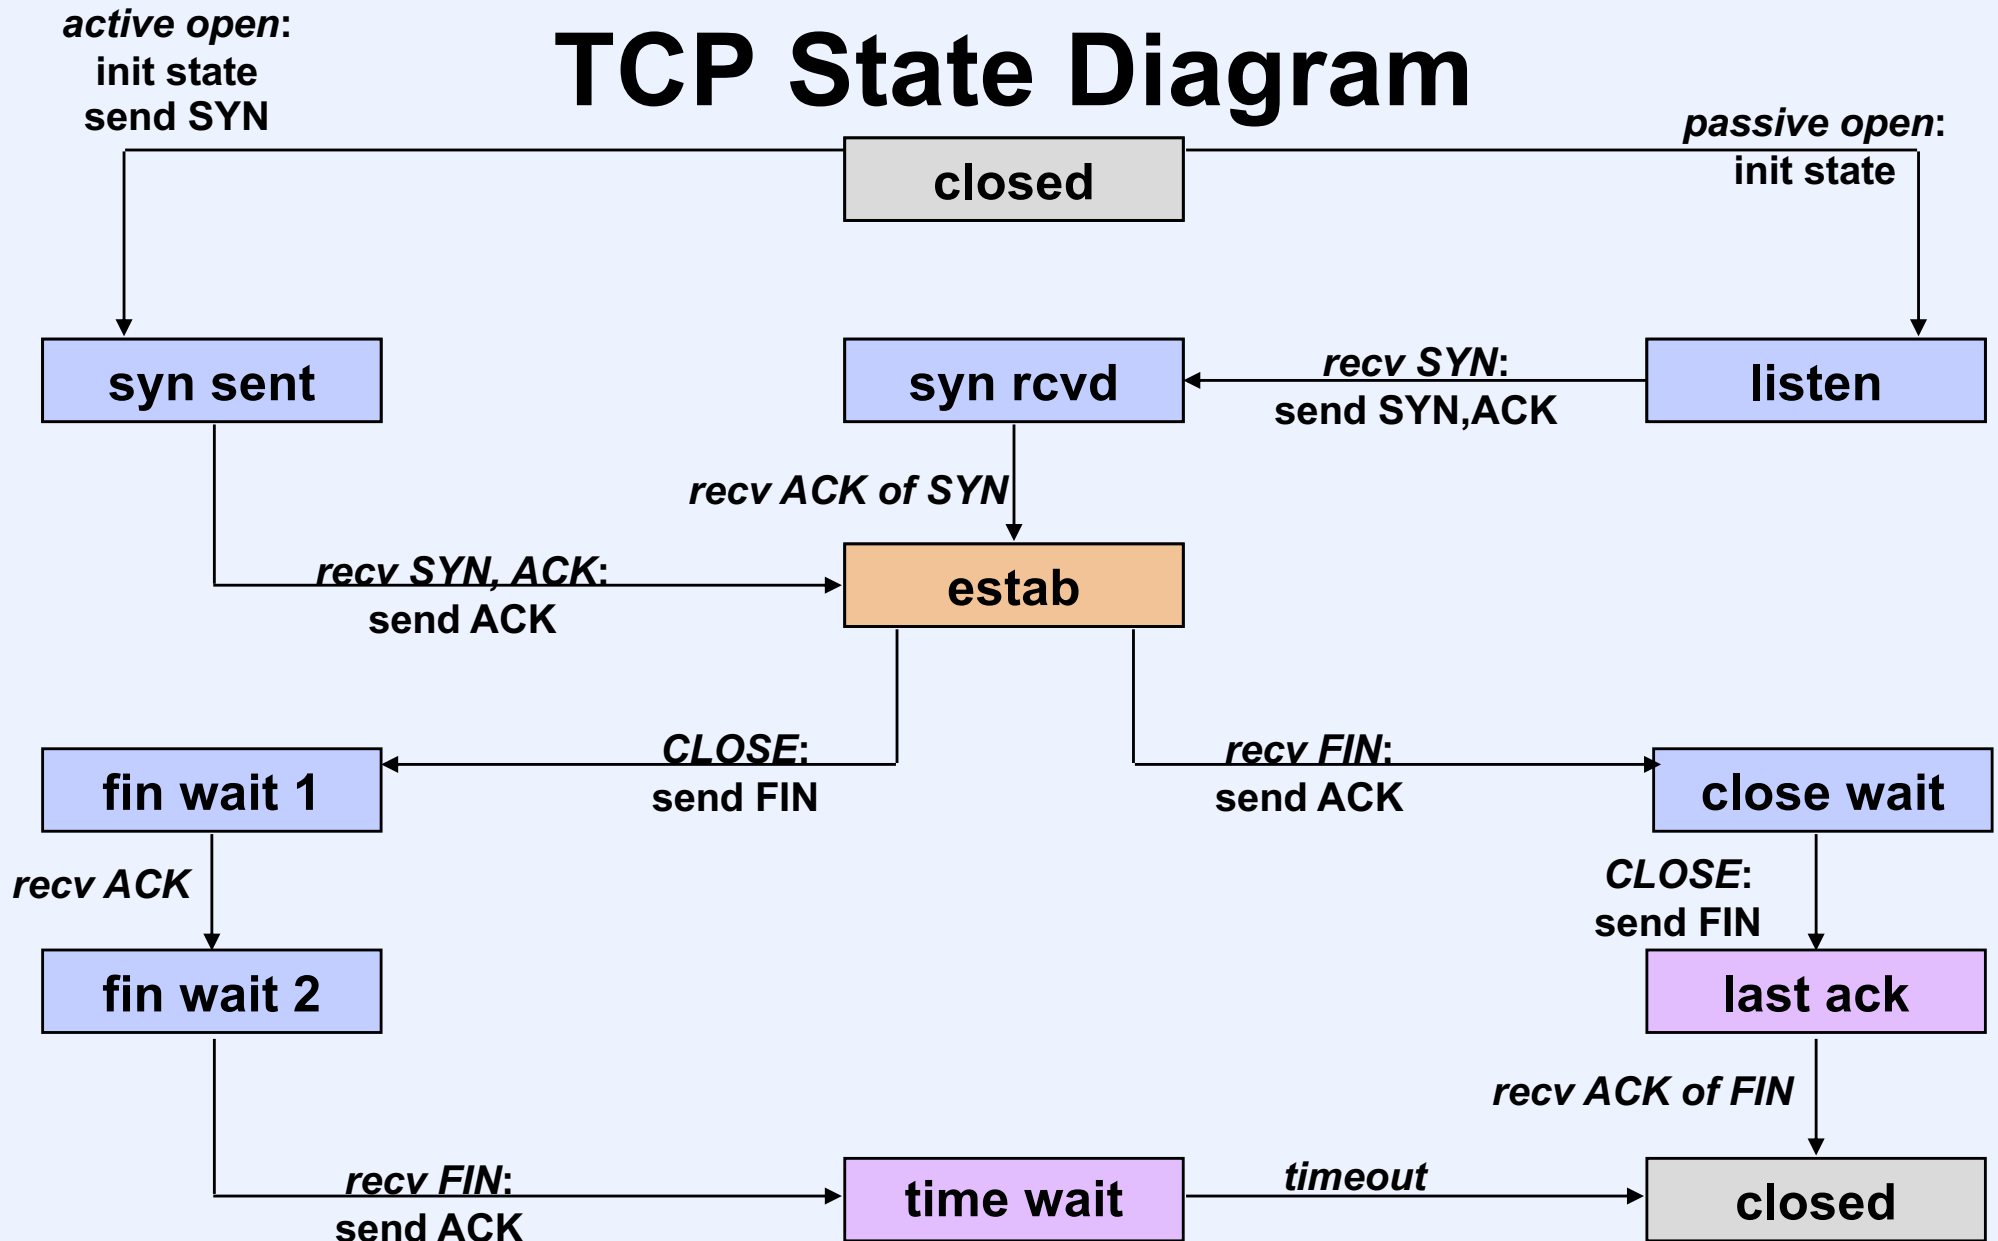
Transmission Control Protocol (TCP)

- **A multiplexing service built on top of IP**
- **A full-duplex reliable stream protocol**
- **Provides reliable data transfer**
- **Packet boundaries are not seen by the application**
 - **it sees a sequence of bytes, not a sequence of packets**
- **A connection must be established for communication to take place**
 - **connections are flow-controlled**
 - **connections are congestion-controlled**

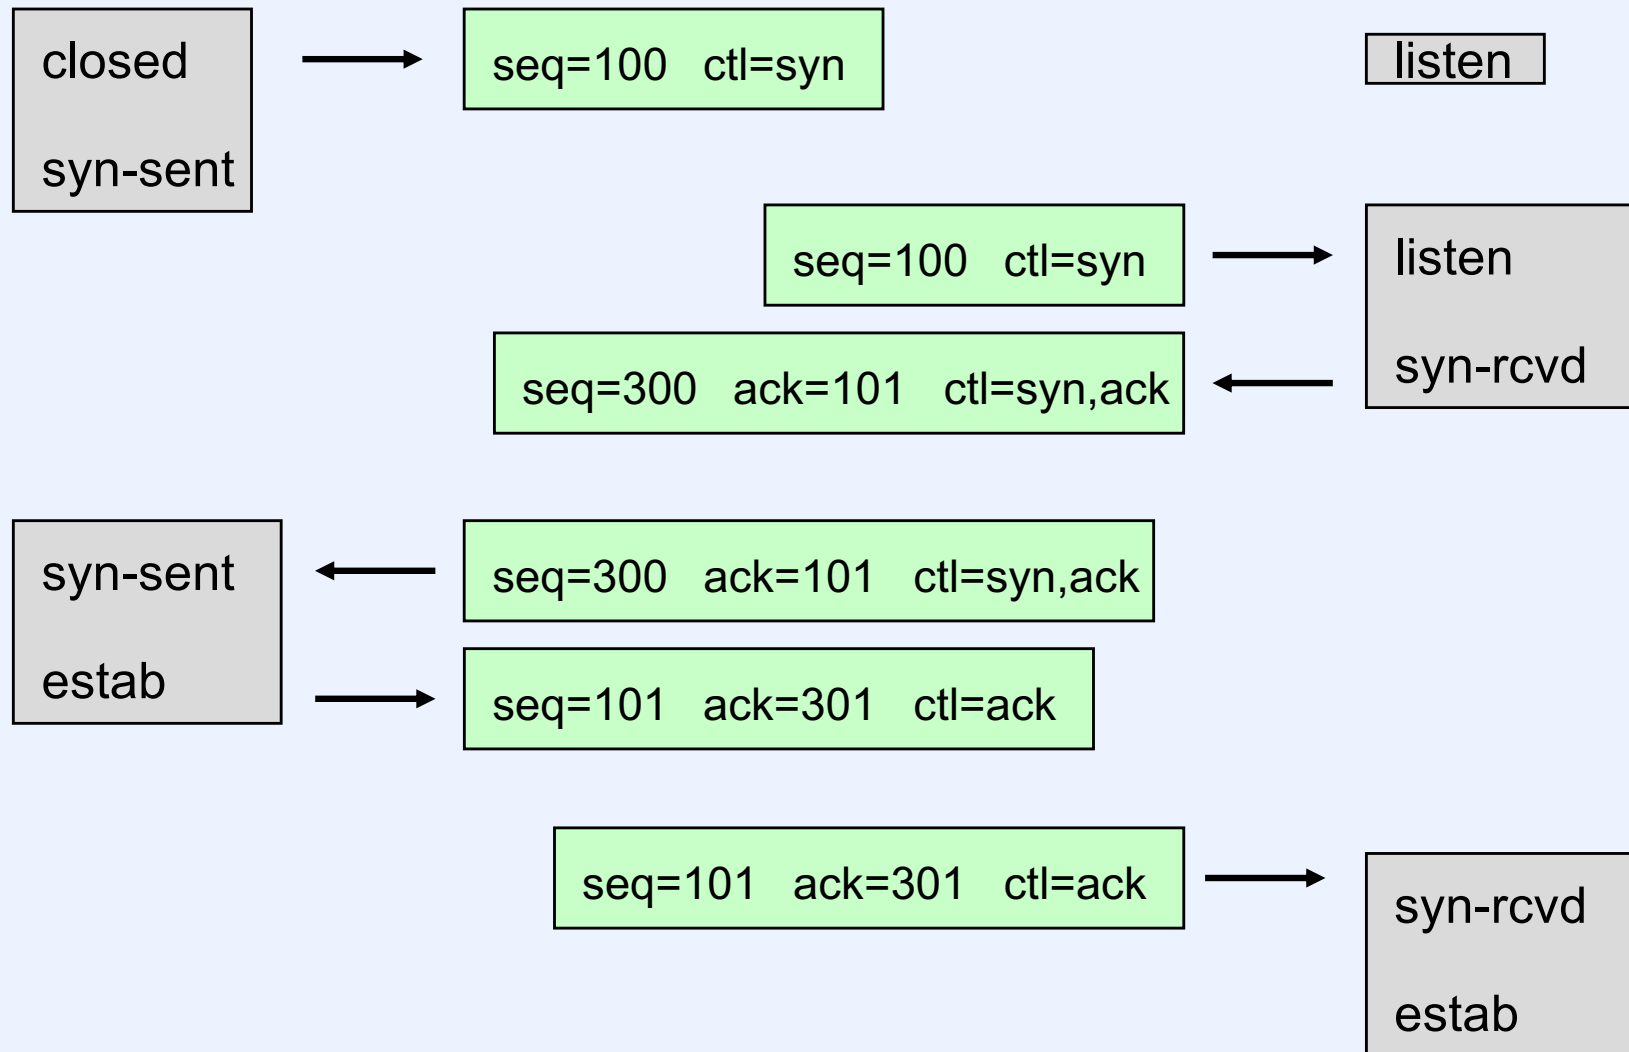
TCP Header Components

source port			destination port		
sequence number					
acknowledgment sequence number					
offset	reserved	flags	window size		
checksum			urgent pointer		
options					padding
data					

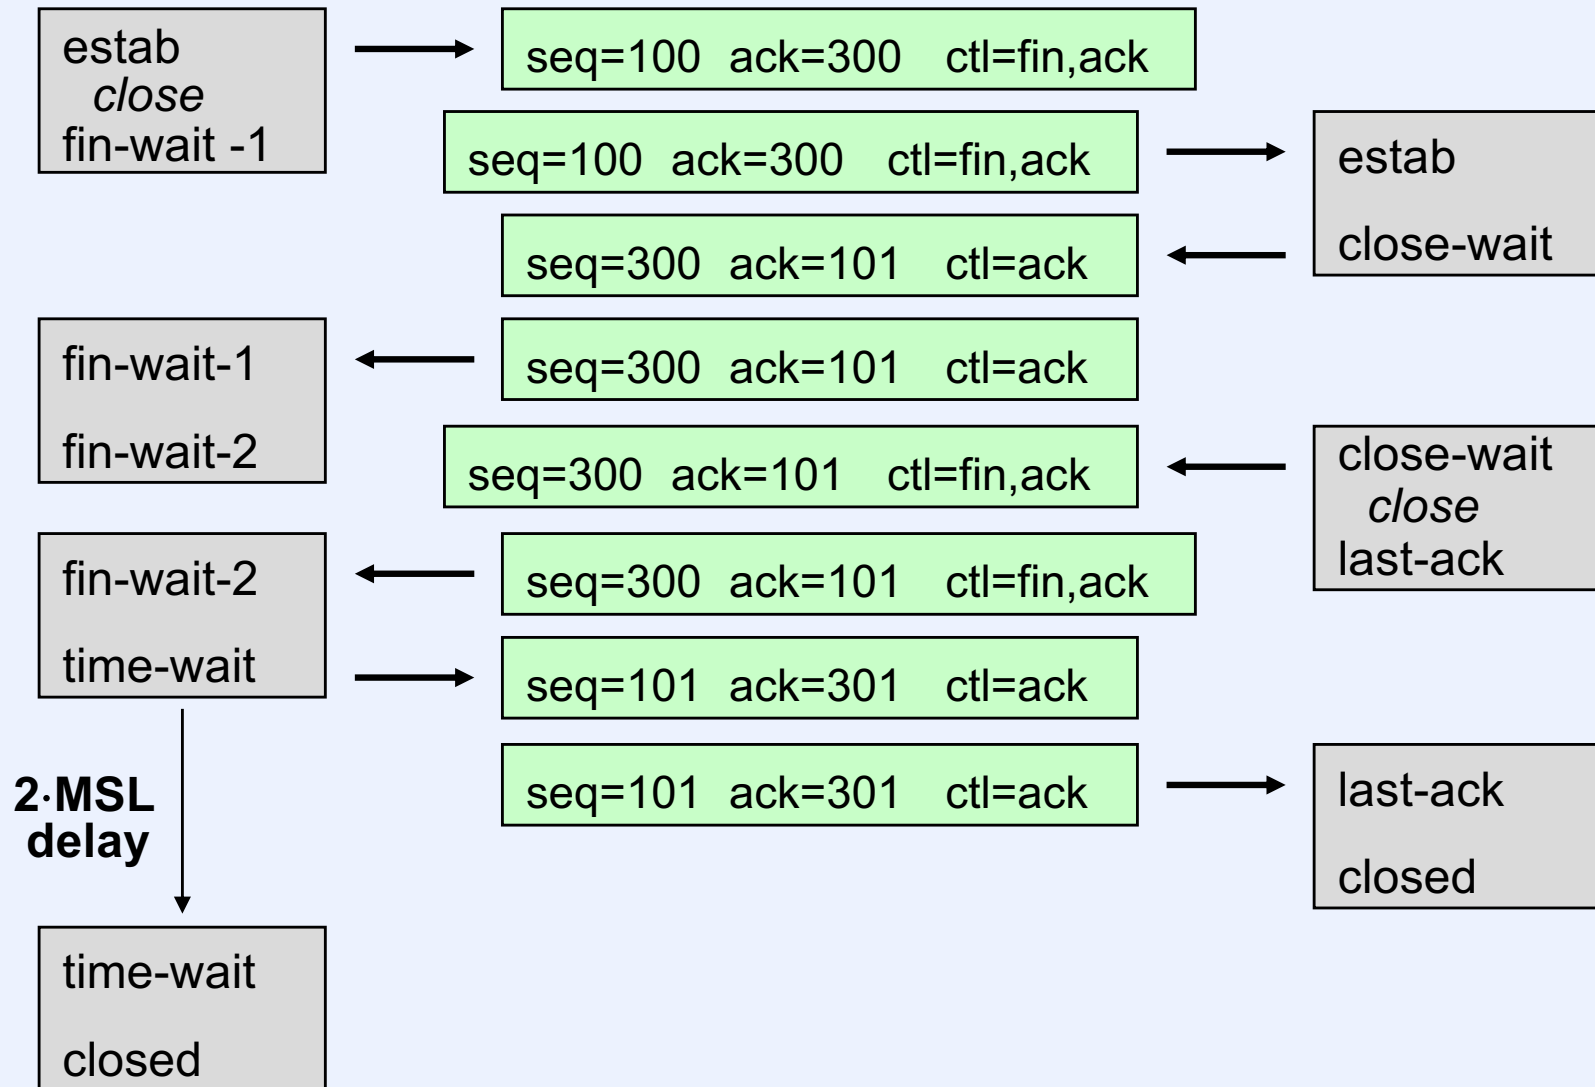
TCP State Diagram



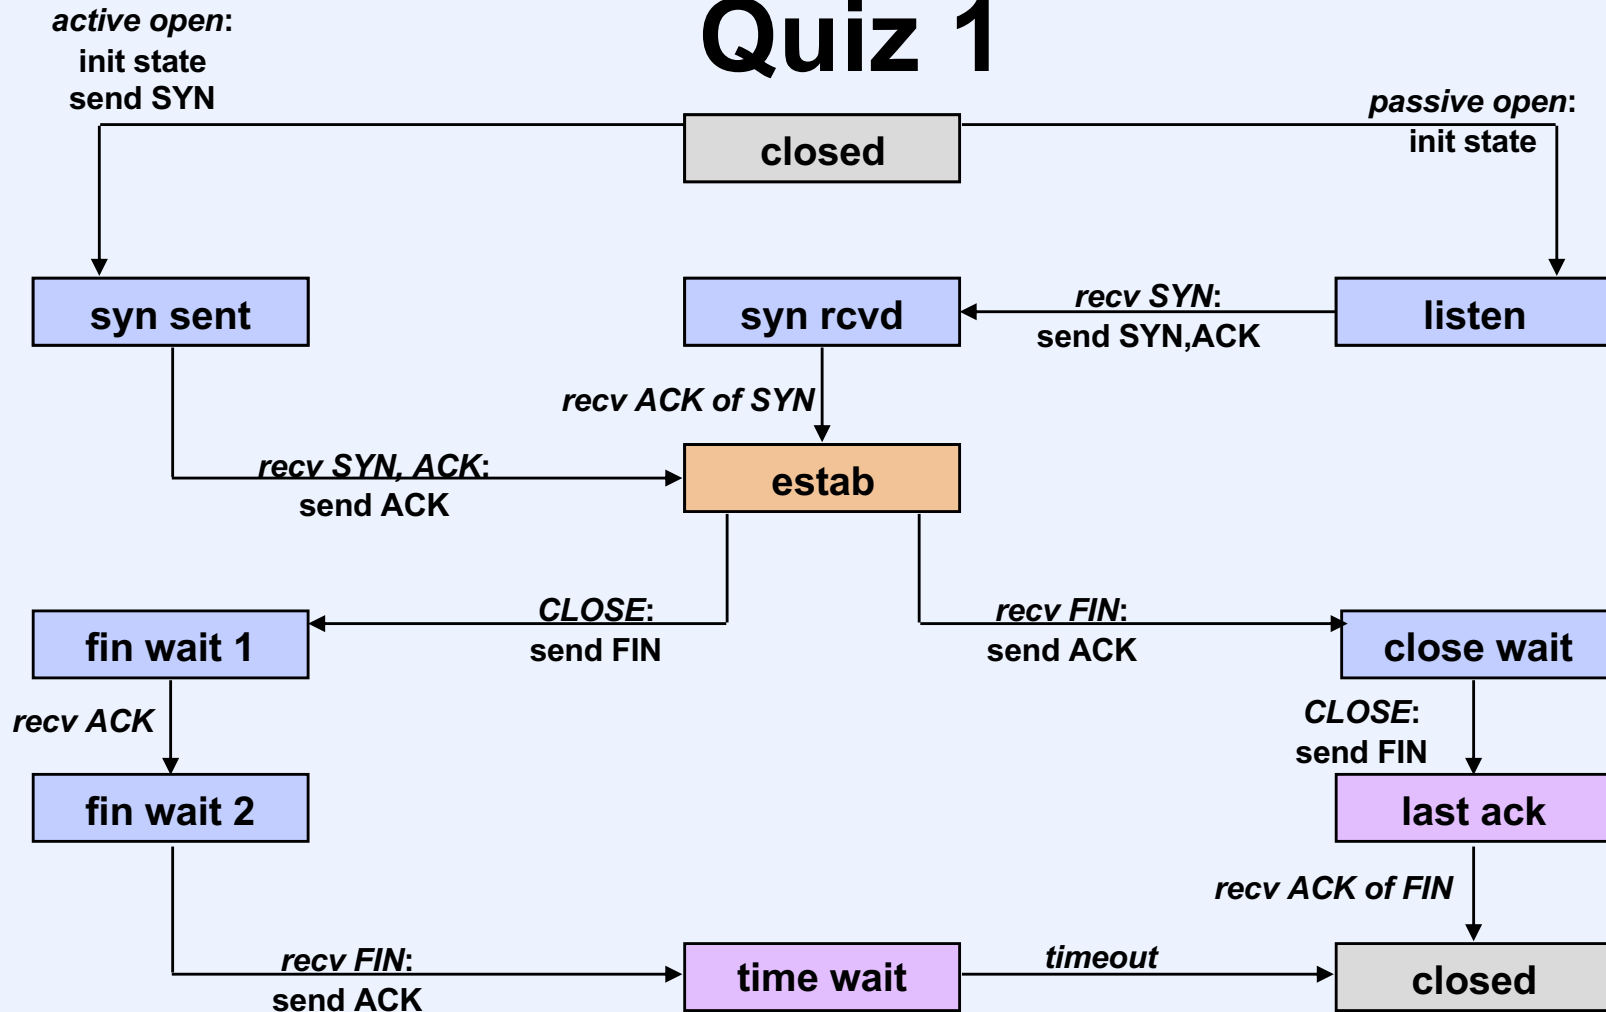
Connection Initiation



Shutdown



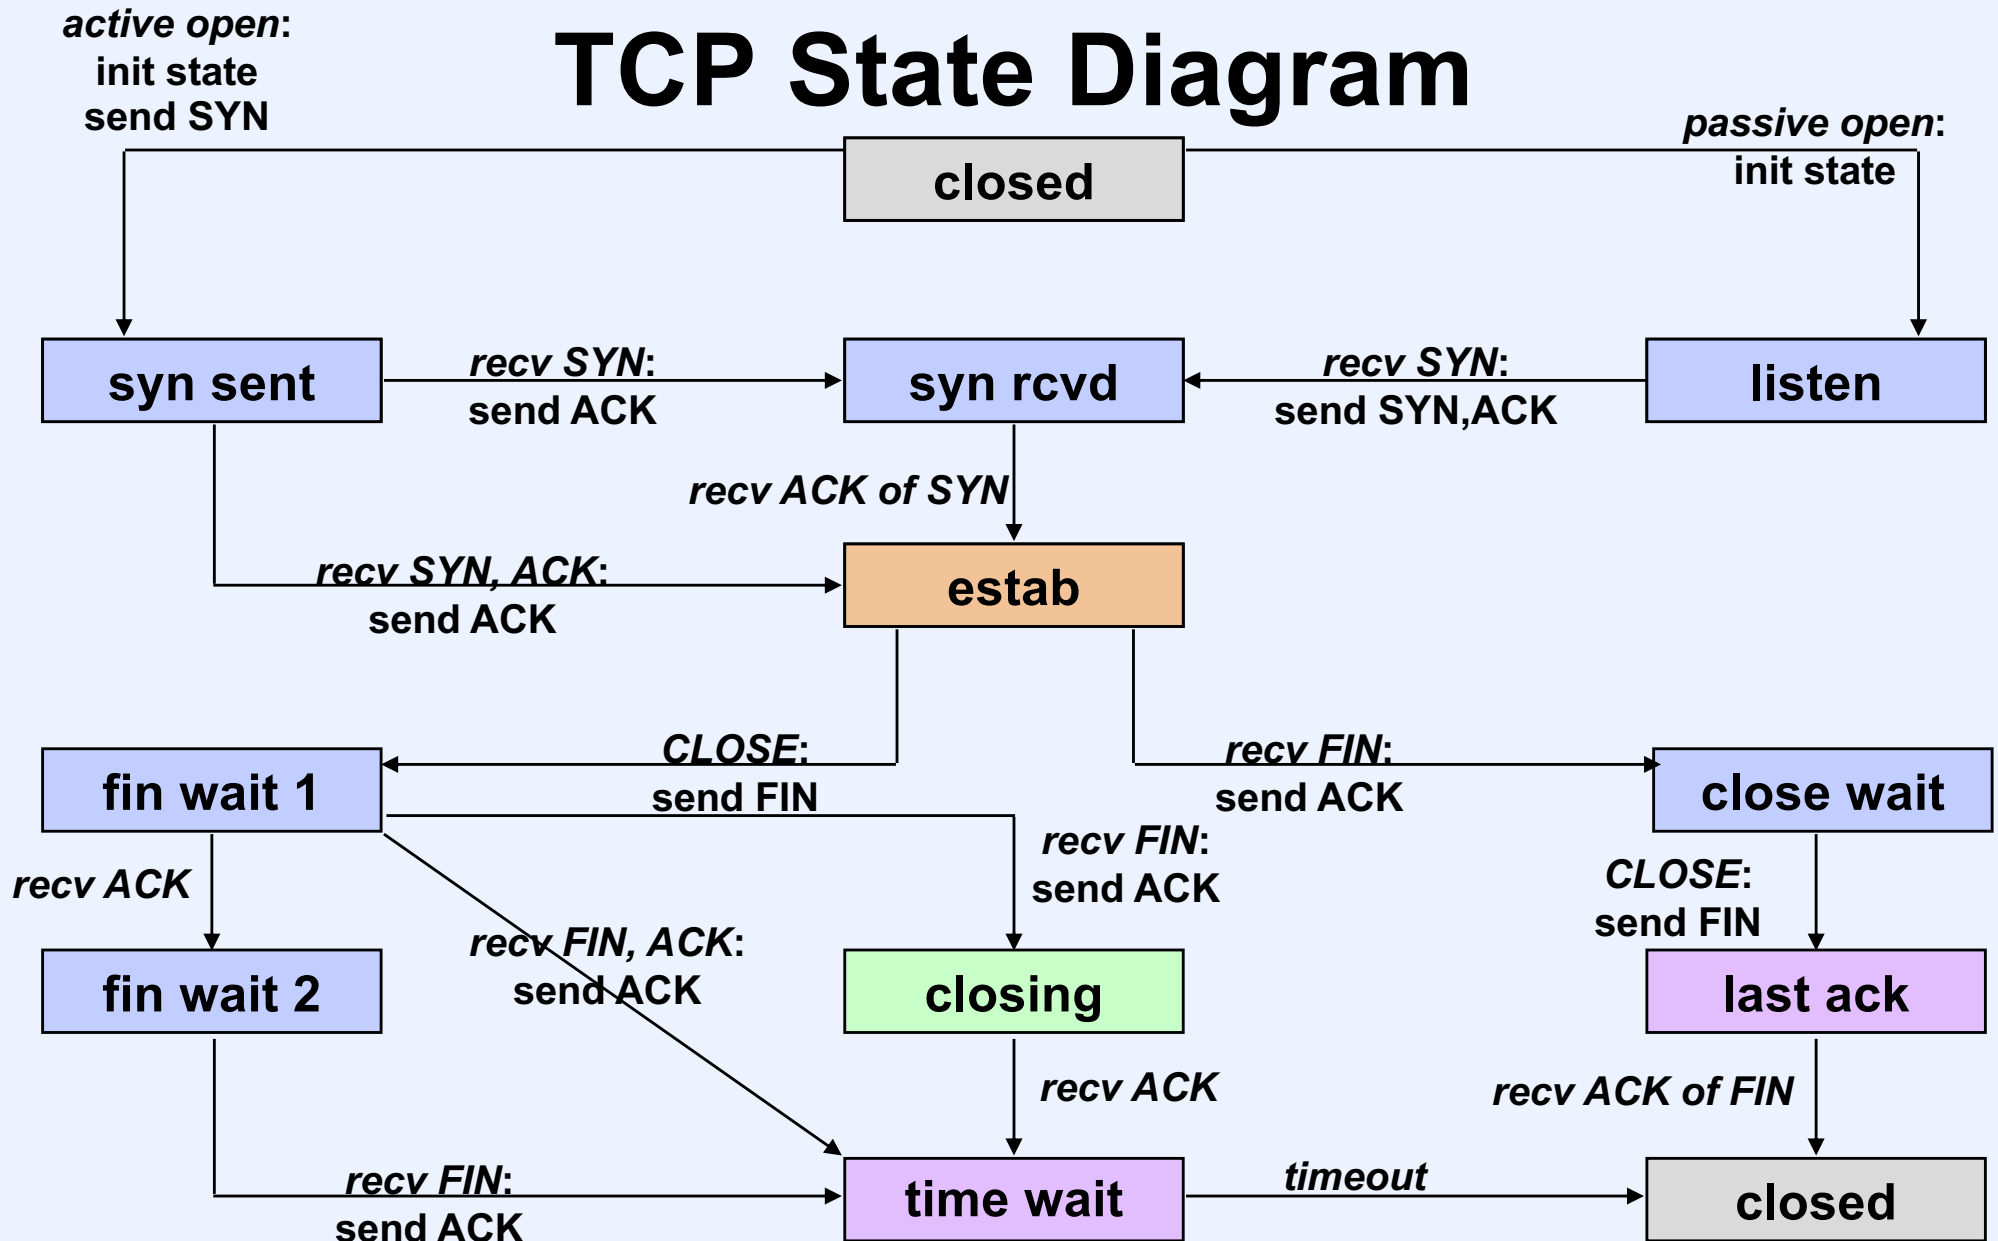
Quiz 1



Suppose both sides send the other a FIN while in the estab state; both then go to fin-wait-1. What happens next?

- a) Both go to fin-wait-2, then time-wait
- b) Both go to a new state that's not in the diagram
- c) Such an event can't happen


TCP State Diagram




Sliding-Window Protocol

- **Used for:**
 - reliable delivery
 - flow control
- **Windows**
 - send
 - receive


TCP Receive Window



TCP Send Window




Coping With Lost Data




When to Retransmit?

- **After waiting significantly longer than the average (“smoothed”) roundtrip time**
- **How much longer?**
 - significantly longer than the average deviation
- **What if one retransmission doesn’t do it?**
 - transmit again ...
- **When?**
 - use exponential backoff
- **When does one give up?**
 - eventually ...


Ack Clocking



Fast Start




Slow Start




Congestion Control (1)


1)





2)



3)



Congestion Control (2)



Acks

- **When a segment is received, the highest permissible Ack is sent back**
 - if data up through i has been received, the ack sequence number is $i+1$
 - if data up through i has been received, as well as $i+100$ through $i+200$, the ack sequence number is $i+1$
 - a higher value would imply that data in the range $[i+1, i+99]$ has been received
- **Every segment sent contains the most up-to-date Ack**

Quiz 2

A TCP sender has sent four hundred-byte segments starting with sequence numbers 1000, 1100, 1200, and 1300, respectively. It receives from the other side three consecutive ACKs, all mentioning sequence number 1100. It may conclude that


- a) The first segment was received, but nothing more**
 - b) The first, third, and fourth segments were received, but not the second**
 - c) The first segment was received, but not the second; nothing is known about the others**
 - d) There's a bug in the receiver**
-

Fast Retransmit and Recovery

- **Waiting an entire RTO before retransmitting causes the “pipeline” to become empty**
 - must slow-start to get going again
- **If one receives three acks that all repeat the same sequence number:**
 - some data is getting through
 - one segment is lost
 - immediately retransmit the lost segment
 - halve the congestion window (i.e., perform congestion control)
 - don't slow-start (there is still data in the pipeline)

Remote Procedure Call Protocols

Local Procedure Calls






```
// Client code
...
result = procedure(arg1, arg2);
...
```

```
// Server code
result_t procedure(a1_t arg1, a2_t arg2) {
    ...
    return(result);
}
```

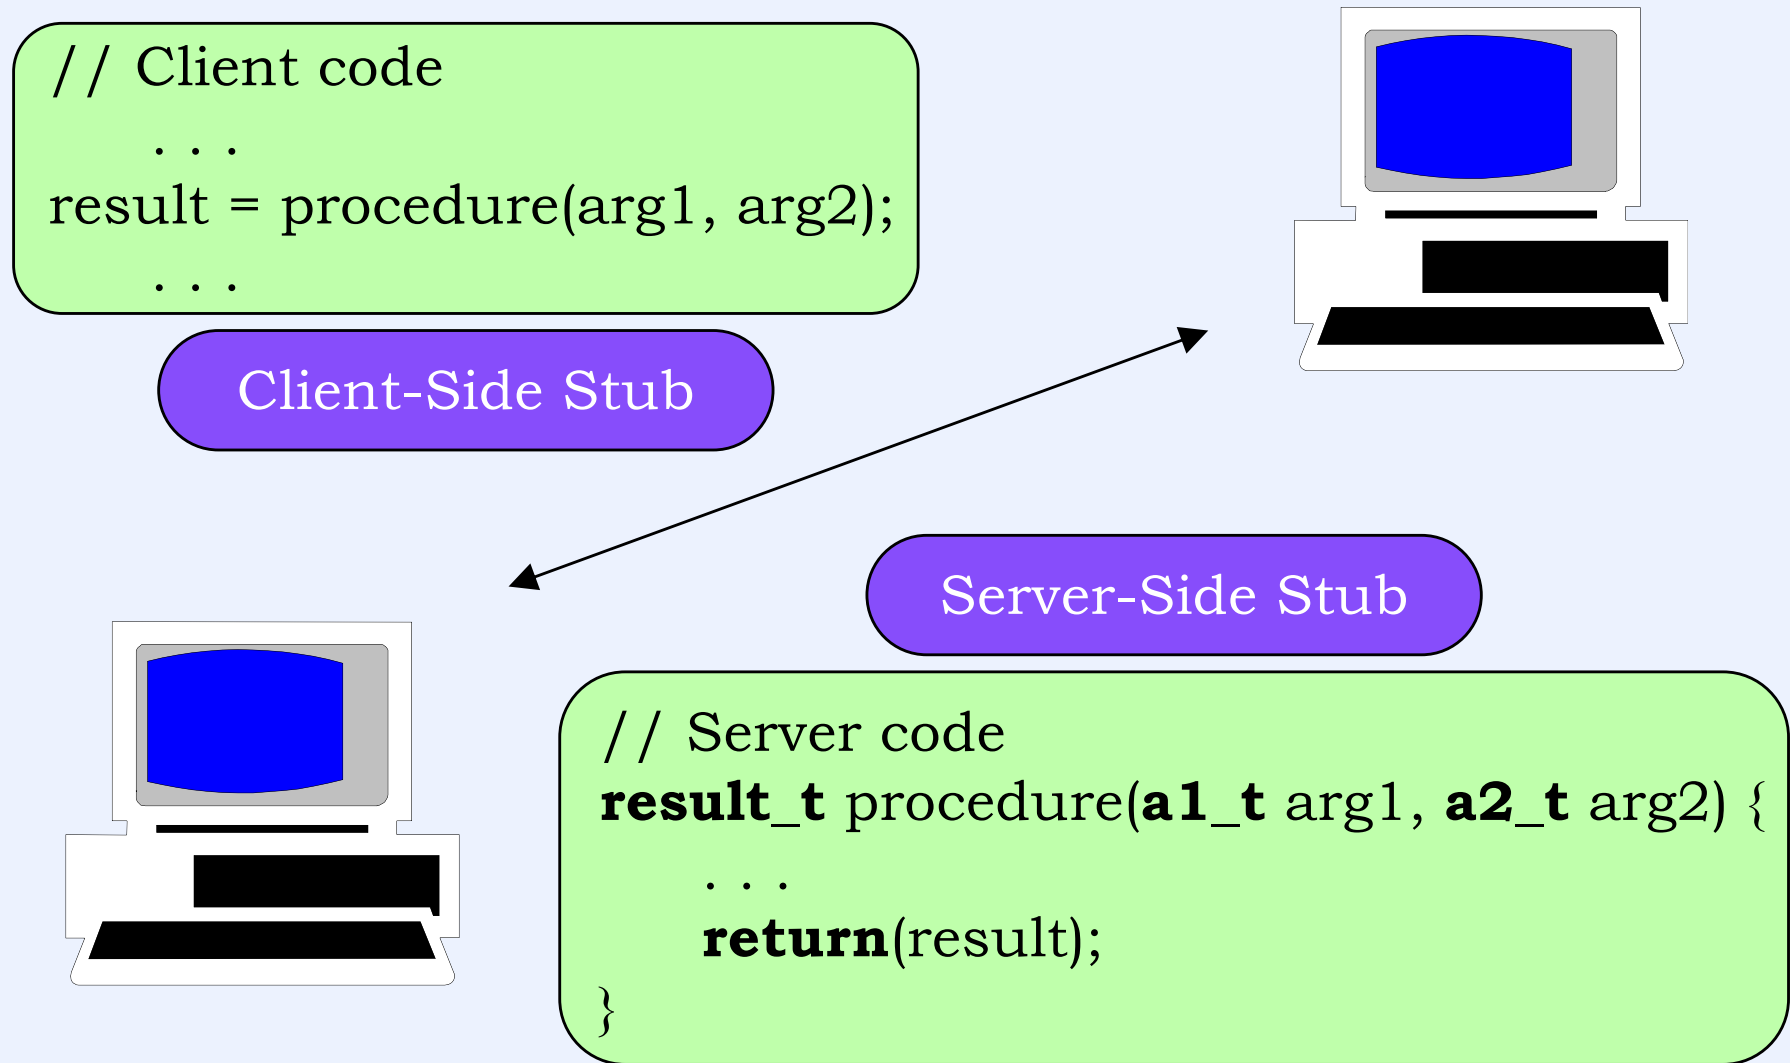

Remote Procedure Calls (1)

```
// Client code  
...  
result = procedure(arg1, arg2);  
...
```

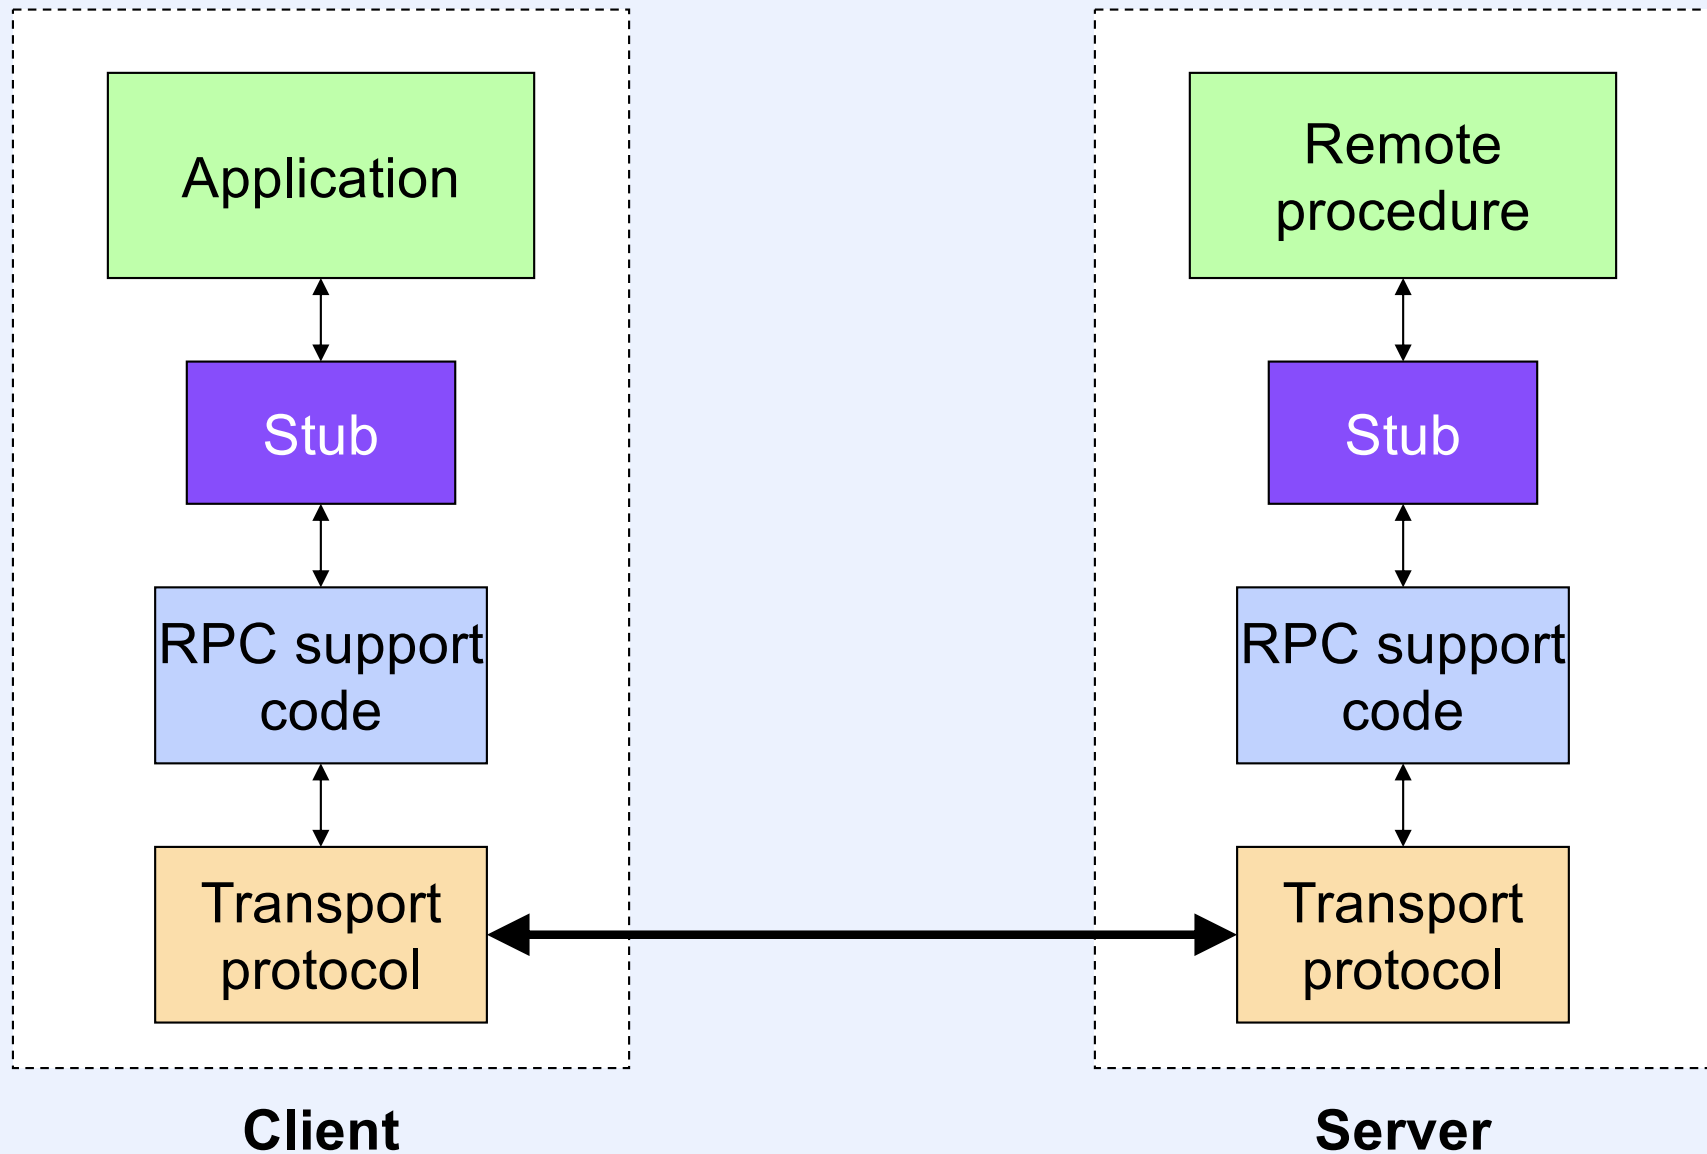


```
// Server code  
result_t procedure(a1_t arg1, a2_t arg2) {  
    ...  
    return(result);  
}
```

Remote Procedure Calls (2)



Block Diagram



ONC RPC

- **Used with NFS**
- **eXternal Data Representation (XDR)**
 - **specification for how data is transmitted**
 - **language for specifying interfaces**

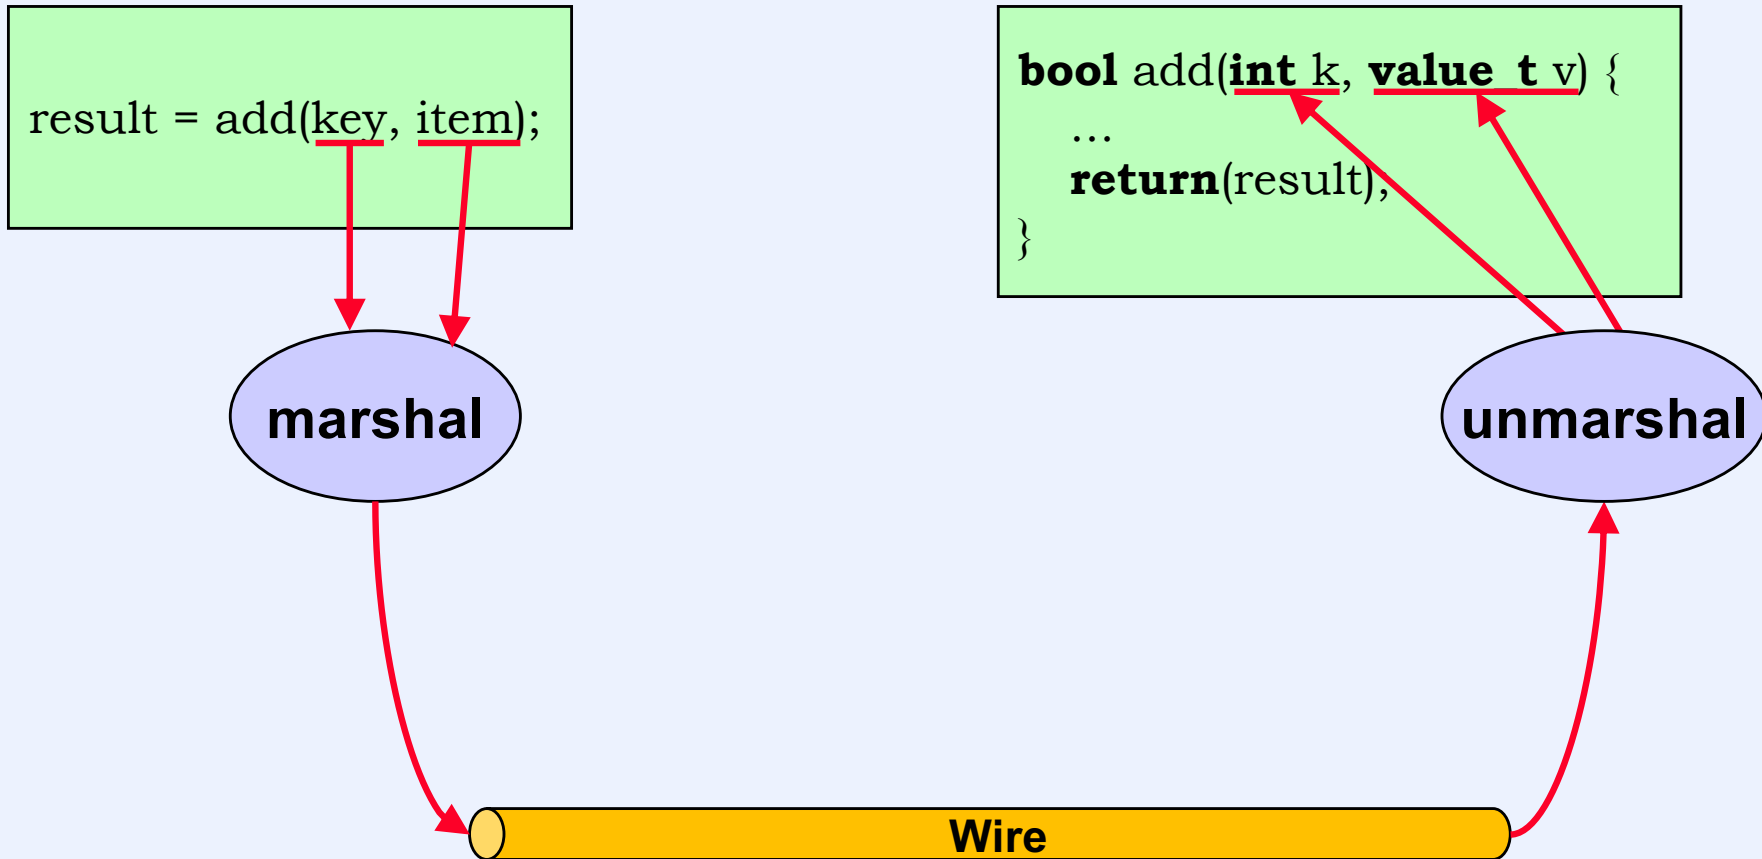
Example

```
typedef struct {  
    int    comp1;  
    float  comp2[6];  
    char   *annotation;  
} value_t;
```

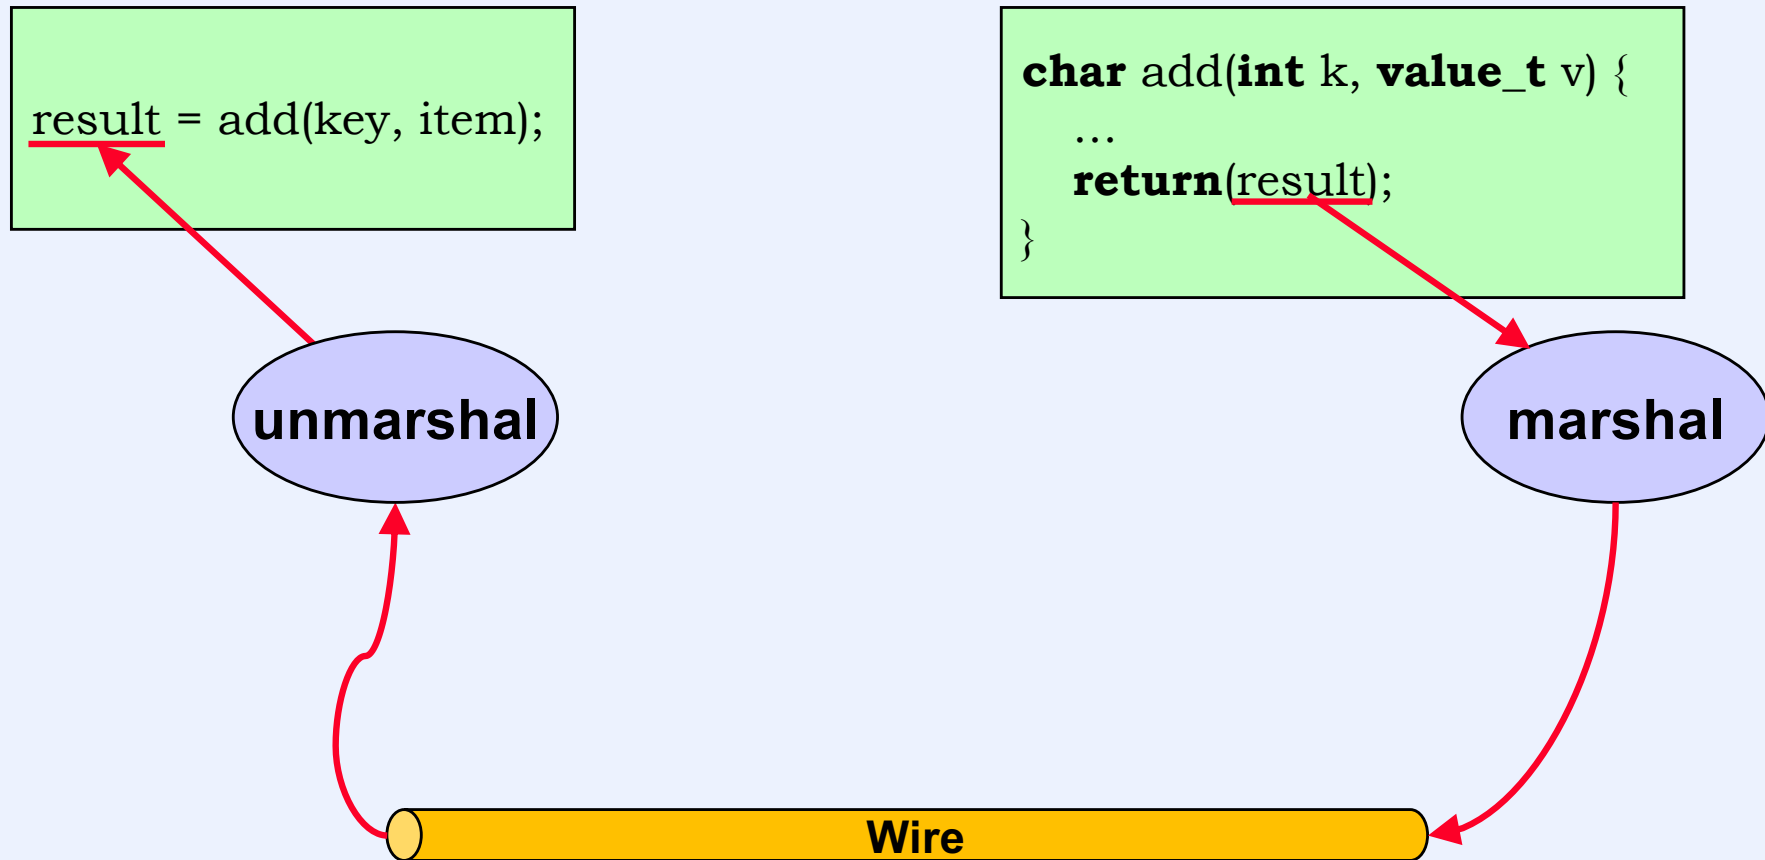
```
typedef struct {  
    value_t  item;  
    list_t   *next;  
} list_t;
```

```
bool add(int key, value_t item);  
bool remove(int key, value_t item);  
list_t query(int key);
```

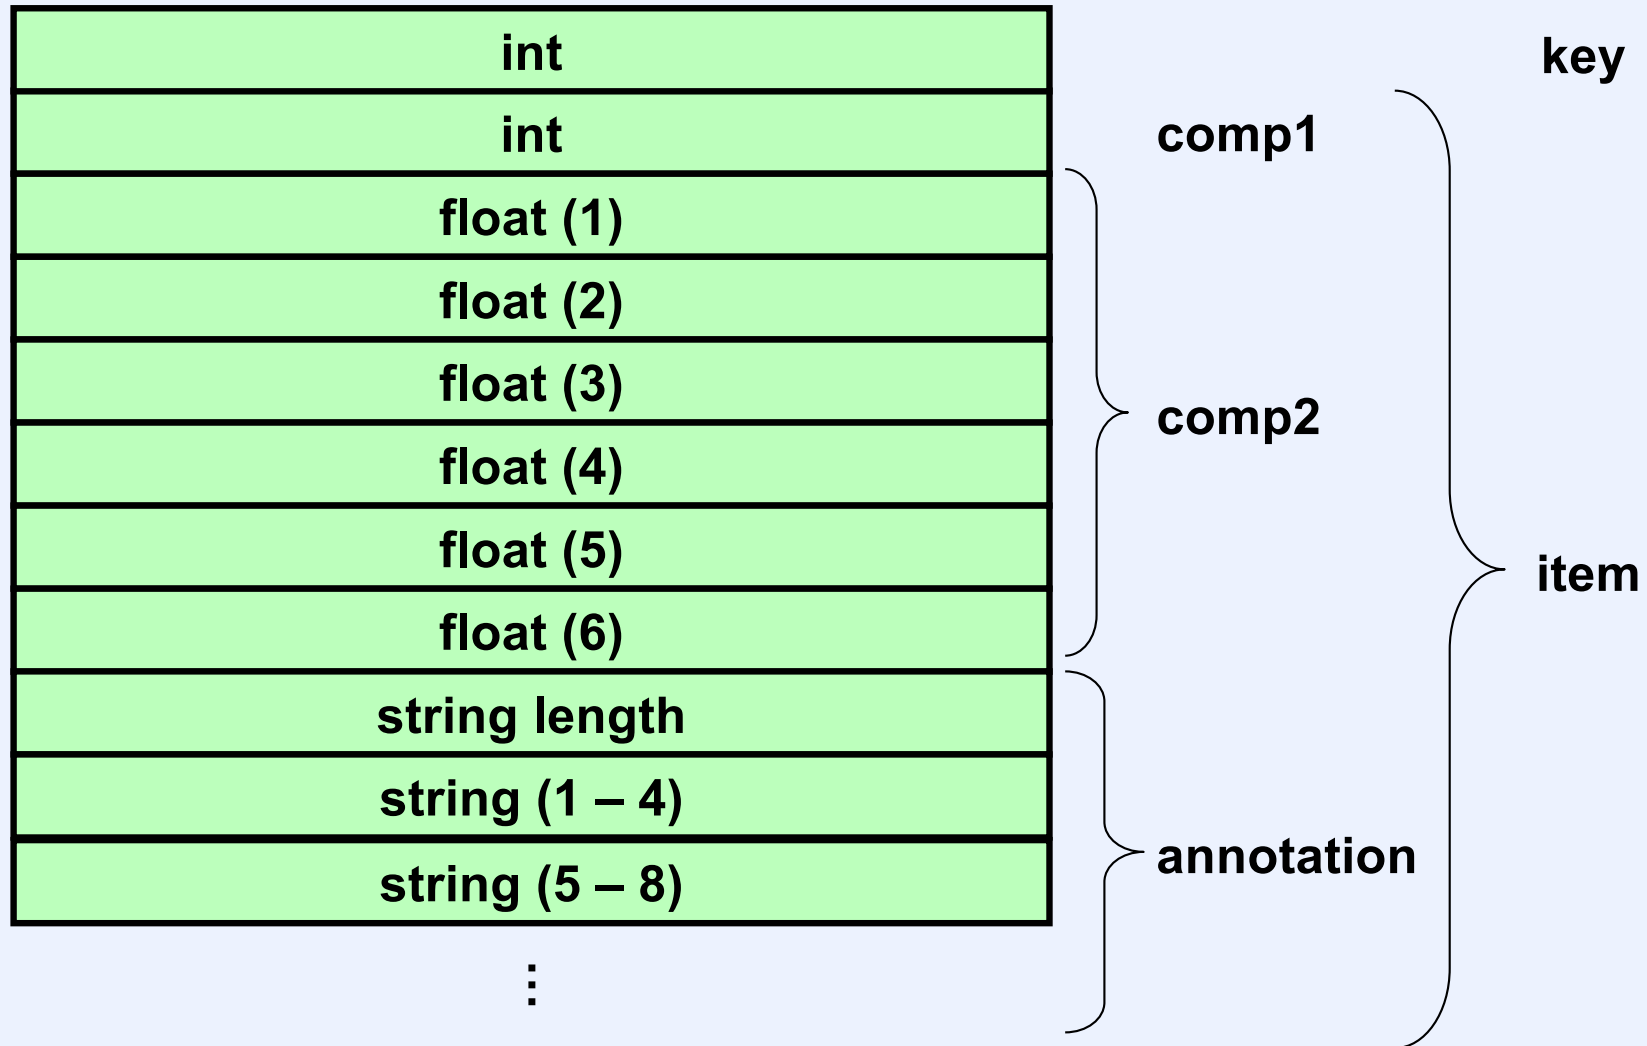
Placing a Call



Returning From the Call



Marshalled Arguments



Marshalled Linked List

		array length
0:	value_t	next: 1
1:	value_t	next: 2
2:	value_t	next: 3
3:	value_t	next: 4
4:	value_t	next: 5
5:	value_t	next: 6
6:	value_t	next: -1